
AWS SDKs and Tools

Reference Guide



AWS SDKs and Tools: Reference Guide

Copyright © Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

Table of Contents

AWS SDKs and Tools Reference Guide	1
Applicable to all SDKs and tools	1
Applicable to some SDKs and tools	1
Configuration	2
Additional topics in this section	2
Shared config and credentials files	2
Profiles	3
Format of the config file	3
Format of the credentials file	4
Example files	4
Location of the shared files	5
Supported SDKs and tools	6
Environment variables	7
How to set environment variables	7
Credentials and access	9
Topics in this section	9
Accounts, users, and roles	9
Create an account	9
Create an administrative user	10
Create an IAM user for development	10
Create an IAM role	11
SSO	12
MFA	12
Step 1: Creating an IAM role to delegate access to IAM users	12
Step 2: Creating an IAM user that assumes the role's permissions	13
Step 3: Managing a virtual MFA device for the IAM user	14
Step 4: Creating profiles to allow MFA	14
IAM roles for Amazon EC2 instances	15
Create an IAM role	15
Launch an Amazon EC2 instance and specify your IAM role	15
Connect to the EC2 instance	16
Run the sample application on the EC2 instance	16
IAM role from an IAM user	16
Scenario description	16
How to configure the profile	17
How to use the profile	17
Config and auth settings reference	19
Creating service clients	19
Precedence of settings	19
Config file settings list	20
Credentials file settings list	21
Environment variables list	21
Standardized credential providers	22
Credential provider chain	23
Assume role credentials	23
Container credentials	27
IMDS credentials	28
Process credentials	30
SSO credentials	32
Static credentials	34
Standardized features	35
Amazon EC2 instance metadata	36
Amazon S3 access points	37
Amazon S3 Multi-Region Access Points	38

AWS Region	39
AWS STS Regionalized endpoints	41
Endpoint discovery	42
General configuration	43
IMDS client	45
Retry behavior	47
Smart configuration defaults	50
Common Runtime	53
CRT dependencies	53
Maintenance and support	54
Maintenance policy	54
Overview	54
Versioning	54
SDK major version life-cycle	54
Dependency life-cycle	55
Communication methods	55
Version support matrix	56
Document history	58
AWS glossary	59

AWS SDKs and Tools Reference Guide

Applicable to all SDKs and tools

[AWS SDKs and Tools maintenance policy \(p. 54\)](#) covers the maintenance policy and versioning for AWS Software Development Kits (SDKs) and tools, including Mobile and Internet of Things (IoT) SDKs, and their underlying dependencies.

Applicable to some SDKs and tools

Many SDKs and tools share some common functionality, either through shared design specifications or through a shared library, such as the AWS Common Runtime (CRT) libraries.

This AWS SDKs and Tools Reference Guide is intended to be a base of information that is applicable to multiple SDKs and tools. The specific SDK or tool guide for the SDK or tool you are using should be used in addition to any information presented here. The following SDKs and tools have additional information represented in this guide:

- [AWS Cloud Development Kit \(CDK\) Developer Guide](#)
- [AWS Command Line Interface User Guide](#)
- [AWS Serverless Application Model Developer Guide](#)
- [AWS SDK for C++ Developer Guide](#)
- [AWS SDK for Go Developer Guide](#)
- [AWS SDK for Java Developer Guide](#)
- [AWS SDK for JavaScript Developer Guide](#)
- [AWS SDK for .NET Developer Guide](#)
- [AWS SDK for PHP Developer Guide](#)
- [AWS SDK for Python \(Boto3\) Getting Started](#)
- [AWS SDK for Ruby Developer Guide](#)
- [AWS Toolkit for Eclipse User Guide](#)
- [AWS Toolkit for JetBrains User Guide](#)
- [AWS Toolkit for Visual Studio User Guide](#)
- [AWS Toolkit for Visual Studio Code User Guide](#)
- [AWS Tools for Windows PowerShell User Guide](#)

This guide includes information regarding:

- [Configuration \(p. 2\)](#) – How to use the shared `config` and `credentials` files or environment variables for the authentication and configuration of an AWS SDK or tool.
- [Configuration and authentication settings reference \(p. 19\)](#) – Reference for all standardized settings available for authentication and configuration.
- [AWS Common Runtime \(CRT\) libraries \(p. 53\)](#) – Overview of the AWS Common Runtime (CRT) libraries that are available to almost all SDKs.

Configuration

With AWS SDKs and other AWS developer tools, such as the AWS Command Line Interface (AWS CLI), you can interact with AWS service APIs. Before attempting that, however, you must configure the SDK or tool with the information that it needs to perform the requested operation.

This information includes the following items:

- **Credentials information** that identifies who is calling the API. The credentials are used to encrypt the request to the AWS servers. Using this information, AWS confirms your identity and can retrieve permissions policies associated with it. Then it can determine what actions you're allowed to perform.
- **Other configuration details** that you use to tell the AWS CLI or SDK how to process the request, where to send the request (to which AWS service endpoint), and how to interpret or display the response.

Each SDK or tool supports multiple sources that you can use to supply the required credential and configuration information. Some sources are unique to the SDK or tool, and you must refer to the documentation for that tool or SDK for the details on how to use that method.

However, most of the AWS SDKs and tools support common settings from two primary sources (beyond the code itself):

- [Shared AWS config and credentials files \(p. 2\)](#) – The shared `config` and `credentials` files are the most common way that you can specify authentication and configuration to an AWS SDK or tool. Use these files to store settings that your tools and applications can use. The primary file is `config`, and you can put all settings into it. However, by default and as a security best practice, sensitive values such as secret keys are stored in a separate `credentials` file. You can protect those settings separately with different permissions. Using these files together, you can configure multiple groups of settings. Each group of settings is called a *profile*. When you use an AWS tool to invoke a command or use an SDK to invoke an AWS API, you can specify which profile, and thus which configuration settings, to use for that action. One of the profiles is designated as the `default` profile and is used automatically when you don't explicitly specify a profile to use. The settings that you can store in these files are documented in this reference guide.
- [Environment variables \(p. 7\)](#) – Some of the settings can alternatively be stored in the environment variables of your operating system. Although you can have only one set of environment variables in effect at a time, they are easily modified dynamically as your program runs and your requirements change.

Additional topics in this section

- [Location of the shared config and credentials files \(p. 5\)](#)
- [Shared AWS config and credentials files \(p. 2\)](#)
- [AWS SDKs and tools that use the shared config and credentials files \(p. 6\)](#)
- [Environment variables support \(p. 7\)](#)

Shared AWS config and credentials files

The shared AWS `config` and `credentials` files contain a set of profiles. A profile is a set of configuration values that can be referenced from the SDK/tool using its profile name. Configuration

values are attached to a profile in order to configure some aspect of the SDK/tool when that profile is used.

As a general rule, any value that you can place in the shared `credentials` file can alternatively be placed in the shared `config` file. The reverse isn't true; only a few settings can be placed in the `credentials` file. However, as a security best practice, we recommend that you keep any sensitive values, such as access key IDs and secret keys, in the separate `credentials` file. This way, you can provide separate permissions for each file, if necessary.

We recommend downloading these files from the AWS Management Console by following the instructions for [Managing access keys](#) in the *IAM User Guide*.

Both the shared `config` and `credentials` files are plaintext files that contain only ASCII characters (UTF-8 encoded). They take the form of what are generally referred to as [INI files](#).

Profiles

Settings within the shared `config` and `credentials` files are associated with a specific profile. With multiple profiles, you can create different settings configurations to apply in different scenarios.

The `[default]` profile contains the values that are used by an SDK or tool operation if a specific named profile is not specified. You can also create separate profiles that you can explicitly reference by name. Each named profile can have a different group of settings.

`[default]` is simply an unnamed profile. This profile is named `default` because it is the default profile used by the SDK if the user does not specify a profile. It does not provide inherited default values to other profiles. For example, if you set something in the `[default]` profile and you don't set it in a named profile, then the value isn't set when you use the named profile.

Optionally, set a named profile that you want to use through your SDK code or AWS CLI commands. Alternatively, you can use the environment variable `AWS_PROFILE` to specify which profile's settings to use.

Linux/macOS example of setting environment variables via command line:

```
export AWS_PROFILE="my_named_profile";
```

Windows example of setting environment variables via command line:

```
setx AWS_PROFILE "my_named_profile"
```

Format of the config file

The `config` file must be a plaintext file that uses the following format:

- Each section begins with the profile name in square brackets `[]`.
- All entries in a section take the general form of `setting-name=value`.

The following example shows a basic `config` file having a `[default]` profile. It sets the [region \(p. 39\)](#) global setting.

```
[default]
region = us-east-2
```

To create a named profile in the `config` file, create a section with a new header, similar to the following example. You must use the word `profile` and follow it with a unique name. You can use letters, numbers, hyphens (`-`), and underscores (`_`), but no spaces.

```
[profile developers]  
...settings...
```

Some settings have their own nested group of subsettings, such as the `s3` setting and subsettings in the following example. Associate the subsettings with the group by indenting them by one or more spaces.

```
[profile testers]  
region = us-west-2  
s3 =  
    max_concurrent_requests=10  
    max_queue_size=1000
```

Format of the credentials file

The following example shows a basic `credentials` file with a `[default]` profile. It sets the `aws_access_key_id` and `aws_secret_access_key` (p. 34) global settings.

```
[default]  
aws_access_key_id=AKIAIOSFODNN7EXAMPLE  
aws_secret_access_key=wJalrXUtnFEMI/K7MDENG/bPxrFcIYEXAMPLEKEY
```

The rules for the `credentials` file are generally identical to those for the `config` file, with the following exceptions:

- The section names don't begin with the word `profile`. Use only the profile name itself between square brackets.

```
[developers]  
...settings...
```

- You can store only a subset of settings and values in the `credentials` file. Generally, it's only those with values that would be considered "secrets" or sensitive, such as access key IDs and secret keys. The page for each setting in this guide states whether it can be stored in the `credentials` file or only in the `config` file.

Example files

In summary, each profile can have some settings in each file. The majority of settings go in the `config` file, while the sensitive information settings go in the `credentials` file.

The following example shows three profiles stored in these two files:

- `default` profile – Provides access by using the long-term credentials of an AWS Identity and Access Management (IAM) user. Tools or code that use this profile send requests to the US West (Oregon) Region (`us-west-2`). AWS CLI commands invoked using this profile output the results as JSON.
- `dev-user` profile – Uses the long-term credentials of a different IAM user. Tools or code that use this profile send requests to the US West (N. California) Region (`us-west-1`). AWS CLI commands invoked using this profile output the results as text.
- `developers` profile – Uses short-term credentials from assuming the specified role. It uses the long-term credentials in the `dev-user` source profile only to assume the role and retrieve the short-term

credentials for the role. Tools or code that use this profile send requests to the US West (Oregon) Region (us-west-2). AWS CLI commands invoked using this profile output the results as JSON. This profile doesn't store any of its values in the `credentials` file.

Contents of the `config` file

```
[default]
region = us-west-2
output = json

[profile dev-user]
region = us-west-1
output = text

[profile developers]
role_arn = arn:aws:iam::123456789012:role/developers
source_profile = dev-user
region = us-west-2
output = json
```

Contents of the `credentials` file

```
[default]
aws_access_key_id = AKIAIOSFODNN7EXAMPLE
aws_secret_access_key = wJalrXUtnFEMI/K7MDENG/bPxRfiCYEXAMPLEKEY

[dev-user]
aws_access_key_id = AKIAI44QH8DHBEXAMPLE
aws_secret_access_key = je7MtGbClwBF/2Zp9Utk/h3yCo8nvbEXAMPLEKEY
```

Location of the shared `config` and `credentials` files

The shared AWS `config` and `credentials` files are plaintext files that reside by default in a folder named `.aws` that is placed in the "home" folder on your computer.

On Linux and macOS, this is typically shown as `~/ .aws`. On Windows, it is `%USERPROFILE%\ .aws`.

Operating system	Default location of files
Linux and macOS	<code>~/ .aws/config</code> <code>~/ .aws/credentials</code>
Windows	<code>%USERPROFILE%\ .aws\config</code> <code>%USERPROFILE%\ .aws\credentials</code>

A `~/` or `~` followed by the file system's default path separator at the start of the path is resolved by checking, in order,

1. (All platforms) The `HOME` environment variable
2. (Windows platforms) The `USERPROFILE` environment variable

3. (Windows platforms) The `HOMEDRIVE` environment variable, prepended to the `HOME` environment variable (for example, `%HOMEDRIVE%HOME`)
4. (Optional per SDK or tool) An SDK or tool-specific home path resolution function or variable

When possible, if a user's home directory is specified at the start of the path (for example, `~username/`), it is resolved to the requested user name's home directory (for example, `/home/username/.aws/config`).

Changing the default location of these files:

The following environment variables can be set to change the location or name of these files from the default to a custom value:

- config file environment variable: `AWS_CONFIG_FILE`
- credentials file environment variable: `AWS_SHARED_CREDENTIALS_FILE`

Linux/macOS

You can specify an alternate location by running the following `export` commands on Linux or macOS.

```
$ export AWS_CONFIG_FILE=/some/file/path/on/the/system/config-file-name
$ export AWS_SHARED_CREDENTIALS_FILE=/some/other/file/path/on/the/system/credentials-file-name
```

Windows

You can specify an alternate location by running the following `setx` commands on Windows.

```
C:\> setx AWS_CONFIG_FILE c:\some\file\path\on\the\system\config-file-name
C:\> setx AWS_SHARED_CREDENTIALS_FILE c:\some\other\file\path\on\the\system\credentials-file-name
```

AWS SDKs and tools that use the shared config and credentials files

The following is a list of the AWS SDKs and other development tools that can use the shared `config` and `credentials` files to retrieve authentication and other configuration settings.

Each entry includes a link to that SDK's or tool's documentation where the use of these files is discussed in detail.

SDK or tool	Shared config and credentials topic
AWS Cloud Development Kit (CDK)	Getting started with AWS Cloud Development Kit (CDK)
AWS Command Line Interface (AWS CLI)	Configuring the AWS CLI
AWS Serverless Application Model	Setting up AWS credentials
AWS SDK for C++	Providing AWS credentials
AWS SDK for Go	Specifying credentials

SDK or tool	Shared config and credentials topic
AWS SDK for Java	Using credentials
AWS SDK for JavaScript	Getting your credentials
AWS SDK for .NET	Configure AWS credentials
AWS SDK for PHP	Using the AWS Credentials File and Credential Profiles
AWS SDK for Python (Boto3)	Configuration
AWS SDK for Ruby	Setting AWS Credentials
AWS Toolkit for Eclipse	Set up AWS credentials
AWS Toolkit for JetBrains	Setting AWS Credentials for the AWS Toolkit for JetBrains
AWS Toolkit for Visual Studio	Providing AWS Credentials
AWS Toolkit for Visual Studio Code	Setting Up Your AWS Credentials
AWS Tools for PowerShell	Getting Started with AWS Tools for PowerShell

Environment variables support

Environment variables provide another way to specify configuration options and credentials, and can be useful for scripting or temporarily setting a named profile as the default.

Precedence of options

- If you specify a setting by using its environment variable, it overrides any value loaded from a profile in the shared `AWS config` and `credentials` files.
- If you specify a setting by using a parameter on the AWS CLI command line, it overrides any value from either the corresponding environment variable or a profile in the configuration file.

How to set environment variables

The following examples show how you can configure environment variables for the default user.

Linux, macOS, or Unix

```
$ export AWS_ACCESS_KEY_ID=AKIAIOSFODNN7EXAMPLE
$ export AWS_SECRET_ACCESS_KEY=wJalrXUtnFEMI/K7MDENG/bPxrFiCYEXAMPLEKEY
$ export AWS_DEFAULT_REGION=us-west-2
```

Setting the environment variable changes the value used until the end of your shell session, or until you set the variable to a different value. You can make the variables persistent across future sessions by setting them in your shell's startup script.

Windows Command Prompt

```
C:\> setx AWS_ACCESS_KEY_ID AKIAIOSFODNN7EXAMPLE
C:\> setx AWS_SECRET_ACCESS_KEY wJalrXUtnFEMI/K7MDENG/bPxrFiCYEXAMPLEKEY
```

```
C:\> setx AWS_DEFAULT_REGION us-west-2
```

Using `set` to set an environment variable changes the value used until the end of the current Command Prompt session, or until you set the variable to a different value. Using `setx` to set an environment variable changes the value used in both the current Command Prompt session and all Command Prompt sessions that you create after running the command. It does **not** affect other command shells that are already running at the time you run the command.

PowerShell

```
PS C:\> $Env:AWS_ACCESS_KEY_ID="AKIAIOSFODNN7EXAMPLE"
PS C:\> $Env:AWS_SECRET_ACCESS_KEY="wJalrXUtnFEMI/K7MDENG/bpɁrfiCYEXAMPLEKEY"
PS C:\> $Env:AWS_DEFAULT_REGION="us-west-2"
```

If you set an environment variable at the PowerShell prompt as shown in the previous examples, it saves the value for only the duration of the current session. To make the environment variable setting persistent across all PowerShell and Command Prompt sessions, store it by using the **System** application in **Control Panel**. Alternatively, you can set the variable for all future PowerShell sessions by adding it to your PowerShell profile. See the [PowerShell documentation](#) for more information about storing environment variables or persisting them across sessions.

Credentials and access

The *Amazon Web Services General Reference* has foundational basics on:

- [AWS account root user credentials and IAM user credentials](#) – Root user credentials vs IAM credentials, tasks that require root user credentials
- [Understanding and getting your AWS credentials](#) – Access key options and management practices for both console and programmatic access
- [Your AWS account identifiers](#) – Finding and understanding your AWS account ID

The *IAM User Guide* is always the best resource for securely controlling access to AWS resources:

- [Security best practices in IAM](#) – Security recommendations to follow when developing AWS applications in keeping with our [shared-responsibility model](#)

Topics in this section

- [Setting up AWS accounts, users, and roles \(p. 9\)](#) for developing with SDKs and tools
- Using [Single sign-on \(SSO\) \(p. 12\)](#) in your application
- Using [multi-factor authentication \(MFA\) \(p. 12\)](#) in your application
- Using [IAM roles for Amazon EC2 instances \(p. 15\)](#)
- Using [an IAM role assumed as an IAM user \(p. 16\)](#)

Setting up AWS accounts, users, and roles

The *IAM User Guide* is always the best resource for securely controlling access to AWS resources. [Identities \(users, groups, and roles\)](#) in the *IAM User Guide* covers these foundational concepts in greater depth.

Roles allow you to gain access to AWS using your existing identities. You can federate in from an external identity provider, like Active Directory or Okta, or you can assume a role from a different user or role within AWS. You can use roles to build your applications securely by assuming different roles for different job functions, and keeping roles with access to sensitive data separated from your other identities.

Users are meant for when you have no other identity to connect to AWS with. If you need an entry point to AWS but you don't have a separate identity provider, you can create a user to sign into AWS with so that you don't have to use the AWS account root user. Users have long-lived credentials that you should rotate regularly. If you need to use a user in AWS, we recommend using the user only to sign in to AWS and then to assume a role that can perform your job function.

Although users are a good way to get started using AWS services, we recommend all developers consider setting up single sign-on for access. For more information, see [Single sign-on \(SSO\) \(p. 12\)](#).

Create an account

To create an AWS account, see [How do I create and activate a new Amazon Web Services account?](#)

Create an administrative user

Avoid using your AWS account root user (the initial account you create) to access the management console and services. Instead, create an administrative user account, as explained in [Creating your first IAM admin user and group](#) in the *IAM User Guide*.

After you create the administrative user account and record the login details, **sign out as root user** and sign back in using the administrative account.

The administrative user account still grants full access to most things - it is not appropriate for doing SDK development on AWS or for running applications on AWS. Thus, you need to create user accounts and service roles that are appropriate for these tasks. The specific user accounts and service roles that you create, and the way in which you use them, depend on the requirements of your applications.

Create an IAM user for development

User accounts are a good way to get started developing applications to use AWS services. [IAM users](#) are detailed in the *IAM User Guide*.

To programmatically interface with AWS, one option is to create an IAM user and obtain credentials for that user. Then make them available to the SDK/tool in your development environment by saving them to the shared AWS `credentials` file.

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. Choose **Users**, and then choose **Add users**.
3. Provide a user name.
4. Under **Select AWS access type**, select **Programmatic access**, and then choose **Next: Permissions**.
5. Choose **Attach existing policies directly**, and then select the [appropriate policies](#) for the AWS services that your application will use.

For example, if you wanted your application to have full access to the Amazon S3 service: In **Search**, enter **s3**, and then select **AmazonS3FullAccess** as the policy to attach to this user. For best practices regarding policies, see [Grant least privilege](#) in the *IAM User Guide*.

Warning

The **AdministratorAccess** policy enables read and write permissions to almost everything in your account. We recommend choosing a more restrictive policy for development.

6. Choose **Next: Tags** and enter any tags you want.

You can find information about tags in [Control access using AWS resource tags](#) in the *IAM User Guide*.

7. Choose **Next: Review**, and then choose **Create user**.
8. On the *Success* screen, choose **Download .csv**.

The downloaded file contains the Access Key ID and the Secret Access Key for this IAM user.

Note

You will not have another opportunity to download or copy the Secret Access Key.

Treat your Secret Access Key as a password; save in a trusted location and do not share it.

Warning

Use appropriate security measures to keep these credentials safe and rotated. See [Security best practices in IAM](#) in the *IAM User Guide*.

You now have the credentials for this user. For an AWS SDK or tool to be able to use these credentials, they can be put into the shared `credentials` file.

9. Create or open the shared `credentials` file. This file is `~/.aws/credentials` on Linux and macOS systems, and `%USERPROFILE%\.aws\credentials` on Windows. See [Location of Credentials Files](#) for more information.
10. Add the following text to the shared `credentials` file, but replace the example ID value and example key value with the ones you obtained earlier.

```
[default]
aws_access_key_id = AKIAIOSFODNN7EXAMPLE
aws_secret_access_key = wJalrXUtnFEMI/K7MDENG/bPxrFiCYEXAMPLEKEY
```

11. Save the file.

Comprehensive coverage of navigating all the options in the **Add users** console wizard is at [Creating IAM users \(console\)](#) in the *IAM User Guide*

Note that there are other ways of storing credentials (such as the [SDK Store](#) available to the AWS SDK for .NET), but the shared `credentials` file is the most common.

Create an IAM role

Roles represent a set of AWS permission policies that are intended to be assumable by anyone who needs it. [IAM roles](#) are detailed in the *IAM User Guide*.

The process for creating a service role varies depending on the situation, but is essentially the following.

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. Choose **Roles**, and then choose **Create role**.
3. Choose **AWS service**.
4. Under **Choose a use case**, select the service you want to allow to call services on your behalf.

For example, if you were working with the Amazon EC2 service, then select **EC2** as your use case.

5. Choose **Next: Permissions**, and select the [appropriate policies](#) for the AWS services that your application will use.

For example, if you wanted your application to have full access to the Amazon EC2 service: In **Search**, enter **ec2**, and then select **AmazonEC2FullAccess** as the policy to attach. For best practices regarding policies, see [Grant least privilege](#) in the *IAM User Guide*.

Warning

The **AdministratorAccess** policy enables read and write permissions to almost everything in your account. We recommend choosing a more restrictive policy for development.

6. Choose **Next: Tags** and enter any tags you want.

You can find information about tags in [Control access using AWS resource tags](#) in the *IAM User Guide*.

7. Choose **Next: Review** and provide a **Role name** and **Role description**. Then choose **Create role**.

Comprehensive coverage of navigating all the options in the **Create role** console wizard is at [Creating a role for an AWS service \(console\)](#) in the *IAM User Guide*

Single sign-on (SSO)

AWS Single Sign-On (AWS SSO) is a cloud-based single sign-on service you can use to centrally manage SSO access to all of your AWS accounts and cloud applications. For full details, see the [AWS Single Sign-On User Guide](#).

To use SSO credentials with your AWS SDK, you'll do the following:

- **Set up AWS SSO** – This includes choosing your identity source and setting up AWS SSO access to your AWS accounts. For steps on this setup, see [Getting started](#) in the *AWS Single Sign-On User Guide*.
- **Set up an AWS SSO profile** – Define a named profile in the shared AWS config file that is used to retrieve temporary credentials for your AWS account. For details on which settings to configure for SSO, see [SSO credentials \(p. 32\)](#).
- **Get an SSO Token** – To use an SSO-enabled profile, you must generate a temporary token. This can be done either programmatically or by running the `aws sso login` command from the AWS CLI.

In the AWS CLI, for example, to sign in with a named profile instead of the default profile, pass the profile name using the `--profile` parameter:

```
aws sso login --profile my-sso-profile
```

To learn how to get and renew tokens programmatically, see your specific AWS SDK developer guide.

Using multi-factor authentication (MFA)

Multi-factor authentication (MFA) offers increased security because it requires users to provide unique authentication from an AWS supported MFA mechanism in addition to their regular sign-in credentials when they access AWS websites or services.

AWS supports a range of both virtual and hardware devices for MFA authentication. The example that's documented here is a virtual MFA device that's enabled by a smartphone application. For more information on MFA device options, see [Using multi-factor authentication \(MFA\) in AWS](#) in the *IAM User Guide*.

Step 1: Creating an IAM role to delegate access to IAM users

This task uses [role delegation](#) to allow an IAM role to delegate permissions to an IAM user. First, you define an IAM role that requires signing in with MFA. You also attach policies to that role that grant permissions to access specific AWS services. Next, you create an IAM user that has no permissions to start with. But you then attach to that user a policy that includes the `AssumeRole` operation, which delegates all the role's permissions to the user.

1. Go to the IAM console at <https://console.aws.amazon.com/iam>.
2. Choose **Roles** in the navigation bar, and then choose **Create Role**.
3. In the **Create role** page, choose **Another AWS account**.
4. Enter your required **Account ID** and mark the **Require MFA** check box.

Note

To find your 12-digit account number (ID), go to the navigation bar in the console, and then choose **Support, Support Center**.

5. Choose **Next: Permissions**.
6. Attach existing policies to your role or create a new policy for it. The policies that you choose on this page determine which AWS services the IAM user can access.
7. After attaching policies, choose **Next: Tags** for the option of adding IAM tags to your role. Then choose **Next: Review** to continue.
8. In the **Review** page, enter a required **Role name** (*my-role*, for example). You can also add an optional **Role description**.
9. Choose **Create role**.
10. When the confirmation message displays ("The role **my-role** has been created", for example), choose the name of the role in the message.
11. In the **Summary** page, choose the copy icon to copy the **Role ARN** and paste it into a file. (You need this ARN when configuring the IAM user to assume the role.).

Step 2: Creating an IAM user that assumes the role's permissions

In this step, you first create the IAM user without permissions. Then you create an in-line policy that allows the user to assume the role (and that role's permissions) that you created in the previous step.

To create the IAM user

1. Go to the IAM console at <https://console.aws.amazon.com/iam>.
2. Choose **Users** in the navigation bar and then choose **Add user**.
3. In the **Add user** page, enter a required **User name** (*my-user*, for example) and mark the **Programmatic access** check box.
4. Choose **Next: Permissions**, **Next: Tags**, and **Next: Review** to move through the next pages. You're not adding permissions at this stage because the user is going to assume the role's permissions.
5. In the **Review** page, you're informed that **This user has no permissions**. Choose **Create user**.
6. In the **Success** page, choose **Download .csv** to download the file containing the access key ID and secret access key. (You need both when defining the user's profile in the shared AWS `credentials` file.)
7. Choose **Close**.

To add a policy to allow the IAM user to assume the role

1. In the **Users** page of the IAM console, choose the IAM user you've just created (*my-user*, for example).
2. In the **Permissions** tab of the **Summary** page, choose **Add inline policy**.
3. In the **Create policy** page, choose **Choose a service**, enter **STS** in **Find a service**, and then choose **STS** from the results.
4. For **Actions**, start entering the term *AssumeRole*. Mark the **AssumeRole** check box when it appears.
5. In the **Resource section**, ensure **Specific** is selected, and click **Add ARN** to restrict access.
6. In the **Add ARN(s)** dialog box, for the **Specify ARN for role** add the ARN of the role you that you created in Step 1.

After you add the role's ARN, the trusted account and role name associated with that role are displayed in **Account** and **Role name with path**.

7. Choose **Add**.

8. Back in the **Create policy** page, choose **Specify request conditions (optional)**, mark the **MFA required** check box, and then choose **close** to confirm..
9. Choose **Review policy**
10. In **Review policy** page, enter a **Name** for the policy, and then choose **Create policy**.

The **Permissions** tab displays the new inline policy attached directly to IAM user.

Step 3: Managing a virtual MFA device for the IAM user

1. Download and install a virtual MFA application to your smartphone.

For a list of supported applications, see the [Multi-factor Authentication](#) resource page.

2. In the IAM console, choose **Users** from the navigation bar and then choose the user that's assuming a role (*my-user*, in this case).
3. In the **Summary** page, choose the **Security credentials** tab, and for **Assigned MFA device** choose **Manage**.
4. In the **Manage MFA device** pane, choose **Virtual MFA device**, and then choose **Continue**.
5. In the **Set up virtual MFA device** pane, choose **Show QR code** and then scan the code using the virtual MFA application that you installed on your smartphone.
6. After you scan the QR code, the virtual MFA application generates one-time MFA codes. Enter two consecutive MFA codes in **MFA code 1** and **MFA code 2**.
7. Choose **Assign MFA**.
8. Back in the **Security credentials** tab for the user, copy the ARN of the new **Assigned MFA device**.

The ARN includes your 12-digit account ID and the format is similar to the following:
`arn:aws:iam::123456789012:mfa/my-user`. You need this ARN when defining the MFA profile in the next step.

Step 4: Creating profiles to allow MFA

In this step, you create the profiles that allow users to use MFA when accessing AWS services.

The profiles that you create include three pieces of information that you've copied and stored during the previous steps:

- Access keys (access key ID and secret access key) for the IAM user
- ARN of the role that's delegating permissions to the IAM user
- ARN of the virtual MFA device that's assigned to the IAM user

In the shared AWS `config` and `credentials` files or SDK Store that contain your AWS credentials, add the following entries:

```
[my-user]
aws_access_key_id = AKIAIOSFODNN7EXAMPLE
aws_secret_access_key = wJalrXUtnFEMI/K7MDENG/bPxRfiCYEXAMPLEKEY

[mfa]
source_profile = my-user
role_arn = arn:aws:iam::111111111111:role/my-role
```

```
mfa_serial = arn:aws:iam::111111111111:mfa/my-user
```

There are two profiles defined in the example provided:

- [my-user] profile includes the access key and secret access key that were generated and saved when you created the IAM user in Step 2.
 - Details on using these two settings can be found at [Static credentials \(p. 34\)](#).
- [mfa] profile defines how multi-factor authentication is supported. There are three entries:
 - source_profile: Specifies the profile whose credentials are used to assume the role specified by this role_arn setting in this profile. In this case, it's the my-user profile.
 - role_arn: Specifies the Amazon Resource Name (ARN) of the IAM role that you want to use to perform operations requested using this profile. In this case, it's the ARN for the role you created in Step 1.
 - mfa_serial: Specifies the identification or serial number of the MFA device that the user must use when assuming a role. In this case, it's the ARN of the virtual device you set up in Step 3.
- Details on using these three settings can be found at [Assume role credentials \(p. 23\)](#).

Using IAM roles for Amazon EC2 instances

This example covers setting up an IAM role with Amazon S3 access to use in your application deployed to an Amazon EC2 instance.

Create an IAM role

Create an IAM role that grants read-only access to Amazon S3.

1. Open the [IAM console](#).
2. In the navigation pane, choose **Roles**, then **Create New Role**.
3. On the **Select Role Type** page, under **AWS-service Roles**, choose **Amazon EC2**.
4. On the **Attach Policy** page, choose **Amazon S3 Read Only Access** from the policy list, then choose **Next Step**. Enter a name for the role, then select **Next Step**. *Remember this name because you'll need it when you launch your Amazon EC2 instance.*
5. On the **Review** page, choose **Create Role**.

Launch an Amazon EC2 instance and specify your IAM role

You can launch an Amazon EC2 instance with an IAM role using the Amazon EC2 console.

Follow the directions to launch an instance in the [Amazon EC2 User Guide for Linux Instances](#) or the [Amazon EC2 User Guide for Windows Instances](#).

When you reach the **Review Instance Launch** page, select **Edit instance details**. In **IAM role**, choose the IAM role that you created previously. Complete the procedure as directed.

Note

You need to create or use an existing security group and key pair to connect to the instance.

With this IAM and Amazon EC2 setup, you can deploy your application to the Amazon EC2 instance and it will have read access to the Amazon S3 service.

Connect to the EC2 instance

Connect to the EC2 instance so that you can transfer the sample application to it and then run the application. You'll need the file that contains the private portion of the key pair you used to launch the instance; that is, the PEM file.

You can do this by following the connect procedure in the [Amazon EC2 User Guide for Linux Instances](#) or the [Amazon EC2 User Guide for Windows Instances](#). When you connect, do so in such a way that you can transfer files from your development machine to your instance.

If you're using an AWS Toolkit, you can often also connect to the instance by using the Toolkit. For more information, see the specific user guide for the Toolkit you use.

Run the sample application on the EC2 instance

1. Copy the application files from your local drive to your instance.

For information about how to transfer files to your instance see the [Amazon EC2 User Guide for Linux Instances](#) or the [Amazon EC2 User Guide for Windows Instances](#).

2. Start the application and verify that it runs with the same results as on your development machine.
3. (Optional) Verify that the application uses the credentials provided by the IAM role.
 - a. Open the [Amazon EC2 console](#).
 - b. Select the instance and detach the IAM role through **Actions, Instance Settings, Attach/Replace IAM Role**.
 - c. Run the application again and confirm that it returns an authorization error.

Using an IAM role assumed as an IAM user

This example shows how to configure the shared AWS `config` and `credentials` files to support logging in to an AWS SDK or developer tool using an IAM role. The SDK or tool assumes the role using the credentials of a separate IAM user.

Scenario description

This scenario requires that you have two entities created in IAM:

- An IAM user, that in this example we call `UserAlpha`. This user has IAM policy permissions that enable it to perform only the `sts:AssumeRole` operation. The following JSON policy document describes this user in IAM policy language. To better understand the components of IAM policy language, see [IAM JSON policy elements reference](#) in the *IAM User Guide*. However, when getting started, it is far more common to simply use the AWS Management Console to graphically create users, roles, etc. rather than writing them manually.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "sts:AssumeRole",
      "Resource": "*"
    }
  ]
}
```

- An IAM role, that in this example we call `RoleBeta`. This role has a trust policy that enables all users in the account to assume the role.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "AWS": "arn:aws:iam::231179739868:root"
      },
      "Action": "sts:AssumeRole",
      "Condition": {}
    }
  ]
}
```

It also has IAM permission policies that enable it to perform any task in Amazon S3. When getting started, it is far more common to simply use the AWS Management Console and assign AWS managed policies.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "s3:*",
      "Resource": "*"
    }
  ]
}
```

How to configure the profile

The following example shared AWS `config` file and shared AWS `credentials` file files show how you can configure SDK or AWS developer tool access using `RoleBeta`.

Contents of shared `config` file

```
[profile UserAlpha]
region = us-west-2
output = json

[profile RoleBeta]
source_profile = UserAlpha
role_arn = arn:aws:iam::123456789012:role/RoleBeta
```

Contents of shared `credentials` file

```
[UserAlpha]
aws_access_key_id = AKIAIOSFODNN7EXAMPLE
aws_secret_access_key = wJalrXUtnFEMI/K7MDENG/bPxRfiCYEXAMPLEKEY
```

How to use the profile

As an example, you can run the following AWS CLI command to list the Amazon S3 buckets available in the account. The AWS CLI sees that `RoleBeta` profile references the `source_profile` `UserAlpha`. It

looks up `UserAlpha`'s access key and secret key, and uses them to call the `sts:AssumeRole` operation on the ARN of `RoleBeta`. That operation returns short-term credentials for `RoleBeta` that the AWS CLI uses to call the `s3:ListBuckets` operation.

```
$ aws s3api list-buckets --profile RoleBeta
{
  "Buckets": [
    {
      "Name": "my-first-bucket",
      "CreationDate": "2018-08-31T07:46:02+00:00"
    },
    {
      "Name": "my-second-bucket",
      "CreationDate": "2019-09-17T19:17:31+00:00"
    },
    {
      "Name": "my-third-bucket",
      "CreationDate": "2018-06-12T23:18:08+00:00"
    }
  ],
  "Owner": { ...truncated... }
}
```

Configuration and authentication settings reference

- [Standardized credential providers \(p. 22\)](#) – Common credential providers standardized across multiple SDKs.
- [Standardized features \(p. 35\)](#) – Common features standardized across multiple SDKs.

SDKs provide language-specific APIs for AWS services. They take care of some of the heavy lifting necessary in successfully making API calls, including authentication, retry behavior, and more. To do this, the SDKs have flexible strategies to obtain credentials to use for your requests, to maintain settings to use with each service, and to obtain values to use for global settings.

Creating service clients

To programmatically access AWS services, SDKs use a client class/object for each AWS service. For example, if your application needs to access Amazon EC2, your application creates an Amazon EC2 client object to interface with that service. You then use the service client to make requests to that AWS service. A service client object is immutable, so you must create a new client for each service to which you make requests and for making requests to the same service using a different configuration.

Precedence of settings

Global settings configure features, credential providers, and other functionality that are supported by most SDKs and have a broad impact across AWS services. All SDKs have a series of places (or sources) that they check in order to find a value for global settings. The following is the setting lookup precedence:

1. Any explicit setting set in the code or on a service client itself takes precedence over anything else.
 - Some settings can be set on a per-operation basis, and can be changed as needed for each operation that you invoke. For the AWS CLI or AWS Tools for PowerShell, these take the form of per-operation parameters that you enter on the command line. For an SDK, explicit assignments can take the form of a parameter that you set when you instantiate an AWS service client or configuration object, or sometimes when you call an individual API.
2. Java/Kotlin only: Sometimes there is a JVM system property associated with the setting. If it's set, that value is used to configure the client.
3. The environment variable is checked. If it's set, that value is used to configure the client.
4. The SDK checks the shared `credentials` file and then the shared `config` file. If the setting is present, the SDK uses it. The `AWS_PROFILE` environment variable or the `aws.profile` system property can be used to specify which profile that the SDK loads.
5. Any default value provided by the SDK code base is used last.

Note

If a setting exists in both the `config` file and the `credentials` file for the same profile, the value in the `credentials` file is used instead of the value in the `config` file.

Note

Some SDKs and tools might check in a different order. Also, some SDKs and tools support other methods of storing and retrieving parameters. For example, the AWS SDK for .NET supports an additional source called the [SDK Store](#). For more information about providers that are unique to a SDK or tool, see the [documentation for that SDK or tool](#) (p. 6).

The order determines which methods take precedence and override others. For example, if you set up a default profile in the shared config file, it's only found and used after the SDK or tool checks the other places first. This means that if you put a setting in the credentials file, it is used instead of one found in the config file. If you configure an environment variable with a setting and value, it would override that setting in both the credentials and config files. And finally, a setting on the individual operation (AWS CLI command-line parameter or API parameter) or in code would override all other values for that one command.

Config file settings list

The settings listed in the following table can be assigned in the shared AWS config file. They are global and affect all AWS services.

Setting name	Details	
api_versions	General configuration settings (p. 43)	
aws_access_key_id	Static credentials (p. 34)	
aws_secret_access_key	Static credentials (p. 34)	
aws_session_token	Static credentials (p. 34)	
ca_bundle	General configuration settings (p. 43)	
credential_process	Process credentials (p. 30)	
credential_source	Assume role credentials (p. 23)	
defaults_mode	Smart configuration defaults (p. 50)	
duration_seconds	Assume role credentials (p. 23)	
ec2_metadata_service_endpoint	IMDS endpoints (p. 28)	
ec2_metadata_service_endpoint_mode	IMDS endpoints (p. 28)	
endpoint_discovery_enabled	Endpoint discovery (p. 42)	
external_id	Assume role credentials (p. 23)	
max_attempts	Retry behavior (p. 47)	
metadata_service_num_attempts	Amazon EC2 instance metadata (p. 36)	
metadata_service_timeout	Amazon EC2 instance metadata (p. 36)	
mfa_serial	Assume role credentials (p. 23)	
parameter_validation	General configuration settings (p. 43)	
region	AWS Region (p. 39)	

Setting name	Details	
retry_mode	Retry behavior (p. 47)	
role_arn	Assume role credentials (p. 23)	
role_session_name	Assume role credentials (p. 23)	
s3_disable_multiregion_access_points	Amazon S3 Multiregion Access Points (p. 38)	
s3_use_arn_region	Amazon S3 access points (p. 37)	
source_profile	Assume role credentials (p. 23)	
sso_account_id	SSO credentials (p. 32)	
sso_region	SSO credentials (p. 32)	
sso_role_name	SSO credentials (p. 32)	
sso_start_url	SSO credentials (p. 32)	
sts_regional_endpoints	STS regionalized endpoints (p. 41)	
web_identity_token_file	Assume role credentials (p. 23)	

Credentials file settings list

The settings listed in the following table can be assigned in the shared AWS credentials file. They are global and affect all AWS services.

Setting name	Details	
aws_access_key_id	Static credentials (p. 34)	
aws_secret_access_key	Static credentials (p. 34)	
aws_session_token	Static credentials (p. 34)	

Environment variables list

Environment variables supported by most SDKs are listed in the following table. They are global and affect all AWS services.

Setting name	Details	
AWS_ACCESS_KEY_ID	Static credentials (p. 34)	
AWS_CA_BUNDLE	General configuration settings (p. 43)	
AWS_CONFIG_FILE	Location of the shared config and credentials files (p. 5)	
AWS_CONTAINER_AUTHORIZATION_TOKEN	Container credentials (p. 27)	

Setting name	Details	
AWS_CONTAINER_CREDENTIALS_FULL_URI	Container credentials (p. 27)	
AWS_CONTAINER_CREDENTIALS_RELATIVE_URI	Container credentials (p. 27)	
AWS_DEFAULTS_MODE	Smart configuration defaults (p. 50)	
AWS_EC2_METADATA_SERVICE_ENDPOINT	IMDS endpoint (p. 28)	
AWS_EC2_METADATA_SERVICE_ENDPOINT_MODE	IMDS endpoint (p. 28)	
AWS_ENABLE_ENDPOINT_DISCOVERY	Endpoint discovery (p. 42)	
AWS_MAX_ATTEMPTS	Retry behavior (p. 47)	
AWS_METADATA_SERVICE_NUM_ATTEMPTS	Amazon EC2 instance metadata (p. 36)	
AWS_METADATA_SERVICE_TIMEOUT	Amazon EC2 instance metadata (p. 36)	
AWS_PROFILE	Shared AWS config and credentials files (p. 2)	
AWS_REGION	AWS Region (p. 39)	
AWS_RETRY_MODE	Retry behavior (p. 47)	
AWS_S3_DISABLE_MULTIREGION_ACCESS_POINTS	Amazon S3 access points (p. 38)	
AWS_S3_USE_ARN_REGION	Amazon S3 access points (p. 37)	
AWS_SECRET_ACCESS_KEY	Static credentials (p. 34)	
AWS_SESSION_TOKEN	Static credentials (p. 34)	
AWS_SHARED_CREDENTIALS_FILE	Location of the shared config and credentials files (p. 5)	
AWS_STS_REGIONAL_ENDPOINTS	Regionalized endpoints (p. 41)	

Standardized credential providers

Many credential providers have been standardized to consistent defaults and to work the same way across many SDKs. This consistency increases productivity and clarity when coding across multiple SDKs. All settings can be overridden in code. For details, see your specific SDK API.

Important

Not all SDKs support all providers, or even all aspects within a provider.

Topics

- [Credential provider chain \(p. 23\)](#)
- [Assume role credentials \(p. 23\)](#)
- [Container credentials \(p. 27\)](#)
- [IMDS credentials \(p. 28\)](#)
- [Process credentials \(p. 30\)](#)
- [SSO credentials \(p. 32\)](#)
- [Static credentials \(p. 34\)](#)

Credential provider chain

All SDKs have a series of places (or sources) that they check in order to find valid credentials to use to make a request to an AWS service. After valid credentials are found, the search is stopped. This systematic search is called the default credential provider chain. Although the distinct chain used by each SDK varies, they most often include sources such as the following:

- Static credentials (such as `AWS_ACCESS_KEY_ID`). For more information, see [Static credentials \(p. 34\)](#).
- Web identity token from AWS Security Token Service (AWS STS). For more information, see [Assume role credentials \(p. 23\)](#).
- AWS Single Sign-On. For more information, see [SSO credentials \(p. 32\)](#).
- Trusted entity provider (such as `AWS_ROLE_ARN`). For more information, see [Assume role credentials \(p. 23\)](#).
- Amazon Elastic Container Service (Amazon ECS) credentials. For more information, see [Container credentials \(p. 27\)](#).
- Custom credential provider. For more information, see [Process credentials \(p. 30\)](#).
- Amazon Elastic Compute Cloud (Amazon EC2) instance profile credentials (IMDS credential provider). For more information, see [IMDS credentials \(p. 28\)](#).

For each step in the chain, there are a variety of ways to assign setting values. Setting values specified in code always take precedence, but there are also [Environment variables \(p. 7\)](#) and the [Shared AWS config and credentials files \(p. 2\)](#). For more information, see [Precedence of settings \(p. 19\)](#).

Assume role credentials

Assuming a role involves using a set of temporary security credentials that you can use to access AWS resources that you might not normally have access to. These temporary credentials consist of an access key ID, a secret access key, and a security token. Typically, you assume a role within your account or for cross-account access.

To learn more about IAM roles, see [Using IAM roles](#) in the *IAM User Guide*.

To learn more about assuming a role, see [AssumeRole](#) in the *AWS Security Token Service API Reference*.

Configure this functionality by using the following:

credential_source - shared AWS config file setting

Used within Amazon EC2 instances or containers to specify where the SDK or development tool can find credentials that have permission to assume the role that you specify with the `role_arn` parameter.

Default value: None

Valid values:

- **Environment** – Specifies that the SDK or tool is to retrieve source credentials from the environment variables `AWS_ACCESS_KEY_ID` and `AWS_SECRET_ACCESS_KEY` ([p. 34](#)).
- **Ec2InstanceMetadata** – Specifies that the SDK or tool is to use the [IAM role attached to the EC2 instance profile](#) to get source credentials.
- **EcsContainer** – Specifies that the SDK or tool is to use the [IAM role attached to the ECS container](#) to get source credentials.

You cannot specify both `credential_source` and `source_profile` in the same profile.

Example of setting this in a config file to indicate that credentials should be sourced from Amazon EC2:

```
credential_source = Ec2InstanceMetadata
role_arn = arn:aws:iam::123456789012:role/my-role-name
```

duration_seconds - shared AWS config file setting

Specifies the maximum duration of the role session, in seconds.

This setting applies only when the profile specifies to assume a role.

Default value: 3600 seconds (one hour)

Valid values: The value can range from 900 seconds (15 minutes) up to the maximum session duration setting configured for the role (which can be a maximum of 43200 seconds, or 12 hours). For more information, see [View the Maximum Session Duration Setting for a Role](#) in the *IAM User Guide*.

Example of setting this in a config file:

```
duration_seconds = 43200
```

external_id - shared AWS config file setting

Specifies a unique identifier that is used by third parties to assume a role in their customers' accounts.

This setting applies only when the profile specifies to assume a role and the trust policy for the role requires a value for ExternalId. The value maps to the ExternalId parameter that is passed to the AssumeRole operation when the profile specifies a role.

Default value: None.

Valid values: See [How to use an External ID When Granting Access to Your AWS Resources to a Third Party](#) in the *IAM User Guide*.

Example of setting this in a config file:

```
external_id = unique_value_assigned_by_3rd_party
```

mfa_serial - shared AWS config file setting

Specifies the identification or serial number of a multi-factor authentication (MFA) device that the user must use when assuming a role.

Required when assuming a role where the trust policy for that role includes a condition that requires MFA authentication.

Default value: None.

Valid values: The value can be either a serial number for a hardware device (such as GAHT12345678), or an Amazon Resource Name (ARN) for a virtual MFA device. For more information about MFA, see [Configuring MFA-Protected API Access](#) in the *IAM User Guide*.

Example of setting this in a config file:

```
mfa_serial = arn:aws:iam::123456789012:mfa/my-user-name
```

role_arn - shared AWS config file setting

Specifies the Amazon Resource Name (ARN) of an IAM role that you want to use to perform operations requested using this profile.

Default value: None.

Valid values: The value must be the ARN of an IAM role, formatted as follows:

arn:aws:iam::*account-id*:role/*role-name*

In addition, you must also specify **one** of the following settings:

- **source_profile** – To identify another profile to use to find credentials that have permission to assume the role in this profile.
- **credential_source** – To use either credentials identified by the current environment variables or credentials attached to an Amazon EC2 instance profile, or an Amazon ECS container instance.

Example of setting this in a config file:

```
role_arn = arn:aws:iam::123456789012:role/my-role-name  
source_profile = profile-with-user-that-can-assume-role
```

```
role_arn = arn:aws:iam::123456789012:role/my-role-name  
credential_source = Ec2InstanceMetadata
```

role_session_name - shared AWS config file setting

Specifies the name to attach to the role session. This name appears in AWS CloudTrail logs for entries associated with this session.

Default value: An optional parameter. If you don't provide this value, a session name is generated automatically if the profile assumes a role.

Valid values: Provided to the RoleSessionName parameter when the AWS CLI calls the AssumeRole operation (or operations such as the AssumeRoleWithWebIdentity operation) on your behalf. The value becomes part of the assumed role user Amazon Resource Name (ARN) that you can query, and shows up as part of the CloudTrail log entries for operations invoked by this profile.

arn:aws:sts::*123456789012*:assumed-role/*my-role-name*/*my-role_session_name*.

Example of setting this in a config file:

```
role_session_name = my-role-session-name
```

source_profile - shared AWS config file setting

Specifies another profile whose credentials are used to assume the role specified by the `role_arn` setting in the original profile. To understand how profiles are used in the shared AWS config and credentials files, see [Shared config and credentials files \(p. 2\)](#).

If you specify a profile that is also an assume role profile, each role will be assumed in sequential order to fully resolve the credentials. This chain is stopped when the SDK encounters a profile with static credentials. Role chaining limits your AWS CLI or AWS API role session to a maximum of one hour and can't be increased. For more information, see [Roles terms and concepts](#) in the *IAM User Guide*.

Default value: None.

Valid values: A text string that consists of the name of a profile defined in the `config` and `credentials` files. You must also specify a value for `role_arn` in the current profile.

Note

This setting is an alternative to `credential_source`. You can't specify both `source_profile` and `credential_source` in the same profile.

Example of setting this in a config file:

```
[profile A]
source_profile = B
role_arn = arn:aws:iam::123456789012:role/RoleA

[profile B]
aws_access_key_id=AKIAIOSFODNN7EXAMPLE
aws_secret_access_key=wJalrXUtnFEMI/K7MDENG/bPxRfiCYEXAMPLEKEY
```

web_identity_token_file - shared AWS config file setting

Specifies the path to a file that contains an access token from a [supported OAuth 2.0 provider](#) or [OpenID Connect ID identity provider](#).

This setting enables authentication by using web identity federation providers, such as [Google](#), [Facebook](#), and [Amazon](#), among many others. The SDK or developer tool loads the contents of this file and passes it as the `WebIdentityToken` argument when it calls the `AssumeRoleWithWebIdentity` operation on your behalf.

Default value: None.

Valid values: This value must be a path and file name. The file must contain an OAuth 2.0 access token or an OpenID Connect token that was provided to you by an identity provider.

Example of setting this in a config file:

```
[profile web-identity]
role_arn=arn:aws:iam::123456789012:role/my-role-name
web_identity_token_file=/path/to/a/token
```

Compatibility with AWS SDKs

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Support	Notes or more information
AWS CLI v2	Yes	
SDK for C++	Partial	<code>credential_source</code> not supported. <code>duration_seconds</code> not supported. <code>mfa_serial</code> not supported.
SDK for Go V2 (1.x)	Yes	
SDK for Go 1.x (V1)	Yes	
SDK for Java 2.x	Partial	<code>mfa_serial</code> not supported.
SDK for Java 1.x	Partial	<code>mfa_serial</code> not supported.

SDK	Support	Notes or more information
SDK for JavaScript 3.x	Partial	credential_source not supported.
SDK for JavaScript 2.x	Partial	credential_source not supported.
SDK for .NET 3.x	Yes	
SDK for PHP 3.x	Yes	
SDK for Python (Boto3)	Yes	
SDK for Ruby 3.x	Yes	

Container credentials

The container credential provider fetches credentials for customer's containerized application. This credential provider is useful for Amazon Elastic Container Service (Amazon ECS) customers. SDKs will attempt to load credentials from the specified HTTP endpoint through a GET request.

Configure this functionality by using the following:

AWS_CONTAINER_CREDENTIALS_FULL_URI - environment variable

Contains the full HTTP URL endpoint for the SDK to use when making a request for credentials. This includes both the scheme and the host.

Default value: None.

Valid values: Valid URI.

Note: This setting is an alternative to AWS_CONTAINER_CREDENTIALS_RELATIVE_URI and will only be used if AWS_CONTAINER_CREDENTIALS_RELATIVE_URI is not set in the same profile.

Linux/macOS example of setting environment variables via command line:

```
export AWS_CONTAINER_CREDENTIALS_FULL_URI=http://localhost/get-credentials
```

or

```
export AWS_CONTAINER_CREDENTIALS_FULL_URI=http://localhost:8080/get-credentials
```

AWS_CONTAINER_CREDENTIALS_RELATIVE_URI - environment variable

Specifies the relative HTTP URL endpoint for the SDK to use when making a request for credentials.

Default value: None.

Valid values: Valid relative URI.

Linux/macOS example of setting environment variables via command line:

```
export AWS_CONTAINER_CREDENTIALS_RELATIVE_URI=/get-credentials?a=1
```

AWS_CONTAINER_AUTHORIZATION_TOKEN - environment variable

If this variable is set, the SDK will set the Authorization header on the HTTP request with the environment variable's value.

Default value: None.

Valid values: String.

Linux/macOS example of setting environment variables via command line:

```
export AWS_CONTAINER_CREDENTIALS_FULL_URI=http://localhost/get-credential
export AWS_CONTAINER_AUTHORIZATION_TOKEN=Basic abcd
```

Compatibility with AWS SDKS

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Support	Notes or more information
AWS CLI v2	Yes	
SDK for C++	No	
SDK for Go V2 (1.x)	Yes	
SDK for Go 1.x (V1)	Yes	
SDK for Java 2.x	Yes	
SDK for Java 1.x	Yes	
SDK for JavaScript 3.x	Yes	
SDK for JavaScript 2.x	Yes	
SDK for .NET 3.x	Yes	
SDK for PHP 3.x	No	
SDK for Python (Boto3)	Yes	
SDK for Ruby 3.x	No	

IMDS credentials

Instance metadata is data about your instance that you can use to configure or manage the running instance. For more information about the data available, see [Instance metadata and user data](#) in the *Amazon EC2 User Guide for Linux Instances* or [Instance metadata and user data](#) in the *Amazon EC2 User Guide for Windows Instances*. Amazon EC2 provides a local endpoint available to instances that can provide various bits of information to the instance. If the instance has a role attached, it can provide a set of credentials that are valid for that role. The SDKs can use that endpoint to resolve credentials as part of their [default credential provider chain](#) (p. 23).

Configure this functionality by using the following:

AWS_EC2_METADATA_DISABLED - environment variable

Whether or not to attempt to use Amazon EC2 Instance Metadata Service (IMDS) to obtain credentials.

Default value: false.

Valid values: true, false.

ec2_metadata_service_endpoint - shared AWS config file setting,
AWS_EC2_METADATA_SERVICE_ENDPOINT - environment variable

The endpoint of IMDS.

Default value: If `ec2_metadata_service_endpoint_mode` equals `IPv4`, then default endpoint is `http://169.254.169.254`. If `ec2_metadata_service_endpoint_mode` equals `IPv6`, then default endpoint is `http://[fd00:ec2::254]`.

Valid values: Valid URI.

ec2_metadata_service_endpoint_mode - shared AWS config file setting,
AWS_EC2_METADATA_SERVICE_ENDPOINT_MODE - environment variable

The endpoint mode of IMDS.

Default value: `IPv4`.

Valid values: `IPv4`, `IPv6`.

Security for IMDS credentials

By default, when the AWS SDK is not configured with valid credentials the SDK will attempt to use the Amazon EC2 Instance Metadata Service (IMDS) to retrieve credentials for an AWS role. This behavior can be disabled by setting the `AWS_EC2_METADATA_DISABLED` environment variable to `true`. This prevents unnecessary network activity and enhances security on untrusted networks where the Amazon EC2 Instance Metadata Service may be impersonated.

Note

AWS SDK clients configured with valid credentials will never use IMDS to retrieve credentials, regardless of any of these settings.

Disabling use of Amazon EC2 IMDS credentials

How you set this environment variable depends on what operating system is in use as well as whether or not you want the change to be persistent.

Linux and macOS

Customers using Linux or macOS can set this environment variable with the following command:

```
$ export AWS_EC2_METADATA_DISABLED=true
```

If you want this setting to be persistent across multiple shell sessions and system restarts, you can add the above command to your shell profile file, such as `.bash_profile`, `.zsh_profile`, or `.profile`.

Windows

Customers using Windows can set this environment variable with the following command:

```
$ set AWS_EC2_METADATA_DISABLED=true
```

If you want this setting to be persistent across multiple shell sessions and system restarts can use the following command instead:

```
$ setx AWS_EC2_METADATA_DISABLED=true
```

Note

The **setx** command does not apply the value to the current shell session, so you will need to reload or reopen the shell for the change to take effect.

Compatibility with AWS SDKS

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Support	Notes or more information
AWS CLI v2	Yes	
SDK for C++	Yes	
SDK for Go V2 (1.x)	Yes	
SDK for Go 1.x (V1)	Yes	
SDK for Java 2.x	Yes	
SDK for Java 1.x	Yes	
SDK for JavaScript 3.x	Yes	
SDK for JavaScript 2.x	Yes	
SDK for .NET 3.x	Yes	
SDK for PHP 3.x	Yes	
SDK for Python (Boto3)	Yes	
SDK for Ruby 3.x	Yes	

Process credentials

SDKs provide a way to extend the credential provider chain for custom use cases.

Warning

The following describes a method of sourcing credentials from an external process. This can potentially be dangerous, so proceed with caution. Other credential providers should be preferred if at all possible. If using this option, you should make sure that the `config` file is as locked down as possible using security best practices for your operating system. Confirm that your custom credential tool does not write any secret information to `stderr`, because the SDKs and AWS CLI can capture and log such information, potentially exposing it to unauthorized users.

Configure this functionality by using the following:

credential_process - shared AWS config file setting

Specifies an external command that the SDK or tool runs on your behalf to generate or retrieve authentication credentials to use. The setting specifies the name of a program/command that the

SDK will invoke. When the SDK invokes the process, it waits for the process to write JSON data to `stdout`. The custom provider must return information in a specific format. That information contains the credentials that the SDK or tool can use to authenticate you.

Specifying the path to the credentials program

The setting's value is a string that contains a path to a program that the SDK or development tool runs on your behalf:

- The path and file name can consist of only these characters: A-Z, a-z, 0-9, hyphen (-), underscore (_), period (.), forward slash (/), backslash (\), and space.
- If the path or file name contains a space, surround the complete path and file name with double-quotation marks (" ").
- If a parameter name or a parameter value contains a space, surround that element with double-quotation marks (" "). Surround only the name or value, not the pair.
- Don't include any environment variables in the strings. For example, don't include `$HOME` or `%USERPROFILE%`.
- Don't specify the home folder as `~`. * You must specify either the full path or a base file name. If there is a base file name, the system attempts to find the program within folders specified by the `PATH` environment variable.

Linux/macOS example of setting environment variables via command line:

```
credential_process = "/path/to/credentials.sh" parameterWithoutSpaces "parameter with spaces"
```

Windows example of setting environment variables via command line:

```
credential_process = "C:\Path\To\credentials.cmd" parameterWithoutSpaces "parameter with spaces"
```

Valid output from the credentials program

The SDK runs the command as specified in the profile and then reads data from the standard output stream. The command you specify, whether a script or binary program, must generate JSON output on `STDOUT` that matches the following syntax.

```
{
  "Version": 1,
  "AccessKeyId": "an AWS access key",
  "SecretAccessKey": "your AWS secret access key",
  "SessionToken": "the AWS session token for temporary credentials",
  "Expiration": "RFC3339 timestamp for when the credentials expire"
}
```

Note

As of this writing, the `Version` key must be set to 1. This might increment over time as the structure evolves.

The `Expiration` key is an RFC3339 formatted timestamp. If the `Expiration` key isn't present in the tool's output, the SDK assumes that the credentials are long-term credentials that don't refresh. Otherwise, the credentials are considered temporary credentials, and they are automatically refreshed by rerunning the `credential_process` command before the credentials expire.

Note

The SDK does *not* cache external process credentials the way it does assume-role credentials. If caching is required, you must implement it in the external process.

The external process can return a non-zero return code to indicate that an error occurred while retrieving the credentials.

Compatibility with AWS SDKS

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Support	Notes or more information
AWS CLI v2	Yes	
SDK for C++	Yes	
SDK for Go V2 (1.x)	Yes	
SDK for Go 1.x (V1)	Yes	
SDK for Java 2.x	Yes	
SDK for Java 1.x	Yes	
SDK for JavaScript 3.x	Yes	
SDK for JavaScript 2.x	Yes	
SDK for .NET 3.x	Yes	
SDK for PHP 3.x	Yes	
SDK for Python (Boto3)	Yes	
SDK for Ruby 3.x	Yes	

SSO credentials

After AWS SSO is [configured](#), you need to define named profiles for SSO in your shared AWS `config` or `credentials` files. These profiles are used to connect to the AWS SSO [user portal](#). After a user successfully authenticates with AWS SSO, the portal returns short-term credentials for the IAM role associated with that user.

The user portal for a profile is specified with the `sso_start_url` and `sso_region` settings. The IAM role is specified with the `sso_account_id` and `sso_role_name` settings. All four settings are required.

The following two guides contain additional information on SSO:

- [AWS Single Sign-On User Guide](#)
- [AWS Single Sign-On Portal API Reference](#)

To learn more about configuring and signing in to AWS SSO through the AWS CLI, see [Configuring the AWS CLI to use AWS Single Sign-On](#) in the *AWS Command Line Interface User Guide*.

Configure this functionality by using the following:

sso_start_url - shared AWS config file setting

The URL that points to your organization's AWS SSO [user portal](#).

sso_region - shared AWS config file setting

The AWS Region that contains your AWS SSO portal host; that is, the Region you selected before enabling AWS SSO. This can be different from your default AWS Region.

For a complete list of the AWS Regions and their codes, see [Regional Endpoints](#) in the *Amazon Web Services General Reference*.

sso_account_id - shared AWS config file setting

The AWS account ID that contains the IAM role with the permissions that you want to grant to AWS SSO users.

sso_role_name - shared AWS config file setting

The role name of the IAM role that defines the user's permissions when using this profile to get credentials through AWS SSO. The role must exist in the AWS account specified by `sso_account_id`. Use the role name, not the role ARN.

Depending on whether you log into the AWS SSO user portal as an administrator or a user, this value may be an [AWS SSO permission set](#) or a role name. Permission sets define the level of access that users have to their assigned AWS accounts.

Example of setting these four required values in the config file:

```
[profile my-sso-profile]
sso_start_url = https://my-sso-portal.awsapps.com/start
sso_region = us-west-2
sso_account_id = 111122223333
sso_role_name = SSORedOnlyRole
```

Compatibility with AWS SDKs

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Support	Notes or more information
AWS CLI v2	Yes	
SDK for C++	Yes	
SDK for Go V2 (1.x)	Yes	
SDK for Go 1.x (V1)	Yes	
SDK for Java 2.x	Yes	
SDK for Java 1.x	No	
SDK for JavaScript 3.x	Yes	
SDK for JavaScript 2.x	No	
SDK for .NET 3.x	Yes	
SDK for PHP 3.x	Yes	

SDK	Support	Notes or more information
SDK for Python (Boto3)	Yes	
SDK for Ruby 3.x	Yes	

Static credentials

For basics on static credentials, see [AWS account root user credentials and IAM user credentials](#) and [Understanding and getting your AWS credentials](#) in the *Amazon Web Services General Reference*.

The AWS SDK automatically uses these AWS credentials to sign API requests to AWS, so that your workloads can access your AWS resources and data securely and conveniently. If you use an IAM role, these temporary AWS credentials are refreshed multiple times a day.

Note

If AWS becomes unable to refresh these temporary credentials, AWS may extend the validity of the credentials so that your workloads are not impacted.

The shared `AWS credentials` file is the recommended location for storing credentials information because it is safely outside of application source directories and separate from the SDK-specific settings of the shared `config` file.

Configure this functionality by using the following:

aws_access_key_id - shared AWS config file setting, **aws_access_key_id** - shared AWS credentials file setting (*recommended method*), **AWS_ACCESS_KEY_ID** - environment variable

Specifies the AWS access key used as part of the credentials to authenticate the user.

aws_secret_access_key - shared AWS config file setting, **aws_secret_access_key** - shared AWS credentials file setting (*recommended method*), **AWS_SECRET_ACCESS_KEY** - environment variable

Specifies the AWS secret key used as part of the credentials to authenticate the user.

aws_session_token - shared AWS config file setting, **aws_session_token** - shared AWS credentials file setting (*recommended method*), **AWS_SESSION_TOKEN** - environment variable

Specifies an AWS session token used as part of the credentials to authenticate the user. A session token is required only if you manually specify temporary security credentials. You receive this value as part of the temporary credentials returned by successful requests to assume a role.

Example of setting these required values in the `config` or `credentials` file:

```
[default]
aws_access_key_id = AKIAIOSFODNN7EXAMPLE
aws_secret_access_key = wJalrXUtnFEMI/K7MDENG/bPxRfiCYEXAMPLEKEY
aws_session_token = AQoEXAMPLEH4aoAH0gNCAPy...truncated...zrkuWJOgQs8IZzaIv2BXIa2R40lgk
```

Linux/macOS example of setting environment variables via command line:

```
export AWS_ACCESS_KEY_ID=AKIAIOSFODNN7EXAMPLE
export AWS_SECRET_ACCESS_KEY=wJalrXUtnFEMI/K7MDENG/bPxRfiCYEXAMPLEKEY
export
  AWS_SESSION_TOKEN=AQoEXAMPLEH4aoAH0gNCAPy...truncated...zrkuWJOgQs8IZzaIv2BXIa2R40lgk
```

Windows example of setting environment variables via command line:

```
setx AWS_ACCESS_KEY_ID AKIAIOSFODNN7EXAMPLE
setx AWS_SECRET_ACCESS_KEY wJalrXUtnFEMI/K7MDENG/bPxrFiCYEXAMPLEKEY
setx AWS_SESSION_TOKEN AQoEXAMPLEH4aoAH0gNCAPy...truncated...zrkuWJOGqs8IZZaIv2BXIa2R4Olgk
```

Compatibility with AWS SDKS

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Support	Notes or more information
AWS CLI v2	Yes	
SDK for C++	Yes	shared config file not supported.
SDK for Go V2 (1.x)	Yes	
SDK for Go 1.x (V1)	Yes	
SDK for Java 2.x	Yes	
SDK for Java 1.x	Yes	
SDK for JavaScript 3.x	Yes	
SDK for JavaScript 2.x	Yes	
SDK for .NET 3.x	Yes	Environment variables not supported.
SDK for PHP 3.x	Yes	
SDK for Python (Boto3)	Yes	
SDK for Ruby 3.x	Yes	

Standardized features

Many features have been standardized to consistent defaults and to work the same way across many SDKs. This consistency increases productivity and clarity when coding across multiple SDKs. All settings can be overridden in code, see your specific SDK API for details.

Important

Not all SDKs support all features, or even all aspects within a feature.

Topics

- [Amazon EC2 instance metadata \(p. 36\)](#)
- [Amazon S3 access points \(p. 37\)](#)
- [Amazon S3 Multi-Region Access Points \(p. 38\)](#)
- [AWS Region \(p. 39\)](#)
- [AWS STS Regionalized endpoints \(p. 41\)](#)
- [Endpoint discovery \(p. 42\)](#)
- [General configuration settings \(p. 43\)](#)
- [IMDS client \(p. 45\)](#)
- [Retry behavior \(p. 47\)](#)
- [Smart configuration defaults \(p. 50\)](#)

Amazon EC2 instance metadata

Amazon EC2 provides a service on instances called the Instance Metadata Service (IMDS). To learn more about this service, see [Instance metadata and user data](#) in the *Amazon EC2 User Guide for Linux Instances* or [Instance metadata and user data](#) in the *Amazon EC2 User Guide for Windows Instances*. When attempting to retrieve credentials on an Amazon EC2 instance that has been configured with an IAM role, the connection to the instance metadata service is adjustable.

Configure this functionality by using the following:

metadata_service_num_attempts - shared AWS config file setting,
AWS_METADATA_SERVICE_NUM_ATTEMPTS - environment variable

This setting specifies the number of total attempts to make before giving up when attempting to retrieve data from the instance metadata service.

Default value: 1

Valid values: Number greater than or equal to 1.

metadata_service_timeout - shared AWS config file setting, **AWS_METADATA_SERVICE_TIMEOUT** - environment variable

Specifies the number of seconds before timing out when attempting to retrieve data from the instance metadata service.

Default value: 1

Valid values: Number greater than or equal to 1.

Example of setting these values in the config file:

```
[default]
metadata_service_num_attempts=10
metadata_service_timeout=10
```

Linux/macOS example of setting environment variables via command line:

```
export AWS_METADATA_SERVICE_NUM_ATTEMPTS=10
export AWS_METADATA_SERVICE_TIMEOUT=10
```

Windows example of setting environment variables via command line:

```
setx AWS_METADATA_SERVICE_NUM_ATTEMPTS 10
setx AWS_METADATA_SERVICE_TIMEOUT 10
```

Compatibility with AWS SDKS

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Support	Notes or more information
AWS CLI v2	Yes	
SDK for C++	No	

SDK	Support	Notes or more information
SDK for Go V2 (1.x)	No	
SDK for Go 1.x (V1)	No	
SDK for Java 2.x	No	
SDK for Java 1.x	Partial	<code>metadata_service_num_attempts</code> not supported.
SDK for JavaScript 3.x	No	
SDK for JavaScript 2.x	No	
SDK for .NET 3.x	No	
SDK for PHP 3.x	Yes	
SDK for Python (Boto3)	Yes	
SDK for Ruby 3.x	No	

Amazon S3 access points

The Amazon S3 service provides access points as an alternative way to interact with Amazon S3 buckets. Access points have unique policies and configurations that can be applied to them instead of directly to the bucket. With AWS SDKs, you can use access point Amazon Resource Names (ARNs) in the bucket field for API operations instead of specifying the bucket name explicitly. They are used for specific operations such as using an access point ARN with `GetObject` to fetch an object from a bucket, or using an access point ARN with `PutObject` to add an object to a bucket.

To learn more about Amazon S3 access points and ARNs, see [Using access points](#) in the *Amazon S3 User Guide*.

Configure this functionality by using the following:

`s3_use_arn_region` - shared AWS config file setting, `AWS_S3_USE_ARN_REGION` - environment variable, To configure value directly in code, consult your specific SDK directly.

This setting controls whether the SDK uses the access point ARN AWS Region to construct the Regional endpoint for the request. The SDK validates that the ARN AWS Region is served by the same AWS partition as the client's configured AWS Region to prevent cross-partition calls that most likely will fail. If multiply defined, the code-configured setting takes precedence, followed by the environment variable setting.

Default value: `false`

Valid values:

- `true` – The SDK uses the ARN's AWS Region when constructing the endpoint instead of the client's configured AWS Region. Exception: If the client's configured AWS Region is a FIPS AWS Region, then it must match the ARN's AWS Region. Otherwise, an error will result.
- `false` – The SDK uses the client's configured AWS Region when constructing the endpoint.

Compatibility with AWS SDKs

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Supported	Notes or more information
AWS CLI v2	Yes	
SDK for C++	Yes	
SDK for Go V2 (1.x)	Yes	
SDK for Go 1.x (V1)	Yes	
SDK for Java 2.x	Yes	
SDK for Java 1.x	Yes	
SDK for JavaScript 3.x	Yes	
SDK for JavaScript 2.x	Yes	
SDK for .NET 3.x	Yes	
SDK for PHP 3.x	Partial	
SDK for Python (Boto3)	Yes	
SDK for Ruby 3.x	Yes	

Amazon S3 Multi-Region Access Points

Amazon S3 Multi-Region Access Points provide a global endpoint that applications can use to fulfill requests from Amazon S3 buckets located in multiple AWS Regions. You can use Multi-Region Access Points to build multi-Region applications with the same architecture used in a single Region, and then run those applications anywhere in the world.

To learn more about Multi-Region Access Points, see [Multi-Region Access Points in Amazon S3](#) in the *Amazon S3 User Guide*.

To learn more about Multi-Region Access Point Amazon Resource Names (ARNs), see [Making requests using a Multi-Region Access Point](#) in the *Amazon S3 User Guide*.

To learn more about creating Multi-Region Access Points, see [Managing Multi-Region Access Points](#) in the *Amazon S3 User Guide*.

The SigV4A algorithm is the signing implementation used to sign the global Region requests. This algorithm is obtained by the SDK through a dependency on the [AWS Common Runtime \(CRT\) libraries](#) (p. 53).

Configure this functionality by using the following:

s3_disable_multiregion_access_points - shared AWS config file setting,
AWS_S3_DISABLE_MULTIREGION_ACCESS_POINTS - environment variable, To configure value directly in code, consult your specific SDK directly.

This setting controls whether the SDK potentially attempts cross-Region requests. If multiply defined, the code-configured setting takes precedence, followed by the environment variable setting.

Default value: `false`

Valid values:

- **true** – Stops the use of cross-Region requests.
- **false** – Enables cross-Region requests using Multi-Region Access Points.

Compatibility with AWS SDKS

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Support	Notes or more information
AWS CLI v2	Yes	
SDK for C++	Yes	
SDK for Go V2 (1.x)	No	
SDK for Go 1.x (V1)	No	
SDK for Java 2.x	Yes	
SDK for Java 1.x	No	
SDK for JavaScript 3.x	Yes	
SDK for JavaScript 2.x	No	
SDK for .NET 3.x	Yes	
SDK for PHP 3.x	Yes	
SDK for Python (Boto3)	Yes	
SDK for Ruby 3.x	Yes	

AWS Region

AWS Regions are an important concept to understand when working with AWS services.

With AWS Regions, you can access AWS services that physically reside in a specific geographic area. This can be useful to keep your data and applications running close to where you and your users will access them. Regions provide fault tolerance, stability, and resilience, and can also reduce latency. With Regions, you can create redundant resources that remain available and unaffected by a Regional outage.

Most AWS service requests are associated with a particular geographic region. The resources that you create in one Region do not exist in any other Region unless you explicitly use a replication feature offered by an AWS service. For example, Amazon S3 and Amazon EC2 support cross-Region replication. Some services, such as IAM, do not have Regional resources.

The *AWS General Reference* contains information on the following:

- To understand the relationship between Regions and endpoints, and to view a list of existing Regional endpoints, see [AWS service endpoints](#).
- To view the current list of all supported Regions and endpoints for each AWS service, see [Service endpoints and quotas](#).

Creating service clients

To programmatically access AWS services, SDKs use a client class/object for each AWS service. If your application needs to access Amazon EC2, for example, your application would create an Amazon EC2 client object to interface with that service.

If no Region is explicitly specified for the client, the client defaults to using the Region set through the following `region` setting. However, the active Region for a client can be explicitly set for any individual client object. Setting the Region in this way takes precedence over any global setting for that particular service client. The alternative Region is specified during instantiation of that client, specific to your SDK (check your specific SDK Guide or your SDK's code base).

Configure this functionality by using the following:

region - shared AWS config file setting, `AWS_REGION` - environment variable

Specifies the default AWS Region to use for AWS requests. This Region is used for SDK service requests that aren't provided with a specific Region to use.

Default value: None. You must specify this value explicitly.

Valid values:

- Any of the Region codes available for the chosen service, as listed in [AWS service endpoints](#) in the *AWS General Reference*. For example, the value `us-east-1` sets the endpoint to the AWS Region US East (N. Virginia).
- `aws-global` specifies the global endpoint for services that support a separate global endpoint in addition to Regional endpoints, such as AWS Security Token Service (AWS STS) and Amazon Simple Storage Service (Amazon S3).

Example of setting this value in the `config` file:

```
[default]
region = us-west-2
```

Linux/macOS example of setting environment variables via command line:

```
export AWS_REGION=us-west-2
```

Windows example of setting environment variables via command line:

```
setx AWS_REGION us-west-2
```

Most SDKs have a "configuration" object that is available for setting the default Region from within the application code. For details, see your specific AWS SDK developer guide.

Compatibility with AWS SDKs

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Support	Notes or more information
AWS CLI v2	Yes	AWS CLI V2 uses any value in <code>AWS_REGION</code> before any value in <code>AWS_DEFAULT_REGION</code> (both variables are checked).
AWS CLI v1	Yes	This SDK uses environment variable named <code>AWS_DEFAULT_REGION</code> for this purpose.

SDK	Supported	Notes or more information
SDK for C++	Yes	
SDK for Go V2 (1.x)	Yes	
SDK for Go 1.x (V1)	Yes	
SDK for Java 2.x	Yes	
SDK for Java 1.x	Yes	
SDK for JavaScript 3.x	Yes	
SDK for JavaScript 2.x	Yes	
SDK for .NET 3.x	Yes	
SDK for PHP 3.x	Yes	
SDK for Python (Boto3)	Yes	This SDK uses environment variable named <code>AWS_DEFAULT_REGION</code> for this purpose.
SDK for Ruby 3.x	Yes	

AWS STS Regionalized endpoints

By default, AWS Security Token Service (AWS STS) is available as a global service, and all AWS STS requests go to a single endpoint at `https://sts.amazonaws.com`. Global requests map to the US East (N. Virginia) Region. AWS recommends using Regional AWS STS endpoints instead of the global endpoint. For more information on AWS STS endpoints, [Endpoints](#) in the *AWS Security Token Service API Reference*.

Configure this functionality by using the following:

sts_regional_endpoints - shared AWS config file setting, **AWS_STS_REGIONAL_ENDPOINTS** - environment variable

This setting specifies how the SDK or tool determines the AWS service endpoint that it uses to talk to the AWS Security Token Service (AWS STS).

Default value: `legacy`

Note: All new SDK major versions will default to `regional` in the future.

Valid values:

- **legacy** – Uses the global AWS STS endpoint, `sts.amazonaws.com`, for the following AWS Regions: `ap-northeast-1`, `ap-south-1`, `ap-southeast-1`, `ap-southeast-2`, `aws-global`, `ca-central-1`, `eu-central-1`, `eu-north-1`, `eu-west-1`, `eu-west-2`, `eu-west-3`, `sa-east-1`, `us-east-1`, `us-east-2`, `us-west-1`, and `us-west-2`. All other Regions automatically use their respective Regional endpoint.
- **regional** – The SDK or tool always uses the AWS STS endpoint for the currently configured Region. For example, if the client is configured to use `us-west-2`, all calls to AWS STS are made to the Regional endpoint `sts.us-west-2.amazonaws.com`, instead of the global `sts.amazonaws.com` endpoint. To send a request to the global endpoint while this setting is enabled, you can set the Region to `aws-global`.

Example of setting these values in the config file:

```
[default]
sts_regional_endpoints = regional
```

Linux/macOS example of setting environment variables via command line:

```
export AWS_STS_REGIONAL_ENDPOINTS=regional
```

Windows example of setting environment variables via command line:

```
setx AWS_STS_REGIONAL_ENDPOINTS regional
```

Compatibility with AWS SDKS

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Support	Notes or more information
AWS CLI v2	Partial	Default value is <code>regional</code> .
SDK for C++	Partial	Environment variable and config file setting not supported. SDK performs with <code>regional</code> setting.
SDK for Go V2 (1.x)	Yes	
SDK for Go 1.x (V1)	Yes	
SDK for Java 2.x	Yes	
SDK for Java 1.x	Yes	
SDK for JavaScript 3.x	Yes	
SDK for JavaScript 2.x	Yes	
SDK for .NET 3.x	Yes	
SDK for PHP 3.x	Yes	
SDK for Python (Boto3)	Yes	
SDK for Ruby 3.x	Yes	

Endpoint discovery

SDKs use endpoint discovery to access service endpoints (URLs to access various resources), while still maintaining flexibility for AWS to alter URLs as needed. This way, your code can automatically detect new endpoints. There are no fixed endpoints for some services. Instead, you get the available endpoints during runtime by making a request to get the endpoints first. After retrieving the available endpoints, the code then uses the endpoint to access other operations. For example, for Amazon Timestream, the SDK makes a `DescribeEndpoints` request to retrieve the available endpoints, and then uses those endpoints to complete specific operations such as `CreateDatabase` or `CreateTable`.

Endpoint discovery is required in some services and optional in others. It defaults to either `true` or `false` depending on whether the service requires endpoint discovery. For example, Timestream

defaults to `true`, and Amazon DynamoDB defaults to `false`. For services where endpoint discovery is not required, endpoint discovery is not enabled. Instead, configuration options are available through environment variables, the shared AWS `config` file, or SDK code constructs (for example, configuration classes). For operations where endpoint discovery is required, the SDK automatically attempts to discover an endpoint.

Configure this functionality by using the following:

endpoint_discovery_enabled - shared AWS `config` file setting,
AWS_ENABLE_ENDPOINT_DISCOVERY - environment variable, To configure value directly in code,
consult your specific SDK directly.

Enables/disables endpoint discovery for services where endpoint discovery is optional. Endpoint discovery is required in some services.

Default value: `false`

Valid values:

- **true** – The SDK should automatically attempt to discover an endpoint for services where endpoint discovery is optional.
- **false** – The SDK should not automatically attempt to discover an endpoint for services where endpoint discovery is optional.

Compatibility with AWS SDKS

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Support	Notes or more information
AWS CLI v2	Yes	
SDK for C++	Yes	
SDK for Go V2 (1.x)	Yes	
SDK for Go 1.x (V1)	Yes	
SDK for Java 2.x	Yes	
SDK for Java 1.x	Yes	
SDK for JavaScript 3.x	Yes	
SDK for JavaScript 2.x	Yes	
SDK for .NET 3.x	Yes	
SDK for PHP 3.x	Yes	
SDK for Python (Boto3)	Yes	
SDK for Ruby 3.x	Yes	

General configuration settings

SDKs support some general settings that configure overall SDK behaviors.

Configure this functionality by using the following:

api_versions - shared AWS config file setting

Some AWS services maintain multiple API versions to support backward compatibility. By default, SDK and AWS CLI operations use the latest available API version. To require a specific API version to use for your requests, include the `api_versions` setting in your profile.

Default value: None. (Latest API version is used by the SDK.)

Valid values: This is a nested setting that's followed by one or more indented lines that each identify one AWS service and the API version to use. See the documentation for the AWS service to understand which API versions are available.

The example sets a specific API version for two AWS services in the `config` file. These API versions are used only for commands that run under the profile that contains these settings. Commands for any other service use the latest version of that service's API.

```
api_versions =  
  ec2 = 2015-03-01  
  cloudfront = 2015-09-017
```

ca_bundle - shared AWS config file setting, AWS_CA_BUNDLE - environment variable

Specifies the path to a custom certificate bundle (a file with a `.pem` extension) to use when establishing SSL/TLS connections.

Default value: none

Valid values: Specify either the full path or a base file name. If there is a base file name, the system attempts to find the program within folders specified by the `PATH` environment variable.

Example of setting this value in the `config` file:

```
[default]  
ca_bundle = dev/apps/ca-certs/cabundle-2019mar05.pem
```

Linux/macOS example of setting environment variables via command line:

```
export AWS_CA_BUNDLE=dev/apps/ca-certs/cabundle-2019mar05.pem
```

Windows example of setting environment variables via command line:

```
setx AWS_CA_BUNDLE C:\dev\apps\ca-certs\cabundle-2019mar05.pem
```

parameter_validation - shared AWS config file setting

Specifies whether the SDK or tool attempts to validate command line parameters before sending them to the AWS service endpoint.

Default value: true

Valid values:

- **true** – The default. The SDK or tool performs client-side validation of command line parameters. This helps the SDK or tool confirm that parameters are valid, and catches some errors. The SDK or tool can reject requests that aren't valid before sending requests to the AWS service endpoint.

- **false** – The SDK or tool doesn't validate command line parameters before sending them to the AWS service endpoint. The AWS service endpoint is responsible for validating all requests and rejecting requests that aren't valid.

Compatibility with AWS SDKS

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Support	Notes or more information
AWS CLI v2	Yes	
SDK for C++	Yes	
SDK for Go V2 (1.x)	Partial	<code>api_versions</code> and <code>parameter_validation</code> not supported.
SDK for Go 1.x (V1)	Partial	<code>api_versions</code> and <code>parameter_validation</code> not supported.
SDK for Java 2.x	No	
SDK for Java 1.x	No	
SDK for JavaScript 3.x	Yes	
SDK for JavaScript 2.x	Yes	
SDK for .NET 3.x	Yes	
SDK for PHP 3.x	Yes	
SDK for Python (Boto3)	Yes	
SDK for Ruby 3.x	Yes	

IMDS client

SDKs implement an Instance Metadata Service Version 2 (IMDSv2) client using session-oriented requests. For more information on IMDSv2, see [Use IMDSv2](#) in the *Amazon EC2 User Guide for Linux Instances* or [Use IMDSv2](#) in the *Amazon EC2 User Guide for Windows Instances*. The IMDS client is configurable via a client configuration object available in the SDK code base.

Configure this functionality by using the following:

retries - client configuration object member

The number of additional retry attempts for any failed request.

Default value: 3

Valid values: Number greater than 0.

port - client configuration object member

The port for the endpoint.

Default value: 80

Valid values: Number.

token_ttl - client configuration object member

The TTL of the token.

Default value: 21,600 seconds (6 hours, the maximum time allotted).

Valid values: Number.

endpoint - client configuration object member

The endpoint of IMDS.

Default value: If `endpoint_mode` equals `IPv4`, then default endpoint is `http://169.254.169.254`. If `endpoint_mode` equals `IPv6`, then default endpoint is `http://[fd00:ec2::254]`.

Valid values: Valid URI.

The following options are supported by most SDKs. See your specific SDK code base for details.

endpoint_mode - client configuration object member

The endpoint mode of IMDS.

Default value: `IPv4`

Valid values: `IPv4`, `IPv6`

http_open_timeout - client configuration object member (name may vary)

The number of seconds to wait for the connection to open.

Default value: 1 second.

Valid values: Number greater than 0.

http_read_timeout - client configuration object member (name may vary)

The number of seconds for one chunk of data to be read.

Default value: 1 second.

Valid values: Number greater than 0.

http_debug_output - client configuration object member (name may vary)

Sets an output stream for debugging.

Default value: None.

Valid values: A valid I/O stream, like `STDOUT`.

backoff - client configuration object member (name may vary)

The number of seconds to sleep in between retries or a customer provided backoff function to call. This overrides the default exponential backoff strategy.

Default value: Varies by SDK.

Valid values: Varies by SDK. Can be either a numeric value or a call out to a custom function.

Compatibility with AWS SDKS

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Support	Notes or more information
AWS CLI v2	Yes	
SDK for C++	No	
SDK for Go V2 (1.x)	No	
SDK for Go 1.x (V1)	No	
SDK for Java 2.x	No	
SDK for Java 1.x	Yes	
SDK for JavaScript 3.x	Yes	
SDK for JavaScript 2.x	Yes	
SDK for .NET 3.x	Yes	
SDK for PHP 3.x	Yes	
SDK for Python (Boto3)	Yes	
SDK for Ruby 3.x	Yes	

Retry behavior

Retry behavior includes settings regarding how the SDKs attempt to recover from failures resulting from requests made to AWS services.

Configure this functionality by using the following:

max_attempts - shared AWS config file setting, **AWS_MAX_ATTEMPTS** - environment variable

Specifies the maximum number attempts to make on a request.

Default value: If this value is not specified, its default depends on the value of the `retry_mode` setting:

- If `retry_mode` is `legacy` – Uses a default value specific to your SDK (check your specific SDK guide or your SDK's code base for `max_attempts` default).
- If `retry_mode` is `standard` – Makes three attempts.
- If `retry_mode` is `adaptive` – Makes three attempts.

Valid values: Number greater than 0.

retry_mode - shared AWS config file setting, **AWS_RETRY_MODE** - environment variable

Specifies how the SDK or developer tool attempts retries.

Default value: `legacy` is the default retry strategy.

Valid values:

- **legacy** – Specific to your SDK (check your specific SDK guide or your SDK's code base).
- **standard** – The standard set of retry rules across AWS SDKs. This mode includes a standard set of errors that are retried, and support for retry quotas. The default maximum number of attempts with this mode is three, unless `max_attempts` is explicitly configured.
- **adaptive** – An experimental retry mode that includes the functionality of standard mode but includes automatic client-side throttling. Because this mode is experimental, it might change behavior in the future.

Following is the high-level pseudocode for both the `standard` and `adaptive` retry modes:

```
MakeSDKRequest() {  
    attempts = 0  
    loop {  
        GetSendToken()  
        response = SendHTTPRequest()  
        RequestBookkeeping(response)  
        if not Retryable(response)  
            return response  
        attempts += 1  
        if attempts >= MAX_ATTEMPTS:  
            return response  
        if not HasRetryQuota(response)  
            return response  
        delay = ExponentialBackoff(attempts)  
        sleep(delay)  
    }  
}
```

Following are more details about the components used in the pseudocode:

GetSendToken:

Token buckets are only used in `adaptive` retry mode. Token buckets enforce a maximum request rate by requiring a token to be available in order to initiate a request. The SDK client is configurable to either fast fail the request or block until a token becomes available.

Client Side Rate Limiting is an algorithm that initially lets requests be made at any rate up to the token allowance. However, after a throttled response is detected, the client rate-of-request is then limited accordingly. The token allowance is also increased accordingly if successful responses are received.

With adaptive rate limiting, SDKs can slow down the rate at which requests are sent in order to better accommodate the capacity of AWS services.

SendHTTPRequest:

Most AWS SDKs use an HTTP library that uses connection pools so that you can reuse an existing connection when making an HTTP request. Generally, connections are reused when retrying requests due to throttling errors. Requests are not reused when retrying due to transient errors.

RequestBookkeeping:

The retry quota should be updated if the request is successful. For `adaptive` retry mode only, the state variable `maxsendrate` is updated based on the type of response received.

Retryable:

This step determines whether a response can be retried based on the following:

- The HTTP status code.

- The error code returned from the service.
- Connection errors, defined as any error received by the SDK in which an HTTP response from the service is not received.

Transient errors (HTTP status codes 400, 408, 500, 502, 503, and 504) and throttling errors (HTTP status codes 400, 403, 429, 502, 503, and 509) can all be retried.

MAX_ATTEMPTS:

Specified by the `config` file setting or the environment variable.

HasRetryQuota

This step throttles retry requests by requiring a token to be available in the retry quota bucket. Retry quota buckets are a mechanism to prevent retries that are unlikely to succeed. These quotas are SDK-dependent, are often client-dependent, and are sometimes even dependent on service endpoints. The available retry quota tokens are removed when requests fail for various reasons, and replenished when they succeed. When no tokens remain, the retry loop is exited.

ExponentialBackoff

For an error that can be retried, the retry delay is calculated using truncated exponential backoff. The SDKs use truncated binary exponential backoff with jitter. The following algorithm shows how the amount of time to sleep, in seconds, is defined for a response for request `i`:

```
seconds_to_sleep_i = min(b*r^i, MAX_BACKOFF)
```

In the preceding algorithm, the following values apply:

`b` = random number within the range of: `0 <= b <= 1`

`r` = 2

`MAX_BACKOFF` = 20 seconds for most SDKs. See your specific SDK guide or source code for confirmation.

Compatibility with AWS SDKS

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Support	Notes or more information
AWS CLI v2	Yes	
SDK for C++	Yes	
SDK for Go V2 (1.x)	Partial	<code>retry_mode</code> value adaptive is not supported.
SDK for Go 1.x (V1)	No	
SDK for Java 2.x	Yes	
SDK for Java 1.x	Yes	
SDK for JavaScript 3.x	Partial	<code>retry_mode</code> value adaptive is not supported.
SDK for JavaScript 2.x	No	

SDK	Supported	Notes or more information
SDK for .NET 3.x	Yes	
SDK for PHP 3.x	Yes	
SDK for Python (Boto3)	Yes	
SDK for Ruby 3.x	Yes	

Smart configuration defaults

With the smart configuration defaults feature, AWS SDKs can provide predefined, optimized default values for other configuration settings.

Configure this functionality by using the following:

defaults_mode - shared AWS config file setting, **AWS_DEFAULTS_MODE** - environment variable

With this setting, you can choose a mode that aligns with your application architecture, which then provides optimized default values for your application. If an AWS SDK setting has a value explicitly set, then that value always takes precedence. If an AWS SDK setting does not have a value explicitly set, and `defaults_mode` is not equal to `legacy`, then this feature can provide different default values for various settings optimized for your application. Settings may include the following: HTTP communication settings, retry behavior, service Regional endpoint settings, and, potentially, any SDK-related configuration. Customers who use this feature can get new configuration defaults tailored to common usage scenarios. If your `defaults_mode` is not equal to `legacy`, we recommend performing tests of your application when you upgrade the SDK, because the default values provided might change as best practices evolve.

Default value: `legacy`

Note: New major versions of SDKs will default to `standard`.

Valid values:

- `legacy` – Provides default settings that vary by SDK and existed before establishment of `defaults_mode`.
- `standard` – Provides the latest recommended default values that should be safe to run in most scenarios.
- `in-region` – Builds on the standard mode and includes optimization tailored for applications that call AWS services from within the same AWS Region.
- `cross-region` – Builds on the standard mode and includes optimization tailored for applications that call AWS services in a different Region.
- `mobile` – Builds on the standard mode and includes optimization tailored for mobile applications.
- `auto` – Builds on the standard mode and includes experimental features. The SDK attempts to discover the runtime environment to determine the appropriate settings automatically. The auto detection is heuristics-based and does not provide 100% accuracy. If the runtime environment can't be determined, `standard` mode is used. The auto detection might query [Instance metadata and user data](#), which might introduce latency. If startup latency is critical to your application, we recommend choosing an explicit `defaults_mode` instead.

Example of setting this value in the config file:

```
[default]
defaults_mode = standard
```

The following parameters might be optimized based on the selection of `defaults_mode`:

- `retryMode` – Specifies how the SDK attempts retries. See [Retry behavior \(p. 47\)](#).
- `stsRegionalEndpoints` – Specifies how the SDK determines the AWS service endpoint that it uses to talk to the AWS Security Token Service (AWS STS). See [AWS STS Regionalized endpoints \(p. 41\)](#).
- `s3UsEast1RegionalEndpoints` – Specifies how the SDK determines the AWS service endpoint that it uses to talk to the Amazon S3 for the `us-east-1` Region.
- `connectTimeoutInMillis` – After making an initial connection attempt on a socket, the amount of time before timing out. If the client does not receive a completion of the connect handshake, the client gives up and fails the operation.
- `tlsNegotiationTimeoutInMillis` – The maximum amount of time that a TLS handshake can take from the time the CLIENT HELLO message is sent to the time the client and server have fully negotiated ciphers and exchanged keys.

The default value for each setting changes depending on the `defaults_mode` selected for your application. These values are currently set as follows (subject to change):

Parameter	standard mode	in-region mode	cross-region mode	mobile mode
<code>retryMode</code>	standard	standard	standard	standard
<code>stsRegionalEndpoints</code>	regional	regional	regional	regional
<code>s3UsEast1RegionalEndpoints</code>	regional	regional	regional	regional
<code>connectTimeoutInMillis</code>	3100	1100	3100	30000
<code>tlsNegotiationTimeoutInMillis</code>	3100	1100	3100	30000

For example, if the `defaults_mode` you selected was `standard`, then the value of `standard` would be assigned for `retry_mode` (from the valid `retry_mode` options) and the value of `regional` would be assigned for `stsRegionalEndpoints` (from the valid `stsRegionalEndpoints` options).

Compatibility with AWS SDKs

The following SDKs support the features and settings described on this page, any partial exceptions are noted:

SDK	Supported	Notes or more information
AWS CLI v2	No	
SDK for C++	Yes	Parameters not optimized: <code>stsRegionalEndpoints</code> , <code>s3UsEast1RegionalEndpoints</code> , <code>tlsNegotiationTimeoutInMillis</code> .
SDK for Go V2 (1.x)	Yes	Parameters not optimized: <code>retryMode</code> , <code>stsRegionalEndpoints</code> , <code>s3UsEast1RegionalEndpoints</code> .
SDK for Go 1.x (V1)	No	

SDK	Supported	Notes or more information
SDK for Java 2.x	Yes	Parameters not optimized: <code>stsRegionalEndpoints</code> .
SDK for Java 1.x	No	
SDK for JavaScript 3.x	Yes	Parameters not optimized: <code>stsRegionalEndpoints</code> , <code>s3UseEast1RegionalEndpoints</code> , <code>tlsNegotiationTimeoutInMillis</code> . <code>connectTimeoutInMillis</code> is called <code>connectionTimeout</code> .
SDK for JavaScript 2.x	No	
SDK for .NET 3.x	Yes	Parameters not optimized: <code>connectTimeoutInMillis</code> , <code>tlsNegotiationTimeoutInMillis</code> .
SDK for PHP 3.x	Yes	Parameters not optimized: <code>tlsNegotiationTimeoutInMillis</code> .
SDK for Python (Boto3)	Yes	Parameters not optimized: <code>tlsNegotiationTimeoutInMillis</code> .
SDK for Ruby 3.x	Yes	

AWS Common Runtime (CRT) libraries

The AWS Common Runtime (CRT) libraries are a base library of the SDKs. The CRT is a modular family of independent packages, written in C. Each package provides good performance and minimal footprint for different required functionalities. These functionalities are common and shared across all SDKs providing better code reuse, optimization, and accuracy. The packages are:

- [awslabs/aws-c-auth](#): AWS client-side authentication (standard credential providers and signing (sigv4))
- [awslabs/aws-c-cal](#): Cryptographic primitive types, hashes (MD5, SHA256, SHA256 HMAC), signers, AES
- [awslabs/aws-c-common](#): Basic data structures, threading/synchronization primitive types, buffer management, stdlib-related functions
- [awslabs/aws-c-compression](#): Compression algorithms (Huffman encoding/decoding)
- [awslabs/aws-c-event-stream](#): Event stream message processing (headers, prelude, payload, crc/trailer), remote procedure call (RPC) implementation over event streams
- [awslabs/aws-c-http](#): C99 implementation of the HTTP/1.1 and HTTP/2 specifications
- [awslabs/aws-c-io](#): Sockets (TCP, UDP), DNS, pipes, event loops, channels, SSL/TLS
- [awslabs/aws-c-iot](#): C99 implementation of AWS IoT cloud services integration with devices
- [awslabs/aws-c-mqtt](#): Standard, lightweight messaging protocol for the Internet of Things (IoT)
- [awslabs/aws-c-s3](#): C99 library implementation for communicating with the Amazon S3 service, designed for maximizing throughput on high bandwidth Amazon EC2 instances
- [awslabs/aws-c-sdkutils](#): A utilities library for parsing and managing AWS profiles
- [awslabs/aws-checksums](#): Cross-platform hardware-accelerated CRC32c and CRC32 with fallback to efficient software implementations
- [awslabs/aws-lc](#): General-purpose cryptographic library maintained by the AWS Cryptography team for AWS and their customers, based on code from the Google BoringSSL project and the OpenSSL project
- [awslabs/s2n](#): C99 implementation of the TLS/SSL protocols, designed to be small and fast with security as a priority

The CRT is available through all SDKs except Go.

CRT dependencies

The CRT libraries form a complex net of relationships and dependencies. Knowing these relationships is helpful if you need to build the CRT directly from source. However, most users access CRT functionality through their language SDK (such as AWS SDK for C++ or AWS SDK for Java) or their language IoT device SDK (such as AWS IoT SDK for C++ or AWS IoT SDK for Java). In the following diagram, the Language CRT Bindings box refers to the package wrapping the CRT libraries for a specific language SDK. This is a collection of packages of the form `aws-crt-*`, where `*` is an SDK language (such as [aws-crt-cpp](#) or [aws-crt-java](#)).

The following is a visual illustration of the hierarchical dependencies of the CRT libraries.

Maintenance and support

For an overview of tools that can help you develop applications on AWS, see [Tools to Build on AWS](#). For information on support, see the [AWS Knowledge Center](#).

The following topics cover the maintenance and version support policies for AWS SDKs.

Topics

- [AWS SDKs and Tools maintenance policy \(p. 54\)](#)
- [AWS SDKs and Tools version support matrix \(p. 56\)](#)

AWS SDKs and Tools maintenance policy

Overview

This document outlines the maintenance policy for AWS Software Development Kits (SDKs) and Tools, including Mobile and IoT SDKs, and their underlying dependencies. AWS regularly provides the AWS SDKs and Tools with updates that may contain support for new or updated AWS APIs, new features, enhancements, bug fixes, security patches, or documentation updates. Updates may also address changes with dependencies, language runtimes, and operating systems. AWS SDK releases are published to package managers (e.g. Maven, NuGet, PyPI), and are available as source code on GitHub.

We recommend users to stay up-to-date with SDK releases to keep up with the latest features, security updates, and underlying dependencies. Continued use of an unsupported SDK version is not recommended and is done at the user's discretion.

Versioning

The AWS SDK release versions are in the form of X.Y.Z where X represents the major version. Increasing the major version of an SDK indicates that this SDK underwent significant and substantial changes to support new idioms and patterns in the language. Major versions are introduced when public interfaces (e.g. classes, methods, types, etc.), behaviors, or semantics have changed. Applications need to be updated in order for them to work with the newest SDK version. It is important to update major versions carefully and in accordance with the upgrade guidelines provided by AWS.

SDK major version life-cycle

The life-cycle for major SDKs and Tools versions consists of 5 phases, which are outlined below.

- *Developer Preview (Phase 0)* - During this phase, SDKs are not supported, should not be used in production environments, and are meant for early access and feedback purposes only. It is possible for future releases to introduce breaking changes. Once AWS identifies a release to be a stable product, it may mark it as a Release Candidate. Release Candidates are ready for GA release unless significant bugs emerge, and will receive full AWS support.
- *General Availability (GA) (Phase 1)* - During this phase, SDKs are fully supported. AWS will provide regular SDK releases that include support for new services, API updates for existing services, as well as bug and security fixes. For Tools, AWS will provide regular releases that include new feature updates and bug fixes. AWS will support the GA version of an SDK for *at least 24 months*.

- *Maintenance Announcement (Phase 2)* - AWS will make a public announcement at least 6 months before an SDK enters maintenance mode. During this period, the SDK will continue to be fully supported. Typically, maintenance mode is announced at the same time as the next major version is transitioned to GA.
- *Maintenance (Phase 3)* - During the maintenance mode, AWS limits SDK releases to address critical bug fixes and security issues only. An SDK will not receive API updates for new or existing services, or be updated to support new regions. Maintenance mode has a *default duration of 12 months*, unless otherwise specified.
- *End-of-Support (Phase 4)* - When an SDK reaches end-of support, it will no longer receive updates or releases. Previously published releases will continue to be available via public package managers and the code will remain on GitHub. The GitHub repository may be archived. Use of an SDK which has reached end-of-support is done at the user's discretion. We recommend users upgrade to the new major version.

The following is a visual illustration of the SDK major version life-cycle. Please note that the timelines shown below are illustrative and not binding.

Dependency life-cycle

Most AWS SDKs have underlying dependencies, such as language runtimes, operating systems, or third party libraries and frameworks. These dependencies are typically tied to the language community or the vendor who owns that particular component. Each community or vendor publishes their own end-of-support schedule for their product.

The following terms are used to classify underlying third party dependencies:

- *Operating System (OS)*: Examples include Amazon Linux AMI, Amazon Linux 2, Windows 2008, Windows 2012, Windows 2016, etc.
- **Language Runtime:** Examples include Java 7, Java 8, Java 11, .NET Core, .NET Standard, .NET PCL, etc.
- **Third party Library / Framework*: Examples include OpenSSL, .NET Framework 4.5, Java EE, etc.

Our policy is to continue supporting SDK dependencies for* at least 6 months* after the community or vendor ends support for the dependency. This policy, however, could vary depending on the specific dependency.

Note

AWS reserves the right to stop support for an underlying dependency without increasing the major SDK version

Communication methods

Maintenance announcements are communicated in several ways:

- An email announcement is sent to affected accounts, announcing our plans to end support for the specific SDK version. The email will outline the path to end-of-support, specify the campaign timelines, and provide upgrade guidance.
- AWS SDK documentation, such as API reference documentation, user guides, SDK product marketing pages, and GitHub readme(s) are updated to indicate the campaign timeline and provide guidance on upgrading affected applications.
- An AWS blog post is published that outlines the path to end-of-support, as well as reiterates the campaign timelines.
- Deprecation warnings are added to the SDKs, outlining the path to end-of-support and linking to the SDK documentation.

To see the list of available major versions of AWS SDKs and Tools and where they are in their maintenance life cycle, see [the section called “Version support matrix” \(p. 56\)](#).

AWS SDKs and Tools version support matrix

The matrix below shows the list of available AWS SDK major versions and where they are in the maintenance life cycle with associated timelines. For detailed information on the life-cycle for the major versions of Software Development Kits (SDKs) and Tools and their underlying dependencies see [the section called “Maintenance policy” \(p. 54\)](#)

SDK	Major version	Current Phase	General Availability Date	Notes
AWS CLI	1.x	General Availability	9/2/2013	
AWS CLI	2.x	General Availability	2/10/2020	
SDK for C++	1.x	General Availability	9/2/2015	
SDK for Go V2	V2 1.x	General Availability	1/19/2021	
SDK for Go	1.x	General Availability	11/19/2015	
SDK for Java	1.x	General Availability	3/25/2010	
SDK for Java	2.x	General Availability	11/20/2018	
SDK for JavaScript	1.x	End-of-Support	5/6/2013	
SDK for JavaScript	2.x	General Availability	6/19/2014	
SDK for JavaScript	3.x	General Availability	12/15/2020	
SDK for Kotlin	1.x	Developer Preview		
SDK for .NET	1.x	End-of-Support	11/2009	
SDK for .NET	2.x	End-of-Support	11/8/2013	
SDK for .NET	3.x	General Availability	7/28/2015	
SDK for PHP	2.x	End-of-Support	11/2/2012	
SDK for PHP	3.x	General Availability	5/27/2015	
SDK for Python (Boto2)	1.x	End-of-Support	7/13/2011	

AWS SDKs and Tools Reference Guide
Version support matrix

SDK	Major version	Current Phase	General Availability Date	Notes
SDK for Python (Boto3)	1.x	General Availability	6/22/2015	
SDK for Python (Botocore)	1.x	General Availability	6/22/2015	
SDK for Ruby	1.x	End-of-Support	7/14/2011	
SDK for Ruby	2.x	End-of-Support	2/15/2015	
SDK for Ruby	3.x	General Availability	8/29/2017	
SDK for Rust	1.x	Developer Preview		
SDK for Swift	1.x	Developer Preview		
Tools for PowerShell	2.x	End-of-Support	11/8/2013	
Tools for PowerShell	3.x	End-of-Support	7/29/2015	
Tools for PowerShell	4.x	General Availability	11/21/2019	

Document history for AWS SDKs and Tools Reference Guide

The following table describes important additions and updates to the *AWS SDKs and Tools Reference Guide*. For notification about updates to this documentation, you can subscribe to the RSS feed.

update-history-change	update-history-description	update-history-date
Update	Massive update of almost all parts of this guide.	February 1, 2022
Initial release	The first release of this guide is released to the public.	March 13, 2020

AWS glossary

For the latest AWS terminology, see the [AWS glossary](#) in the *AWS General Reference*.