
Amazon ElastiCache

API Reference

API Version 2015-02-02



Amazon ElastiCache: API Reference

Copyright © Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

Table of Contents

Welcome	1
Actions	2
AddTagsToResource	4
Request Parameters	4
Response Elements	4
Errors	4
Examples	6
See Also	6
AuthorizeCacheSecurityGroupIngress	7
Request Parameters	7
Response Elements	7
Errors	7
Examples	8
See Also	9
BatchApplyUpdateAction	10
Request Parameters	10
Response Elements	10
Errors	10
See Also	11
BatchStopUpdateAction	12
Request Parameters	12
Response Elements	12
Errors	12
See Also	13
CompleteMigration	14
Request Parameters	14
Response Elements	14
Errors	14
See Also	15
CopySnapshot	16
Request Parameters	16
Response Elements	17
Errors	17
Examples	18
See Also	19
CreateCacheCluster	20
Request Parameters	20
Response Elements	26
Errors	26
Examples	27
See Also	28
CreateCacheParameterGroup	30
Request Parameters	30
Response Elements	30
Errors	31
Examples	31
See Also	32
CreateCacheSecurityGroup	33
Request Parameters	33
Response Elements	33
Errors	33
Examples	34
See Also	35
CreateCacheSubnetGroup	36

Request Parameters	36
Response Elements	36
Errors	37
Examples	37
See Also	38
CreateGlobalReplicationGroup	39
Request Parameters	39
Response Elements	39
Errors	40
See Also	40
CreateReplicationGroup	41
Request Parameters	41
Response Elements	49
Errors	49
Examples	51
See Also	53
CreateSnapshot	54
Request Parameters	54
Response Elements	54
Errors	55
Examples	56
See Also	56
CreateUser	58
Request Parameters	58
Response Elements	59
Errors	60
See Also	60
CreateUserGroup	62
Request Parameters	62
Response Elements	62
Errors	63
See Also	64
DecreaseNodeGroupsInGlobalReplicationGroup	65
Request Parameters	65
Response Elements	65
Errors	66
See Also	66
DecreaseReplicaCount	67
Request Parameters	67
Response Elements	68
Errors	68
Examples	69
See Also	70
DeleteCacheCluster	71
Request Parameters	71
Response Elements	71
Errors	71
Examples	72
See Also	73
DeleteCacheParameterGroup	74
Request Parameters	74
Errors	74
Examples	74
See Also	75
DeleteCacheSecurityGroup	76
Request Parameters	76
Errors	76

Examples	76
See Also	77
DeleteCacheSubnetGroup	78
Request Parameters	78
Errors	78
Examples	78
See Also	79
DeleteGlobalReplicationGroup	80
Request Parameters	80
Response Elements	80
Errors	80
See Also	81
DeleteReplicationGroup	82
Request Parameters	82
Response Elements	82
Errors	82
Examples	83
See Also	84
DeleteSnapshot	85
Request Parameters	85
Response Elements	85
Errors	85
Examples	86
See Also	86
DeleteUser	88
Request Parameters	88
Response Elements	88
Errors	89
See Also	89
DeleteUserGroup	91
Request Parameters	91
Response Elements	91
Errors	92
See Also	92
DescribeCacheClusters	93
Request Parameters	93
Response Elements	94
Errors	94
Examples	94
See Also	95
DescribeCacheEngineVersions	97
Request Parameters	97
Response Elements	98
Errors	98
Examples	98
See Also	99
DescribeCacheParameterGroups	100
Request Parameters	100
Response Elements	100
Errors	101
Examples	101
See Also	102
DescribeCacheParameters	103
Request Parameters	103
Response Elements	103
Errors	104
Examples	104

See Also	105
DescribeCacheSecurityGroups	107
Request Parameters	107
Response Elements	107
Errors	108
Examples	108
See Also	109
DescribeCacheSubnetGroups	110
Request Parameters	110
Response Elements	110
Errors	111
Examples	111
See Also	112
DescribeEngineDefaultParameters	113
Request Parameters	113
Response Elements	113
Errors	113
Examples	114
See Also	115
DescribeEvents	116
Request Parameters	116
Response Elements	117
Errors	117
Examples	117
See Also	118
DescribeGlobalReplicationGroups	119
Request Parameters	119
Response Elements	119
Errors	120
See Also	120
DescribeReplicationGroups	121
Request Parameters	121
Response Elements	121
Errors	122
Examples	122
See Also	123
DescribeReservedCacheNodes	124
Request Parameters	124
Response Elements	126
Errors	126
Examples	127
See Also	128
DescribeReservedCacheNodesOfferings	129
Request Parameters	129
Response Elements	131
Errors	131
Examples	132
See Also	132
DescribeServiceUpdates	134
Request Parameters	134
Response Elements	134
Errors	135
See Also	135
DescribeSnapshots	136
Request Parameters	136
Response Elements	137
Errors	137

Examples	138
See Also	139
DescribeUpdateActions	140
Request Parameters	140
Response Elements	141
Errors	141
See Also	142
DescribeUserGroups	143
Request Parameters	143
Response Elements	143
Errors	143
See Also	144
DescribeUsers	145
Request Parameters	145
Response Elements	145
Errors	146
See Also	146
DisassociateGlobalReplicationGroup	147
Request Parameters	147
Response Elements	147
Errors	147
See Also	148
FailoverGlobalReplicationGroup	149
Request Parameters	149
Response Elements	149
Errors	149
See Also	150
IncreaseNodeGroupsInGlobalReplicationGroup	151
Request Parameters	151
Response Elements	151
Errors	152
See Also	152
IncreaseReplicaCount	153
Request Parameters	153
Response Elements	153
Errors	154
Examples	155
See Also	156
ListAllowedNodeTypeModifications	157
Request Parameters	157
Response Elements	157
Errors	158
Examples	158
See Also	159
ListTagsForResource	160
Request Parameters	160
Response Elements	160
Errors	160
Examples	161
See Also	162
ModifyCacheCluster	163
Request Parameters	163
Response Elements	168
Errors	168
Examples	169
See Also	170
ModifyCacheParameterGroup	171

Request Parameters	171
Response Elements	171
Errors	171
Examples	172
See Also	172
ModifyCacheSubnetGroup	174
Request Parameters	174
Response Elements	174
Errors	174
Examples	175
See Also	176
ModifyGlobalReplicationGroup	177
Request Parameters	177
Response Elements	178
Errors	178
See Also	178
ModifyReplicationGroup	179
Request Parameters	179
Response Elements	183
Errors	183
Examples	185
See Also	186
ModifyReplicationGroupShardConfiguration	187
Request Parameters	187
Response Elements	188
Errors	188
Examples	189
See Also	191
ModifyUser	192
Request Parameters	192
Response Elements	192
Errors	193
See Also	194
ModifyUserGroup	195
Request Parameters	195
Response Elements	195
Errors	196
See Also	197
PurchaseReservedCacheNodesOffering	198
Request Parameters	198
Response Elements	198
Errors	199
Examples	199
See Also	200
RebalanceSlotsInGlobalReplicationGroup	201
Request Parameters	201
Response Elements	201
Errors	201
See Also	202
RebootCacheCluster	203
Request Parameters	203
Response Elements	203
Errors	203
Examples	204
See Also	205
RemoveTagsFromResource	206
Request Parameters	206

Response Elements	206
Errors	206
Examples	207
See Also	208
ResetCacheParameterGroup	209
Request Parameters	209
Response Elements	209
Errors	209
Examples	210
See Also	210
RevokeCacheSecurityGroupIngress	212
Request Parameters	212
Response Elements	212
Errors	212
Examples	213
See Also	214
StartMigration	215
Request Parameters	215
Response Elements	215
Errors	215
See Also	216
TestFailover	217
Request Parameters	217
Response Elements	218
Errors	218
Examples	219
See Also	219
Data Types	220
Authentication	222
Contents	222
See Also	222
AvailabilityZone	223
Contents	223
See Also	223
CacheCluster	224
Contents	224
See Also	229
CacheEngineVersion	230
Contents	230
See Also	230
CacheNode	232
Contents	233
See Also	234
CacheNodeTypeSpecificParameter	235
Contents	235
See Also	236
CacheNodeTypeSpecificValue	237
Contents	237
See Also	237
CacheNodeUpdateStatus	238
Contents	238
See Also	239
CacheParameterGroup	240
Contents	240
See Also	240
CacheParameterGroupStatus	242
Contents	242

See Also	242
CacheSecurityGroup	243
Contents	243
See Also	243
CacheSecurityGroupMembership	245
Contents	245
See Also	245
CacheSubnetGroup	246
Contents	246
See Also	246
CloudWatchLogsDestinationDetails	248
Contents	248
See Also	248
ConfigureShard	249
Contents	249
See Also	250
CustomerNodeEndpoint	251
Contents	251
See Also	251
DestinationDetails	252
Contents	252
See Also	252
EC2SecurityGroup	253
Contents	253
See Also	253
Endpoint	254
Contents	254
See Also	254
EngineDefaults	255
Contents	255
See Also	255
Event	256
Contents	256
See Also	256
Filter	257
Contents	257
See Also	257
GlobalNodeGroup	258
Contents	258
See Also	258
GlobalReplicationGroup	259
Contents	259
See Also	261
GlobalReplicationGroupInfo	262
Contents	262
See Also	262
GlobalReplicationGroupMember	263
Contents	263
See Also	263
KinesisFirehoseDestinationDetails	265
Contents	265
See Also	265
LogDeliveryConfiguration	266
Contents	266
See Also	267
LogDeliveryConfigurationRequest	268
Contents	268

See Also	268
NodeGroup	270
Contents	270
See Also	271
NodeGroupConfiguration	272
Contents	272
See Also	273
NodeGroupMember	274
Contents	274
See Also	274
NodeGroupMemberUpdateStatus	276
Contents	276
See Also	277
NodeGroupUpdateStatus	278
Contents	278
See Also	278
NodeSnapshot	279
Contents	279
See Also	280
NotificationConfiguration	281
Contents	281
See Also	281
Parameter	282
Contents	282
See Also	283
ParameterNameValue	284
Contents	284
See Also	284
PendingLogDeliveryConfiguration	285
Contents	285
See Also	285
PendingModifiedValues	286
Contents	286
See Also	287
ProcessedUpdateAction	288
Contents	288
See Also	288
RecurringCharge	289
Contents	289
See Also	289
RegionalConfiguration	290
Contents	290
See Also	290
ReplicationGroup	291
Contents	291
See Also	294
ReplicationGroupPendingModifiedValues	296
Contents	296
See Also	297
ReservedCacheNode	298
Contents	298
See Also	300
ReservedCacheNodesOffering	302
Contents	302
See Also	304
ReshardingConfiguration	305
Contents	305

See Also	305
ReshardingStatus	306
Contents	306
See Also	306
SecurityGroupMembership	307
Contents	307
See Also	307
ServiceUpdate	308
Contents	308
See Also	309
SlotMigration	310
Contents	310
See Also	310
Snapshot	311
Contents	311
See Also	316
Subnet	317
Contents	317
See Also	317
SubnetOutpost	318
Contents	318
See Also	318
Tag	319
Contents	319
See Also	319
TimeRangeFilter	320
Contents	320
See Also	320
UnprocessedUpdateAction	321
Contents	321
See Also	321
UpdateAction	322
Contents	322
See Also	324
User	325
Contents	325
See Also	326
UserGroup	327
Contents	327
See Also	328
UserGroupPendingChanges	329
Contents	329
See Also	329
UserGroupsUpdateStatus	330
Contents	330
See Also	330
Common Parameters	331
Common Errors	333

Welcome

Amazon ElastiCache is a web service that makes it easier to set up, operate, and scale a distributed cache in the cloud.

With ElastiCache, customers get all of the benefits of a high-performance, in-memory cache with less of the administrative burden involved in launching and managing a distributed cache. The service makes setup, scaling, and cluster failure handling much simpler than in a self-managed cache deployment.

In addition, through integration with Amazon CloudWatch, customers get enhanced visibility into the key performance statistics associated with their cache and can receive alarms if a part of their cache runs hot.

This document was last published on October 6, 2021.

Actions

The following actions are supported:

- [AddTagsToResource](#) (p. 4)
- [AuthorizeCacheSecurityGroupIngress](#) (p. 7)
- [BatchApplyUpdateAction](#) (p. 10)
- [BatchStopUpdateAction](#) (p. 12)
- [CompleteMigration](#) (p. 14)
- [CopySnapshot](#) (p. 16)
- [CreateCacheCluster](#) (p. 20)
- [CreateCacheParameterGroup](#) (p. 30)
- [CreateCacheSecurityGroup](#) (p. 33)
- [CreateCacheSubnetGroup](#) (p. 36)
- [CreateGlobalReplicationGroup](#) (p. 39)
- [CreateReplicationGroup](#) (p. 41)
- [CreateSnapshot](#) (p. 54)
- [CreateUser](#) (p. 58)
- [CreateUserGroup](#) (p. 62)
- [DecreaseNodeGroupsInGlobalReplicationGroup](#) (p. 65)
- [DecreaseReplicaCount](#) (p. 67)
- [DeleteCacheCluster](#) (p. 71)
- [DeleteCacheParameterGroup](#) (p. 74)
- [DeleteCacheSecurityGroup](#) (p. 76)
- [DeleteCacheSubnetGroup](#) (p. 78)
- [DeleteGlobalReplicationGroup](#) (p. 80)
- [DeleteReplicationGroup](#) (p. 82)
- [DeleteSnapshot](#) (p. 85)
- [DeleteUser](#) (p. 88)
- [DeleteUserGroup](#) (p. 91)
- [DescribeCacheClusters](#) (p. 93)
- [DescribeCacheEngineVersions](#) (p. 97)
- [DescribeCacheParameterGroups](#) (p. 100)
- [DescribeCacheParameters](#) (p. 103)
- [DescribeCacheSecurityGroups](#) (p. 107)
- [DescribeCacheSubnetGroups](#) (p. 110)
- [DescribeEngineDefaultParameters](#) (p. 113)
- [DescribeEvents](#) (p. 116)
- [DescribeGlobalReplicationGroups](#) (p. 119)
- [DescribeReplicationGroups](#) (p. 121)
- [DescribeReservedCacheNodes](#) (p. 124)
- [DescribeReservedCacheNodesOfferings](#) (p. 129)
- [DescribeServiceUpdates](#) (p. 134)
- [DescribeSnapshots](#) (p. 136)

- [DescribeUpdateActions](#) (p. 140)
- [DescribeUserGroups](#) (p. 143)
- [DescribeUsers](#) (p. 145)
- [DisassociateGlobalReplicationGroup](#) (p. 147)
- [FailoverGlobalReplicationGroup](#) (p. 149)
- [IncreaseNodeGroupsInGlobalReplicationGroup](#) (p. 151)
- [IncreaseReplicaCount](#) (p. 153)
- [ListAllowedNodeTypeModifications](#) (p. 157)
- [ListTagsForResource](#) (p. 160)
- [ModifyCacheCluster](#) (p. 163)
- [ModifyCacheParameterGroup](#) (p. 171)
- [ModifyCacheSubnetGroup](#) (p. 174)
- [ModifyGlobalReplicationGroup](#) (p. 177)
- [ModifyReplicationGroup](#) (p. 179)
- [ModifyReplicationGroupShardConfiguration](#) (p. 187)
- [ModifyUser](#) (p. 192)
- [ModifyUserGroup](#) (p. 195)
- [PurchaseReservedCacheNodesOffering](#) (p. 198)
- [RebalanceSlotsInGlobalReplicationGroup](#) (p. 201)
- [RebootCacheCluster](#) (p. 203)
- [RemoveTagsFromResource](#) (p. 206)
- [ResetCacheParameterGroup](#) (p. 209)
- [RevokeCacheSecurityGroupIngress](#) (p. 212)
- [StartMigration](#) (p. 215)
- [TestFailover](#) (p. 217)

AddTagsToResource

A tag is a key-value pair where the key and value are case-sensitive. You can use tags to categorize and track all your ElastiCache resources, with the exception of global replication group. When you add or remove tags on replication groups, those actions will be replicated to all nodes in the replication group. For more information, see [Resource-level permissions](#).

For example, you can use cost-allocation tags to your ElastiCache resources, Amazon generates a cost allocation report as a comma-separated value (CSV) file with your usage and costs aggregated by your tags. You can apply tags that represent business categories (such as cost centers, application names, or owners) to organize your costs across multiple services.

For more information, see [Using Cost Allocation Tags in Amazon ElastiCache](#) in the *ElastiCache User Guide*.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

ResourceName

The Amazon Resource Name (ARN) of the resource to which the tags are to be added, for example `arn:aws:elasticache:us-west-2:0123456789:cluster:myCluster` or `arn:aws:elasticache:us-west-2:0123456789:snapshot:mySnapshot`. ElastiCache resources are *cluster* and *snapshot*.

For more information about ARNs, see [Amazon Resource Names \(ARNs\) and Amazon Service Namespaces](#).

Type: String

Required: Yes

Tags.Tag.N

A list of tags to be added to this resource. A tag is a key-value pair. A tag key must be accompanied by a tag value, although null is accepted.

Type: Array of [Tag](#) (p. 319) objects

Required: Yes

Response Elements

The following element is returned by the service.

TagList.Tag.N

A list of tags as key-value pairs.

Type: Array of [Tag](#) (p. 319) objects

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

CacheClusterNotFound

The requested cluster ID does not refer to an existing cluster.

HTTP Status Code: 404

CacheParameterGroupNotFound

The requested cache parameter group name does not refer to an existing cache parameter group.

HTTP Status Code: 404

CacheSecurityGroupNotFound

The requested cache security group name does not refer to an existing cache security group.

HTTP Status Code: 404

CacheSubnetGroupNotFoundFault

The requested cache subnet group name does not refer to an existing cache subnet group.

HTTP Status Code: 400

InvalidARN

The requested Amazon Resource Name (ARN) does not refer to an existing resource.

HTTP Status Code: 400

InvalidReplicationGroupState

The requested replication group is not in the available state.

HTTP Status Code: 400

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

ReservedCacheNodeNotFound

The requested reserved cache node was not found.

HTTP Status Code: 404

SnapshotNotFoundFault

The requested snapshot name does not refer to an existing snapshot.

HTTP Status Code: 404

TagQuotaPerResourceExceeded

The request cannot be processed because it would cause the resource to have more than the allowed number of tags. The maximum number of tags permitted on a resource is 50.

HTTP Status Code: 400

UserGroupNotFound

The user group was not found or does not exist

HTTP Status Code: 404

UserNotFound

The user does not exist or could not be found.

HTTP Status Code: 404

Examples

AddTagsToResource

This example illustrates one usage of AddTagsToResource.

Sample Request

```
https://elasticache.us-east-1.amazonaws.com/  
?Action=AddTagsToResource  
&ResourceName=arn:aws:elasticache:us-west-2:0123456789:cluster:myCluster  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Tags.Tag.1.Key=Service  
&Tags.Tag.1.Value=elasticache  
&Tags.Tag.2.Key=Region  
&Tags.Tag.2.Value=us-west-2  
&Version=2015-02-02  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

AuthorizeCacheSecurityGroupIngress

Allows network ingress to a cache security group. Applications using ElastiCache must be running on Amazon EC2, and Amazon EC2 security groups are used as the authorization mechanism.

Note

You cannot authorize ingress from an Amazon EC2 security group in one region to an ElastiCache cluster in another region.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheSecurityGroupName

The cache security group that allows network ingress.

Type: String

Required: Yes

EC2SecurityGroupName

The Amazon EC2 security group to be authorized for ingress to the cache security group.

Type: String

Required: Yes

EC2SecurityGroupOwnerId

The Amazon account number of the Amazon EC2 security group owner. Note that this is not the same thing as an Amazon access key ID - you must provide a valid Amazon account number for this parameter.

Type: String

Required: Yes

Response Elements

The following element is returned by the service.

CacheSecurityGroup

Represents the output of one of the following operations:

- `AuthorizeCacheSecurityGroupIngress`
- `CreateCacheSecurityGroup`
- `RevokeCacheSecurityGroupIngress`

Type: [CacheSecurityGroup](#) (p. 243) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

AuthorizationAlreadyExists

The specified Amazon EC2 security group is already authorized for the specified cache security group.

HTTP Status Code: 400

CacheSecurityGroupNotFound

The requested cache security group name does not refer to an existing cache security group.

HTTP Status Code: 404

InvalidCacheSecurityGroupState

The current state of the cache security group does not allow deletion.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

Examples

AuthorizeCacheSecurityGroupIngress

This example illustrates one usage of AuthorizeCacheSecurityGroupIngress.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=AuthorizeCacheSecurityGroupIngress  
&EC2SecurityGroupName=default  
&CacheSecurityGroupName=mygroup  
&EC2SecurityGroupOwnerId=1234-5678-1234  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Example

This example illustrates one usage of AuthorizeCacheSecurityGroupIngress.

Sample Response

```
<AuthorizeCacheSecurityGroupIngressResponse xmlns="http://elasticache.amazonaws.com/  
doc/2015-02-02/">  
  <AuthorizeCacheSecurityGroupIngressResult>  
    <CacheSecurityGroup>
```

```
<EC2SecurityGroups>
  <EC2SecurityGroup>
    <Status>authorizing</Status>
    <EC2SecurityGroupName>default</EC2SecurityGroupName>
    <EC2SecurityGroupOwnerId>565419523791</EC2SecurityGroupOwnerId>
  </EC2SecurityGroup>
</EC2SecurityGroups>
<CacheSecurityGroupName>mygroup</CacheSecurityGroupName>
<OwnerId>123456781234</OwnerId>
<Description>My security group</Description>
</CacheSecurityGroup>
</AuthorizeCacheSecurityGroupIngress>
<ResponseMetadata>
  <RequestId>817fa999-3647-11e0-ae57-f96cfe56749c</RequestId>
</ResponseMetadata>
</AuthorizeCacheSecurityGroupIngressResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

BatchApplyUpdateAction

Apply the service update. For more information on service updates and applying them, see [Applying Service Updates](#).

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

ServiceUpdateName

The unique ID of the service update

Type: String

Required: Yes

CacheClusterIds.member.N

The cache cluster IDs

Type: Array of strings

Array Members: Maximum number of 20 items.

Required: No

ReplicationGroupIds.member.N

The replication group IDs

Type: Array of strings

Array Members: Maximum number of 20 items.

Required: No

Response Elements

The following elements are returned by the service.

ProcessedUpdateActions.ProcessedUpdateAction.N

Update actions that have been processed successfully

Type: Array of [ProcessedUpdateAction](#) (p. 288) objects

UnprocessedUpdateActions.UnprocessedUpdateAction.N

Update actions that haven't been processed successfully

Type: Array of [UnprocessedUpdateAction](#) (p. 321) objects

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

ServiceUpdateNotFoundFault

The service update doesn't exist

HTTP Status Code: 404

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

BatchStopUpdateAction

Stop the service update. For more information on service updates and stopping them, see [Stopping Service Updates](#).

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

ServiceUpdateName

The unique ID of the service update

Type: String

Required: Yes

CacheClusterIds.member.N

The cache cluster IDs

Type: Array of strings

Array Members: Maximum number of 20 items.

Required: No

ReplicationGroupIds.member.N

The replication group IDs

Type: Array of strings

Array Members: Maximum number of 20 items.

Required: No

Response Elements

The following elements are returned by the service.

ProcessedUpdateActions.ProcessedUpdateAction.N

Update actions that have been processed successfully

Type: Array of [ProcessedUpdateAction \(p. 288\)](#) objects

UnprocessedUpdateActions.UnprocessedUpdateAction.N

Update actions that haven't been processed successfully

Type: Array of [UnprocessedUpdateAction \(p. 321\)](#) objects

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

ServiceUpdateNotFoundFault

The service update doesn't exist

HTTP Status Code: 404

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

CompleteMigration

Complete the migration of data.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

ReplicationGroupId

The ID of the replication group to which data is being migrated.

Type: String

Required: Yes

Force

Forces the migration to stop without ensuring that data is in sync. It is recommended to use this option only to abort the migration and not recommended when application wants to continue migration to ElastiCache.

Type: Boolean

Required: No

Response Elements

The following element is returned by the service.

ReplicationGroup

Contains all of the attributes of a specific Redis replication group.

Type: [ReplicationGroup \(p. 291\)](#) object

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

InvalidReplicationGroupState

The requested replication group is not in the `available` state.

HTTP Status Code: 400

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

ReplicationGroupNotUnderMigrationFault

The designated replication group is not available for data migration.

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

CopySnapshot

Makes a copy of an existing snapshot.

Note

This operation is valid for Redis only.

Important

Users or groups that have permissions to use the `CopySnapshot` operation can create their own Amazon S3 buckets and copy snapshots to it. To control access to your snapshots, use an IAM policy to control who has the ability to use the `CopySnapshot` operation. For more information about using IAM to control the use of ElastiCache operations, see [Exporting Snapshots and Authentication & Access Control](#).

You could receive the following error messages.

Error Messages

- **Error Message:** The S3 bucket %s is outside of the region.
Solution: Create an Amazon S3 bucket in the same region as your snapshot. For more information, see [Step 1: Create an Amazon S3 Bucket](#) in the ElastiCache User Guide.
- **Error Message:** The S3 bucket %s does not exist.
Solution: Create an Amazon S3 bucket in the same region as your snapshot. For more information, see [Step 1: Create an Amazon S3 Bucket](#) in the ElastiCache User Guide.
- **Error Message:** The S3 bucket %s is not owned by the authenticated user.
Solution: Create an Amazon S3 bucket in the same region as your snapshot. For more information, see [Step 1: Create an Amazon S3 Bucket](#) in the ElastiCache User Guide.
- **Error Message:** The authenticated user does not have sufficient permissions to perform the desired activity.
Solution: Contact your system administrator to get the needed permissions.
- **Error Message:** The S3 bucket %s already contains an object with key %s.
Solution: Give the `TargetSnapshotName` a new and unique value. If exporting a snapshot, you could alternatively create a new Amazon S3 bucket and use this same value for `TargetSnapshotName`.
- **Error Message:** ElastiCache has not been granted READ permissions %s on the S3 Bucket.
Solution: Add List and Read permissions on the bucket. For more information, see [Step 2: Grant ElastiCache Access to Your Amazon S3 Bucket](#) in the ElastiCache User Guide.
- **Error Message:** ElastiCache has not been granted WRITE permissions %s on the S3 Bucket.
Solution: Add Upload/Delete permissions on the bucket. For more information, see [Step 2: Grant ElastiCache Access to Your Amazon S3 Bucket](#) in the ElastiCache User Guide.
- **Error Message:** ElastiCache has not been granted READ_ACP permissions %s on the S3 Bucket.
Solution: Add View Permissions on the bucket. For more information, see [Step 2: Grant ElastiCache Access to Your Amazon S3 Bucket](#) in the ElastiCache User Guide.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

SourceSnapshotName

The name of an existing snapshot from which to make a copy.

Type: String

Required: Yes

TargetSnapshotName

A name for the snapshot copy. ElastiCache does not permit overwriting a snapshot, therefore this name must be unique within its context - ElastiCache or an Amazon S3 bucket if exporting.

Type: String

Required: Yes

KmsKeyId

The ID of the KMS key used to encrypt the target snapshot.

Type: String

Required: No

Tags.Tag.N

A list of tags to be added to this resource. A tag is a key-value pair. A tag key must be accompanied by a tag value, although null is accepted.

Type: Array of [Tag \(p. 319\)](#) objects

Required: No

TargetBucket

The Amazon S3 bucket to which the snapshot is exported. This parameter is used only when exporting a snapshot for external access.

When using this parameter to export a snapshot, be sure Amazon ElastiCache has the needed permissions to this S3 bucket. For more information, see [Step 2: Grant ElastiCache Access to Your Amazon S3 Bucket](#) in the *Amazon ElastiCache User Guide*.

For more information, see [Exporting a Snapshot](#) in the *Amazon ElastiCache User Guide*.

Type: String

Required: No

Response Elements

The following element is returned by the service.

Snapshot

Represents a copy of an entire Redis cluster as of the time when the snapshot was taken.

Type: [Snapshot \(p. 311\)](#) object

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidSnapshotState

The current state of the snapshot does not allow the requested operation to occur.

HTTP Status Code: 400

SnapshotAlreadyExistsFault

You already have a snapshot with the given name.

HTTP Status Code: 400

SnapshotNotFoundFault

The requested snapshot name does not refer to an existing snapshot.

HTTP Status Code: 404

SnapshotQuotaExceededFault

The request cannot be processed because it would exceed the maximum number of snapshots.

HTTP Status Code: 400

TagQuotaPerResourceExceeded

The request cannot be processed because it would cause the resource to have more than the allowed number of tags. The maximum number of tags permitted on a resource is 50.

HTTP Status Code: 400

Examples

Snapshot copy

The following example makes a copy of the snapshot `automatic.my-redis-primary-2016-04-27-03-15` named `my-snapshot-copy`.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/
?Action=CopySnapshot
&SourceSnapshotName=automatic.my-redis-primary-2016-04-27-03-15
&TargetSnapshotName=my-snapshot-copy
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20141201T220302Z
&Version=2015-02-02
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Date=20141201T220302Z
&X-Amz-SignedHeaders=Host
```

```
&X-Amz-Expires=20141201T220302Z  
&X-Amz-Credential=<credential>  
&X-Amz-Signature=<signature>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

CreateCacheCluster

Creates a cluster. All nodes in the cluster run the same protocol-compliant cache engine software, either Memcached or Redis.

This operation is not supported for Redis (cluster mode enabled) clusters.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheClusterId

The node group (shard) identifier. This parameter is stored as a lowercase string.

Constraints:

- A name must contain from 1 to 50 alphanumeric characters or hyphens.
- The first character must be a letter.
- A name cannot end with a hyphen or contain two consecutive hyphens.

Type: String

Required: Yes

AuthToken

Reserved parameter. The password used to access a password protected server.

Password constraints:

- Must be only printable ASCII characters.
- Must be at least 16 characters and no more than 128 characters in length.
- The only permitted printable special characters are !, &, #, \$, ^, <, >, and -. Other printable special characters cannot be used in the AUTH token.

For more information, see [AUTH password](http://redis.io/commands/AUTH) at <http://redis.io/commands/AUTH>.

Type: String

Required: No

AutoMinorVersionUpgrade

This parameter is currently disabled.

Type: Boolean

Required: No

AZMode

Specifies whether the nodes in this Memcached cluster are created in a single Availability Zone or created across multiple Availability Zones in the cluster's region.

This parameter is only supported for Memcached clusters.

If the `AZMode` and `PreferredAvailabilityZones` are not specified, ElastiCache assumes `single-az` mode.

Type: String

Valid Values: `single-az` | `cross-az`

Required: No

CacheNodeType

The compute and memory capacity of the nodes in the node group (shard).

The following node types are supported by ElastiCache. Generally speaking, the current generation types provide more memory and computational power at lower cost when compared to their equivalent previous generation counterparts.

- General purpose:
 - Current generation:

M6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

cache.m6g.large, cache.m6g.xlarge, cache.m6g.2xlarge, cache.m6g.4xlarge, cache.m6g.8xlarge, cache.m6g.12xlarge, cache.m6g.16xlarge

Note

For region availability, see [Supported Node Types](#)

M5 node types: cache.m5.large, cache.m5.xlarge, cache.m5.2xlarge, cache.m5.4xlarge, cache.m5.12xlarge, cache.m5.24xlarge

M4 node types: cache.m4.large, cache.m4.xlarge, cache.m4.2xlarge, cache.m4.4xlarge, cache.m4.10xlarge

T3 node types: cache.t3.micro, cache.t3.small, cache.t3.medium

T2 node types: cache.t2.micro, cache.t2.small, cache.t2.medium

- Previous generation: (not recommended)

T1 node types: cache.t1.micro

M1 node types: cache.m1.small, cache.m1.medium, cache.m1.large, cache.m1.xlarge

M3 node types: cache.m3.medium, cache.m3.large, cache.m3.xlarge, cache.m3.2xlarge

- Compute optimized:
 - Previous generation: (not recommended)

C1 node types: cache.c1.xlarge

- Memory optimized:
 - Current generation:

R6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

cache.r6g.large, cache.r6g.xlarge, cache.r6g.2xlarge, cache.r6g.4xlarge, cache.r6g.8xlarge, cache.r6g.12xlarge, cache.r6g.16xlarge

Note

For region availability, see [Supported Node Types](#)

R5 node types: cache.r5.large, cache.r5.xlarge, cache.r5.2xlarge, cache.r5.4xlarge, cache.r5.12xlarge, cache.r5.24xlarge

R4 node types: cache.r4.large, cache.r4.xlarge, cache.r4.2xlarge, cache.r4.4xlarge, cache.r4.8xlarge, cache.r4.16xlarge

- Previous generation: (not recommended)

M2 node types: `cache.m2.xlarge`, `cache.m2.2xlarge`, `cache.m2.4xlarge`

R3 node types: `cache.r3.large`, `cache.r3.xlarge`, `cache.r3.2xlarge`,
`cache.r3.4xlarge`, `cache.r3.8xlarge`

Additional node type info

- All current generation instance types are created in Amazon VPC by default.
- Redis append-only files (AOF) are not supported for T1 or T2 instances.
- Redis Multi-AZ with automatic failover is not supported on T1 instances.
- Redis configuration variables `appendonly` and `appendfsync` are not supported on Redis version 2.8.22 and later.

Type: String

Required: No

CacheParameterGroupName

The name of the parameter group to associate with this cluster. If this argument is omitted, the default parameter group for the specified engine is used. You cannot use any parameter group which has `cluster-enabled='yes'` when creating a cluster.

Type: String

Required: No

CacheSecurityGroupNames.CacheSecurityGroupName.N

A list of security group names to associate with this cluster.

Use this parameter only when you are creating a cluster outside of an Amazon Virtual Private Cloud (Amazon VPC).

Type: Array of strings

Required: No

CacheSubnetGroupName

The name of the subnet group to be used for the cluster.

Use this parameter only when you are creating a cluster in an Amazon Virtual Private Cloud (Amazon VPC).

Important

If you're going to launch your cluster in an Amazon VPC, you need to create a subnet group before you start creating a cluster. For more information, see [Subnets and Subnet Groups](#).

Type: String

Required: No

Engine

The name of the cache engine to be used for this cluster.

Valid values for this parameter are: `memcached` | `redis`

Type: String

Required: No

EngineVersion

The version number of the cache engine to be used for this cluster. To view the supported cache engine versions, use the DescribeCacheEngineVersions operation.

Important: You can upgrade to a newer engine version (see [Selecting a Cache Engine and Version](#)), but you cannot downgrade to an earlier engine version. If you want to use an earlier engine version, you must delete the existing cluster or replication group and create it anew with the earlier engine version.

Type: String

Required: No

LogDeliveryConfigurations.LogDeliveryConfigurationRequest.N

Specifies the destination, format and type of the logs.

Type: Array of [LogDeliveryConfigurationRequest](#) (p. 268) objects

Required: No

NotificationTopicArn

The Amazon Resource Name (ARN) of the Amazon Simple Notification Service (SNS) topic to which notifications are sent.

Note

The Amazon SNS topic owner must be the same as the cluster owner.

Type: String

Required: No

NumCacheNodes

The initial number of cache nodes that the cluster has.

For clusters running Redis, this value must be 1. For clusters running Memcached, this value must be between 1 and 40.

If you need more than 40 nodes for your Memcached cluster, please fill out the ElastiCache Limit Increase Request form at <http://aws.amazon.com/contact-us/elasticache-node-limit-request/>.

Type: Integer

Required: No

OutpostMode

Specifies whether the nodes in the cluster are created in a single outpost or across multiple outposts.

Type: String

Valid Values: `single-outpost` | `cross-outpost`

Required: No

Port

The port number on which each of the cache nodes accepts connections.

Type: Integer

Required: No

PreferredAvailabilityZone

The EC2 Availability Zone in which the cluster is created.

All nodes belonging to this cluster are placed in the preferred Availability Zone. If you want to create your nodes across multiple Availability Zones, use `PreferredAvailabilityZones`.

Default: System chosen Availability Zone.

Type: String

Required: No

PreferredAvailabilityZones.PreferredAvailabilityZone.N

A list of the Availability Zones in which cache nodes are created. The order of the zones in the list is not important.

This option is only supported on Memcached.

Note

If you are creating your cluster in an Amazon VPC (recommended) you can only locate nodes in Availability Zones that are associated with the subnets in the selected subnet group.

The number of Availability Zones listed must equal the value of `NumCacheNodes`.

If you want all the nodes in the same Availability Zone, use `PreferredAvailabilityZone` instead, or repeat the Availability Zone multiple times in the list.

Default: System chosen Availability Zones.

Type: Array of strings

Required: No

PreferredMaintenanceWindow

Specifies the weekly time range during which maintenance on the cluster is performed. It is specified as a range in the format `ddd:hh24:mi-ddd:hh24:mi` (24H Clock UTC). The minimum maintenance window is a 60 minute period. Valid values for `ddd` are:

Type: String

Required: No

PreferredOutpostArn

The outpost ARN in which the cache cluster is created.

Type: String

Required: No

PreferredOutpostArns.PreferredOutpostArn.N

The outpost ARNs in which the cache cluster is created.

Type: Array of strings

Required: No

ReplicationGroupId

The ID of the replication group to which this cluster should belong. If this parameter is specified, the cluster is added to the specified replication group as a read replica; otherwise, the cluster is a standalone primary that is not part of any replication group.

If the specified replication group is Multi-AZ enabled and the Availability Zone is not specified, the cluster is created in Availability Zones that provide the best spread of read replicas across Availability Zones.

Note

This parameter is only valid if the `Engine` parameter is `redis`.

Type: String

Required: No

SecurityGroupIds.SecurityGroupId.N

One or more VPC security groups associated with the cluster.

Use this parameter only when you are creating a cluster in an Amazon Virtual Private Cloud (Amazon VPC).

Type: Array of strings

Required: No

SnapshotArns.SnapshotArn.N

A single-element string list containing an Amazon Resource Name (ARN) that uniquely identifies a Redis RDB snapshot file stored in Amazon S3. The snapshot file is used to populate the node group (shard). The Amazon S3 object name in the ARN cannot contain any commas.

Note

This parameter is only valid if the `Engine` parameter is `redis`.

Example of an Amazon S3 ARN: `arn:aws:s3:::my_bucket/snapshot1.rdb`

Type: Array of strings

Required: No

SnapshotName

The name of a Redis snapshot from which to restore data into the new node group (shard). The snapshot status changes to `restoring` while the new node group (shard) is being created.

Note

This parameter is only valid if the `Engine` parameter is `redis`.

Type: String

Required: No

SnapshotRetentionLimit

The number of days for which ElastiCache retains automatic snapshots before deleting them. For example, if you set `SnapshotRetentionLimit` to 5, a snapshot taken today is retained for 5 days before being deleted.

Note

This parameter is only valid if the `Engine` parameter is `redis`.

Default: 0 (i.e., automatic backups are disabled for this cache cluster).

Type: Integer

Required: No

SnapshotWindow

The daily time range (in UTC) during which ElastiCache begins taking a daily snapshot of your node group (shard).

Example: 05:00–09:00

If you do not specify this parameter, ElastiCache automatically chooses an appropriate time range.

Note

This parameter is only valid if the `Engine` parameter is `redis`.

Type: String

Required: No

Tags.Tag.N

A list of tags to be added to this resource.

Type: Array of [Tag \(p. 319\)](#) objects

Required: No

Response Elements

The following element is returned by the service.

CacheCluster

Contains all of the attributes of a specific cluster.

Type: [CacheCluster \(p. 224\)](#) object

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

CacheClusterAlreadyExists

You already have a cluster with the given identifier.

HTTP Status Code: 400

CacheParameterGroupNotFound

The requested cache parameter group name does not refer to an existing cache parameter group.

HTTP Status Code: 404

CacheSecurityGroupNotFound

The requested cache security group name does not refer to an existing cache security group.

HTTP Status Code: 404

CacheSubnetGroupNotFoundFault

The requested cache subnet group name does not refer to an existing cache subnet group.

HTTP Status Code: 400

ClusterQuotaForCustomerExceeded

The request cannot be processed because it would exceed the allowed number of clusters per customer.

HTTP Status Code: 400

InsufficientCacheClusterCapacity

The requested cache node type is not available in the specified Availability Zone. For more information, see [InsufficientCacheClusterCapacity](#) in the ElastiCache User Guide.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidReplicationGroupState

The requested replication group is not in the available state.

HTTP Status Code: 400

InvalidVPCNetworkStateFault

The VPC network is in an invalid state.

HTTP Status Code: 400

NodeQuotaForClusterExceeded

The request cannot be processed because it would exceed the allowed number of cache nodes in a single cluster.

HTTP Status Code: 400

NodeQuotaForCustomerExceeded

The request cannot be processed because it would exceed the allowed number of cache nodes per customer.

HTTP Status Code: 400

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

TagQuotaPerResourceExceeded

The request cannot be processed because it would cause the resource to have more than the allowed number of tags. The maximum number of tags permitted on a resource is 50.

HTTP Status Code: 400

Examples

CreateCacheCluster

This example illustrates one usage of CreateCacheCluster.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/
?Action=CreateCacheCluster
&CacheClusterId=myMemcachedCluster
&CacheNodeType=cache.m1.small
&CacheSecurityGroupNames.CacheSecurityGroupName.1=default
&Engine=memcached
&NumCacheNodes=3
&PreferredAvailabilityZones.PreferredAvailabilityZone.1=us-west-2a
&PreferredAvailabilityZones.PreferredAvailabilityZone.2=us-west-2b
&PreferredAvailabilityZones.PreferredAvailabilityZone.3=us-west-2c
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2015-02-02
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=[your-access-key-id]/20150202/us-west-2/elasticache/AWS4_request
&X-Amz-Date=20150202T170651Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=[signature-value]
```

Sample Response

```
<CreateCacheClusterResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <CreateCacheClusterResult>
    <CacheCluster>
      <CacheClusterId>myMemcachedCluster</CacheClusterId>
      <CacheClusterStatus>creating</CacheClusterStatus>
      <CacheParameterGroup>
        <CacheParameterGroupName>default.memcached1.4</CacheParameterGroupName>
        <ParameterApplyStatus>in-sync</ParameterApplyStatus>
        <CacheNodeIdsToReboot/>
      </CacheParameterGroup>
      <CacheNodeType>cache.m1.small</CacheNodeType>
      <Engine>memcached</Engine>
      <PendingModifiedValues/>
      <EngineVersion>1.4.14</EngineVersion>
      <AutoMinorVersionUpgrade>true</AutoMinorVersionUpgrade>
      <PreferredMaintenanceWindow>sat:09:00-sat:10:00</PreferredMaintenanceWindow>
      <ClientDownloadLandingPage>https://console.aws.amazon.com/elasticache/home#client-
download:</ClientDownloadLandingPage>
      <CacheSecurityGroups>
        <CacheSecurityGroup>
          <CacheSecurityGroupName>default</CacheSecurityGroupName>
          <Status>active</Status>
        </CacheSecurityGroup>
      </CacheSecurityGroups>
      <NumCacheNodes>3</NumCacheNodes>
    </CacheCluster>
  </CreateCacheClusterResult>
  <ResponseMetadata>
    <RequestId>69134921-10f9-11e4-81bb-d76bad68b8fd</RequestId>
  </ResponseMetadata>
</CreateCacheClusterResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

CreateCacheParameterGroup

Creates a new Amazon ElastiCache cache parameter group. An ElastiCache cache parameter group is a collection of parameters and their values that are applied to all of the nodes in any cluster or replication group using the CacheParameterGroup.

A newly created CacheParameterGroup is an exact duplicate of the default parameter group for the CacheParameterGroupFamily. To customize the newly created CacheParameterGroup you can change the values of specific parameters. For more information, see:

- [ModifyCacheParameterGroup](#) in the ElastiCache API Reference.
- [Parameters and Parameter Groups](#) in the ElastiCache User Guide.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheParameterGroupFamily

The name of the cache parameter group family that the cache parameter group can be used with.

Valid values are: memcached1.4 | memcached1.5 | memcached1.6 | redis2.6 | redis2.8 | redis3.2 | redis4.0 | redis5.0 | redis6.x |

Type: String

Required: Yes

CacheParameterGroupName

A user-specified name for the cache parameter group.

Type: String

Required: Yes

Description

A user-specified description for the cache parameter group.

Type: String

Required: Yes

Tags.Tag.N

A list of tags to be added to this resource. A tag is a key-value pair. A tag key must be accompanied by a tag value, although null is accepted.

Type: Array of [Tag](#) (p. 319) objects

Required: No

Response Elements

The following element is returned by the service.

CacheParameterGroup

Represents the output of a `CreateCacheParameterGroup` operation.

Type: [CacheParameterGroup](#) (p. 240) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

CacheParameterGroupAlreadyExists

A cache parameter group with the requested name already exists.

HTTP Status Code: 400

CacheParameterGroupQuotaExceeded

The request cannot be processed because it would exceed the maximum number of cache security groups.

HTTP Status Code: 400

InvalidCacheParameterGroupState

The current state of the cache parameter group does not allow the requested operation to occur.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

TagQuotaPerResourceExceeded

The request cannot be processed because it would cause the resource to have more than the allowed number of tags. The maximum number of tags permitted on a resource is 50.

HTTP Status Code: 400

Examples

CreateCacheParameterGroup

This example illustrates one usage of `CreateCacheParameterGroup`.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=CreateCacheParameterGroup  
&CacheParameterGroupFamily=memcached1.4  
&CacheParameterGroupName=mycacheparametergroup1
```

```
&Description=My%20custom%20Redis%20cache%20parameter%20group
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20150202T192317Z
&Version=2015-02-02
&X-Amz-Credential=<credential>
```

Sample Response

```
<CreateCacheParameterGroupResponse xmlns="http://elasticache.amazonaws.com/
doc/2015-02-02/">
  <CreateCacheParameterGroupResult>
    <CacheParameterGroup>
      <CacheParameterGroupName>mycacheparametergroup1</CacheParameterGroupName>
      <CacheParameterGroupFamily>memcached1.4</CacheParameterGroupFamily>
      <Description>My first cache parameter group</Description>
    </CacheParameterGroup>
  </CreateCacheParameterGroupResult>
  <ResponseMetadata>
    <RequestId>05699541-b7f9-11e0-9326-b7275b9d4a6c</RequestId>
  </ResponseMetadata>
</CreateCacheParameterGroupResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

CreateCacheSecurityGroup

Creates a new cache security group. Use a cache security group to control access to one or more clusters.

Cache security groups are only used when you are creating a cluster outside of an Amazon Virtual Private Cloud (Amazon VPC). If you are creating a cluster inside of a VPC, use a cache subnet group instead. For more information, see [CreateCacheSubnetGroup](#).

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheSecurityGroupName

A name for the cache security group. This value is stored as a lowercase string.

Constraints: Must contain no more than 255 alphanumeric characters. Cannot be the word "Default".

Example: mysecuritygroup

Type: String

Required: Yes

Description

A description for the cache security group.

Type: String

Required: Yes

Tags.Tag.N

A list of tags to be added to this resource. A tag is a key-value pair. A tag key must be accompanied by a tag value, although null is accepted.

Type: Array of [Tag](#) (p. 319) objects

Required: No

Response Elements

The following element is returned by the service.

CacheSecurityGroup

Represents the output of one of the following operations:

- [AuthorizeCacheSecurityGroupIngress](#)
- [CreateCacheSecurityGroup](#)
- [RevokeCacheSecurityGroupIngress](#)

Type: [CacheSecurityGroup](#) (p. 243) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

CacheSecurityGroupAlreadyExists

A cache security group with the specified name already exists.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

QuotaExceeded.CacheSecurityGroup

The request cannot be processed because it would exceed the allowed number of cache security groups.

HTTP Status Code: 400

TagQuotaPerResourceExceeded

The request cannot be processed because it would cause the resource to have more than the allowed number of tags. The maximum number of tags permitted on a resource is 50.

HTTP Status Code: 400

Examples

CreateCacheSecurityGroup

This example illustrates one usage of CreateCacheSecurityGroup.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=CreateCacheSecurityGroup  
&CacheSecurityGroupName=mycachesecuritygroup  
&Description=My%20cache%20security%20group  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<CreateCacheSecurityGroupResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">  
  <CreateCacheSecurityGroupResult>  
    <CacheSecurityGroup>  
      <EC2SecurityGroups/>  
      <CacheSecurityGroupName>mycachesecuritygroup</CacheSecurityGroupName>  
      <OwnerId>123456789012</OwnerId>  
      <Description>My cache security group</Description>  
    </CacheSecurityGroup>  
  </CreateCacheSecurityGroupResult>
```

```
<ResponseMetadata>  
  <RequestId>2b1c8035-b7fa-11e0-9326-b7275b9d4a6c</RequestId>  
</ResponseMetadata>  
</CreateCacheSecurityGroupResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

CreateCacheSubnetGroup

Creates a new cache subnet group.

Use this parameter only when you are creating a cluster in an Amazon Virtual Private Cloud (Amazon VPC).

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheSubnetGroupDescription

A description for the cache subnet group.

Type: String

Required: Yes

CacheSubnetGroupName

A name for the cache subnet group. This value is stored as a lowercase string.

Constraints: Must contain no more than 255 alphanumeric characters or hyphens.

Example: mysubnetgroup

Type: String

Required: Yes

SubnetIds.SubnetIdentifier.N

A list of VPC subnet IDs for the cache subnet group.

Type: Array of strings

Required: Yes

Tags.Tag.N

A list of tags to be added to this resource. A tag is a key-value pair. A tag key must be accompanied by a tag value, although null is accepted.

Type: Array of [Tag](#) (p. 319) objects

Required: No

Response Elements

The following element is returned by the service.

CacheSubnetGroup

Represents the output of one of the following operations:

- `CreateCacheSubnetGroup`
- `ModifyCacheSubnetGroup`

Type: [CacheSubnetGroup](#) (p. 246) object

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

CacheSubnetGroupAlreadyExists

The requested cache subnet group name is already in use by an existing cache subnet group.

HTTP Status Code: 400

CacheSubnetGroupQuotaExceeded

The request cannot be processed because it would exceed the allowed number of cache subnet groups.

HTTP Status Code: 400

CacheSubnetQuotaExceededFault

The request cannot be processed because it would exceed the allowed number of subnets in a cache subnet group.

HTTP Status Code: 400

InvalidSubnet

An invalid subnet identifier was specified.

HTTP Status Code: 400

SubnetNotAllowedFault

At least one subnet ID does not match the other subnet IDs. This mismatch typically occurs when a user sets one subnet ID to a regional Availability Zone and a different one to an outpost. Or when a user sets the subnet ID to an Outpost when not subscribed on this service.

HTTP Status Code: 400

TagQuotaPerResourceExceeded

The request cannot be processed because it would cause the resource to have more than the allowed number of tags. The maximum number of tags permitted on a resource is 50.

HTTP Status Code: 400

Examples

CreateCacheSubnetGroup

This example illustrates one usage of CreateCacheSubnetGroup.

Sample Request

```
https://elasticache.amazonaws.com/  
?Action=CreateCacheSubnetGroup  
&CacheSubnetGroupName=myCachesubnetgroup  
&CacheSubnetGroupDescription=My%20new%20CacheSubnetGroup  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z
```

```
&X-Amz-Credential=<credential>
```

Sample Response

```
<CreateCacheSubnetGroupResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <CreateCacheSubnetGroupResult>
    <CacheSubnetGroup>
      <VpcId>990524496922</VpcId>
      <CacheSubnetGroupDescription>My new CacheSubnetGroup</
CacheSubnetGroupDescription>
      <CacheSubnetGroupName>myCachesubnetgroup</CacheSubnetGroupName>
      <Subnets>
        <Subnet>
          <SubnetStatus>Active</SubnetStatus>
          <SubnetIdentifier>subnet-7c5b4115</SubnetIdentifier>
          <SubnetAvailabilityZone>
            <Name>us-west-2c</Name>
          </SubnetAvailabilityZone>
        </Subnet>
        <Subnet>
          <SubnetStatus>Active</SubnetStatus>
          <SubnetIdentifier>subnet-7b5b4112</SubnetIdentifier>
          <SubnetAvailabilityZone>
            <Name>us-west-2b</Name>
          </SubnetAvailabilityZone>
        </Subnet>
        <Subnet>
          <SubnetStatus>Active</SubnetStatus>
          <SubnetIdentifier>subnet-3ea6bd57</SubnetIdentifier>
          <SubnetAvailabilityZone>
            <Name>us-west-2c</Name>
          </SubnetAvailabilityZone>
        </Subnet>
      </Subnets>
    </CacheSubnetGroup>
  </CreateCacheSubnetGroupResult>
  <ResponseMetadata>
    <RequestId>ed662948-a57b-11df-9e38-7ffab86c801f</RequestId>
  </ResponseMetadata>
</CreateCacheSubnetGroupResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

CreateGlobalReplicationGroup

Global Datastore for Redis offers fully managed, fast, reliable and secure cross-region replication. Using Global Datastore for Redis, you can create cross-region read replica clusters for ElastiCache for Redis to enable low-latency reads and disaster recovery across regions. For more information, see [Replication Across Regions Using Global Datastore](#).

- The **GlobalReplicationGroupIdSuffix** is the name of the Global datastore.
- The **PrimaryReplicationGroupId** represents the name of the primary cluster that accepts writes and will replicate updates to the secondary cluster.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

GlobalReplicationGroupIdSuffix

The suffix name of a Global datastore. Amazon ElastiCache automatically applies a prefix to the Global datastore ID when it is created. Each Amazon Region has its own prefix. For instance, a Global datastore ID created in the US-West-1 region will begin with "dsdfu" along with the suffix name you provide. The suffix, combined with the auto-generated prefix, guarantees uniqueness of the Global datastore name across multiple regions.

For a full list of Amazon Regions and their respective Global datastore ID prefixes, see [Using the Amazon CLI with Global datastores](#).

Type: String

Required: Yes

PrimaryReplicationGroupId

The name of the primary cluster that accepts writes and will replicate updates to the secondary cluster.

Type: String

Required: Yes

GlobalReplicationGroupDescription

Provides details of the Global datastore

Type: String

Required: No

Response Elements

The following element is returned by the service.

GlobalReplicationGroup

Consists of a primary cluster that accepts writes and an associated secondary cluster that resides in a different Amazon region. The secondary cluster accepts only reads. The primary cluster automatically replicates updates to the secondary cluster.

- The **GlobalReplicationGroupIdSuffix** represents the name of the Global datastore, which is what you use to associate a secondary cluster.

Type: [GlobalReplicationGroup](#) (p. 259) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

GlobalReplicationGroupAlreadyExistsFault

The Global datastore name already exists.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidReplicationGroupState

The requested replication group is not in the available state.

HTTP Status Code: 400

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

ServiceLinkedRoleNotFoundFault

The specified service linked role (SLR) was not found.

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

CreateReplicationGroup

Creates a Redis (cluster mode disabled) or a Redis (cluster mode enabled) replication group.

This API can be used to create a standalone regional replication group or a secondary replication group associated with a Global datastore.

A Redis (cluster mode disabled) replication group is a collection of clusters, where one of the clusters is a read/write primary and the others are read-only replicas. Writes to the primary are asynchronously propagated to the replicas.

A Redis cluster-mode enabled cluster is comprised of from 1 to 90 shards (API/CLI: node groups). Each shard has a primary node and up to 5 read-only replica nodes. The configuration can range from 90 shards and 0 replicas to 15 shards and 5 replicas, which is the maximum number of replicas allowed.

The node or shard limit can be increased to a maximum of 500 per cluster if the Redis engine version is 5.0.6 or higher. For example, you can choose to configure a 500 node cluster that ranges between 83 shards (one primary and 5 replicas per shard) and 500 shards (single primary and no replicas). Make sure there are enough available IP addresses to accommodate the increase. Common pitfalls include the subnets in the subnet group have too small a CIDR range or the subnets are shared and heavily used by other clusters. For more information, see [Creating a Subnet Group](#). For versions below 5.0.6, the limit is 250 per cluster.

To request a limit increase, see [Amazon Service Limits](#) and choose the limit type **Nodes per cluster per instance type**.

When a Redis (cluster mode disabled) replication group has been successfully created, you can add one or more read replicas to it, up to a total of 5 read replicas. If you need to increase or decrease the number of node groups (console: shards), you can avail yourself of ElastiCache for Redis' scaling. For more information, see [Scaling ElastiCache for Redis Clusters](#) in the *ElastiCache User Guide*.

Note

This operation is valid for Redis only.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

ReplicationGroupDescription

A user-created description for the replication group.

Type: String

Required: Yes

ReplicationGroupId

The replication group identifier. This parameter is stored as a lowercase string.

Constraints:

- A name must contain from 1 to 40 alphanumeric characters or hyphens.
- The first character must be a letter.
- A name cannot end with a hyphen or contain two consecutive hyphens.

Type: String

Required: Yes

AtRestEncryptionEnabled

A flag that enables encryption at rest when set to `true`.

You cannot modify the value of `AtRestEncryptionEnabled` after the replication group is created. To enable encryption at rest on a replication group you must set `AtRestEncryptionEnabled` to `true` when you create the replication group.

Required: Only available when creating a replication group in an Amazon VPC using redis version 3.2.6, 4.x or later.

Default: `false`

Type: Boolean

Required: No

AuthToken

Reserved parameter. The password used to access a password protected server.

`AuthToken` can be specified only on replication groups where `TransitEncryptionEnabled` is `true`.

Important

For HIPAA compliance, you must specify `TransitEncryptionEnabled` as `true`, an `AuthToken`, and a `CacheSubnetGroup`.

Password constraints:

- Must be only printable ASCII characters.
- Must be at least 16 characters and no more than 128 characters in length.
- The only permitted printable special characters are `!`, `&`, `#`, `$`, `^`, `<`, `>`, and `-`. Other printable special characters cannot be used in the AUTH token.

For more information, see [AUTH password](http://redis.io/commands/AUTH) at <http://redis.io/commands/AUTH>.

Type: String

Required: No

AutomaticFailoverEnabled

Specifies whether a read-only replica is automatically promoted to read/write primary if the existing primary fails.

`AutomaticFailoverEnabled` must be enabled for Redis (cluster mode enabled) replication groups.

Default: `false`

Type: Boolean

Required: No

AutoMinorVersionUpgrade

This parameter is currently disabled.

Type: Boolean

Required: No

CacheNodeType

The compute and memory capacity of the nodes in the node group (shard).

The following node types are supported by ElastiCache. Generally speaking, the current generation types provide more memory and computational power at lower cost when compared to their equivalent previous generation counterparts.

- General purpose:
 - Current generation:

M6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

cache.m6g.large, cache.m6g.xlarge, cache.m6g.2xlarge, cache.m6g.4xlarge, cache.m6g.8xlarge, cache.m6g.12xlarge, cache.m6g.16xlarge

Note

For region availability, see [Supported Node Types](#)

M5 node types: cache.m5.large, cache.m5.xlarge, cache.m5.2xlarge, cache.m5.4xlarge, cache.m5.12xlarge, cache.m5.24xlarge

M4 node types: cache.m4.large, cache.m4.xlarge, cache.m4.2xlarge, cache.m4.4xlarge, cache.m4.10xlarge

T3 node types: cache.t3.micro, cache.t3.small, cache.t3.medium

T2 node types: cache.t2.micro, cache.t2.small, cache.t2.medium

- Previous generation: (not recommended)

T1 node types: cache.t1.micro

M1 node types: cache.m1.small, cache.m1.medium, cache.m1.large, cache.m1.xlarge

M3 node types: cache.m3.medium, cache.m3.large, cache.m3.xlarge, cache.m3.2xlarge

- Compute optimized:
 - Previous generation: (not recommended)

C1 node types: cache.c1.xlarge

- Memory optimized:
 - Current generation:

R6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

cache.r6g.large, cache.r6g.xlarge, cache.r6g.2xlarge, cache.r6g.4xlarge, cache.r6g.8xlarge, cache.r6g.12xlarge, cache.r6g.16xlarge

Note

For region availability, see [Supported Node Types](#)

R5 node types: cache.r5.large, cache.r5.xlarge, cache.r5.2xlarge, cache.r5.4xlarge, cache.r5.12xlarge, cache.r5.24xlarge

R4 node types: cache.r4.large, cache.r4.xlarge, cache.r4.2xlarge, cache.r4.4xlarge, cache.r4.8xlarge, cache.r4.16xlarge

- Previous generation: (not recommended)

M2 node types: `cache.m2.xlarge`, `cache.m2.2xlarge`, `cache.m2.4xlarge`

R3 node types: `cache.r3.large`, `cache.r3.xlarge`, `cache.r3.2xlarge`,
`cache.r3.4xlarge`, `cache.r3.8xlarge`

Additional node type info

- All current generation instance types are created in Amazon VPC by default.
- Redis append-only files (AOF) are not supported for T1 or T2 instances.
- Redis Multi-AZ with automatic failover is not supported on T1 instances.
- Redis configuration variables `appendonly` and `appendfsync` are not supported on Redis version 2.8.22 and later.

Type: String

Required: No

CacheParameterGroupName

The name of the parameter group to associate with this replication group. If this argument is omitted, the default cache parameter group for the specified engine is used.

If you are running Redis version 3.2.4 or later, only one node group (shard), and want to use a default parameter group, we recommend that you specify the parameter group by name.

- To create a Redis (cluster mode disabled) replication group, use
`CacheParameterGroupName=default.redis3.2.`
- To create a Redis (cluster mode enabled) replication group, use
`CacheParameterGroupName=default.redis3.2.cluster.on.`

Type: String

Required: No

CacheSecurityGroupNames.CacheSecurityGroupName.N

A list of cache security group names to associate with this replication group.

Type: Array of strings

Required: No

CacheSubnetGroupName

The name of the cache subnet group to be used for the replication group.

Important

If you're going to launch your cluster in an Amazon VPC, you need to create a subnet group before you start creating a cluster. For more information, see [Subnets and Subnet Groups](#).

Type: String

Required: No

Engine

The name of the cache engine to be used for the clusters in this replication group. Must be Redis.

Type: String

Required: No

EngineVersion

The version number of the cache engine to be used for the clusters in this replication group. To view the supported cache engine versions, use the `DescribeCacheEngineVersions` operation.

Important: You can upgrade to a newer engine version (see [Selecting a Cache Engine and Version](#)) in the *ElastiCache User Guide*, but you cannot downgrade to an earlier engine version. If you want to use an earlier engine version, you must delete the existing cluster or replication group and create it anew with the earlier engine version.

Type: String

Required: No

GlobalReplicationGroupId

The name of the Global datastore

Type: String

Required: No

KmsKeyId

The ID of the KMS key used to encrypt the disk in the cluster.

Type: String

Required: No

LogDeliveryConfigurations.LogDeliveryConfigurationRequest.N

Specifies the destination, format and type of the logs.

Type: Array of [LogDeliveryConfigurationRequest](#) (p. 268) objects

Required: No

MultiAZEnabled

A flag indicating if you have Multi-AZ enabled to enhance fault tolerance. For more information, see [Minimizing Downtime: Multi-AZ](#).

Type: Boolean

Required: No

NodeGroupConfiguration.NodeGroupConfiguration.N

A list of node group (shard) configuration options. Each node group (shard) configuration has the following members: `PrimaryAvailabilityZone`, `ReplicaAvailabilityZones`, `ReplicaCount`, and `Slots`.

If you're creating a Redis (cluster mode disabled) or a Redis (cluster mode enabled) replication group, you can use this parameter to individually configure each node group (shard), or you can omit this parameter. However, it is required when seeding a Redis (cluster mode enabled) cluster from a S3 rdb file. You must configure each node group (shard) using this parameter because you must specify the slots for each node group.

Type: Array of [NodeGroupConfiguration](#) (p. 272) objects

Required: No

NotificationTopicArn

The Amazon Resource Name (ARN) of the Amazon Simple Notification Service (SNS) topic to which notifications are sent.

Note

The Amazon SNS topic owner must be the same as the cluster owner.

Type: String

Required: No

NumCacheClusters

The number of clusters this replication group initially has.

This parameter is not used if there is more than one node group (shard). You should use `ReplicasPerNodeGroup` instead.

If `AutomaticFailoverEnabled` is `true`, the value of this parameter must be at least 2. If `AutomaticFailoverEnabled` is `false` you can omit this parameter (it will default to 1), or you can explicitly set it to a value between 2 and 6.

The maximum permitted value for `NumCacheClusters` is 6 (1 primary plus 5 replicas).

Type: Integer

Required: No

NumNodeGroups

An optional parameter that specifies the number of node groups (shards) for this Redis (cluster mode enabled) replication group. For Redis (cluster mode disabled) either omit this parameter or set it to 1.

Default: 1

Type: Integer

Required: No

Port

The port number on which each member of the replication group accepts connections.

Type: Integer

Required: No

PreferredCacheClusterAZs.AvailabilityZone.N

A list of EC2 Availability Zones in which the replication group's clusters are created. The order of the Availability Zones in the list is the order in which clusters are allocated. The primary cluster is created in the first AZ in the list.

This parameter is not used if there is more than one node group (shard). You should use `NodeGroupConfiguration` instead.

Note

If you are creating your replication group in an Amazon VPC (recommended), you can only locate clusters in Availability Zones associated with the subnets in the selected subnet group.

The number of Availability Zones listed must equal the value of `NumCacheClusters`.

Default: system chosen Availability Zones.

Type: Array of strings

Required: No

PreferredMaintenanceWindow

Specifies the weekly time range during which maintenance on the cluster is performed. It is specified as a range in the format ddd:hh24:mi-ddd:hh24:mi (24H Clock UTC). The minimum maintenance window is a 60 minute period. Valid values for ddd are:

Specifies the weekly time range during which maintenance on the cluster is performed. It is specified as a range in the format ddd:hh24:mi-ddd:hh24:mi (24H Clock UTC). The minimum maintenance window is a 60 minute period.

Valid values for ddd are:

- sun
- mon
- tue
- wed
- thu
- fri
- sat

Example: sun:23:00-mon:01:30

Type: String

Required: No

PrimaryClusterId

The identifier of the cluster that serves as the primary for this replication group. This cluster must already exist and have a status of available.

This parameter is not required if NumCacheClusters, NumNodeGroups, or ReplicasPerNodeGroup is specified.

Type: String

Required: No

ReplicasPerNodeGroup

An optional parameter that specifies the number of replica nodes in each node group (shard). Valid values are 0 to 5.

Type: Integer

Required: No

SecurityGroupIds.SecurityGroupId.N

One or more Amazon VPC security groups associated with this replication group.

Use this parameter only when you are creating a replication group in an Amazon Virtual Private Cloud (Amazon VPC).

Type: Array of strings

Required: No

SnapshotArns.SnapshotArn.N

A list of Amazon Resource Names (ARN) that uniquely identify the Redis RDB snapshot files stored in Amazon S3. The snapshot files are used to populate the new replication group. The Amazon

S3 object name in the ARN cannot contain any commas. The new replication group will have the number of node groups (console: shards) specified by the parameter *NumNodeGroups* or the number of node groups configured by *NodeGroupConfiguration* regardless of the number of ARNs specified here.

Example of an Amazon S3 ARN: `arn:aws:s3:::my_bucket/snapshot1.rdb`

Type: Array of strings

Required: No

SnapshotName

The name of a snapshot from which to restore data into the new replication group. The snapshot status changes to `restoring` while the new replication group is being created.

Type: String

Required: No

SnapshotRetentionLimit

The number of days for which ElastiCache retains automatic snapshots before deleting them. For example, if you set `SnapshotRetentionLimit` to 5, a snapshot that was taken today is retained for 5 days before being deleted.

Default: 0 (i.e., automatic backups are disabled for this cluster).

Type: Integer

Required: No

SnapshotWindow

The daily time range (in UTC) during which ElastiCache begins taking a daily snapshot of your node group (shard).

Example: `05:00-09:00`

If you do not specify this parameter, ElastiCache automatically chooses an appropriate time range.

Type: String

Required: No

Tags.Tag.N

A list of tags to be added to this resource. Tags are comma-separated key,value pairs (e.g. `Key=myKey, Value=myKeyValue`). You can include multiple tags as shown following: `Key=myKey, Value=myKeyValue Key=mySecondKey, Value=mySecondKeyValue`. Tags on replication groups will be replicated to all nodes.

Type: Array of [Tag](#) (p. 319) objects

Required: No

TransitEncryptionEnabled

A flag that enables in-transit encryption when set to `true`.

You cannot modify the value of `TransitEncryptionEnabled` after the cluster is created. To enable in-transit encryption on a cluster you must set `TransitEncryptionEnabled` to `true` when you create a cluster.

This parameter is valid only if the `Engine` parameter is `redis`, the `EngineVersion` parameter is `3.2.6`, `4.x` or later, and the cluster is being created in an Amazon VPC.

If you enable in-transit encryption, you must also specify a value for `CacheSubnetGroup`.

Required: Only available when creating a replication group in an Amazon VPC using redis version `3.2.6`, `4.x` or later.

Default: `false`

Important

For HIPAA compliance, you must specify `TransitEncryptionEnabled` as `true`, an `AuthToken`, and a `CacheSubnetGroup`.

Type: Boolean

Required: No

UserGroupIds.member.N

The user group to associate with the replication group.

Type: Array of strings

Array Members: Minimum number of 1 item.

Length Constraints: Minimum length of 1.

Pattern: `[a-zA-Z][a-zA-Z0-9\-\]*`

Required: No

Response Elements

The following element is returned by the service.

ReplicationGroup

Contains all of the attributes of a specific Redis replication group.

Type: [ReplicationGroup](#) (p. 291) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

CacheClusterNotFound

The requested cluster ID does not refer to an existing cluster.

HTTP Status Code: 404

CacheParameterGroupNotFound

The requested cache parameter group name does not refer to an existing cache parameter group.

HTTP Status Code: 404

CacheSecurityGroupNotFound

The requested cache security group name does not refer to an existing cache security group.

HTTP Status Code: 404

CacheSubnetGroupNotFoundFault

The requested cache subnet group name does not refer to an existing cache subnet group.

HTTP Status Code: 400

ClusterQuotaForCustomerExceeded

The request cannot be processed because it would exceed the allowed number of clusters per customer.

HTTP Status Code: 400

GlobalReplicationGroupNotFoundFault

The Global datastore does not exist

HTTP Status Code: 404

InsufficientCacheClusterCapacity

The requested cache node type is not available in the specified Availability Zone. For more information, see [InsufficientCacheClusterCapacity](#) in the ElastiCache User Guide.

HTTP Status Code: 400

InvalidCacheClusterState

The requested cluster is not in the `available` state.

HTTP Status Code: 400

InvalidGlobalReplicationGroupState

The Global datastore is not available or in `primary-only` state.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidUserGroupState

The user group is not in an `active` state.

HTTP Status Code: 400

InvalidVPCNetworkStateFault

The VPC network is in an `invalid` state.

HTTP Status Code: 400

NodeGroupsPerReplicationGroupQuotaExceeded

The request cannot be processed because it would exceed the maximum allowed number of node groups (shards) in a single replication group. The default maximum is 90

HTTP Status Code: 400

NodeQuotaForClusterExceeded

The request cannot be processed because it would exceed the allowed number of cache nodes in a single cluster.

HTTP Status Code: 400

NodeQuotaForCustomerExceeded

The request cannot be processed because it would exceed the allowed number of cache nodes per customer.

HTTP Status Code: 400

ReplicationGroupAlreadyExists

The specified replication group already exists.

HTTP Status Code: 400

TagQuotaPerResourceExceeded

The request cannot be processed because it would cause the resource to have more than the allowed number of tags. The maximum number of tags permitted on a resource is 50.

HTTP Status Code: 400

UserGroupNotFound

The user group was not found or does not exist

HTTP Status Code: 404

Examples

CreateReplicationGroup - Redis (cluster mode disabled) Replication Group

The following example creates a Redis (cluster mode disabled) replication group with three nodes (NumCacheClusters=3), a primary and two read replicas. Because a single node group (shard) replication group technically could be either clustered or non-clustered, the parameter group default.redis3.2 is specified, making this a non-clustered replication group.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=CreateReplicationGroup  
&CacheParameterGroup=default.redis3.2  
&Engine=redis  
&EngineVersion=3.2.4  
&NumCacheClusters=3  
&ReplicationGroupDescription=My%20replication%20group  
&ReplicationGroupId=my-repgroup  
&PrimaryClusterId=my-redis-primary  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<CreateReplicationGroupResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <CreateReplicationGroupResult>
    <ReplicationGroup>
      <SnapshottingClusterId>my-redis-primary</SnapshottingClusterId>
      <MemberClusters>
        <ClusterId>my-redis-primary</ClusterId>
      </MemberClusters>
      <ReplicationGroupId>my-repgroup</ReplicationGroupId>
      <Status>creating</Status>
      <PendingModifiedValues />
      <Description>My replication group</Description>
    </ReplicationGroup>
  </CreateReplicationGroupResult>
  <ResponseMetadata>
    <RequestId>f3b7b32d-b9d2-11e3-8a16-7978bb24ffdf</RequestId>
  </ResponseMetadata>
</CreateReplicationGroupResponse>
```

Redis (cluster mode enabled) Replication Group - all shards same profile

The following example creates a Redis (cluster mode enabled) replication group with three node groups (shards) and four replica nodes in each node group (shard). Note the following parameters and their values.

- EngineVersion=3.2.4
- CacheParameterGroup=default.redis3.2.cluster.on
- NumNodeGroups=3
- ReplicasPerNodeGroup=4

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/
?Action=CreateReplicationGroup
&CacheParameterGroup=default.redis3.2.cluster.on
&Engine=redis
&EngineVersion=3.2.4
&ReplicationGroupDescription=My%20replication%20group
&ReplicationGroupId=my-repgroup
&NumNodeGroups=3
&PrimaryClusterId=my-redis-primary
&ReplicasPerNodeGroup=4
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20150202T192317Z
&X-Amz-Credential=<credential>
```

Redis (cluster mode enabled) Replication Group - each shard configured separately

The following example creates a Redis (cluster mode enabled) replication group with two node groups (shards). The first shard has two replica nodes and slots 0-8192. The second shard has one replica and slots 8193-16383. Note the following parameters and their values.

- EngineVersion
- CacheParameterGroup
- NodeGroupConfiguration.NodeGroupConfiguration.n.PrimaryAvailabilityZone
- NodeGroupConfiguration.NodeGroupConfiguration.n.ReplicaAvailabilityZones.AvailabilityZone.n
- NodeGroupConfiguration.NodeGroupConfiguration.n.ReplicaCount
- NodeGroupConfiguration.NodeGroupConfiguration.n.Slots

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/
?Action=CreateReplicationGroup
&CacheParameterGroup=default.redis3.2.cluster.on
&Engine=redis
&EngineVersion=3.2.4
&ReplicationGroupDescription=My%20replication%20group
&ReplicationGroupId=my-repgroup
&NodeGroupConfiguration.NodeGroupConfiguration.1.PrimaryAvailabilityZone=us-east-2a

&NodeGroupConfiguration.NodeGroupConfiguration.1.ReplicaAvailabilityZones.AvailabilityZone.1=us-
east-2b

&NodeGroupConfiguration.NodeGroupConfiguration.1.ReplicaAvailabilityZones.AvailabilityZone.2=us-
east-2c
&NodeGroupConfiguration.NodeGroupConfiguration.1.ReplicaCount=2
&NodeGroupConfiguration.NodeGroupConfiguration.1.Slots=0-8192
&NodeGroupConfiguration.NodeGroupConfiguration.2.PrimaryAvailabilityZone=us-east-2b

&NodeGroupConfiguration.NodeGroupConfiguration.2.ReplicaAvailabilityZones.AvailabilityZone.1=us-
east-2d
&NodeGroupConfiguration.NodeGroupConfiguration.2.ReplicaCount=1
&NodeGroupConfiguration.NodeGroupConfiguration.2.Slots=8193-16383
&PrimaryClusterId=my-redis-primary
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20150202T192317Z
&X-Amz-Credential=<credential>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

CreateSnapshot

Creates a copy of an entire cluster or replication group at a specific moment in time.

Note

This operation is valid for Redis only.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

SnapshotName

A name for the snapshot being created.

Type: String

Required: Yes

CacheClusterId

The identifier of an existing cluster. The snapshot is created from this cluster.

Type: String

Required: No

KmsKeyId

The ID of the KMS key used to encrypt the snapshot.

Type: String

Required: No

ReplicationGroupId

The identifier of an existing replication group. The snapshot is created from this replication group.

Type: String

Required: No

Tags.Tag.N

A list of tags to be added to this resource. A tag is a key-value pair. A tag key must be accompanied by a tag value, although null is accepted.

Type: Array of [Tag](#) (p. 319) objects

Required: No

Response Elements

The following element is returned by the service.

Snapshot

Represents a copy of an entire Redis cluster as of the time when the snapshot was taken.

Type: [Snapshot](#) (p. 311) object

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

CacheClusterNotFound

The requested cluster ID does not refer to an existing cluster.

HTTP Status Code: 404

InvalidCacheClusterState

The requested cluster is not in the available state.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidReplicationGroupState

The requested replication group is not in the available state.

HTTP Status Code: 400

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

SnapshotAlreadyExistsFault

You already have a snapshot with the given name.

HTTP Status Code: 400

SnapshotFeatureNotSupportedFault

You attempted one of the following operations:

- Creating a snapshot of a Redis cluster running on a `cache.t1.micro` cache node.
- Creating a snapshot of a cluster that is running Memcached rather than Redis.

Neither of these are supported by ElastiCache.

HTTP Status Code: 400

SnapshotQuotaExceededFault

The request cannot be processed because it would exceed the maximum number of snapshots.

HTTP Status Code: 400

TagQuotaPerResourceExceeded

The request cannot be processed because it would cause the resource to have more than the allowed number of tags. The maximum number of tags permitted on a resource is 50.

HTTP Status Code: 400

Examples

CreateSnapshot

This example illustrates one usage of CreateSnapshot.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=CreateSnapshot  
&CacheClusterId=my-redis-primary  
&SnapshotName=my-manual-snapshot  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<CreateSnapshotResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">  
  <CreateSnapshotResult>  
    <Snapshot>  
      <CacheClusterId>my-redis-primary</CacheClusterId>  
      <Port>6379</Port>  
      <CacheNodeType>cache.m1.small</CacheNodeType>  
      <CacheParameterGroupName>default.redis2.8</CacheParameterGroupName>  
      <Engine>redis</Engine>  
      <PreferredAvailabilityZone>us-west-2c</PreferredAvailabilityZone>  
      <CacheClusterCreateTime>2015-02-02T18:46:57.972Z</CacheClusterCreateTime>  
      <EngineVersion>2.8.6</EngineVersion>  
      <SnapshotSource>manual</SnapshotSource>  
      <AutoMinorVersionUpgrade>true</AutoMinorVersionUpgrade>  
      <PreferredMaintenanceWindow>wed:09:00-wed:10:00</PreferredMaintenanceWindow>  
      <SnapshotName>my-manual-snapshot</SnapshotName>  
      <SnapshotRetentionLimit>5</SnapshotRetentionLimit>  
      <NodeSnapshots>  
        <NodeSnapshot>  
          <CacheNodeCreateTime>2015-02-02T18:46:57.972Z</CacheNodeCreateTime>  
          <CacheNodeId>0001</CacheNodeId>  
          <CacheSize />  
        </NodeSnapshot>  
      </NodeSnapshots>  
      <SnapshotStatus>creating</SnapshotStatus>  
      <NumCacheNodes>1</NumCacheNodes>  
      <SnapshotWindow>07:30-08:30</SnapshotWindow>  
    </Snapshot>  
  </CreateSnapshotResult>  
  <ResponseMetadata>  
    <RequestId>faf5a232-b9ce-11e3-8a16-7978bb24ffdf</RequestId>  
  </ResponseMetadata>  
</CreateSnapshotResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

CreateUser

For Redis engine version 6.x onwards: Creates a Redis user. For more information, see [Using Role Based Access Control \(RBAC\)](#).

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

AccessString

Access permissions string used for this user.

Type: String

Pattern: .*\S.*

Required: Yes

Engine

The current supported value is Redis.

Type: String

Pattern: [a-zA-Z]*

Required: Yes

UserId

The ID of the user.

Type: String

Length Constraints: Minimum length of 1.

Pattern: [a-zA-Z][a-zA-Z0-9\-*]

Required: Yes

UserName

The username of the user.

Type: String

Length Constraints: Minimum length of 1.

Required: Yes

NoPasswordRequired

Indicates a password is not required for this user.

Type: Boolean

Required: No

Passwords.member.N

Passwords used for this user. You can create up to two passwords for each user.

Type: Array of strings

Array Members: Minimum number of 1 item.

Required: No

Tags.Tag.N

A list of tags to be added to this resource. A tag is a key-value pair. A tag key must be accompanied by a tag value, although null is accepted.

Type: Array of [Tag \(p. 319\)](#) objects

Required: No

Response Elements

The following elements are returned by the service.

AccessString

Access permissions string used for this user.

Type: String

ARN

The Amazon Resource Name (ARN) of the user.

Type: String

Authentication

Denotes whether the user requires a password to authenticate.

Type: [Authentication \(p. 222\)](#) object

Engine

The current supported value is Redis.

Type: String

Pattern: [a-zA-Z]*

Status

Indicates the user status. Can be "active", "modifying" or "deleting".

Type: String

UserGroupIds.member.N

Returns a list of the user group IDs the user belongs to.

Type: Array of strings

Length Constraints: Minimum length of 1.

Pattern: [a-zA-Z][a-zA-Z0-9\-*]

UserId

The ID of the user.

Type: String

UserName

The username of the user.

Type: String

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

DuplicateUserName

A user with this username already exists.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

TagQuotaPerResourceExceeded

The request cannot be processed because it would cause the resource to have more than the allowed number of tags. The maximum number of tags permitted on a resource is 50.

HTTP Status Code: 400

UserAlreadyExists

A user with this ID already exists.

HTTP Status Code: 400

UserQuotaExceeded

The quota of users has been exceeded.

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)

- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

CreateUserGroup

For Redis engine version 6.x onwards: Creates a Redis user group. For more information, see [Using Role Based Access Control \(RBAC\)](#)

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

Engine

The current supported value is Redis.

Type: String

Pattern: [a-zA-Z]*

Required: Yes

UserGroupId

The ID of the user group.

Type: String

Required: Yes

Tags.Tag.N

A list of tags to be added to this resource. A tag is a key-value pair. A tag key must be accompanied by a tag value, although null is accepted.

Type: Array of [Tag \(p. 319\)](#) objects

Required: No

UserIds.member.N

The list of user IDs that belong to the user group.

Type: Array of strings

Array Members: Minimum number of 1 item.

Length Constraints: Minimum length of 1.

Pattern: [a-zA-Z][a-zA-Z0-9\-*]

Required: No

Response Elements

The following elements are returned by the service.

ARN

The Amazon Resource Name (ARN) of the user group.

Type: String

Engine

The current supported value is Redis.

Type: String

Pattern: [a-zA-Z]*

PendingChanges

A list of updates being applied to the user group.

Type: [UserGroupPendingChanges](#) (p. 329) object

ReplicationGroups.member.N

A list of replication groups that the user group can access.

Type: Array of strings

Status

Indicates user group status. Can be "creating", "active", "modifying", "deleting".

Type: String

UserGroupId

The ID of the user group.

Type: String

UserIds.member.N

The list of user IDs that belong to the user group.

Type: Array of strings

Length Constraints: Minimum length of 1.

Pattern: [a-zA-Z][a-zA-Z0-9\-\-]*

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

DefaultUserRequired

You must add default user to a user group.

HTTP Status Code: 400

DuplicateUserName

A user with this username already exists.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

TagQuotaPerResourceExceeded

The request cannot be processed because it would cause the resource to have more than the allowed number of tags. The maximum number of tags permitted on a resource is 50.

HTTP Status Code: 400

UserGroupAlreadyExists

The user group with this ID already exists.

HTTP Status Code: 400

UserGroupQuotaExceeded

The number of users exceeds the user group limit.

HTTP Status Code: 400

UserNotFound

The user does not exist or could not be found.

HTTP Status Code: 404

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DecreaseNodeGroupsInGlobalReplicationGroup

Decreases the number of node groups in a Global datastore

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

ApplyImmediately

Indicates that the shard reconfiguration process begins immediately. At present, the only permitted value for this parameter is true.

Type: Boolean

Required: Yes

GlobalReplicationGroupId

The name of the Global datastore

Type: String

Required: Yes

NodeGroupCount

The number of node groups (shards) that results from the modification of the shard configuration

Type: Integer

Required: Yes

GlobalNodeGroupsToRemove.GlobalNodeId.N

If the value of NodeGroupCount is less than the current number of node groups (shards), then either NodeGroupsToRemove or NodeGroupsToRetain is required. GlobalNodeGroupsToRemove is a list of NodeGroupIds to remove from the cluster. ElastiCache for Redis will attempt to remove all node groups listed by GlobalNodeGroupsToRemove from the cluster.

Type: Array of strings

Required: No

GlobalNodeGroupsToRetain.GlobalNodeId.N

If the value of NodeGroupCount is less than the current number of node groups (shards), then either NodeGroupsToRemove or NodeGroupsToRetain is required. GlobalNodeGroupsToRetain is a list of NodeGroupIds to retain from the cluster. ElastiCache for Redis will attempt to retain all node groups listed by GlobalNodeGroupsToRetain from the cluster.

Type: Array of strings

Required: No

Response Elements

The following element is returned by the service.

GlobalReplicationGroup

Consists of a primary cluster that accepts writes and an associated secondary cluster that resides in a different Amazon region. The secondary cluster accepts only reads. The primary cluster automatically replicates updates to the secondary cluster.

- The **GlobalReplicationGroupIdSuffix** represents the name of the Global datastore, which is what you use to associate a secondary cluster.

Type: [GlobalReplicationGroup](#) (p. 259) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

GlobalReplicationGroupNotFoundFault

The Global datastore does not exist

HTTP Status Code: 404

InvalidGlobalReplicationGroupState

The Global datastore is not available or in primary-only state.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DecreaseReplicaCount

Dynamically decreases the number of replicas in a Redis (cluster mode disabled) replication group or the number of replica nodes in one or more node groups (shards) of a Redis (cluster mode enabled) replication group. This operation is performed with no cluster down time.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

ApplyImmediately

If `True`, the number of replica nodes is decreased immediately. `ApplyImmediately=False` is not currently supported.

Type: Boolean

Required: Yes

ReplicationGroupId

The id of the replication group from which you want to remove replica nodes.

Type: String

Required: Yes

NewReplicaCount

The number of read replica nodes you want at the completion of this operation. For Redis (cluster mode disabled) replication groups, this is the number of replica nodes in the replication group. For Redis (cluster mode enabled) replication groups, this is the number of replica nodes in each of the replication group's node groups.

The minimum number of replicas in a shard or replication group is:

- Redis (cluster mode disabled)
 - If Multi-AZ is enabled: 1
 - If Multi-AZ is not enabled: 0
- Redis (cluster mode enabled): 0 (though you will not be able to failover to a replica if your primary node fails)

Type: Integer

Required: No

ReplicaConfiguration.ConfigureShard.N

A list of `ConfigureShard` objects that can be used to configure each shard in a Redis (cluster mode enabled) replication group. The `ConfigureShard` has three members: `NewReplicaCount`, `NodeGroupId`, and `PreferredAvailabilityZones`.

Type: Array of [ConfigureShard](#) (p. 249) objects

Required: No

ReplicasToRemove.member.N

A list of the node ids to remove from the replication group or node group (shard).

Type: Array of strings

Required: No

Response Elements

The following element is returned by the service.

ReplicationGroup

Contains all of the attributes of a specific Redis replication group.

Type: [ReplicationGroup](#) (p. 291) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

ClusterQuotaForCustomerExceeded

The request cannot be processed because it would exceed the allowed number of clusters per customer.

HTTP Status Code: 400

InsufficientCacheClusterCapacity

The requested cache node type is not available in the specified Availability Zone. For more information, see [InsufficientCacheClusterCapacity](#) in the ElastiCache User Guide.

HTTP Status Code: 400

InvalidCacheClusterState

The requested cluster is not in the `available` state.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidReplicationGroupState

The requested replication group is not in the `available` state.

HTTP Status Code: 400

InvalidVPCNetworkStateFault

The VPC network is in an invalid state.

HTTP Status Code: 400

NodeGroupsPerReplicationGroupQuotaExceeded

The request cannot be processed because it would exceed the maximum allowed number of node groups (shards) in a single replication group. The default maximum is 90

HTTP Status Code: 400

NodeQuotaForCustomerExceeded

The request cannot be processed because it would exceed the allowed number of cache nodes per customer.

HTTP Status Code: 400

NoOperationFault

The operation was not performed because no changes were required.

HTTP Status Code: 400

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

ServiceLinkedRoleNotFoundFault

The specified service linked role (SLR) was not found.

HTTP Status Code: 400

Examples

Example

The following example removes two replicas from each node group in the replication group `sample-repl-group`.

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=DecreaseReplicaCount  
&ApplyImmediately=True  
&NewReplicaCount=2  
&ReplicasToRemove.ReplicaToRemove.1=0001  
&ReplicasToRemove.ReplicaToRemove.2=0003  
&ReplicationGroupId=sample-repl-group  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Example

The following example removes replicas from two node groups. Because there are multiple node groups, this example is for a Redis (cluster mode enabled) replication group.

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=DecreaseReplicaCount  
&ApplyImmediately=True  
&ReplicaConfiguration.ConfigureShard.1.NodeGroupId=0001
```

```
&ReplicaConfiguration.ConfigureShard.1.NewReplicaCount=1

&ReplicaConfiguration.ConfigureShard.1.PreferredAvailabilityZones.PreferredAvailabilityZone.1=us-east-1a

&ReplicaConfiguration.ConfigureShard.1.PreferredAvailabilityZones.PreferredAvailabilityZone.2=us-east-1c
  &ReplicaConfiguration.ConfigureShard.2.NodeGroupId=0003
  &ReplicaConfiguration.ConfigureShard.2.NewReplicaCount=2

&ReplicaConfiguration.ConfigureShard.2.PreferredAvailabilityZones.PreferredAvailabilityZone.1=us-east-1a

&ReplicaConfiguration.ConfigureShard.2.PreferredAvailabilityZones.PreferredAvailabilityZone.2=us-east-1b

&ReplicaConfiguration.ConfigureShard.2.PreferredAvailabilityZones.PreferredAvailabilityZone.4=us-east-1c
  &ReplicationGroupId=samplem--repl-group
  &Version=2015-02-02
  &SignatureVersion=4
  &SignatureMethod=HmacSHA256
  &Timestamp=20150202T192317Z
  &X-Amz-Credential=<credential>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DeleteCacheCluster

Deletes a previously provisioned cluster. `DeleteCacheCluster` deletes all associated cache nodes, node endpoints and the cluster itself. When you receive a successful response from this operation, Amazon ElastiCache immediately begins deleting the cluster; you cannot cancel or revert this operation.

This operation is not valid for:

- Redis (cluster mode enabled) clusters
- Redis (cluster mode disabled) clusters
- A cluster that is the last read replica of a replication group
- A cluster that is the primary node of a replication group
- A node group (shard) that has Multi-AZ mode enabled
- A cluster from a Redis (cluster mode enabled) replication group
- A cluster that is not in the `available` state

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

CacheClusterId

The cluster identifier for the cluster to be deleted. This parameter is not case sensitive.

Type: String

Required: Yes

FinalSnapshotIdentifier

The user-supplied name of a final cluster snapshot. This is the unique name that identifies the snapshot. ElastiCache creates the snapshot, and then deletes the cluster immediately afterward.

Type: String

Required: No

Response Elements

The following element is returned by the service.

CacheCluster

Contains all of the attributes of a specific cluster.

Type: [CacheCluster \(p. 224\)](#) object

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

CacheClusterNotFound

The requested cluster ID does not refer to an existing cluster.

HTTP Status Code: 404

InvalidCacheClusterState

The requested cluster is not in the available state.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

SnapshotAlreadyExistsFault

You already have a snapshot with the given name.

HTTP Status Code: 400

SnapshotFeatureNotSupportedFault

You attempted one of the following operations:

- Creating a snapshot of a Redis cluster running on a `cache.t1.micro` cache node.
- Creating a snapshot of a cluster that is running Memcached rather than Redis.

Neither of these are supported by ElastiCache.

HTTP Status Code: 400

SnapshotQuotaExceededFault

The request cannot be processed because it would exceed the maximum number of snapshots.

HTTP Status Code: 400

Examples

DeleteCacheCluster

This example illustrates one usage of DeleteCacheCluster.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=DeleteCacheCluster  
&CacheClusterId=simcoprod43  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<DeleteCacheClusterResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <DeleteCacheClusterResult>
    <CacheCluster>
      <CacheParameterGroup>
        <ParameterApplyStatus>in-sync</ParameterApplyStatus>
        <CacheParameterGroupName>default.memcached1.4</CacheParameterGroupName>
        <CacheNodeIdsToReboot/>
      </CacheParameterGroup>
      <CacheClusterId>simcoprod43</CacheClusterId>
      <CacheClusterStatus>deleting</CacheClusterStatus>
      <ConfigurationEndpoint>
        <Port>11211</Port>
        <Address>simcoprod43.m2st2p.cfg.cache.amazonaws.com</Address>
      </ConfigurationEndpoint>
      <CacheNodeType>cache.m1.large</CacheNodeType>
      <Engine>memcached</Engine>
      <PendingModifiedValues/>
      <PreferredAvailabilityZone>us-west-2b</PreferredAvailabilityZone>
      <CacheClusterCreateTime>2015-02-02T02:18:26.497Z</CacheClusterCreateTime>
      <EngineVersion>1.4.5</EngineVersion>
      <AutoMinorVersionUpgrade>true</AutoMinorVersionUpgrade>
      <PreferredMaintenanceWindow>mon:05:00-mon:06:00</PreferredMaintenanceWindow>
      <CacheSecurityGroups>
        <CacheSecurityGroup>
          <CacheSecurityGroupName>default</CacheSecurityGroupName>
          <Status>active</Status>
        </CacheSecurityGroup>
      </CacheSecurityGroups>
      <NumCacheNodes>3</NumCacheNodes>
    </CacheCluster>
  </DeleteCacheClusterResult>
  <ResponseMetadata>
    <RequestId>ab84aa7e-b7fa-11e0-9b0b-a9261be2b354</RequestId>
  </ResponseMetadata>
</DeleteCacheClusterResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DeleteCacheParameterGroup

Deletes the specified cache parameter group. You cannot delete a cache parameter group if it is associated with any cache clusters. You cannot delete the default cache parameter groups in your account.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheParameterGroupName

The name of the cache parameter group to delete.

Note

The specified cache security group must not be associated with any clusters.

Type: String

Required: Yes

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

CacheParameterGroupNotFound

The requested cache parameter group name does not refer to an existing cache parameter group.

HTTP Status Code: 404

InvalidCacheParameterGroupState

The current state of the cache parameter group does not allow the requested operation to occur.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

Examples

DeleteCacheParameterGroup

This example illustrates one usage of DeleteCacheParameterGroup.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=DeleteCacheParameterGroup  
&CacheParameterGroupName=myparametergroup  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<DeleteCacheParameterGroupResponse xmlns="http://elasticache.amazonaws.com/  
doc/2015-02-02/">  
  <ResponseMetadata>  
    <RequestId>d0a417cb-575b-11e0-8869-cd22b4f9d96f</RequestId>  
  </ResponseMetadata>  
</DeleteCacheParameterGroupResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DeleteCacheSecurityGroup

Deletes a cache security group.

Note

You cannot delete a cache security group if it is associated with any clusters.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheSecurityGroupName

The name of the cache security group to delete.

Note

You cannot delete the default security group.

Type: String

Required: Yes

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

CacheSecurityGroupNotFound

The requested cache security group name does not refer to an existing cache security group.

HTTP Status Code: 404

InvalidCacheSecurityGroupState

The current state of the cache security group does not allow deletion.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

Examples

DeleteCacheSecurityGroup

This example illustrates one usage of DeleteCacheSecurityGroup.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=DeleteCacheSecurityGroup  
&CacheSecurityGroupName=mycachesecuritygroup3  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<DeleteCacheSecurityGroupResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">  
  <ResponseMetadata>  
    <RequestId>c130cfb7-3650-11e0-ae57-f96cfe56749c</RequestId>  
  </ResponseMetadata>  
</DeleteCacheSecurityGroupResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DeleteCacheSubnetGroup

Deletes a cache subnet group.

Note

You cannot delete a default cache subnet group or one that is associated with any clusters.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheSubnetGroupName

The name of the cache subnet group to delete.

Constraints: Must contain no more than 255 alphanumeric characters or hyphens.

Type: String

Required: Yes

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

CacheSubnetGroupInUse

The requested cache subnet group is currently in use.

HTTP Status Code: 400

CacheSubnetGroupNotFoundFault

The requested cache subnet group name does not refer to an existing cache subnet group.

HTTP Status Code: 400

Examples

DeleteCacheSubnetGroup

This example illustrates one usage of DeleteCacheSubnetGroup.

Sample Request

```
https://elasticache.amazonaws.com/  
?Action=DeleteCacheSubnetGroup  
&CacheSubnetGroupName=mysubnetgroup  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<DeleteCacheSubnetGroupResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <ResponseMetadata>
    <RequestId>5d013245-4172-11df-8520-e7e1e602a915</RequestId>
  </ResponseMetadata>
</DeleteCacheSubnetGroupResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DeleteGlobalReplicationGroup

Deleting a Global datastore is a two-step process:

- First, you must [DisassociateGlobalReplicationGroup](#) (p. 147) to remove the secondary clusters in the Global datastore.
- Once the Global datastore contains only the primary cluster, you can use the `DeleteGlobalReplicationGroup` API to delete the Global datastore while retaining the primary cluster using `RetainPrimaryReplicationGroup=true`.

Since the Global Datastore has only a primary cluster, you can delete the Global Datastore while retaining the primary by setting `RetainPrimaryReplicationGroup=true`. The primary cluster is never deleted when deleting a Global Datastore. It can only be deleted when it no longer is associated with any Global Datastore.

When you receive a successful response from this operation, Amazon ElastiCache immediately begins deleting the selected resources; you cannot cancel or revert this operation.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

GlobalReplicationGroupId

The name of the Global datastore

Type: String

Required: Yes

RetainPrimaryReplicationGroup

The primary replication group is retained as a standalone replication group.

Type: Boolean

Required: Yes

Response Elements

The following element is returned by the service.

GlobalReplicationGroup

Consists of a primary cluster that accepts writes and an associated secondary cluster that resides in a different Amazon region. The secondary cluster accepts only reads. The primary cluster automatically replicates updates to the secondary cluster.

- The **GlobalReplicationGroupIdSuffix** represents the name of the Global datastore, which is what you use to associate a secondary cluster.

Type: [GlobalReplicationGroup](#) (p. 259) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

GlobalReplicationGroupNotFoundFault

The Global datastore does not exist

HTTP Status Code: 404

InvalidGlobalReplicationGroupState

The Global datastore is not available or in primary-only state.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DeleteReplicationGroup

Deletes an existing replication group. By default, this operation deletes the entire replication group, including the primary/primaries and all of the read replicas. If the replication group has only one primary, you can optionally delete only the read replicas, while retaining the primary by setting `RetainPrimaryCluster=true`.

When you receive a successful response from this operation, Amazon ElastiCache immediately begins deleting the selected resources; you cannot cancel or revert this operation.

Note

This operation is valid for Redis only.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

ReplicationGroupId

The identifier for the cluster to be deleted. This parameter is not case sensitive.

Type: String

Required: Yes

FinalSnapshotIdentifier

The name of a final node group (shard) snapshot. ElastiCache creates the snapshot from the primary node in the cluster, rather than one of the replicas; this is to ensure that it captures the freshest data. After the final snapshot is taken, the replication group is immediately deleted.

Type: String

Required: No

RetainPrimaryCluster

If set to `true`, all of the read replicas are deleted, but the primary node is retained.

Type: Boolean

Required: No

Response Elements

The following element is returned by the service.

ReplicationGroup

Contains all of the attributes of a specific Redis replication group.

Type: [ReplicationGroup](#) (p. 291) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidReplicationGroupState

The requested replication group is not in the available state.

HTTP Status Code: 400

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

SnapshotAlreadyExistsFault

You already have a snapshot with the given name.

HTTP Status Code: 400

SnapshotFeatureNotSupportedFault

You attempted one of the following operations:

- Creating a snapshot of a Redis cluster running on a `cache.t1.micro` cache node.
- Creating a snapshot of a cluster that is running Memcached rather than Redis.

Neither of these are supported by ElastiCache.

HTTP Status Code: 400

SnapshotQuotaExceededFault

The request cannot be processed because it would exceed the maximum number of snapshots.

HTTP Status Code: 400

Examples

DeleteReplicationGroup

This example illustrates one usage of DeleteReplicationGroup.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=DeleteReplicationGroup &RetainPrimaryCluster=false  
&FinalSnapshotIdentifier=my-final-snapshot  
&ReplicationGroupId=my-repgroup  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z
```

```
&X-Amz-Credential=<credential>
```

Sample Response

```
<DeleteReplicationGroupResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <DeleteReplicationGroupResult>
    <ReplicationGroup>
      <SnapshottingClusterId>my-redis-primary</SnapshottingClusterId>
      decrease-replica-count      <ReplicationGroupId>my-repgroup</ReplicationGroupId>
      <Status>deleting</Status>
      <PendingModifiedValues />
      <Description>My replication group</Description>
    </ReplicationGroup>
  </DeleteReplicationGroupResult>
  <ResponseMetadata>
    <RequestId>93eb37db-b9d7-11e3-8a16-7978bb24ffdf</RequestId>
  </ResponseMetadata>
</DeleteReplicationGroupResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DeleteSnapshot

Deletes an existing snapshot. When you receive a successful response from this operation, ElastiCache immediately begins deleting the snapshot; you cannot cancel or revert this operation.

Note

This operation is valid for Redis only.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

SnapshotName

The name of the snapshot to be deleted.

Type: String

Required: Yes

Response Elements

The following element is returned by the service.

Snapshot

Represents a copy of an entire Redis cluster as of the time when the snapshot was taken.

Type: [Snapshot \(p. 311\)](#) object

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidSnapshotState

The current state of the snapshot does not allow the requested operation to occur.

HTTP Status Code: 400

SnapshotNotFoundFault

The requested snapshot name does not refer to an existing snapshot.

HTTP Status Code: 404

Examples

DeleteSnapshot

This example illustrates one usage of DeleteSnapshot.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/
?Action=DeleteSnapshot
&SnapshotName=my-manual-snapshot
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20150202T192317Z
&X-Amz-Credential=<credential>
```

Sample Response

```
<DeleteSnapshotResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <DeleteSnapshotResult>
    <Snapshot>
      <CacheClusterId>my-redis-primary</CacheClusterId>
      <Port>6379</Port>
      <CacheNodeType>cache.m1.small</CacheNodeType>
      <CacheParameterGroupName>default.redis2.8</CacheParameterGroupName>
      <Engine>redis</Engine>
      <PreferredAvailabilityZone>us-west-2c</PreferredAvailabilityZone>
      <CacheClusterCreateTime>2015-02-02T18:46:57.972Z</CacheClusterCreateTime>
      <EngineVersion>2.8.6</EngineVersion>
      <SnapshotSource>manual</SnapshotSource>
      <AutoMinorVersionUpgrade>true</AutoMinorVersionUpgrade>
      <PreferredMaintenanceWindow>wed:09:00-wed:10:00</PreferredMaintenanceWindow>
      <SnapshotName>my-manual-snapshot</SnapshotName>
      <SnapshotRetentionLimit>5</SnapshotRetentionLimit>
      <NodeSnapshots>
        <NodeSnapshot>
          <SnapshotCreateTime>2015-02-02T18:54:12Z</SnapshotCreateTime>
          <CacheNodeCreateTime>2015-02-02T18:46:57.972Z</CacheNodeCreateTime>
          <CacheNodeId>0001</CacheNodeId>
          <CacheSize>3 MB</CacheSize>
        </NodeSnapshot>
      </NodeSnapshots>
      <SnapshotStatus>deleting</SnapshotStatus>
      <NumCacheNodes>1</NumCacheNodes>
      <SnapshotWindow>07:30-08:30</SnapshotWindow>
    </Snapshot>
  </DeleteSnapshotResult>
  <ResponseMetadata>
    <RequestId>694d7017-b9d2-11e3-8a16-7978bb24ffdf</RequestId>
  </ResponseMetadata>
</DeleteSnapshotResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DeleteUser

For Redis engine version 6.x onwards: Deletes a user. The user will be removed from all user groups and in turn removed from all replication groups. For more information, see [Using Role Based Access Control \(RBAC\)](#).

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

UserId

The ID of the user.

Type: String

Length Constraints: Minimum length of 1.

Pattern: [a-zA-Z][a-zA-Z0-9\-*]

Required: Yes

Response Elements

The following elements are returned by the service.

AccessString

Access permissions string used for this user.

Type: String

ARN

The Amazon Resource Name (ARN) of the user.

Type: String

Authentication

Denotes whether the user requires a password to authenticate.

Type: [Authentication](#) (p. 222) object

Engine

The current supported value is Redis.

Type: String

Pattern: [a-zA-Z]*

Status

Indicates the user status. Can be "active", "modifying" or "deleting".

Type: String

UserGroupIds.member.N

Returns a list of the user group IDs the user belongs to.

Type: Array of strings

Length Constraints: Minimum length of 1.

Pattern: [a-zA-Z][a-zA-Z0-9\-*]

UserId

The ID of the user.

Type: String

UserName

The username of the user.

Type: String

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

DefaultUserAssociatedToUserGroup

The default user assigned to the user group.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidUserState

The user is not in active state.

HTTP Status Code: 400

UserNotFound

The user does not exist or could not be found.

HTTP Status Code: 404

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DeleteUserGroup

For Redis engine version 6.x onwards: Deletes a user group. The user group must first be disassociated from the replication group before it can be deleted. For more information, see [Using Role Based Access Control \(RBAC\)](#).

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

UserGroupId

The ID of the user group.

Type: String

Required: Yes

Response Elements

The following elements are returned by the service.

ARN

The Amazon Resource Name (ARN) of the user group.

Type: String

Engine

The current supported value is Redis.

Type: String

Pattern: [a-zA-Z]*

PendingChanges

A list of updates being applied to the user group.

Type: [UserGroupPendingChanges](#) (p. 329) object

ReplicationGroups.member.N

A list of replication groups that the user group can access.

Type: Array of strings

Status

Indicates user group status. Can be "creating", "active", "modifying", "deleting".

Type: String

UserGroupId

The ID of the user group.

Type: String

UsersIds.member.N

The list of user IDs that belong to the user group.

Type: Array of strings

Length Constraints: Minimum length of 1.

Pattern: `[a-zA-Z][a-zA-Z0-9\-\]*`

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidUserGroupState

The user group is not in an active state.

HTTP Status Code: 400

UserGroupNotFound

The user group was not found or does not exist

HTTP Status Code: 404

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeCacheClusters

Returns information about all provisioned clusters if no cluster identifier is specified, or about a specific cache cluster if a cluster identifier is supplied.

By default, abbreviated information about the clusters is returned. You can use the optional *ShowCacheNodeInfo* flag to retrieve detailed information about the cache nodes associated with the clusters. These details include the DNS address and port for the cache node endpoint.

If the cluster is in the *creating* state, only cluster-level information is displayed until all of the nodes are successfully provisioned.

If the cluster is in the *deleting* state, only cluster-level information is displayed.

If cache nodes are currently being added to the cluster, node endpoint information and creation time for the additional nodes are not displayed until they are completely provisioned. When the cluster state is *available*, the cluster is ready for use.

If cache nodes are currently being removed from the cluster, no endpoint information for the removed nodes is displayed.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

CacheClusterId

The user-supplied cluster identifier. If this parameter is specified, only information about that specific cluster is returned. This parameter isn't case sensitive.

Type: String

Required: No

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by *MaxRecords*.

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response. If more records exist than the specified *MaxRecords* value, a marker is included in the response so that the remaining results can be retrieved.

Default: 100

Constraints: minimum 20; maximum 100.

Type: Integer

Required: No

ShowCacheClustersNotInReplicationGroups

An optional flag that can be included in the *DescribeCacheCluster* request to show only nodes (API/CLI: clusters) that are not members of a replication group. In practice, this means Memcached and single node Redis clusters.

Type: Boolean

Required: No

ShowCacheNodeInfo

An optional flag that can be included in the `DescribeCacheCluster` request to retrieve information about the individual cache nodes.

Type: Boolean

Required: No

Response Elements

The following elements are returned by the service.

CacheClusters.CacheCluster.N

A list of clusters. Each item in the list contains detailed information about one cluster.

Type: Array of [CacheCluster](#) (p. 224) objects

Marker

Provides an identifier to allow retrieval of paginated results.

Type: String

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

CacheClusterNotFound

The requested cluster ID does not refer to an existing cluster.

HTTP Status Code: 404

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

Examples

DescribeCacheClusters

This example illustrates one usage of `DescribeCacheClusters`.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/
```

```
?Action=DescribeCacheClusters
&MaxRecords=100
&ShowCacheNodeInfo=false
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20150202T192317Z
&X-Amz-Credential=<credential>
```

Sample Response

```
<DescribeCacheClustersResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <DescribeCacheClustersResult>
    <CacheClusters>
      <CacheCluster>
        <CacheParameterGroup>
          <ParameterApplyStatus>in-sync</ParameterApplyStatus>
          <CacheParameterGroupName>default.memcached1.4</CacheParameterGroupName>
          <CacheNodeIdsToReboot/>
        </CacheParameterGroup>
        <CacheClusterId>simcoprod42</CacheClusterId>
        <CacheClusterStatus>available</CacheClusterStatus>
        <ConfigurationEndpoint>
          <Port>11211</Port>
          <Address>simcoprod42.m2st2p.cfg.cache.amazonaws.com</Address>
        </ConfigurationEndpoint>
        <ClientDownloadLandingPage>
          https://console.aws.amazon.com/elasticache/home#client-download:
        </ClientDownloadLandingPage>
        <CacheNodeType>cache.m1.large</CacheNodeType>
        <Engine>memcached</Engine>
        <PendingModifiedValues/>
        <PreferredAvailabilityZone>us-west-2c</PreferredAvailabilityZone>
        <CacheClusterCreateTime>2015-02-02T01:21:46.607Z</CacheClusterCreateTime>
        <EngineVersion>1.4.5</EngineVersion>
        <AutoMinorVersionUpgrade>true</AutoMinorVersionUpgrade>
        <PreferredMaintenanceWindow>fri:08:30-fri:09:30</PreferredMaintenanceWindow>
        <CacheSecurityGroups>
          <CacheSecurityGroup>
            <CacheSecurityGroupName>default</CacheSecurityGroupName>
            <Status>active</Status>
          </CacheSecurityGroup>
        </CacheSecurityGroups>
        <NotificationConfiguration>
          <TopicStatus>active</TopicStatus>
          <TopicArn>arn:aws:sns:us-west-2:123456789012:ElastiCacheNotifications</TopicArn>
        </NotificationConfiguration>
        <NumCacheNodes>6</NumCacheNodes>
      </CacheCluster>
    </CacheClusters>
  </DescribeCacheClustersResult>
  <ResponseMetadata>
    <RequestId>f270d58f-b7fb-11e0-9326-b7275b9d4a6c</RequestId>
  </ResponseMetadata>
</DescribeCacheClustersResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)

- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeCacheEngineVersions

Returns a list of the available cache engines and their versions.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheParameterGroupFamily

The name of a specific cache parameter group family to return details for.

Valid values are: `memcached1.4` | `memcached1.5` | `memcached1.6` | `redis2.6` | `redis2.8` | `redis3.2` | `redis4.0` | `redis5.0` | `redis6.x` |

Constraints:

- Must be 1 to 255 alphanumeric characters
- First character must be a letter
- Cannot end with a hyphen or contain two consecutive hyphens

Type: String

Required: No

DefaultOnly

If `true`, specifies that only the default version of the specified engine or engine and major version combination is to be returned.

Type: Boolean

Required: No

Engine

The cache engine to return. Valid values: `memcached` | `redis`

Type: String

Required: No

EngineVersion

The cache engine version to return.

Example: `1.4.14`

Type: String

Required: No

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response. If more records exist than the specified `MaxRecords` value, a marker is included in the response so that the remaining results can be retrieved.

Default: 100

Constraints: minimum 20; maximum 100.

Type: Integer

Required: No

Response Elements

The following elements are returned by the service.

CacheEngineVersions.CacheEngineVersion.N

A list of cache engine version details. Each element in the list contains detailed information about one cache engine version.

Type: Array of [CacheEngineVersion](#) (p. 230) objects

Marker

Provides an identifier to allow retrieval of paginated results.

Type: String

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

Examples

DescribeCacheEngineVersions

This example illustrates one usage of `DescribeCacheEngineVersions`.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=DescribeCacheEngineVersions  
&MaxRecords=100  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<DescribeCacheEngineVersionsResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <DescribeCacheEngineVersionsResult>
    <CacheEngineVersions>
      <CacheEngineVersion>
        <CacheParameterGroupFamily>memcached1.4</CacheParameterGroupFamily>
        <Engine>memcached</Engine>
        <CacheEngineVersionDescription>memcached version 1.4.14</
CacheEngineVersionDescription>
        <CacheEngineDescription>memcached</CacheEngineDescription>
        <EngineVersion>1.4.14</EngineVersion>
      </CacheEngineVersion>
      <CacheEngineVersion>
        <CacheParameterGroupFamily>memcached1.4</CacheParameterGroupFamily>
        <Engine>memcached</Engine>
        <CacheEngineVersionDescription>memcached version 1.4.5</
CacheEngineVersionDescription>
        <CacheEngineDescription>memcached</CacheEngineDescription>
        <EngineVersion>1.4.5</EngineVersion>
      </CacheEngineVersion>
    </CacheEngineVersions>
  </DescribeCacheEngineVersionsResult>
  <ResponseMetadata>
    <RequestId>a6ac9ad2-f8a4-11e1-a4d1-a345e5370093</RequestId>
  </ResponseMetadata>
</DescribeCacheEngineVersionsResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeCacheParameterGroups

Returns a list of cache parameter group descriptions. If a cache parameter group name is specified, the list contains only the descriptions for that group.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheParameterGroupName

The name of a specific cache parameter group to return details for.

Type: String

Required: No

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response. If more records exist than the specified `MaxRecords` value, a marker is included in the response so that the remaining results can be retrieved.

Default: 100

Constraints: minimum 20; maximum 100.

Type: Integer

Required: No

Response Elements

The following elements are returned by the service.

CacheParameterGroups.CacheParameterGroup.N

A list of cache parameter groups. Each element in the list contains detailed information about one cache parameter group.

Type: Array of [CacheParameterGroup](#) (p. 240) objects

Marker

Provides an identifier to allow retrieval of paginated results.

Type: String

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

CacheParameterGroupNotFound

The requested cache parameter group name does not refer to an existing cache parameter group.

HTTP Status Code: 404

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

Examples

DescribeCacheParameterGroups

This example illustrates one usage of DescribeCacheParameterGroups.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=DescribeCacheParameterGroups  
&MaxRecords=100  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<DescribeCacheParameterGroupsResponse xmlns="http://elasticache.amazonaws.com/  
doc/2015-02-02/">  
  <DescribeCacheParameterGroupsResult>  
    <CacheParameterGroups>  
      <CacheParameterGroup>  
        <CacheParameterGroupName>default.memcached1.4</CacheParameterGroupName>  
        <CacheParameterGroupFamily>memcached1.4</CacheParameterGroupFamily>  
        <Description>Default parameter group for memcached1.4</Description>  
      </CacheParameterGroup>  
      <CacheParameterGroup>  
        <CacheParameterGroupName>mycacheparametergroup</CacheParameterGroupName>  
        <CacheParameterGroupFamily>memcached1.4</CacheParameterGroupFamily>  
        <Description>My cache parameter group</Description>  
      </CacheParameterGroup>  
      <CacheParameterGroup>  
        <CacheParameterGroupName>mycacheparametergroup1</CacheParameterGroupName>  
        <CacheParameterGroupFamily>memcached1.4</CacheParameterGroupFamily>  
        <Description>My first cache parameter group</Description>  
      </CacheParameterGroup>  
    </CacheParameterGroups>  
  </DescribeCacheParameterGroupsResult>  
</DescribeCacheParameterGroupsResponse>
```

```
</CacheParameterGroup>
<CacheParameterGroup>
  <CacheParameterGroupName>mycacheparametergroup3</CacheParameterGroupName>
  <CacheParameterGroupFamily>memcached1.4</CacheParameterGroupFamily>
  <Description>My first cache parameter group</Description>
</CacheParameterGroup>
</CacheParameterGroups>
</DescribeCacheParameterGroupsResult>
<ResponseMetadata>
  <RequestId>7193fbb8-b7fc-11e0-9b0b-a9261be2b354</RequestId>
</ResponseMetadata>
</DescribeCacheParameterGroupsResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeCacheParameters

Returns the detailed parameter list for a particular cache parameter group.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheParameterGroupName

The name of a specific cache parameter group to return details for.

Type: String

Required: Yes

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response. If more records exist than the specified `MaxRecords` value, a marker is included in the response so that the remaining results can be retrieved.

Default: 100

Constraints: minimum 20; maximum 100.

Type: Integer

Required: No

Source

The parameter types to return.

Valid values: `user` | `system` | `engine-default`

Type: String

Required: No

Response Elements

The following elements are returned by the service.

CacheNodeTypeSpecificParameters.CacheNodeTypeSpecificParameter.N

A list of parameters specific to a particular cache node type. Each element in the list contains detailed information about one parameter.

Type: Array of [CacheNodeTypeSpecificParameter](#) (p. 235) objects

Marker

Provides an identifier to allow retrieval of paginated results.

Type: String

Parameters.Parameter.N

A list of [Parameter](#) (p. 282) instances.

Type: Array of [Parameter](#) (p. 282) objects

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

CacheParameterGroupNotFound

The requested cache parameter group name does not refer to an existing cache parameter group.

HTTP Status Code: 404

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

Examples

DescribeCacheParameters

Some of the output has been omitted for brevity.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/
?Action=DescribeCacheParameters
&CacheParameterGroupName=default.memcached1.4
&MaxRecords=100
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20150202T192317Z
&X-Amz-Credential=<credential>
```

Sample Response

```
<DescribeCacheParametersResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <DescribeCacheParametersResult>
```

```
<CacheNodeTypeSpecificParameters>
  <CacheNodeTypeSpecificParameter>
    <CacheNodeTypeSpecificValues>
      <CacheNodeTypeSpecificValue>
        <CacheNodeType>cache.cl.xlarge</CacheNodeType>
        <Value>6000</Value>
      </CacheNodeTypeSpecificValue>
    </CacheNodeTypeSpecificValues>
    <DataType>integer</DataType>
    <Source>system</Source>
    <IsModifiable>false</IsModifiable>
    <Description>The maximum configurable amount of memory to use to store items, in megabytes.</Description>
    <AllowedValues>1-100000</AllowedValues>
    <ParameterName>max_cache_memory</ParameterName>
    <MinimumEngineVersion>1.4.5</MinimumEngineVersion>
  </CacheNodeTypeSpecificParameter>
</CacheNodeTypeSpecificParameters>
<Parameters>
  <Parameter>
    <ParameterValue>1024</ParameterValue>
    <DataType>integer</DataType>
    <Source>system</Source>
    <IsModifiable>false</IsModifiable>
    <Description>The backlog queue limit.</Description>
    <AllowedValues>1-10000</AllowedValues>
    <ParameterName>backlog_queue_limit</ParameterName>
    <MinimumEngineVersion>1.4.5</MinimumEngineVersion>
  </Parameter>
</Parameters>
</DescribeCacheParametersResult>
<ResponseMetadata>
  <RequestId>0c507368-b7fe-11e0-9326-b7275b9d4a6c</RequestId>
</ResponseMetadata>
</DescribeCacheParametersResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- AWS Command Line Interface
- AWS SDK for .NET
- AWS SDK for C++
- AWS SDK for Go
- AWS SDK for Java V2
- AWS SDK for JavaScript
- AWS SDK for PHP V3
- AWS SDK for Python
- AWS SDK for Ruby V3

DescribeCacheSecurityGroups

Returns a list of cache security group descriptions. If a cache security group name is specified, the list contains only the description of that group. This applicable only when you have ElastiCache in Classic setup

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

CacheSecurityGroupName

The name of the cache security group to return details for.

Type: String

Required: No

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response. If more records exist than the specified `MaxRecords` value, a marker is included in the response so that the remaining results can be retrieved.

Default: 100

Constraints: minimum 20; maximum 100.

Type: Integer

Required: No

Response Elements

The following elements are returned by the service.

CacheSecurityGroups.CacheSecurityGroup.N

A list of cache security groups. Each element in the list contains detailed information about one group.

Type: Array of [CacheSecurityGroup \(p. 243\)](#) objects

Marker

Provides an identifier to allow retrieval of paginated results.

Type: String

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

CacheSecurityGroupNotFound

The requested cache security group name does not refer to an existing cache security group.

HTTP Status Code: 404

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

Examples

DescribeCacheSecurityGroups

This example illustrates one usage of DescribeCacheSecurityGroups.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=DescribeCacheSecurityGroups  
&MaxRecords=100  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<DescribeCacheSecurityGroupsResponse xmlns="http://elasticache.amazonaws.com/  
doc/2015-02-02/">  
  <DescribeCacheSecurityGroupsResult>  
    <CacheSecurityGroups>  
      <CacheSecurityGroup>  
        <EC2SecurityGroups/>  
        <CacheSecurityGroupName>default</CacheSecurityGroupName>  
        <OwnerId>123456789012</OwnerId>  
        <Description>default</Description>  
      </CacheSecurityGroup>  
      <CacheSecurityGroup>  
        <EC2SecurityGroups/>  
        <CacheSecurityGroupName>mycachesecuritygroup</CacheSecurityGroupName>  
        <OwnerId>123456789012</OwnerId>  
        <Description>My Security Group</Description>  
      </CacheSecurityGroup>  
    </CacheSecurityGroups>
```



```
</DescribeCacheSecurityGroupsResult>  
<ResponseMetadata>  
  <RequestId>a95360ae-b7fc-11e0-9326-b7275b9d4a6c</RequestId>  
</ResponseMetadata>  
</DescribeCacheSecurityGroupsResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeCacheSubnetGroups

Returns a list of cache subnet group descriptions. If a subnet group name is specified, the list contains only the description of that group. This is applicable only when you have ElastiCache in VPC setup. All ElastiCache clusters now launch in VPC by default.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

CacheSubnetGroupName

The name of the cache subnet group to return details for.

Type: String

Required: No

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response. If more records exist than the specified `MaxRecords` value, a marker is included in the response so that the remaining results can be retrieved.

Default: 100

Constraints: minimum 20; maximum 100.

Type: Integer

Required: No

Response Elements

The following elements are returned by the service.

CacheSubnetGroups.CacheSubnetGroup.N

A list of cache subnet groups. Each element in the list contains detailed information about one group.

Type: Array of [CacheSubnetGroup \(p. 246\)](#) objects

Marker

Provides an identifier to allow retrieval of paginated results.

Type: String

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

CacheSubnetGroupNotFoundFault

The requested cache subnet group name does not refer to an existing cache subnet group.

HTTP Status Code: 400

Examples

DescribeCacheSubnetGroups

Some of the output has been omitted for brevity.

Sample Request

```
https://elasticache.amazonaws.com/  
?Action=DescribeCacheSubnetGroups  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<DescribeCacheSubnetGroupsResponse xmlns="http://elasticache.amazonaws.com/  
doc/2015-02-02/">  
  <DescribeCacheSubnetGroupsResult>  
    <CacheSubnetGroups>  
      <CacheSubnetGroup>  
        <VpcId>990524496922</VpcId>  
        <CacheSubnetGroupDescription>description</CacheSubnetGroupDescription>  
        <CacheSubnetGroupName>subnet_grp1</CacheSubnetGroupName>  
        <Subnets>  
          <Subnet>  
            <SubnetStatus>Active</SubnetStatus>  
            <SubnetIdentifier>subnet-7c5b4115</SubnetIdentifier>  
            <SubnetAvailabilityZone>  
              <Name>us-west-2c</Name>  
            </SubnetAvailabilityZone>  
          </Subnet>  
          <Subnet>  
            <SubnetStatus>Active</SubnetStatus>  
            <SubnetIdentifier>subnet-7b5b4112</SubnetIdentifier>  
            <SubnetAvailabilityZone>  
              <Name>us-west-2b</Name>  
            </SubnetAvailabilityZone>  
          </Subnet>  
          <Subnet>  
            <SubnetStatus>Active</SubnetStatus>  
            <SubnetIdentifier>subnet-3ea6bd57</SubnetIdentifier>  
            <SubnetAvailabilityZone>  
              <Name>us-west-2c</Name>  
            </SubnetAvailabilityZone>  
          </Subnet>  
        </Subnets>  
      </CacheSubnetGroup>  
    </CacheSubnetGroups>  
  </DescribeCacheSubnetGroupsResult>  
</DescribeCacheSubnetGroupsResponse>
```

```
        </CacheSubnetGroup>
    (...output omitted...)
    </CacheSubnetGroups>
</DescribeCacheSubnetGroupsResult>
<ResponseMetadata>
    <RequestId>31d0faee-229b-11e1-81f1-df3a2a803dad</RequestId>
</ResponseMetadata>
</DescribeCacheSubnetGroupsResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeEngineDefaultParameters

Returns the default engine and system parameter information for the specified cache engine.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheParameterGroupFamily

The name of the cache parameter group family.

Valid values are: `memcached1.4` | `memcached1.5` | `memcached1.6` | `redis2.6` | `redis2.8` | `redis3.2` | `redis4.0` | `redis5.0` | `redis6.x` |

Type: String

Required: Yes

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response. If more records exist than the specified `MaxRecords` value, a marker is included in the response so that the remaining results can be retrieved.

Default: 100

Constraints: minimum 20; maximum 100.

Type: Integer

Required: No

Response Elements

The following element is returned by the service.

EngineDefaults

Represents the output of a `DescribeEngineDefaultParameters` operation.

Type: [EngineDefaults](#) (p. 255) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

Examples

DescribeEngineDefaultParameters

Some of the output has been omitted for brevity.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=DescribeEngineDefaultParameters  
&CacheParameterGroupFamily=memcached1.4  
&MaxRecords=100  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<DescribeEngineDefaultParametersResponse xmlns="http://elasticache.amazonaws.com/  
doc/2015-02-02/">  
  <DescribeEngineDefaultParametersResult>  
    <EngineDefaults>  
      <CacheParameterGroupFamily>memcached1.4</CacheParameterGroupFamily>  
      <Parameters>  
        <Parameter>  
          <ParameterValue>1024</ParameterValue>  
          <DataType>integer</DataType>  
          <Source>system</Source>  
          <IsModifiable>false</IsModifiable>  
          <Description>The backlog queue limit.</Description>  
          <AllowedValues>1-10000</AllowedValues>  
          <ParameterName>backlog_queue_limit</ParameterName>  
          <MinimumEngineVersion>1.4.5</MinimumEngineVersion>  
        </Parameter>  
        <Parameter>  
          (...output omitted...)  
        </Parameter>  
      </Parameters>  
      <CacheNodeTypeSpecificParameters>  
        <CacheNodeTypeSpecificParameter>  
          <CacheNodeTypeSpecificValues>  
            <CacheNodeTypeSpecificValue>  
              <CacheNodeType>cache.c1.xlarge</CacheNodeType>
```

```
        <Value>6000</Value>
      </CacheNodeTypeSpecificValue>

(...output omitted...)

    </CacheNodeTypeSpecificValues>
    <DataType>integer</DataType>
    <Source>system</Source>
    <IsModifiable>false</IsModifiable>
    <Description>The maximum configurable amount of memory to use to store items, in
megabytes.</Description>
    <AllowedValues>1-100000</AllowedValues>
    <ParameterName>max_cache_memory</ParameterName>
    <MinimumEngineVersion>1.4.5</MinimumEngineVersion>
  </CacheNodeTypeSpecificParameter>

(...output omitted...)

</CacheNodeTypeSpecificParameters>
</EngineDefaults>
</DescribeEngineDefaultParametersResult>
<ResponseMetadata>
  <RequestId>061282fe-b7fd-11e0-9326-b7275b9d4a6c</RequestId>
</ResponseMetadata>
</DescribeEngineDefaultParametersResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeEvents

Returns events related to clusters, cache security groups, and cache parameter groups. You can obtain events specific to a particular cluster, cache security group, or cache parameter group by providing the name as a parameter.

By default, only the events occurring within the last hour are returned; however, you can retrieve up to 14 days' worth of events if necessary.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

Duration

The number of minutes worth of events to retrieve.

Type: Integer

Required: No

EndTime

The end of the time interval for which to retrieve events, specified in ISO 8601 format.

Example: 2017-03-30T07:03:49.555Z

Type: Timestamp

Required: No

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response. If more records exist than the specified `MaxRecords` value, a marker is included in the response so that the remaining results can be retrieved.

Default: 100

Constraints: minimum 20; maximum 100.

Type: Integer

Required: No

SourceIdentifier

The identifier of the event source for which events are returned. If not specified, all sources are included in the response.

Type: String

Required: No

SourceType

The event source to retrieve events for. If no value is specified, all events are returned.

Type: String

Valid Values: `cache-cluster` | `cache-parameter-group` | `cache-security-group` | `cache-subnet-group` | `replication-group` | `user` | `user-group`

Required: No

StartTime

The beginning of the time interval to retrieve events for, specified in ISO 8601 format.

Example: 2017-03-30T07:03:49.555Z

Type: Timestamp

Required: No

Response Elements

The following elements are returned by the service.

Events.Event.N

A list of events. Each element in the list contains detailed information about one event.

Type: Array of [Event](#) (p. 256) objects

Marker

Provides an identifier to allow retrieval of paginated results.

Type: String

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

Examples

DescribeEvents

Some of the output has been omitted for brevity.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=DescribeEvents  
&MaxRecords=100  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<DescribeEventsResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">  
  <DescribeEventsResult>  
    <Events>  
      <Event>  
        <Message>Cache cluster created</Message>  
        <SourceType>cache-cluster</SourceType>  
        <Date>2015-02-02T18:22:18.202Z</Date>  
        <SourceIdentifier>my-redis-primary</SourceIdentifier>  
      </Event>  
  
      (...output omitted...)  
  
    </Events>  
  </DescribeEventsResult>  
  <ResponseMetadata>  
    <RequestId>e21c81b4-b9cd-11e3-8a16-7978bb24ffdf</RequestId>  
  </ResponseMetadata>  
</DescribeEventsResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeGlobalReplicationGroups

Returns information about a particular global replication group. If no identifier is specified, returns information about all Global datastores.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

GlobalReplicationGroupId

The name of the Global datastore

Type: String

Required: No

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response. If more records exist than the specified `MaxRecords` value, a marker is included in the response so that the remaining results can be retrieved.

Type: Integer

Required: No

ShowMemberInfo

Returns the list of members that comprise the Global datastore.

Type: Boolean

Required: No

Response Elements

The following elements are returned by the service.

GlobalReplicationGroups.GlobalReplicationGroup.N

Indicates the slot configuration and global identifier for each slice group.

Type: Array of [GlobalReplicationGroup](#) (p. 259) objects

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`. >

Type: String

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

GlobalReplicationGroupNotFoundFault

The Global datastore does not exist

HTTP Status Code: 404

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeReplicationGroups

Returns information about a particular replication group. If no identifier is specified, `DescribeReplicationGroups` returns information about all replication groups.

Note

This operation is valid for Redis only.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response. If more records exist than the specified `MaxRecords` value, a marker is included in the response so that the remaining results can be retrieved.

Default: 100

Constraints: minimum 20; maximum 100.

Type: Integer

Required: No

ReplicationGroupId

The identifier for the replication group to be described. This parameter is not case sensitive.

If you do not specify this parameter, information about all replication groups is returned.

Type: String

Required: No

Response Elements

The following elements are returned by the service.

Marker

Provides an identifier to allow retrieval of paginated results.

Type: String

ReplicationGroups.ReplicationGroup.N

A list of replication groups. Each item in the list contains detailed information about one replication group.

Type: Array of [ReplicationGroup](#) (p. 291) objects

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

Examples

DescribeReplicationGroups

This example illustrates one usage of DescribeReplicationGroups.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=DescribeReplicationGroups  
&MaxRecords=100  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<DescribeReplicationGroupsResponse xmlns="http://elasticache.amazonaws.com/  
doc/2015-02-02/">  
  <DescribeReplicationGroupsResult>  
    <ReplicationGroups>  
      <ReplicationGroup>  
        <SnapshottingClusterId>my-redis-primary</SnapshottingClusterId>  
        <MemberClusters>  
          <ClusterId>my-redis-primary</ClusterId>  
        </MemberClusters>  
        <NodeGroups>  
          <NodeGroup>  
            <NodeGroupId>0001</NodeGroupId>  
            <PrimaryEndpoint>  
              <Port>6379</Port>  
              <Address>my-repgroup.q68zge.ng.0001.use1devo.elmo-dev.amazonaws.com</  
Address>
```

```
</PrimaryEndpoint>
<Status>available</Status>
<NodeGroupMembers>
  <NodeGroupMember>
    <CacheClusterId>my-redis-primary</CacheClusterId>
    <ReadEndpoint>
      <Port>6379</Port>
      <Address>my-redis-primary.q68zge.0001.useldevo.elmo-
dev.amazonaws.com</Address>
    </ReadEndpoint>
    <PreferredAvailabilityZone>us-west-2c</PreferredAvailabilityZone>
    <CacheNodeId>0001</CacheNodeId>
    <CurrentRole>primary</CurrentRole>
  </NodeGroupMember>
</NodeGroupMembers>
</NodeGroup>
</NodeGroups>
<ReplicationGroupId>my-repgroup</ReplicationGroupId>
<Status>available</Status>
<PendingModifiedValues />
<Description>My replication group</Description>
</ReplicationGroup>
</ReplicationGroups>
</DescribeReplicationGroupsResult>
<ResponseMetadata>
  <RequestId>144745b0-b9d3-11e3-8a16-7978bb24ffdf</RequestId>
</ResponseMetadata>
</DescribeReplicationGroupsResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeReservedCacheNodes

Returns information about reserved cache nodes for this account, or about a specified reserved cache node.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

CacheNodeType

The cache node type filter value. Use this parameter to show only those reservations matching the specified cache node type.

The following node types are supported by ElastiCache. Generally speaking, the current generation types provide more memory and computational power at lower cost when compared to their equivalent previous generation counterparts.

- General purpose:

- Current generation:

M6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

`cache.m6g.large`, `cache.m6g.xlarge`, `cache.m6g.2xlarge`, `cache.m6g.4xlarge`,
`cache.m6g.8xlarge`, `cache.m6g.12xlarge`, `cache.m6g.16xlarge`

Note

For region availability, see [Supported Node Types](#)

M5 node types: `cache.m5.large`, `cache.m5.xlarge`, `cache.m5.2xlarge`,
`cache.m5.4xlarge`, `cache.m5.12xlarge`, `cache.m5.24xlarge`

M4 node types: `cache.m4.large`, `cache.m4.xlarge`, `cache.m4.2xlarge`,
`cache.m4.4xlarge`, `cache.m4.10xlarge`

T3 node types: `cache.t3.micro`, `cache.t3.small`, `cache.t3.medium`

T2 node types: `cache.t2.micro`, `cache.t2.small`, `cache.t2.medium`

- Previous generation: (not recommended)

T1 node types: `cache.t1.micro`

M1 node types: `cache.m1.small`, `cache.m1.medium`, `cache.m1.large`,
`cache.m1.xlarge`

M3 node types: `cache.m3.medium`, `cache.m3.large`, `cache.m3.xlarge`,
`cache.m3.2xlarge`

- Compute optimized:

- Previous generation: (not recommended)

C1 node types: `cache.c1.xlarge`

- Memory optimized:

- Current generation:

R6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

`cache.r6g.large`, `cache.r6g.xlarge`, `cache.r6g.2xlarge`, `cache.r6g.4xlarge`,
`cache.r6g.8xlarge`, `cache.r6g.12xlarge`, `cache.r6g.16xlarge`

Note

For region availability, see [Supported Node Types](#)

R5 node types: `cache.r5.large`, `cache.r5.xlarge`, `cache.r5.2xlarge`,
`cache.r5.4xlarge`, `cache.r5.12xlarge`, `cache.r5.24xlarge`

R4 node types: `cache.r4.large`, `cache.r4.xlarge`, `cache.r4.2xlarge`,
`cache.r4.4xlarge`, `cache.r4.8xlarge`, `cache.r4.16xlarge`

- Previous generation: (not recommended)

M2 node types: `cache.m2.xlarge`, `cache.m2.2xlarge`, `cache.m2.4xlarge`

R3 node types: `cache.r3.large`, `cache.r3.xlarge`, `cache.r3.2xlarge`,
`cache.r3.4xlarge`, `cache.r3.8xlarge`

Additional node type info

- All current generation instance types are created in Amazon VPC by default.
- Redis append-only files (AOF) are not supported for T1 or T2 instances.
- Redis Multi-AZ with automatic failover is not supported on T1 instances.
- Redis configuration variables `appendonly` and `appendfsync` are not supported on Redis version 2.8.22 and later.

Type: String

Required: No

Duration

The duration filter value, specified in years or seconds. Use this parameter to show only reservations for this duration.

Valid Values: 1 | 3 | 31536000 | 94608000

Type: String

Required: No

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response. If more records exist than the specified `MaxRecords` value, a marker is included in the response so that the remaining results can be retrieved.

Default: 100

Constraints: minimum 20; maximum 100.

Type: Integer

Required: No

OfferingType

The offering type filter value. Use this parameter to show only the available offerings matching the specified offering type.

Valid values: "Light Utilization"|"Medium Utilization"|"Heavy Utilization"|"All Upfront"|"Partial Upfront"|"No Upfront"

Type: String

Required: No

ProductDescription

The product description filter value. Use this parameter to show only those reservations matching the specified product description.

Type: String

Required: No

ReservedCacheNodeId

The reserved cache node identifier filter value. Use this parameter to show only the reservation that matches the specified reservation ID.

Type: String

Required: No

ReservedCacheNodesOfferingId

The offering identifier filter value. Use this parameter to show only purchased reservations matching the specified offering identifier.

Type: String

Required: No

Response Elements

The following elements are returned by the service.

Marker

Provides an identifier to allow retrieval of paginated results.

Type: String

ReservedCacheNodes.ReservedCacheNode.N

A list of reserved cache nodes. Each element in the list contains detailed information about one node.

Type: Array of [ReservedCacheNode](#) (p. 298) objects

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

ReservedCacheNodeNotFound

The requested reserved cache node was not found.

HTTP Status Code: 404

Examples

DescribeReservedCacheNodes

This example illustrates one usage of DescribeReservedCacheNodes.

Sample Request

```
https://elasticache.amazonaws.com/  
?Action=DescribeReservedCacheNodes  
&ReservedCacheNodeId=customerSpecifiedID  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<DescribeReservedCacheNodesResponse xmlns="http://elasticache.amazonaws.com/  
doc/2015-02-02/">  
  <DescribeReservedCacheNodesResult>  
    <ReservedCacheNodes>  
      <ReservedCacheNode>  
        <OfferingType>Medium Utilization</OfferingType>  
        <RecurringCharges/>  
        <ProductDescription>memcached</ProductDescription>  
        <ReservedCacheNodesOfferingId>649fd0c8-cf6d-47a0-bfa6-060f8e75e95f</  
ReservedCacheNodesOfferingId>  
        <State>payment-failed</State>  
        <ReservedCacheNodeId>myreservationid</ReservedCacheNodeId>  
        <CacheNodeCount>1</CacheNodeCount>  
        <StartTime>2010-12-15T00:25:14.131Z</StartTime>  
        <Duration>31536000</Duration>  
        <FixedPrice>227.5</FixedPrice>  
        <UsagePrice>0.046</UsagePrice>  
        <CacheNodeType>cache.m1.small</CacheNodeType>  
      </ReservedCacheNode>  
    </DescribeReservedCacheNodesResult>  
    <ResponseMetadata>  
      <RequestId>c695119b-2961-11e1-bd06-6fe008f046c3</RequestId>  
    </ResponseMetadata>
```

</DescribeReservedCacheNodesResponse>

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeReservedCacheNodesOfferings

Lists available reserved cache node offerings.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheNodeType

The cache node type filter value. Use this parameter to show only the available offerings matching the specified cache node type.

The following node types are supported by ElastiCache. Generally speaking, the current generation types provide more memory and computational power at lower cost when compared to their equivalent previous generation counterparts.

- General purpose:
 - Current generation:

M6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

`cache.m6g.large`, `cache.m6g.xlarge`, `cache.m6g.2xlarge`, `cache.m6g.4xlarge`,
`cache.m6g.8xlarge`, `cache.m6g.12xlarge`, `cache.m6g.16xlarge`

Note

For region availability, see [Supported Node Types](#)

M5 node types: `cache.m5.large`, `cache.m5.xlarge`, `cache.m5.2xlarge`,
`cache.m5.4xlarge`, `cache.m5.12xlarge`, `cache.m5.24xlarge`

M4 node types: `cache.m4.large`, `cache.m4.xlarge`, `cache.m4.2xlarge`,
`cache.m4.4xlarge`, `cache.m4.10xlarge`

T3 node types: `cache.t3.micro`, `cache.t3.small`, `cache.t3.medium`

T2 node types: `cache.t2.micro`, `cache.t2.small`, `cache.t2.medium`

- Previous generation: (not recommended)

T1 node types: `cache.t1.micro`

M1 node types: `cache.m1.small`, `cache.m1.medium`, `cache.m1.large`,
`cache.m1.xlarge`

M3 node types: `cache.m3.medium`, `cache.m3.large`, `cache.m3.xlarge`,
`cache.m3.2xlarge`

- Compute optimized:
 - Previous generation: (not recommended)

C1 node types: `cache.c1.xlarge`

- Memory optimized:
 - Current generation:

R6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

`cache.r6g.large`, `cache.r6g.xlarge`, `cache.r6g.2xlarge`, `cache.r6g.4xlarge`,
`cache.r6g.8xlarge`, `cache.r6g.12xlarge`, `cache.r6g.16xlarge`

Note

For region availability, see [Supported Node Types](#)

R5 node types: `cache.r5.large`, `cache.r5.xlarge`, `cache.r5.2xlarge`,
`cache.r5.4xlarge`, `cache.r5.12xlarge`, `cache.r5.24xlarge`

R4 node types: `cache.r4.large`, `cache.r4.xlarge`, `cache.r4.2xlarge`,
`cache.r4.4xlarge`, `cache.r4.8xlarge`, `cache.r4.16xlarge`

- Previous generation: (not recommended)

M2 node types: `cache.m2.xlarge`, `cache.m2.2xlarge`, `cache.m2.4xlarge`

R3 node types: `cache.r3.large`, `cache.r3.xlarge`, `cache.r3.2xlarge`,
`cache.r3.4xlarge`, `cache.r3.8xlarge`

Additional node type info

- All current generation instance types are created in Amazon VPC by default.
- Redis append-only files (AOF) are not supported for T1 or T2 instances.
- Redis Multi-AZ with automatic failover is not supported on T1 instances.
- Redis configuration variables `appendonly` and `appendfsync` are not supported on Redis version 2.8.22 and later.

Type: String

Required: No

Duration

Duration filter value, specified in years or seconds. Use this parameter to show only reservations for a given duration.

Valid Values: 1 | 3 | 31536000 | 94608000

Type: String

Required: No

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response. If more records exist than the specified `MaxRecords` value, a marker is included in the response so that the remaining results can be retrieved.

Default: 100

Constraints: minimum 20; maximum 100.

Type: Integer

Required: No

OfferingType

The offering type filter value. Use this parameter to show only the available offerings matching the specified offering type.

Valid Values: "Light Utilization"|"Medium Utilization"|"Heavy Utilization"|"All Upfront"|"Partial Upfront"|"No Upfront"

Type: String

Required: No

ProductDescription

The product description filter value. Use this parameter to show only the available offerings matching the specified product description.

Type: String

Required: No

ReservedCacheNodesOfferingId

The offering identifier filter value. Use this parameter to show only the available offering that matches the specified reservation identifier.

Example: 438012d3-4052-4cc7-b2e3-8d3372e0e706

Type: String

Required: No

Response Elements

The following elements are returned by the service.

Marker

Provides an identifier to allow retrieval of paginated results.

Type: String

ReservedCacheNodesOfferings.ReservedCacheNodesOffering.N

A list of reserved cache node offerings. Each element in the list contains detailed information about one offering.

Type: Array of [ReservedCacheNodesOffering](#) (p. 302) objects

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

ReservedCacheNodesOfferingNotFound

The requested cache node offering does not exist.

HTTP Status Code: 404

Examples

DescribeReservedCacheNodesOfferings

This example illustrates one usage of DescribeReservedCacheNodesOfferings.

Sample Request

```
https://elasticache.amazonaws.com/  
?Action=DescribeReservedCacheNodesOfferings  
&ReservedCacheNodesOfferingId=438012d3-4052-4cc7-b2e3-8d3372e0e706  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<DescribeReservedCacheNodesOfferingsResponse xmlns="http://elasticache.amazonaws.com/  
doc/2015-02-02/">  
  <DescribeReservedCacheNodesOfferingsResult>  
    <ReservedCacheNodesOfferings>  
      <ReservedCacheNodesOffering>  
        <Duration>31536000</Duration>  
        <OfferingType>Heavy Utilization</OfferingType>  
        <RecurringCharges>  
          <RecurringCharge>  
            <RecurringChargeFrequency>Hourly</RecurringChargeFrequency>  
            <RecurringChargeAmount>0.123</RecurringChargeAmount>  
          </RecurringCharge>  
        </RecurringCharges>  
        <FixedPrice>162.0</FixedPrice>  
        <ProductDescription>memcached</ProductDescription>  
        <UsagePrice>0.0</UsagePrice>  
        <ReservedCacheNodesOfferingId>SampleOfferingId</ReservedCacheNodesOfferingId>  
        <CacheNodeType>cache.m1.small</CacheNodeType>  
      </ReservedCacheNodesOffering>  
    </ReservedCacheNodesOfferings>  
  </DescribeReservedCacheNodesOfferingsResult>  
  <ResponseMetadata>  
    <RequestId>521b420a-2961-11e1-bd06-6fe008f046c3</RequestId>  
  </ResponseMetadata>  
</DescribeReservedCacheNodesOfferingsResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)

- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeServiceUpdates

Returns details of the service updates

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response

Type: Integer

Required: No

ServiceUpdateName

The unique ID of the service update

Type: String

Required: No

ServiceUpdateStatus.member.N

The status of the service update

Type: Array of strings

Array Members: Maximum number of 3 items.

Valid Values: `available` | `cancelled` | `expired`

Required: No

Response Elements

The following elements are returned by the service.

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

ServiceUpdates.ServiceUpdate.N

A list of service updates

Type: Array of [ServiceUpdate](#) (p. 308) objects

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

ServiceUpdateNotFoundFault

The service update doesn't exist

HTTP Status Code: 404

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeSnapshots

Returns information about cluster or replication group snapshots. By default, `DescribeSnapshots` lists all of your snapshots; it can optionally describe a single snapshot, or just the snapshots associated with a particular cache cluster.

Note

This operation is valid for Redis only.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheClusterId

A user-supplied cluster identifier. If this parameter is specified, only snapshots associated with that specific cluster are described.

Type: String

Required: No

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response. If more records exist than the specified `MaxRecords` value, a marker is included in the response so that the remaining results can be retrieved.

Default: 50

Constraints: minimum 20; maximum 50.

Type: Integer

Required: No

ReplicationGroupId

A user-supplied replication group identifier. If this parameter is specified, only snapshots associated with that specific replication group are described.

Type: String

Required: No

ShowNodeGroupConfig

A Boolean value which if true, the node group (shard) configuration is included in the snapshot description.

Type: Boolean

Required: No

SnapshotName

A user-supplied name of the snapshot. If this parameter is specified, only this snapshot are described.

Type: String

Required: No

SnapshotSource

If set to `system`, the output shows snapshots that were automatically created by ElastiCache. If set to `user` the output shows snapshots that were manually created. If omitted, the output shows both automatically and manually created snapshots.

Type: String

Required: No

Response Elements

The following elements are returned by the service.

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

Snapshots.Snapshot.N

A list of snapshots. Each item in the list contains detailed information about one snapshot.

Type: Array of [Snapshot](#) (p. 311) objects

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

CacheClusterNotFound

The requested cluster ID does not refer to an existing cluster.

HTTP Status Code: 404

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

SnapshotNotFoundFault

The requested snapshot name does not refer to an existing snapshot.

HTTP Status Code: 404

Examples

DescribeSnapshots

This example illustrates one usage of DescribeSnapshots.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/
?Action=DescribeSnapshots
&MaxRecords=50
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20150202T192317Z
&X-Amz-Credential=<credential>
```

Sample Response

```
<DescribeSnapshotsResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <DescribeSnapshotsResult>
    <Snapshots>
      <Snapshot>
        <CacheClusterId>my-redis-primary</CacheClusterId>
        <Port>6379</Port>
        <CacheNodeType>cache.m1.small</CacheNodeType>
        <CacheParameterGroupName>default.redis2.8</CacheParameterGroupName>
        <Engine>redis</Engine>
        <PreferredAvailabilityZone>us-west-2c</PreferredAvailabilityZone>
        <CacheClusterCreateTime>2015-02-02T18:46:57.972Z</CacheClusterCreateTime>
        <EngineVersion>2.8.6</EngineVersion>
        <SnapshotSource>manual</SnapshotSource>
        <AutoMinorVersionUpgrade>true</AutoMinorVersionUpgrade>
        <PreferredMaintenanceWindow>wed:09:00-wed:10:00</PreferredMaintenanceWindow>
        <SnapshotName>my-manual-snapshot</SnapshotName>
        <SnapshotRetentionLimit>5</SnapshotRetentionLimit>
        <NodeSnapshots>
          <NodeSnapshot>
            <SnapshotCreateTime>2015-02-02T18:54:12Z</SnapshotCreateTime>
            <CacheNodeCreateTime>2015-02-02T18:46:57.972Z</CacheNodeCreateTime>
            <CacheNodeId>0001</CacheNodeId>
            <CacheSize>3 MB</CacheSize>
          </NodeSnapshot>
        </NodeSnapshots>
        <SnapshotStatus>creating</SnapshotStatus>
        <NumCacheNodes>1</NumCacheNodes>
        <SnapshotWindow>07:30-08:30</SnapshotWindow>
      </Snapshot>
    </Snapshots>
  </DescribeSnapshotsResult>
  <ResponseMetadata>
    <RequestId>51b0b25e-b9cf-11e3-8a16-7978bb24ffdf</RequestId>
```

```
</ResponseMetadata>  
</DescribeSnapshotsResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeUpdateActions

Returns details of the update actions

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheClusterIds.member.N

The cache cluster IDs

Type: Array of strings

Array Members: Maximum number of 20 items.

Required: No

Engine

The ElastiCache engine to which the update applies. Either Redis or Memcached

Type: String

Required: No

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response

Type: Integer

Required: No

ReplicationGroupIds.member.N

The replication group IDs

Type: Array of strings

Array Members: Maximum number of 20 items.

Required: No

ServiceUpdateName

The unique ID of the service update

Type: String

Required: No

ServiceUpdateStatus.member.N

The status of the service update

Type: Array of strings

Array Members: Maximum number of 3 items.

Valid Values: `available` | `cancelled` | `expired`

Required: No

ServiceUpdateTimeRange

The range of time specified to search for service updates that are in available status

Type: [TimeRangeFilter](#) (p. 320) object

Required: No

ShowNodeLevelUpdateStatus

Dictates whether to include node level update status in the response

Type: Boolean

Required: No

UpdateActionStatus.member.N

The status of the update action.

Type: Array of strings

Array Members: Maximum number of 9 items.

Valid Values: `not-applied` | `waiting-to-start` | `in-progress` | `stopping` | `stopped` | `complete` | `scheduling` | `scheduled` | `not-applicable`

Required: No

Response Elements

The following elements are returned by the service.

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by `MaxRecords`.

Type: String

UpdateActions.UpdateAction.N

Returns a list of update actions

Type: Array of [UpdateAction](#) (p. 322) objects

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeUserGroups

Returns a list of user groups.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by MaxRecords. >

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response. If more records exist than the specified MaxRecords value, a marker is included in the response so that the remaining results can be retrieved.

Type: Integer

Required: No

UserGroupId

The ID of the user group.

Type: String

Required: No

Response Elements

The following elements are returned by the service.

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by MaxRecords. >

Type: String

UserGroups.member.N

Returns a list of user groups.

Type: Array of [UserGroup](#) (p. 327) objects

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

UserGroupNotFound

The user group was not found or does not exist

HTTP Status Code: 404

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeUsers

Returns a list of users.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

Engine

The Redis engine.

Type: String

Pattern: [a-zA-Z]*

Required: No

Filters.member.N

Filter to determine the list of User IDs to return.

Type: Array of [Filter \(p. 257\)](#) objects

Required: No

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by MaxRecords. >

Type: String

Required: No

MaxRecords

The maximum number of records to include in the response. If more records exist than the specified MaxRecords value, a marker is included in the response so that the remaining results can be retrieved.

Type: Integer

Required: No

UserId

The ID of the user.

Type: String

Length Constraints: Minimum length of 1.

Pattern: [a-zA-Z][a-zA-Z0-9\-*]

Required: No

Response Elements

The following elements are returned by the service.

Marker

An optional marker returned from a prior request. Use this marker for pagination of results from this operation. If this parameter is specified, the response includes only records beyond the marker, up to the value specified by MaxRecords. >

Type: String

Users.member.N

A list of users.

Type: Array of [User](#) (p. 325) objects

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

UserNotFound

The user does not exist or could not be found.

HTTP Status Code: 404

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DisassociateGlobalReplicationGroup

Remove a secondary cluster from the Global datastore using the Global datastore name. The secondary cluster will no longer receive updates from the primary cluster, but will remain as a standalone cluster in that Amazon region.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

GlobalReplicationGroupId

The name of the Global datastore

Type: String

Required: Yes

ReplicationGroupId

The name of the secondary cluster you wish to remove from the Global datastore

Type: String

Required: Yes

ReplicationGroupRegion

The Amazon region of secondary cluster you wish to remove from the Global datastore

Type: String

Required: Yes

Response Elements

The following element is returned by the service.

GlobalReplicationGroup

Consists of a primary cluster that accepts writes and an associated secondary cluster that resides in a different Amazon region. The secondary cluster accepts only reads. The primary cluster automatically replicates updates to the secondary cluster.

- The **GlobalReplicationGroupIdSuffix** represents the name of the Global datastore, which is what you use to associate a secondary cluster.

Type: [GlobalReplicationGroup \(p. 259\)](#) object

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

GlobalReplicationGroupNotFoundFault

The Global datastore does not exist

HTTP Status Code: 404

InvalidGlobalReplicationGroupState

The Global datastore is not available or in primary-only state.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

FailoverGlobalReplicationGroup

Used to failover the primary region to a selected secondary region. The selected secondary region will become primary, and all other clusters will become secondary.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

GlobalReplicationGroupId

The name of the Global datastore

Type: String

Required: Yes

PrimaryRegion

The Amazon region of the primary cluster of the Global datastore

Type: String

Required: Yes

PrimaryReplicationGroupId

The name of the primary replication group

Type: String

Required: Yes

Response Elements

The following element is returned by the service.

GlobalReplicationGroup

Consists of a primary cluster that accepts writes and an associated secondary cluster that resides in a different Amazon region. The secondary cluster accepts only reads. The primary cluster automatically replicates updates to the secondary cluster.

- The **GlobalReplicationGroupIdSuffix** represents the name of the Global datastore, which is what you use to associate a secondary cluster.

Type: [GlobalReplicationGroup \(p. 259\)](#) object

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

GlobalReplicationGroupNotFoundFault

The Global datastore does not exist

HTTP Status Code: 404

InvalidGlobalReplicationGroupState

The Global datastore is not available or in primary-only state.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

IncreaseNodeGroupsInGlobalReplicationGroup

Increase the number of node groups in the Global datastore

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

ApplyImmediately

Indicates that the process begins immediately. At present, the only permitted value for this parameter is true.

Type: Boolean

Required: Yes

GlobalReplicationGroupId

The name of the Global datastore

Type: String

Required: Yes

NodeGroupCount

The number of node groups you wish to add

Type: Integer

Required: Yes

RegionalConfigurations.RegionalConfiguration.N

Describes the replication group IDs, the Amazon regions where they are stored and the shard configuration for each that comprise the Global datastore

Type: Array of [RegionalConfiguration](#) (p. 290) objects

Required: No

Response Elements

The following element is returned by the service.

GlobalReplicationGroup

Consists of a primary cluster that accepts writes and an associated secondary cluster that resides in a different Amazon region. The secondary cluster accepts only reads. The primary cluster automatically replicates updates to the secondary cluster.

- The **GlobalReplicationGroupIdSuffix** represents the name of the Global datastore, which is what you use to associate a secondary cluster.

Type: [GlobalReplicationGroup](#) (p. 259) object

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

GlobalReplicationGroupNotFoundFault

The Global datastore does not exist

HTTP Status Code: 404

InvalidGlobalReplicationGroupState

The Global datastore is not available or in primary-only state.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

IncreaseReplicaCount

Dynamically increases the number of replicas in a Redis (cluster mode disabled) replication group or the number of replica nodes in one or more node groups (shards) of a Redis (cluster mode enabled) replication group. This operation is performed with no cluster down time.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

ApplyImmediately

If `True`, the number of replica nodes is increased immediately. `ApplyImmediately=False` is not currently supported.

Type: Boolean

Required: Yes

ReplicationGroupId

The id of the replication group to which you want to add replica nodes.

Type: String

Required: Yes

NewReplicaCount

The number of read replica nodes you want at the completion of this operation. For Redis (cluster mode disabled) replication groups, this is the number of replica nodes in the replication group. For Redis (cluster mode enabled) replication groups, this is the number of replica nodes in each of the replication group's node groups.

Type: Integer

Required: No

ReplicaConfiguration.ConfigureShard.N

A list of `ConfigureShard` objects that can be used to configure each shard in a Redis (cluster mode enabled) replication group. The `ConfigureShard` has three members: `NewReplicaCount`, `NodeId`, and `PreferredAvailabilityZones`.

Type: Array of [ConfigureShard](#) (p. 249) objects

Required: No

Response Elements

The following element is returned by the service.

ReplicationGroup

Contains all of the attributes of a specific Redis replication group.

Type: [ReplicationGroup](#) (p. 291) object

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

ClusterQuotaForCustomerExceeded

The request cannot be processed because it would exceed the allowed number of clusters per customer.

HTTP Status Code: 400

InsufficientCacheClusterCapacity

The requested cache node type is not available in the specified Availability Zone. For more information, see [InsufficientCacheClusterCapacity](#) in the ElastiCache User Guide.

HTTP Status Code: 400

InvalidCacheClusterState

The requested cluster is not in the `available` state.

HTTP Status Code: 400

InvalidKMSKeyFault

The KMS key supplied is not valid.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidReplicationGroupState

The requested replication group is not in the `available` state.

HTTP Status Code: 400

InvalidVPCNetworkStateFault

The VPC network is in an invalid state.

HTTP Status Code: 400

NodeGroupsPerReplicationGroupQuotaExceeded

The request cannot be processed because it would exceed the maximum allowed number of node groups (shards) in a single replication group. The default maximum is 90

HTTP Status Code: 400

NodeQuotaForCustomerExceeded

The request cannot be processed because it would exceed the allowed number of cache nodes per customer.

HTTP Status Code: 400

NoOperationFault

The operation was not performed because no changes were required.

HTTP Status Code: 400

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

Examples

Example

The following example increases the replica count to 3 in all node groups of `sample-repl-group`.

```
https://elasticache.us-west-2.amazonaws.com/
?Action=IncreaseReplicaCount
&ApplyImmediately=True
&NewReplicaCount=3
&ReplicationGroupId=sample-repl-group
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20150202T192317Z
&X-Amz-Credential=<credential>
```

Example

The following example increases the replica count in two node groups. Because there are multiple node groups, this code only works for Redis (cluster mode enabled) replication groups.

```
https://elasticache.us-west-2.amazonaws.com/
?Action=IncreaseReplicaCount
&ApplyImmediately=True
&ReplicaConfiguration.ConfigureShard.1.NodeGroupId=0001
&ReplicaConfiguration.ConfigureShard.1.NewReplicaCount=2

&ReplicaConfiguration.ConfigureShard.1.PreferredAvailabilityZones.PreferredAvailabilityZone.1=us-east-1a

&ReplicaConfiguration.ConfigureShard.1.PreferredAvailabilityZones.PreferredAvailabilityZone.2=us-east-1c

&ReplicaConfiguration.ConfigureShard.1.PreferredAvailabilityZones.PreferredAvailabilityZone.3=us-east-1b
&ReplicaConfiguration.ConfigureShard.2.NodeGroupId=0003
&ReplicaConfiguration.ConfigureShard.2.NewReplicaCount=3

&ReplicaConfiguration.ConfigureShard.2.PreferredAvailabilityZones.PreferredAvailabilityZone.1=us-east-1a

&ReplicaConfiguration.ConfigureShard.2.PreferredAvailabilityZones.PreferredAvailabilityZone.2=us-east-1b

&ReplicaConfiguration.ConfigureShard.2.PreferredAvailabilityZones.PreferredAvailabilityZone.3=us-east-1c
```

```
&ReplicaConfiguration.ConfigureShard.2.PreferredAvailabilityZones.PreferredAvailabilityZone.4=us-east-1c
&ReplicationGroupId=sample-repl-group
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20150202T192317Z
&X-Amz-Credential=<credential>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ListAllowedNodeTypeModifications

Lists all available node types that you can scale your Redis cluster's or replication group's current node type.

When you use the `ModifyCacheCluster` or `ModifyReplicationGroup` operations to scale your cluster or replication group, the value of the `CacheNodeType` parameter must be one of the node types returned by this operation.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheClusterId

The name of the cluster you want to scale up to a larger node instanced type. ElastiCache uses the cluster id to identify the current node type of this cluster and from that to create a list of node types you can scale up to.

Important

You must provide a value for either the `CacheClusterId` or the `ReplicationGroupId`.

Type: String

Required: No

ReplicationGroupId

The name of the replication group want to scale up to a larger node type. ElastiCache uses the replication group id to identify the current node type being used by this replication group, and from that to create a list of node types you can scale up to.

Important

You must provide a value for either the `CacheClusterId` or the `ReplicationGroupId`.

Type: String

Required: No

Response Elements

The following elements are returned by the service.

ScaleDownModifications.member.N

A string list, each element of which specifies a cache node type which you can use to scale your cluster or replication group. When scaling down a Redis cluster or replication group using `ModifyCacheCluster` or `ModifyReplicationGroup`, use a value from this list for the `CacheNodeType` parameter.

Type: Array of strings

ScaleUpModifications.member.N

A string list, each element of which specifies a cache node type which you can use to scale your cluster or replication group.

When scaling up a Redis cluster or replication group using `ModifyCacheCluster` or `ModifyReplicationGroup`, use a value from this list for the `CacheNodeType` parameter.

Type: Array of strings

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

CacheClusterNotFound

The requested cluster ID does not refer to an existing cluster.

HTTP Status Code: 404

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

Examples

ListAllowedNodeTypeModifications for a Cluster

The following example request a list of node types you can use to scale `myCluster` up.

Sample Request

```
https://elasticache.us-east-1.amazonaws.com/  
?Action=ListAllowedNodeTypeModifications  
&CacheClusterId=mycachecluster  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Version=2015-02-02  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

ListAllowedNodeTypeModifications for a Replication Group

The following example requests a list of node types you can use to scale `myRep1Group` up.

Sample Request

```
https://elasticache.us-east-1.amazonaws.com/  
?Action=ListAllowedNodeTypeModifications  
&ReplicationGroupId=myreplgroup  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256
```

```
&Version=2015-02-02  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ListTagsForResource

Lists all tags currently on a named resource.

A tag is a key-value pair where the key and value are case-sensitive. You can use tags to categorize and track all your ElastiCache resources, with the exception of global replication group. When you add or remove tags on replication groups, those actions will be replicated to all nodes in the replication group. For more information, see [Resource-level permissions](#).

If the cluster is not in the *available* state, `ListTagsForResource` returns an error.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

ResourceName

The Amazon Resource Name (ARN) of the resource for which you want the list of tags, for example `arn:aws:elasticache:us-west-2:0123456789:cluster:myCluster` or `arn:aws:elasticache:us-west-2:0123456789:snapshot:mySnapshot`.

For more information about ARNs, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Required: Yes

Response Elements

The following element is returned by the service.

TagList.Tag.N

A list of tags as key-value pairs.

Type: Array of [Tag \(p. 319\)](#) objects

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

CacheClusterNotFound

The requested cluster ID does not refer to an existing cluster.

HTTP Status Code: 404

CacheParameterGroupNotFound

The requested cache parameter group name does not refer to an existing cache parameter group.

HTTP Status Code: 404

CacheSecurityGroupNotFound

The requested cache security group name does not refer to an existing cache security group.

HTTP Status Code: 404

CacheSubnetGroupNotFoundFault

The requested cache subnet group name does not refer to an existing cache subnet group.

HTTP Status Code: 400

InvalidARN

The requested Amazon Resource Name (ARN) does not refer to an existing resource.

HTTP Status Code: 400

InvalidReplicationGroupState

The requested replication group is not in the available state.

HTTP Status Code: 400

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

ReservedCacheNodeNotFound

The requested reserved cache node was not found.

HTTP Status Code: 404

SnapshotNotFoundFault

The requested snapshot name does not refer to an existing snapshot.

HTTP Status Code: 404

UserGroupNotFound

The user group was not found or does not exist

HTTP Status Code: 404

UserNotFound

The user does not exist or could not be found.

HTTP Status Code: 404

Examples

ListTagsForResource

This example illustrates one usage of ListTagsForResource.

Sample Request

```
https://elasticache.us-east-1.amazonaws.com/  
?Action=ListTagsForResource  
&ResourceName=arn:aws:elasticache:us-west-2:0123456789:cluster:myCluster  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Version=2015-02-02
```

```
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ModifyCacheCluster

Modifies the settings for a cluster. You can use this operation to change one or more cluster configuration parameters by specifying the parameters and the new values.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheClusterId

The cluster identifier. This value is stored as a lowercase string.

Type: String

Required: Yes

ApplyImmediately

If `true`, this parameter causes the modifications in this request and any pending modifications to be applied, asynchronously and as soon as possible, regardless of the `PreferredMaintenanceWindow` setting for the cluster.

If `false`, changes to the cluster are applied on the next maintenance reboot, or the next failure reboot, whichever occurs first.

Important

If you perform a `ModifyCacheCluster` before a pending modification is applied, the pending modification is replaced by the newer modification.

Valid values: `true` | `false`

Default: `false`

Type: Boolean

Required: No

AuthToken

Reserved parameter. The password used to access a password protected server. This parameter must be specified with the `auth-token-update` parameter. Password constraints:

- Must be only printable ASCII characters
- Must be at least 16 characters and no more than 128 characters in length
- Cannot contain any of the following characters: `/`, `''`, or `@`, `%`

For more information, see AUTH password at [AUTH](#).

Type: String

Required: No

AuthTokenUpdateStrategy

Specifies the strategy to use to update the AUTH token. This parameter must be specified with the `auth-token` parameter. Possible values:

- Rotate
- Set

For more information, see [Authenticating Users with Redis AUTH](#)

Type: String

Valid Values: `SET` | `ROTATE` | `DELETE`

Required: No

AutoMinorVersionUpgrade

This parameter is currently disabled.

Type: Boolean

Required: No

AZMode

Specifies whether the new nodes in this Memcached cluster are all created in a single Availability Zone or created across multiple Availability Zones.

Valid values: `single-az` | `cross-az`.

This option is only supported for Memcached clusters.

Note

You cannot specify `single-az` if the Memcached cluster already has cache nodes in different Availability Zones. If `cross-az` is specified, existing Memcached nodes remain in their current Availability Zone.

Only newly created nodes are located in different Availability Zones.

Type: String

Valid Values: `single-az` | `cross-az`

Required: No

CacheNodeIdsToRemove.CacheNodeId.N

A list of cache node IDs to be removed. A node ID is a numeric identifier (0001, 0002, etc.). This parameter is only valid when `NumCacheNodes` is less than the existing number of cache nodes. The number of cache node IDs supplied in this parameter must match the difference between the existing number of cache nodes in the cluster or pending cache nodes, whichever is greater, and the value of `NumCacheNodes` in the request.

For example: If you have 3 active cache nodes, 7 pending cache nodes, and the number of cache nodes in this `ModifyCacheCluster` call is 5, you must list 2 (7 - 5) cache node IDs to remove.

Type: Array of strings

Required: No

CacheNodeType

A valid cache node type that you want to scale this cluster up to.

Type: String

Required: No

CacheParameterGroupName

The name of the cache parameter group to apply to this cluster. This change is asynchronously applied as soon as possible for parameters when the `ApplyImmediately` parameter is specified as `true` for this request.

Type: String

Required: No

CacheSecurityGroupNames.CacheSecurityGroupName.N

A list of cache security group names to authorize on this cluster. This change is asynchronously applied as soon as possible.

You can use this parameter only with clusters that are created outside of an Amazon Virtual Private Cloud (Amazon VPC).

Constraints: Must contain no more than 255 alphanumeric characters. Must not be "Default".

Type: Array of strings

Required: No

EngineVersion

The upgraded version of the cache engine to be run on the cache nodes.

Important: You can upgrade to a newer engine version (see [Selecting a Cache Engine and Version](#)), but you cannot downgrade to an earlier engine version. If you want to use an earlier engine version, you must delete the existing cluster and create it anew with the earlier engine version.

Type: String

Required: No

LogDeliveryConfigurations.LogDeliveryConfigurationRequest.N

Specifies the destination, format and type of the logs.

Type: Array of [LogDeliveryConfigurationRequest](#) (p. 268) objects

Required: No

NewAvailabilityZones.PreferredAvailabilityZone.N

Note

This option is only supported on Memcached clusters.

The list of Availability Zones where the new Memcached cache nodes are created.

This parameter is only valid when `NumCacheNodes` in the request is greater than the sum of the number of active cache nodes and the number of cache nodes pending creation (which may be zero). The number of Availability Zones supplied in this list must match the cache nodes being added in this request.

Scenarios:

- **Scenario 1:** You have 3 active nodes and wish to add 2 nodes. Specify `NumCacheNodes=5` ($3 + 2$) and optionally specify two Availability Zones for the two new nodes.
- **Scenario 2:** You have 3 active nodes and 2 nodes pending creation (from the scenario 1 call) and want to add 1 more node. Specify `NumCacheNodes=6` ($(3 + 2) + 1$) and optionally specify an Availability Zone for the new node.
- **Scenario 3:** You want to cancel all pending operations. Specify `NumCacheNodes=3` to cancel all pending operations.

The Availability Zone placement of nodes pending creation cannot be modified. If you wish to cancel any nodes pending creation, add 0 nodes by setting `NumCacheNodes` to the number of current nodes.

If `cross-az` is specified, existing Memcached nodes remain in their current Availability Zone. Only newly created nodes can be located in different Availability Zones. For guidance on how to move existing Memcached nodes to different Availability Zones, see the **Availability Zone Considerations** section of [Cache Node Considerations for Memcached](#).

Impact of new add/remove requests upon pending requests

- Scenario-1
 - Pending Action: Delete
 - New Request: Delete
 - Result: The new delete, pending or immediate, replaces the pending delete.
- Scenario-2
 - Pending Action: Delete
 - New Request: Create
 - Result: The new create, pending or immediate, replaces the pending delete.
- Scenario-3
 - Pending Action: Create
 - New Request: Delete
 - Result: The new delete, pending or immediate, replaces the pending create.
- Scenario-4
 - Pending Action: Create
 - New Request: Create
 - Result: The new create is added to the pending create.

Important

Important: If the new create request is **Apply Immediately - Yes**, all creates are performed immediately. If the new create request is **Apply Immediately - No**, all creates are pending.

Type: Array of strings

Required: No

NotificationTopicArn

The Amazon Resource Name (ARN) of the Amazon SNS topic to which notifications are sent.

Note

The Amazon SNS topic owner must be same as the cluster owner.

Type: String

Required: No

NotificationTopicStatus

The status of the Amazon SNS notification topic. Notifications are sent only if the status is `active`.

Valid values: `active` | `inactive`

Type: String

Required: No

NumCacheNodes

The number of cache nodes that the cluster should have. If the value for `NumCacheNodes` is greater than the sum of the number of current cache nodes and the number of cache nodes pending creation (which may be zero), more nodes are added. If the value is less than the number of existing

cache nodes, nodes are removed. If the value is equal to the number of current cache nodes, any pending add or remove requests are canceled.

If you are removing cache nodes, you must use the `CacheNodeIdsToRemove` parameter to provide the IDs of the specific cache nodes to remove.

For clusters running Redis, this value must be 1. For clusters running Memcached, this value must be between 1 and 40.

Note

Adding or removing Memcached cache nodes can be applied immediately or as a pending operation (see `ApplyImmediately`).

A pending operation to modify the number of cache nodes in a cluster during its maintenance window, whether by adding or removing nodes in accordance with the scale out architecture, is not queued. The customer's latest request to add or remove nodes to the cluster overrides any previous pending operations to modify the number of cache nodes in the cluster. For example, a request to remove 2 nodes would override a previous pending operation to remove 3 nodes. Similarly, a request to add 2 nodes would override a previous pending operation to remove 3 nodes and vice versa. As Memcached cache nodes may now be provisioned in different Availability Zones with flexible cache node placement, a request to add nodes does not automatically override a previous pending operation to add nodes. The customer can modify the previous pending operation to add more nodes or explicitly cancel the pending request and retry the new request. To cancel pending operations to modify the number of cache nodes in a cluster, use the `ModifyCacheCluster` request and set `NumCacheNodes` equal to the number of cache nodes currently in the cluster.

Type: Integer

Required: No

PreferredMaintenanceWindow

Specifies the weekly time range during which maintenance on the cluster is performed. It is specified as a range in the format `ddd:hh24:mi-ddd:hh24:mi` (24H Clock UTC). The minimum maintenance window is a 60 minute period.

Valid values for `ddd` are:

- `sun`
- `mon`
- `tue`
- `wed`
- `thu`
- `fri`
- `sat`

Example: `sun:23:00-mon:01:30`

Type: String

Required: No

SecurityGroupIds.SecurityGroupId.N

Specifies the VPC Security Groups associated with the cluster.

This parameter can be used only with clusters that are created in an Amazon Virtual Private Cloud (Amazon VPC).

Type: Array of strings

Required: No

SnapshotRetentionLimit

The number of days for which ElastiCache retains automatic cluster snapshots before deleting them. For example, if you set `SnapshotRetentionLimit` to 5, a snapshot that was taken today is retained for 5 days before being deleted.

Note

If the value of `SnapshotRetentionLimit` is set to zero (0), backups are turned off.

Type: Integer

Required: No

SnapshotWindow

The daily time range (in UTC) during which ElastiCache begins taking a daily snapshot of your cluster.

Type: String

Required: No

Response Elements

The following element is returned by the service.

CacheCluster

Contains all of the attributes of a specific cluster.

Type: [CacheCluster](#) (p. 224) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

CacheClusterNotFound

The requested cluster ID does not refer to an existing cluster.

HTTP Status Code: 404

CacheParameterGroupNotFound

The requested cache parameter group name does not refer to an existing cache parameter group.

HTTP Status Code: 404

CacheSecurityGroupNotFound

The requested cache security group name does not refer to an existing cache security group.

HTTP Status Code: 404

InsufficientCacheClusterCapacity

The requested cache node type is not available in the specified Availability Zone. For more information, see [InsufficientCacheClusterCapacity](#) in the ElastiCache User Guide.

HTTP Status Code: 400

InvalidCacheClusterState

The requested cluster is not in the available state.

HTTP Status Code: 400

InvalidCacheSecurityGroupState

The current state of the cache security group does not allow deletion.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidVPCNetworkStateFault

The VPC network is in an invalid state.

HTTP Status Code: 400

NodeQuotaForClusterExceeded

The request cannot be processed because it would exceed the allowed number of cache nodes in a single cluster.

HTTP Status Code: 400

NodeQuotaForCustomerExceeded

The request cannot be processed because it would exceed the allowed number of cache nodes per customer.

HTTP Status Code: 400

Examples

ModifyCacheCluster

This example illustrates one usage of ModifyCacheCluster.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=ModifyCacheCluster  
  NumCacheNodes=5  
  &CacheClusterId=simcoprod01  
  &ApplyImmediately=true  
  &Version=2015-02-02  
  &SignatureVersion=4  
  &SignatureMethod=HmacSHA256  
  &Timestamp=20150202T192317Z  
  &X-Amz-Credential=<credential>
```

Sample Response

```
<ModifyCacheClusterResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <ModifyCacheClusterResult>
    <CacheCluster>
      <CacheParameterGroup>
        <ParameterApplyStatus>in-sync</ParameterApplyStatus>
        <CacheParameterGroupName>default.memcached1.4</CacheParameterGroupName>
        <CacheNodeIdsToReboot/>
      </CacheParameterGroup>
      <CacheClusterId>simcoprod01</CacheClusterId>
      <CacheClusterStatus>available</CacheClusterStatus>
      <ConfigurationEndpoint>
        <Port>11211</Port>
        <Address>simcoprod01.m2st2p.cfg.cache.amazonaws.com</Address>
      </ConfigurationEndpoint>
      <CacheNodeType>cache.m1.large</CacheNodeType>
      <Engine>memcached</Engine>
      <PendingModifiedValues>
        <NumCacheNodes>5</NumCacheNodes>
      </PendingModifiedValues>
      <PreferredAvailabilityZone>us-west-2b</PreferredAvailabilityZone>
      <CacheClusterCreateTime>2015-02-02T23:45:20.937Z</CacheClusterCreateTime>
      <EngineVersion>1.4.5</EngineVersion>
      <AutoMinorVersionUpgrade>true</AutoMinorVersionUpgrade>
      <PreferredMaintenanceWindow>fri:04:30-fri:05:30</PreferredMaintenanceWindow>
      <CacheSecurityGroups>
        <CacheSecurityGroup>
          <CacheSecurityGroupName>default</CacheSecurityGroupName>
          <Status>active</Status>
        </CacheSecurityGroup>
      </CacheSecurityGroups>
      <NumCacheNodes>3</NumCacheNodes>
    </CacheCluster>
  </ModifyCacheClusterResult>
  <ResponseMetadata>
    <RequestId>d5786c6d-b7fe-11e0-9326-b7275b9d4a6c</RequestId>
  </ResponseMetadata>
</ModifyCacheClusterResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ModifyCacheParameterGroup

Modifies the parameters of a cache parameter group. You can modify up to 20 parameters in a single request by submitting a list parameter name and value pairs.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheParameterGroupName

The name of the cache parameter group to modify.

Type: String

Required: Yes

ParameterNameValues.ParameterNameValue.N

An array of parameter names and values for the parameter update. You must supply at least one parameter name and value; subsequent arguments are optional. A maximum of 20 parameters may be modified per request.

Type: Array of [ParameterNameValue](#) (p. 284) objects

Required: Yes

Response Elements

The following element is returned by the service.

CacheParameterGroupName

The name of the cache parameter group.

Type: String

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

CacheParameterGroupNotFound

The requested cache parameter group name does not refer to an existing cache parameter group.

HTTP Status Code: 404

InvalidCacheParameterGroupState

The current state of the cache parameter group does not allow the requested operation to occur.

HTTP Status Code: 400

InvalidGlobalReplicationGroupState

The Global datastore is not available or in primary-only state.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

Examples

ModifyCacheParameterGroup

This example illustrates one usage of ModifyCacheParameterGroup.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/
?Action=ModifyCacheParameterGroup
&CacheParameterGroupName=mycacheparametergroup
&ParameterNameValues.ParameterNameValue.1.ParameterName=chunk_size_growth_factor
&ParameterNameValues.ParameterNameValue.1.ParameterValue=1.02
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20150202T192317Z
&X-Amz-Credential=<credential>
```

Sample Response

```
<ModifyCacheParameterGroupResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <ModifyCacheParameterGroupResult>
    <CacheParameterGroupName>mycacheparametergroup</CacheParameterGroupName>
  </ModifyCacheParameterGroupResult>
  <ResponseMetadata>
    <RequestId>fcedeef2-b7ff-11e0-9326-b7275b9d4a6c</RequestId>
  </ResponseMetadata>
</ModifyCacheParameterGroupResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)

- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ModifyCacheSubnetGroup

Modifies an existing cache subnet group.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

CacheSubnetGroupName

The name for the cache subnet group. This value is stored as a lowercase string.

Constraints: Must contain no more than 255 alphanumeric characters or hyphens.

Example: mysubnetgroup

Type: String

Required: Yes

CacheSubnetGroupDescription

A description of the cache subnet group.

Type: String

Required: No

SubnetIds.SubnetIdentifier.N

The EC2 subnet IDs for the cache subnet group.

Type: Array of strings

Required: No

Response Elements

The following element is returned by the service.

CacheSubnetGroup

Represents the output of one of the following operations:

- `CreateCacheSubnetGroup`
- `ModifyCacheSubnetGroup`

Type: [CacheSubnetGroup \(p. 246\)](#) object

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

CacheSubnetGroupNotFoundFault

The requested cache subnet group name does not refer to an existing cache subnet group.

HTTP Status Code: 400

CacheSubnetQuotaExceededFault

The request cannot be processed because it would exceed the allowed number of subnets in a cache subnet group.

HTTP Status Code: 400

InvalidSubnet

An invalid subnet identifier was specified.

HTTP Status Code: 400

SubnetInUse

The requested subnet is being used by another cache subnet group.

HTTP Status Code: 400

SubnetNotAllowedFault

At least one subnet ID does not match the other subnet IDs. This mismatch typically occurs when a user sets one subnet ID to a regional Availability Zone and a different one to an outpost. Or when a user sets the subnet ID to an Outpost when not subscribed on this service.

HTTP Status Code: 400

Examples

ModifyCacheSubnetGroup

This example illustrates one usage of ModifyCacheSubnetGroup.

Sample Request

```
https://elasticache.amazonaws.com/
?Action=ModifyCacheSubnetGroup
&CacheSubnetGroupName=myCachesubnetgroup
&CacheSubnetGroupDescription=My%20modified%20CacheSubnetGroup
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20150202T192317Z
&X-Amz-Credential=<credential>
```

Sample Response

```
<ModifyCacheSubnetGroupResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <ModifyCacheSubnetGroupResult>
    <CacheSubnetGroup>
      <VpcId>990524496922</VpcId>
      <CacheSubnetGroupDescription>My modified CacheSubnetGroup</
CacheSubnetGroupDescription>
      <CacheSubnetGroupName>myCachesubnetgroup</CacheSubnetGroupName>
      <Subnets>
        <Subnet>
          <SubnetStatus>Active</SubnetStatus>
          <SubnetIdentifier>subnet-7c5b4115</SubnetIdentifier>
          <SubnetAvailabilityZone>
```

```
        <Name>us-west-2c</Name>
      </SubnetAvailabilityZone>
    </Subnet>
    <Subnet>
      <SubnetStatus>Active</SubnetStatus>
      <SubnetIdentifier>subnet-7b5b4112</SubnetIdentifier>
      <SubnetAvailabilityZone>
        <Name>us-west-2b</Name>
      </SubnetAvailabilityZone>
    </Subnet>
    <Subnet>
      <SubnetStatus>Active</SubnetStatus>
      <SubnetIdentifier>subnet-3ea6bd57</SubnetIdentifier>
      <SubnetAvailabilityZone>
        <Name>us-west-2c</Name>
      </SubnetAvailabilityZone>
    </Subnet>
  </Subnets>
</CacheSubnetGroup>
</ModifyCacheSubnetGroupResult>
<ResponseMetadata>
  <RequestId>ed662948-a57b-11df-9e38-7ffab86c801f</RequestId>
</ResponseMetadata>
</ModifyCacheSubnetGroupResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ModifyGlobalReplicationGroup

Modifies the settings for a Global datastore.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

ApplyImmediately

This parameter causes the modifications in this request and any pending modifications to be applied, asynchronously and as soon as possible. Modifications to Global Replication Groups cannot be requested to be applied in PreferredMaintenanceWindow.

Type: Boolean

Required: Yes

GlobalReplicationGroupId

The name of the Global datastore

Type: String

Required: Yes

AutomaticFailoverEnabled

Determines whether a read replica is automatically promoted to read/write primary if the existing primary encounters a failure.

Type: Boolean

Required: No

CacheNodeType

A valid cache node type that you want to scale this Global datastore to.

Type: String

Required: No

CacheParameterGroupName

The name of the cache parameter group to use with the Global datastore. It must be compatible with the major engine version used by the Global datastore.

Type: String

Required: No

EngineVersion

The upgraded version of the cache engine to be run on the clusters in the Global datastore.

Type: String

Required: No

GlobalReplicationGroupDescription

A description of the Global datastore

Type: String

Required: No

Response Elements

The following element is returned by the service.

GlobalReplicationGroup

Consists of a primary cluster that accepts writes and an associated secondary cluster that resides in a different Amazon region. The secondary cluster accepts only reads. The primary cluster automatically replicates updates to the secondary cluster.

- The **GlobalReplicationGroupIdSuffix** represents the name of the Global datastore, which is what you use to associate a secondary cluster.

Type: [GlobalReplicationGroup](#) (p. 259) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

GlobalReplicationGroupNotFoundFault

The Global datastore does not exist

HTTP Status Code: 404

InvalidGlobalReplicationGroupState

The Global datastore is not available or in primary-only state.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ModifyReplicationGroup

Modifies the settings for a replication group.

- [Scaling for Amazon ElastiCache for Redis \(cluster mode enabled\)](#) in the ElastiCache User Guide
- [ModifyReplicationGroupShardConfiguration](#) in the ElastiCache API Reference

Note

This operation is valid for Redis only.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

ReplicationGroupId

The identifier of the replication group to modify.

Type: String

Required: Yes

ApplyImmediately

If `true`, this parameter causes the modifications in this request and any pending modifications to be applied, asynchronously and as soon as possible, regardless of the `PreferredMaintenanceWindow` setting for the replication group.

If `false`, changes to the nodes in the replication group are applied on the next maintenance reboot, or the next failure reboot, whichever occurs first.

Valid values: `true` | `false`

Default: `false`

Type: Boolean

Required: No

AuthToken

Reserved parameter. The password used to access a password protected server. This parameter must be specified with the `auth-token-update-strategy` parameter. Password constraints:

- Must be only printable ASCII characters
- Must be at least 16 characters and no more than 128 characters in length
- Cannot contain any of the following characters: `'/'`, `'\"'`, or `'@'`, `'%'`

For more information, see AUTH password at [AUTH](#).

Type: String

Required: No

AuthTokenUpdateStrategy

Specifies the strategy to use to update the AUTH token. This parameter must be specified with the `auth-token` parameter. Possible values:

- Rotate
- Set

For more information, see [Authenticating Users with Redis AUTH](#)

Type: String

Valid Values: SET | ROTATE | DELETE

Required: No

AutomaticFailoverEnabled

Determines whether a read replica is automatically promoted to read/write primary if the existing primary encounters a failure.

Valid values: true | false

Type: Boolean

Required: No

AutoMinorVersionUpgrade

This parameter is currently disabled.

Type: Boolean

Required: No

CacheNodeType

A valid cache node type that you want to scale this replication group to.

Type: String

Required: No

CacheParameterGroupName

The name of the cache parameter group to apply to all of the clusters in this replication group. This change is asynchronously applied as soon as possible for parameters when the `ApplyImmediately` parameter is specified as `true` for this request.

Type: String

Required: No

CacheSecurityGroupNames.CacheSecurityGroupName.N

A list of cache security group names to authorize for the clusters in this replication group. This change is asynchronously applied as soon as possible.

This parameter can be used only with replication group containing clusters running outside of an Amazon Virtual Private Cloud (Amazon VPC).

Constraints: Must contain no more than 255 alphanumeric characters. Must not be `Default`.

Type: Array of strings

Required: No

EngineVersion

The upgraded version of the cache engine to be run on the clusters in the replication group.

Important: You can upgrade to a newer engine version (see [Selecting a Cache Engine and Version](#)), but you cannot downgrade to an earlier engine version. If you want to use an earlier engine version, you must delete the existing replication group and create it anew with the earlier engine version.

Type: String

Required: No

LogDeliveryConfigurations.LogDeliveryConfigurationRequest.N

Specifies the destination, format and type of the logs.

Type: Array of [LogDeliveryConfigurationRequest](#) (p. 268) objects

Required: No

MultiAZEnabled

A flag to indicate MultiAZ is enabled.

Type: Boolean

Required: No

NodeId

Deprecated. This parameter is not used.

Type: String

Required: No

NotificationTopicArn

The Amazon Resource Name (ARN) of the Amazon SNS topic to which notifications are sent.

Note

The Amazon SNS topic owner must be same as the replication group owner.

Type: String

Required: No

NotificationTopicStatus

The status of the Amazon SNS notification topic for the replication group. Notifications are sent only if the status is `active`.

Valid values: `active` | `inactive`

Type: String

Required: No

PreferredMaintenanceWindow

Specifies the weekly time range during which maintenance on the cluster is performed. It is specified as a range in the format `ddd:hh24:mi-ddd:hh24:mi` (24H Clock UTC). The minimum maintenance window is a 60 minute period.

Valid values for `ddd` are:

- `sun`
- `mon`
- `tue`

- wed
- thu
- fri
- sat

Example: sun:23:00-mon:01:30

Type: String

Required: No

PrimaryClusterId

For replication groups with a single primary, if this parameter is specified, ElastiCache promotes the specified cluster in the specified replication group to the primary role. The nodes of all other clusters in the replication group are read replicas.

Type: String

Required: No

RemoveUserGroups

Removes the user group associated with this replication group.

Type: Boolean

Required: No

ReplicationGroupDescription

A description for the replication group. Maximum length is 255 characters.

Type: String

Required: No

SecurityGroupIds.SecurityGroupId.N

Specifies the VPC Security Groups associated with the clusters in the replication group.

This parameter can be used only with replication group containing clusters running in an Amazon Virtual Private Cloud (Amazon VPC).

Type: Array of strings

Required: No

SnapshotRetentionLimit

The number of days for which ElastiCache retains automatic node group (shard) snapshots before deleting them. For example, if you set `SnapshotRetentionLimit` to 5, a snapshot that was taken today is retained for 5 days before being deleted.

Important If the value of `SnapshotRetentionLimit` is set to zero (0), backups are turned off.

Type: Integer

Required: No

SnapshottingClusterId

The cluster ID that is used as the daily snapshot source for the replication group. This parameter cannot be set for Redis (cluster mode enabled) replication groups.

Type: String

Required: No

SnapshotWindow

The daily time range (in UTC) during which ElastiCache begins taking a daily snapshot of the node group (shard) specified by `SnapshottingClusterId`.

Example: 05:00-09:00

If you do not specify this parameter, ElastiCache automatically chooses an appropriate time range.

Type: String

Required: No

UserGroupIdsToAdd.member.N

The ID of the user group you are associating with the replication group.

Type: Array of strings

Length Constraints: Minimum length of 1.

Pattern: `[a-zA-Z][a-zA-Z0-9\-*]`

Required: No

UserGroupIdsToRemove.member.N

The ID of the user group to disassociate from the replication group, meaning the users in the group no longer can access the replication group.

Type: Array of strings

Length Constraints: Minimum length of 1.

Pattern: `[a-zA-Z][a-zA-Z0-9\-*]`

Required: No

Response Elements

The following element is returned by the service.

ReplicationGroup

Contains all of the attributes of a specific Redis replication group.

Type: [ReplicationGroup](#) (p. 291) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

CacheClusterNotFound

The requested cluster ID does not refer to an existing cluster.

HTTP Status Code: 404

CacheParameterGroupNotFound

The requested cache parameter group name does not refer to an existing cache parameter group.

HTTP Status Code: 404

CacheSecurityGroupNotFound

The requested cache security group name does not refer to an existing cache security group.

HTTP Status Code: 404

InsufficientCacheClusterCapacity

The requested cache node type is not available in the specified Availability Zone. For more information, see [InsufficientCacheClusterCapacity](#) in the ElastiCache User Guide.

HTTP Status Code: 400

InvalidCacheClusterState

The requested cluster is not in the `available` state.

HTTP Status Code: 400

InvalidCacheSecurityGroupState

The current state of the cache security group does not allow deletion.

HTTP Status Code: 400

InvalidKMSKeyFault

The KMS key supplied is not valid.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidReplicationGroupState

The requested replication group is not in the `available` state.

HTTP Status Code: 400

InvalidUserGroupState

The user group is not in an active state.

HTTP Status Code: 400

InvalidVPCNetworkStateFault

The VPC network is in an invalid state.

HTTP Status Code: 400

NodeQuotaForClusterExceeded

The request cannot be processed because it would exceed the allowed number of cache nodes in a single cluster.

HTTP Status Code: 400

NodeQuotaForCustomerExceeded

The request cannot be processed because it would exceed the allowed number of cache nodes per customer.

HTTP Status Code: 400

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

UserGroupNotFound

The user group was not found or does not exist

HTTP Status Code: 404

Examples

ModifyReplicationGroup

This example illustrates one usage of ModifyReplicationGroup.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/
?Action=ModifyReplicationGroup
&ApplyImmediately=false
&ReplicationGroupId=my-repgroup
&PrimaryClusterId=my-replica-1
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20150202T192317Z
&X-Amz-Credential=<credential>
```

Sample Response

```
<ModifyReplicationGroupResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <ModifyReplicationGroupResult>
    <ReplicationGroup>
      <SnapshottingClusterId>my-redis-primary</SnapshottingClusterId>
      <MemberClusters>
        <ClusterId>my-redis-primary</ClusterId>
        <ClusterId>my-replica-1</ClusterId>
      </MemberClusters>
      <NodeGroups>
        <NodeGroup>
          <NodeGroupId>0001</NodeGroupId>
          <PrimaryEndpoint>
            <Port>6379</Port>
```

```

        <Address>my-repgroup.q68zge.ng.0001.use1devo.elmo-dev.amazonaws.com</
Address>
        </PrimaryEndpoint>
        <Status>available</Status>
        <NodeGroupMembers>
            <NodeGroupMember>
                <CacheClusterId>my-redis-primary</CacheClusterId>
                <ReadEndpoint>
                    <Port>6379</Port>
                    <Address>my-redis-primary.q68zge.0001.use1devo.elmo-
dev.amazonaws.com</Address>
                </ReadEndpoint>
                <PreferredAvailabilityZone>us-west-2c</PreferredAvailabilityZone>
                <CacheNodeId>0001</CacheNodeId>
                <CurrentRole>primary</CurrentRole>
            </NodeGroupMember>
            <NodeGroupMember>
                <CacheClusterId>my-replica-1</CacheClusterId>
                <ReadEndpoint>
                    <Port>6379</Port>
                    <Address>my-replica-1.q68zge.0001.use1devo.elmo-dev.amazonaws.com</
Address>
                </ReadEndpoint>
                <PreferredAvailabilityZone>us-west-2b</PreferredAvailabilityZone>
                <CacheNodeId>0001</CacheNodeId>
                <CurrentRole>replica</CurrentRole>
            </NodeGroupMember>
        </NodeGroupMembers>
    </NodeGroup>
</NodeGroups>
    <ReplicationGroupId>my-repgroup</ReplicationGroupId>
    <Status>available</Status>
    <PendingModifiedValues>
        <PrimaryClusterId>my-replica-1</PrimaryClusterId>
    </PendingModifiedValues>
    <Description>My replication group</Description>
</ReplicationGroup>
</ModifyReplicationGroupResult>
<ResponseMetadata>
    <RequestId>6fd0aad6-b9d7-11e3-8a16-7978bb24ffdf</RequestId>
</ResponseMetadata>
</ModifyReplicationGroupResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ModifyReplicationGroupShardConfiguration

Modifies a replication group's shards (node groups) by allowing you to add shards, remove shards, or rebalance the keyspace among existing shards.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

ApplyImmediately

Indicates that the shard reconfiguration process begins immediately. At present, the only permitted value for this parameter is `true`.

Value: `true`

Type: Boolean

Required: Yes

NodeGroupCount

The number of node groups (shards) that results from the modification of the shard configuration.

Type: Integer

Required: Yes

ReplicationGroupId

The name of the Redis (cluster mode enabled) cluster (replication group) on which the shards are to be configured.

Type: String

Required: Yes

NodeGroupsToRemove.NodeGroupToRemove.N

If the value of `NodeGroupCount` is less than the current number of node groups (shards), then either `NodeGroupsToRemove` or `NodeGroupsToRetain` is required. `NodeGroupsToRemove` is a list of `NodeGroupIds` to remove from the cluster.

ElastiCache for Redis will attempt to remove all node groups listed by `NodeGroupsToRemove` from the cluster.

Type: Array of strings

Length Constraints: Minimum length of 1. Maximum length of 4.

Pattern: `\d+`

Required: No

NodeGroupsToRetain.NodeGroupToRetain.N

If the value of `NodeGroupCount` is less than the current number of node groups (shards), then either `NodeGroupsToRemove` or `NodeGroupsToRetain` is required. `NodeGroupsToRetain` is a list of `NodeGroupIds` to retain in the cluster.

ElastiCache for Redis will attempt to remove all node groups except those listed by `NodeGroupsToRetain` from the cluster.

Type: Array of strings

Length Constraints: Minimum length of 1. Maximum length of 4.

Pattern: `\d+`

Required: No

ReshardingConfiguration.ReshardingConfiguration.N

Specifies the preferred availability zones for each node group in the cluster. If the value of `NodeGroupCount` is greater than the current number of node groups (shards), you can use this parameter to specify the preferred availability zones of the cluster's shards. If you omit this parameter ElastiCache selects availability zones for you.

You can specify this parameter only if the value of `NodeGroupCount` is greater than the current number of node groups (shards).

Type: Array of [ReshardingConfiguration](#) (p. 305) objects

Required: No

Response Elements

The following element is returned by the service.

ReplicationGroup

Contains all of the attributes of a specific Redis replication group.

Type: [ReplicationGroup](#) (p. 291) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

InsufficientCacheClusterCapacity

The requested cache node type is not available in the specified Availability Zone. For more information, see [InsufficientCacheClusterCapacity](#) in the ElastiCache User Guide.

HTTP Status Code: 400

InvalidCacheClusterState

The requested cluster is not in the available state.

HTTP Status Code: 400

InvalidKMSKeyFault

The KMS key supplied is not valid.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidReplicationGroupState

The requested replication group is not in the available state.

HTTP Status Code: 400

InvalidVPCNetworkStateFault

The VPC network is in an invalid state.

HTTP Status Code: 400

NodeGroupsPerReplicationGroupQuotaExceeded

The request cannot be processed because it would exceed the maximum allowed number of node groups (shards) in a single replication group. The default maximum is 90

HTTP Status Code: 400

NodeQuotaForCustomerExceeded

The request cannot be processed because it would exceed the allowed number of cache nodes per customer.

HTTP Status Code: 400

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

Examples

Add shards

The following example adds shards to the replication group `my-cluster` so that at the completion of the call there are 4 shards. The availability zones for the nodes in the shards are specified by the parameter `ReshardingConfiguration.ReshardingConfiguration.N.PreferredAvailabilityZones.AvailabilityZone.N`. If there are already 4 shards in this replication group, the call fails.

```
https://elasticache.us-east-2.amazonaws.com/
?Action=ModifyReplicationGroupShardConfiguration
  &ApplyImmediately=true
  &NodeGroupCount=4
  &ReplicationGroupId=my-cluster

  &ReshardingConfiguration.ReshardingConfiguration.1.PreferredAvailabilityZones.AvailabilityZone.1=us-east-2a

  &ReshardingConfiguration.ReshardingConfiguration.1.PreferredAvailabilityZones.AvailabilityZone.2=us-east-2c
```

```
&ReshardingConfiguration.ReshardingConfiguration.2.PreferredAvailabilityZones.AvailabilityZone.1=us-east-2b
&ReshardingConfiguration.ReshardingConfiguration.2.PreferredAvailabilityZones.AvailabilityZone.2=us-east-2a
&ReshardingConfiguration.ReshardingConfiguration.3.PreferredAvailabilityZones.AvailabilityZone.1=us-east-2c
&ReshardingConfiguration.ReshardingConfiguration.3.PreferredAvailabilityZones.AvailabilityZone.2=us-east-2d
&ReshardingConfiguration.ReshardingConfiguration.4.PreferredAvailabilityZones.AvailabilityZone.1=us-east-2d
&ReshardingConfiguration.ReshardingConfiguration.4.PreferredAvailabilityZones.AvailabilityZone.2=us-east-2c
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20171002T192317Z
&X-Amz-Credential=<credential>
```

Remove shards

The following example removes two shards from the replication group `my-cluster`, leaving 2 shards. When removing shards, the parameter `NodeGroupsToRemove.NodeGroupToRemove` is required.

You cannot remove the last shard leaving zero shards.

```
https://elasticache.us-east-2.amazonaws.com/
?Action=ModifyReplicationGroupShardConfiguration
&ApplyImmediately=true
&NodeGroupCount=2
&ReplicationGroupId=my-cluster
&NodeGroupsToRemove.NodeGroupToRemove.1=0002
&NodeGroupsToRemove.NodeGroupToRemove.2=0003
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20171002T192317Z
&X-Amz-Credential=<credential>
```

Rebalance shards

The following rebalances the keyspaces among the existing shards in the replication group `my-cluster`. The value specified by `NodeGroupCount` must be the existing number of shards. If the keyspaces are already balanced the call fails.

```
https://elasticache.us-east-2.amazonaws.com/
?Action=ModifyReplicationGroupShardConfiguration
&ApplyImmediately=true
&NodeGroupCount=4
&ReplicationGroupId=my-cluster
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20171002T192317Z
&X-Amz-Credential=<credential>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ModifyUser

Changes user password(s) and/or access string.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

UserId

The ID of the user.

Type: String

Length Constraints: Minimum length of 1.

Pattern: `[a-zA-Z][a-zA-Z0-9\-\]*`

Required: Yes

AccessString

Access permissions string used for this user.

Type: String

Pattern: `.*\S.*`

Required: No

AppendAccessString

Adds additional user permissions to the access string.

Type: String

Pattern: `.*\S.*`

Required: No

NoPasswordRequired

Indicates no password is required for the user.

Type: Boolean

Required: No

Passwords.member.N

The passwords belonging to the user. You are allowed up to two.

Type: Array of strings

Array Members: Minimum number of 1 item.

Required: No

Response Elements

The following elements are returned by the service.

AccessString

Access permissions string used for this user.

Type: String

ARN

The Amazon Resource Name (ARN) of the user.

Type: String

Authentication

Denotes whether the user requires a password to authenticate.

Type: [Authentication](#) (p. 222) object

Engine

The current supported value is Redis.

Type: String

Pattern: [a-zA-Z]*

Status

Indicates the user status. Can be "active", "modifying" or "deleting".

Type: String

UserGroupIds.member.N

Returns a list of the user group IDs the user belongs to.

Type: Array of strings

Length Constraints: Minimum length of 1.

Pattern: [a-zA-Z][a-zA-Z0-9\-*]

UserId

The ID of the user.

Type: String

UserName

The username of the user.

Type: String

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidUserState

The user is not in active state.

HTTP Status Code: 400

UserNotFound

The user does not exist or could not be found.

HTTP Status Code: 404

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ModifyUserGroup

Changes the list of users that belong to the user group.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

UserGroupId

The ID of the user group.

Type: String

Required: Yes

UserIdsToAdd.member.N

The list of user IDs to add to the user group.

Type: Array of strings

Array Members: Minimum number of 1 item.

Length Constraints: Minimum length of 1.

Pattern: `[a-zA-Z][a-zA-Z0-9\-*]`

Required: No

UserIdsToRemove.member.N

The list of user IDs to remove from the user group.

Type: Array of strings

Array Members: Minimum number of 1 item.

Length Constraints: Minimum length of 1.

Pattern: `[a-zA-Z][a-zA-Z0-9\-*]`

Required: No

Response Elements

The following elements are returned by the service.

ARN

The Amazon Resource Name (ARN) of the user group.

Type: String

Engine

The current supported value is Redis.

Type: String

Pattern: [a-zA-Z]*

PendingChanges

A list of updates being applied to the user group.

Type: [UserGroupPendingChanges](#) (p. 329) object

ReplicationGroups.member.N

A list of replication groups that the user group can access.

Type: Array of strings

Status

Indicates user group status. Can be "creating", "active", "modifying", "deleting".

Type: String

UserGroupId

The ID of the user group.

Type: String

UserIds.member.N

The list of user IDs that belong to the user group.

Type: Array of strings

Length Constraints: Minimum length of 1.

Pattern: [a-zA-Z][a-zA-Z0-9\-*]

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

DefaultUserRequired

You must add default user to a user group.

HTTP Status Code: 400

DuplicateUserName

A user with this username already exists.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidUserGroupState

The user group is not in an active state.

HTTP Status Code: 400

UserGroupNotFound

The user group was not found or does not exist

HTTP Status Code: 404

UserNotFound

The user does not exist or could not be found.

HTTP Status Code: 404

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

PurchaseReservedCacheNodesOffering

Allows you to purchase a reserved cache node offering. Reserved nodes are not eligible for cancellation and are non-refundable. For more information, see [Managing Costs with Reserved Nodes](#) for Redis or [Managing Costs with Reserved Nodes](#) for Memcached.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

ReservedCacheNodesOfferingId

The ID of the reserved cache node offering to purchase.

Example: 438012d3-4052-4cc7-b2e3-8d3372e0e706

Type: String

Required: Yes

CacheNodeCount

The number of cache node instances to reserve.

Default: 1

Type: Integer

Required: No

ReservedCacheNodeId

A customer-specified identifier to track this reservation.

Note

The Reserved Cache Node ID is an unique customer-specified identifier to track this reservation. If this parameter is not specified, ElastiCache automatically generates an identifier for the reservation.

Example: myreservationID

Type: String

Required: No

Tags.Tag.N

A list of tags to be added to this resource. A tag is a key-value pair. A tag key must be accompanied by a tag value, although null is accepted.

Type: Array of [Tag](#) (p. 319) objects

Required: No

Response Elements

The following element is returned by the service.

ReservedCacheNode

Represents the output of a `PurchaseReservedCacheNodesOffering` operation.

Type: [ReservedCacheNode](#) (p. 298) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

ReservedCacheNodeAlreadyExists

You already have a reservation with the given identifier.

HTTP Status Code: 404

ReservedCacheNodeQuotaExceeded

The request cannot be processed because it would exceed the user's cache node quota.

HTTP Status Code: 400

ReservedCacheNodesOfferingNotFound

The requested cache node offering does not exist.

HTTP Status Code: 404

TagQuotaPerResourceExceeded

The request cannot be processed because it would cause the resource to have more than the allowed number of tags. The maximum number of tags permitted on a resource is 50.

HTTP Status Code: 400

Examples

PurchaseReservedCacheNodesOffering

This example illustrates one usage of `PurchaseReservedCacheNodesOffering`.

Sample Request

```
https://elasticache.amazonaws.com/
?Action=PurchaseReservedCacheNodesOffering
&ReservedCacheNodeId=myreservationID
&ReservedCacheNodesOfferingId=438012d3-4052-4cc7-b2e3-8d3372e0e706
&CacheNodeCount=1
```

```
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20150202T192317Z
&X-Amz-Credential=<credential>
```

Sample Response

```
<PurchaseReservedCacheNodesOfferingResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <PurchaseReservedCacheNodesOfferingResult>
    <ReservedCacheNode>
      <OfferingType>Medium Utilization</OfferingType>
      <RecurringCharges/>
      <ProductDescription>memcached</ProductDescription>
      <ReservedCacheNodesOfferingId> 438012d3-4052-4cc7-b2e3-8d3372e0e706</ReservedCacheNodesOfferingId>
      <State>payment-pending</State>
      <ReservedCacheNodeId>myreservationID</ReservedCacheNodeId>
      <CacheNodeCount>10</CacheNodeCount>
      <StartTime>2015-02-02T23:24:56.577Z</StartTime>
      <Duration>31536000</Duration>
      <FixedPrice>123.0</FixedPrice>
      <UsagePrice>0.123</UsagePrice>
      <CacheNodeType>cache.m1.small</CacheNodeType>
    </ReservedCacheNode>
  </PurchaseReservedCacheNodesOfferingResult>
  <ResponseMetadata>
    <RequestId>7f099901-29cf-11e1-bd06-6fe008f046c3</RequestId>
  </ResponseMetadata>
</PurchaseReservedCacheNodesOfferingResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

RebalanceSlotsInGlobalReplicationGroup

Redistribute slots to ensure uniform distribution across existing shards in the cluster.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

ApplyImmediately

If `True`, redistribution is applied immediately.

Type: Boolean

Required: Yes

GlobalReplicationGroupId

The name of the Global datastore

Type: String

Required: Yes

Response Elements

The following element is returned by the service.

GlobalReplicationGroup

Consists of a primary cluster that accepts writes and an associated secondary cluster that resides in a different Amazon region. The secondary cluster accepts only reads. The primary cluster automatically replicates updates to the secondary cluster.

- The **GlobalReplicationGroupIdSuffix** represents the name of the Global datastore, which is what you use to associate a secondary cluster.

Type: [GlobalReplicationGroup](#) (p. 259) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

GlobalReplicationGroupNotFoundFault

The Global datastore does not exist

HTTP Status Code: 404

InvalidGlobalReplicationGroupState

The Global datastore is not available or in primary-only state.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

RebootCacheCluster

Reboots some, or all, of the cache nodes within a provisioned cluster. This operation applies any modified cache parameter groups to the cluster. The reboot operation takes place as soon as possible, and results in a momentary outage to the cluster. During the reboot, the cluster status is set to REBOOTING.

The reboot causes the contents of the cache (for each cache node being rebooted) to be lost.

When the reboot is complete, a cluster event is created.

Rebooting a cluster is currently supported on Memcached and Redis (cluster mode disabled) clusters. Rebooting is not supported on Redis (cluster mode enabled) clusters.

If you make changes to parameters that require a Redis (cluster mode enabled) cluster reboot for the changes to be applied, see [Rebooting a Cluster](#) for an alternate process.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CacheClusterId

The cluster identifier. This parameter is stored as a lowercase string.

Type: String

Required: Yes

CacheNodeIdsToReboot.CacheNodeId.N

A list of cache node IDs to reboot. A node ID is a numeric identifier (0001, 0002, etc.). To reboot an entire cluster, specify all of the cache node IDs.

Type: Array of strings

Required: Yes

Response Elements

The following element is returned by the service.

CacheCluster

Contains all of the attributes of a specific cluster.

Type: [CacheCluster](#) (p. 224) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

CacheClusterNotFound

The requested cluster ID does not refer to an existing cluster.

HTTP Status Code: 404

InvalidCacheClusterState

The requested cluster is not in the available state.

HTTP Status Code: 400

Examples

RebootCacheCluster

This example illustrates one usage of RebootCacheCluster.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/
?Action=RebootCacheCluster
&CacheClusterId=mycache
&CacheNodeIdsToReboot.CacheNodeId.1=0001
&CacheNodeIdsToReboot.CacheNodeId.2=0002
&CacheNodeIdsToReboot.CacheNodeId.3=0003
&Version=2015-02-02
&SignatureVersion=4
&SignatureMethod=HmacSHA256
&Timestamp=20150202T192317Z
&X-Amz-Credential=<credential>
```

Sample Response

```
<RebootCacheClusterResponse xmlns="http://elasticache.amazonaws.com/doc/2015-02-02/">
  <RebootCacheClusterResult>
    <CacheCluster>
      <CacheClusterStatus>rebooting cluster nodes</CacheClusterStatus>
      <CacheParameterGroup>
        <CacheParameterGroupName>default.memcached1.4</CacheParameterGroupName>
        <ParameterApplyStatus>in-sync</ParameterApplyStatus>
        <CacheNodeIdsToReboot />
      </CacheParameterGroup>
      <CacheClusterId>mycache</CacheClusterId>
      <ConfigurationEndpoint>
        <Port>11211</Port>
        <Address>mycache.q68zge.cfg.use1devo.elmo-dev.amazonaws.com</Address>
      </ConfigurationEndpoint>
      <CacheNodeType>cache.m1.small</CacheNodeType>
      <Engine>memcached</Engine>
      <PendingModifiedValues />
      <PreferredAvailabilityZone>us-west-2b</PreferredAvailabilityZone>
      <CacheClusterCreateTime>2015-02-02T19:04:12.812Z</CacheClusterCreateTime>
      <EngineVersion>1.4.17</EngineVersion>
      <AutoMinorVersionUpgrade>true</AutoMinorVersionUpgrade>
      <PreferredMaintenanceWindow>wed:09:00-wed:10:00</PreferredMaintenanceWindow>
      <ClientDownloadLandingPage>https://console.aws.amazon.com/elasticache/home#client-
download:</ClientDownloadLandingPage>
      <CacheSecurityGroups>
        <CacheSecurityGroup>
          <CacheSecurityGroupName>default</CacheSecurityGroupName>
          <Status>active</Status>
        </CacheSecurityGroup>
      </CacheSecurityGroups>
      <NumCacheNodes>3</NumCacheNodes>
    </CacheCluster>
  </RebootCacheClusterResult>
</RebootCacheClusterResponse>
```



```
</RebootCacheClusterResult>
<ResponseMetadata>
  <RequestId>cf7e6fc4-b9d1-11e3-8a16-7978bb24ffdf</RequestId>
</ResponseMetadata>
</RebootCacheClusterResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

RemoveTagsFromResource

Removes the tags identified by the `TagKeys` list from the named resource. A tag is a key-value pair where the key and value are case-sensitive. You can use tags to categorize and track all your ElastiCache resources, with the exception of global replication group. When you add or remove tags on replication groups, those actions will be replicated to all nodes in the replication group. For more information, see [Resource-level permissions](#).

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

ResourceName

The Amazon Resource Name (ARN) of the resource from which you want the tags removed, for example `arn:aws:elasticache:us-west-2:0123456789:cluster:myCluster` or `arn:aws:elasticache:us-west-2:0123456789:snapshot:mySnapshot`.

For more information about ARNs, see [Amazon Resource Names \(ARNs\) and Amazon Service Namespaces](#).

Type: String

Required: Yes

TagKeys.member.N

A list of `TagKeys` identifying the tags you want removed from the named resource.

Type: Array of strings

Required: Yes

Response Elements

The following element is returned by the service.

TagList.Tag.N

A list of tags as key-value pairs.

Type: Array of [Tag](#) (p. 319) objects

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

CacheClusterNotFound

The requested cluster ID does not refer to an existing cluster.

HTTP Status Code: 404

CacheParameterGroupNotFound

The requested cache parameter group name does not refer to an existing cache parameter group.

HTTP Status Code: 404

CacheSecurityGroupNotFound

The requested cache security group name does not refer to an existing cache security group.

HTTP Status Code: 404

CacheSubnetGroupNotFoundFault

The requested cache subnet group name does not refer to an existing cache subnet group.

HTTP Status Code: 400

InvalidARN

The requested Amazon Resource Name (ARN) does not refer to an existing resource.

HTTP Status Code: 400

InvalidReplicationGroupState

The requested replication group is not in the available state.

HTTP Status Code: 400

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

ReservedCacheNodeNotFound

The requested reserved cache node was not found.

HTTP Status Code: 404

SnapshotNotFoundFault

The requested snapshot name does not refer to an existing snapshot.

HTTP Status Code: 404

TagNotFound

The requested tag was not found on this resource.

HTTP Status Code: 404

UserGroupNotFound

The user group was not found or does not exist

HTTP Status Code: 404

UserNotFound

The user does not exist or could not be found.

HTTP Status Code: 404

Examples

RemoveTagsFromResource

This example illustrates one usage of RemoveTagsFromResource.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=RemoveTagsFromResource  
&ResourceName=arn:aws:elasticache:us-west-2:0123456789:cluster:myCluster  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&TagKeys.TagKey.1=service  
&TagKeys.TagKey.2=organization  
&Version=2015-02-02  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ResetCacheParameterGroup

Modifies the parameters of a cache parameter group to the engine or system default value. You can reset specific parameters by submitting a list of parameter names. To reset the entire cache parameter group, specify the `ResetAllParameters` and `CacheParameterGroupName` parameters.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

CacheParameterGroupName

The name of the cache parameter group to reset.

Type: String

Required: Yes

ParameterNameValues.ParameterNameValue.N

An array of parameter names to reset to their default values. If `ResetAllParameters` is `true`, do not use `ParameterNameValues`. If `ResetAllParameters` is `false`, you must specify the name of at least one parameter to reset.

Type: Array of [ParameterNameValue \(p. 284\)](#) objects

Required: No

ResetAllParameters

If `true`, all parameters in the cache parameter group are reset to their default values. If `false`, only the parameters listed by `ParameterNameValues` are reset to their default values.

Valid values: `true` | `false`

Type: Boolean

Required: No

Response Elements

The following element is returned by the service.

CacheParameterGroupName

The name of the cache parameter group.

Type: String

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

CacheParameterGroupNotFound

The requested cache parameter group name does not refer to an existing cache parameter group.

HTTP Status Code: 404

InvalidCacheParameterGroupState

The current state of the cache parameter group does not allow the requested operation to occur.

HTTP Status Code: 400

InvalidGlobalReplicationGroupState

The Global datastore is not available or in primary-only state.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

Examples

ResetCacheParameterGroup

This example illustrates one usage of ResetCacheParameterGroup.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=ResetCacheParameterGroup  
&ResetAllParameters=true  
&CacheParameterGroupName=mycacheparametergroup1  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<ResetCacheParameterGroupResponse xmlns="http://elasticache.amazonaws.com/  
doc/2015-02-02/">  
  <ResetCacheParameterGroupResult>  
    <CacheParameterGroupName>mycacheparametergroup1</CacheParameterGroupName>  
  </ResetCacheParameterGroupResult>  
  <ResponseMetadata>  
    <RequestId>cb7cc855-b9d2-11e3-8a16-7978bb24ffdf</RequestId>  
  </ResponseMetadata>  
</ResetCacheParameterGroupResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

RevokeCacheSecurityGroupIngress

Revokes ingress from a cache security group. Use this operation to disallow access from an Amazon EC2 security group that had been previously authorized.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 331\)](#).

CacheSecurityGroupName

The name of the cache security group to revoke ingress from.

Type: String

Required: Yes

EC2SecurityGroupName

The name of the Amazon EC2 security group to revoke access from.

Type: String

Required: Yes

EC2SecurityGroupOwnerId

The Amazon account number of the Amazon EC2 security group owner. Note that this is not the same thing as an Amazon access key ID - you must provide a valid Amazon account number for this parameter.

Type: String

Required: Yes

Response Elements

The following element is returned by the service.

CacheSecurityGroup

Represents the output of one of the following operations:

- `AuthorizeCacheSecurityGroupIngress`
- `CreateCacheSecurityGroup`
- `RevokeCacheSecurityGroupIngress`

Type: [CacheSecurityGroup \(p. 243\)](#) object

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 333\)](#).

AuthorizationNotFound

The specified Amazon EC2 security group is not authorized for the specified cache security group.

HTTP Status Code: 404

CacheSecurityGroupNotFound

The requested cache security group name does not refer to an existing cache security group.

HTTP Status Code: 404

InvalidCacheSecurityGroupState

The current state of the cache security group does not allow deletion.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

Examples

RevokeCacheSecurityGroupIngress

This example illustrates one usage of RevokeCacheSecurityGroupIngress.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=RevokeCacheSecurityGroupIngress  
&EC2SecurityGroupName=default  
&CacheSecurityGroupName=mygroup  
&EC2SecurityGroupOwnerId=1234-5678-1234  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20150202T192317Z  
&X-Amz-Credential=<credential>
```

Sample Response

```
<RevokeCacheSecurityGroupIngressResponse xmlns="http://elasticache.amazonaws.com/  
doc/2015-02-02/">  
  <RevokeCacheSecurityGroupIngressResult>  
    <CacheSecurityGroup>  
      <EC2SecurityGroups>  
        <EC2SecurityGroup>  
          <Status>revoking</Status>  
          <EC2SecurityGroupName>default</EC2SecurityGroupName>  
          <EC2SecurityGroupOwnerId>123456781234</EC2SecurityGroupOwnerId>  
        </EC2SecurityGroup>  
      </EC2SecurityGroups>  
      <CacheSecurityGroupName>mygroup</CacheSecurityGroupName>  
      <OwnerId>123456789012</OwnerId>
```

```
<Description>My security group</Description>
</CacheSecurityGroup>
</RevokeCacheSecurityGroupIngressResult>
<ResponseMetadata>
  <RequestId>02ae3699-3650-11e0-a564-8f11342c56b0</RequestId>
</ResponseMetadata>
</RevokeCacheSecurityGroupIngressResponse>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

StartMigration

Start the migration of data.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

CustomerNodeEndpointList.member.N

List of endpoints from which data should be migrated. For Redis (cluster mode disabled), list should have only one element.

Type: Array of [CustomerNodeEndpoint](#) (p. 251) objects

Required: Yes

ReplicationGroupId

The ID of the replication group to which data should be migrated.

Type: String

Required: Yes

Response Elements

The following element is returned by the service.

ReplicationGroup

Contains all of the attributes of a specific Redis replication group.

Type: [ReplicationGroup](#) (p. 291) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidReplicationGroupState

The requested replication group is not in the available state.

HTTP Status Code: 400

ReplicationGroupAlreadyUnderMigrationFault

The targeted replication group is not available.

HTTP Status Code: 400

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

TestFailover

Represents the input of a `TestFailover` operation which test automatic failover on a specified node group (called shard in the console) in a replication group (called cluster in the console).

Note the following

- A customer can use this operation to test automatic failover on up to 5 shards (called node groups in the ElastiCache API and Amazon CLI) in any rolling 24-hour period.
- If calling this operation on shards in different clusters (called replication groups in the API and CLI), the calls can be made concurrently.
- If calling this operation multiple times on different shards in the same Redis (cluster mode enabled) replication group, the first node replacement must complete before a subsequent call can be made.
- To determine whether the node replacement is complete you can check Events using the Amazon ElastiCache console, the Amazon CLI, or the ElastiCache API. Look for the following automatic failover related events, listed here in order of occurrence:
 1. Replication group message: `Test Failover API called for node group <node-group-id>`
 2. Cache cluster message: `Failover from primary node <primary-node-id> to replica node <node-id> completed`
 3. Replication group message: `Failover from primary node <primary-node-id> to replica node <node-id> completed`
 4. Cache cluster message: `Recovering cache nodes <node-id>`
 5. Cache cluster message: `Finished recovery for cache nodes <node-id>`

For more information see:

- [Viewing ElastiCache Events](#) in the *ElastiCache User Guide*
- [DescribeEvents](#) in the ElastiCache API Reference

Also see, [Testing Multi-AZ](#) in the *ElastiCache User Guide*.

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 331).

NodeGroupId

The name of the node group (called shard in the console) in this replication group on which automatic failover is to be tested. You may test automatic failover on up to 5 node groups in any rolling 24-hour period.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 4.

Pattern: `\d+`

Required: Yes

ReplicationGroupId

The name of the replication group (console: cluster) whose automatic failover is being tested by this operation.

Type: String

Required: Yes

Response Elements

The following element is returned by the service.

ReplicationGroup

Contains all of the attributes of a specific Redis replication group.

Type: [ReplicationGroup](#) (p. 291) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 333).

APICallRateForCustomerExceeded

The customer has exceeded the allowed rate of API calls.

HTTP Status Code: 400

InvalidCacheClusterState

The requested cluster is not in the `available` state.

HTTP Status Code: 400

InvalidKMSKeyFault

The KMS key supplied is not valid.

HTTP Status Code: 400

InvalidParameterCombination

Two or more incompatible parameters were specified.

HTTP Status Code: 400

InvalidParameterValue

The value for a parameter is invalid.

HTTP Status Code: 400

InvalidReplicationGroupState

The requested replication group is not in the `available` state.

HTTP Status Code: 400

NodeGroupNotFoundFault

The node group specified by the `NodeGroupId` parameter could not be found. Please verify that the node group exists and that you spelled the `NodeGroupId` value correctly.

HTTP Status Code: 404

ReplicationGroupNotFoundFault

The specified replication group does not exist.

HTTP Status Code: 404

TestFailoverNotAvailableFault

The `TestFailover` action is not available.

HTTP Status Code: 400

Examples

Example

The following example tests automatic failover on the Redis (cluster mode disabled) replication group (console: cluster) `redis00`. Since there is only one node group in Redis (cluster mode disabled) clusters, the *NodeGroupId* will always be `<cluster-name>-0001`.

Sample Request

```
https://elasticache.us-west-2.amazonaws.com/  
?Action=TestFailover  
&NodeGroupId=redis00-0001  
&ReplicationGroupId=redis00  
&Version=2015-02-02  
&SignatureVersion=4  
&SignatureMethod=HmacSHA256  
&Timestamp=20170418T192317Z  
&X-Amz-Credential=<credential>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

Data Types

The Amazon ElastiCache API contains several data types that various actions use. This section describes each data type in detail.

Note

The order of each element in a data type structure is not guaranteed. Applications should not assume a particular order.

The following data types are supported:

- [Authentication](#) (p. 222)
- [AvailabilityZone](#) (p. 223)
- [CacheCluster](#) (p. 224)
- [CacheEngineVersion](#) (p. 230)
- [CacheNode](#) (p. 232)
- [CacheNodeTypeSpecificParameter](#) (p. 235)
- [CacheNodeTypeSpecificValue](#) (p. 237)
- [CacheNodeUpdateStatus](#) (p. 238)
- [CacheParameterGroup](#) (p. 240)
- [CacheParameterGroupStatus](#) (p. 242)
- [CacheSecurityGroup](#) (p. 243)
- [CacheSecurityGroupMembership](#) (p. 245)
- [CacheSubnetGroup](#) (p. 246)
- [CloudWatchLogsDestinationDetails](#) (p. 248)
- [ConfigureShard](#) (p. 249)
- [CustomerNodeEndpoint](#) (p. 251)
- [DestinationDetails](#) (p. 252)
- [EC2SecurityGroup](#) (p. 253)
- [Endpoint](#) (p. 254)
- [EngineDefaults](#) (p. 255)
- [Event](#) (p. 256)
- [Filter](#) (p. 257)
- [GlobalNodeGroup](#) (p. 258)
- [GlobalReplicationGroup](#) (p. 259)
- [GlobalReplicationGroupInfo](#) (p. 262)
- [GlobalReplicationGroupMember](#) (p. 263)
- [KinesisFirehoseDestinationDetails](#) (p. 265)
- [LogDeliveryConfiguration](#) (p. 266)
- [LogDeliveryConfigurationRequest](#) (p. 268)
- [NodeGroup](#) (p. 270)
- [NodeGroupConfiguration](#) (p. 272)
- [NodeGroupMember](#) (p. 274)
- [NodeGroupMemberUpdateStatus](#) (p. 276)
- [NodeGroupUpdateStatus](#) (p. 278)
- [NodeSnapshot](#) (p. 279)

- [NotificationConfiguration](#) (p. 281)
- [Parameter](#) (p. 282)
- [ParameterNameValue](#) (p. 284)
- [PendingLogDeliveryConfiguration](#) (p. 285)
- [PendingModifiedValues](#) (p. 286)
- [ProcessedUpdateAction](#) (p. 288)
- [RecurringCharge](#) (p. 289)
- [RegionalConfiguration](#) (p. 290)
- [ReplicationGroup](#) (p. 291)
- [ReplicationGroupPendingModifiedValues](#) (p. 296)
- [ReservedCacheNode](#) (p. 298)
- [ReservedCacheNodesOffering](#) (p. 302)
- [ReshardingConfiguration](#) (p. 305)
- [ReshardingStatus](#) (p. 306)
- [SecurityGroupMembership](#) (p. 307)
- [ServiceUpdate](#) (p. 308)
- [SlotMigration](#) (p. 310)
- [Snapshot](#) (p. 311)
- [Subnet](#) (p. 317)
- [SubnetOutpost](#) (p. 318)
- [Tag](#) (p. 319)
- [TimeRangeFilter](#) (p. 320)
- [UnprocessedUpdateAction](#) (p. 321)
- [UpdateAction](#) (p. 322)
- [User](#) (p. 325)
- [UserGroup](#) (p. 327)
- [UserGroupPendingChanges](#) (p. 329)
- [UserGroupsUpdateStatus](#) (p. 330)

Authentication

Indicates whether the user requires a password to authenticate.

Contents

Note

In the following list, the required parameters are described first.

PasswordCount

The number of passwords belonging to the user. The maximum is two.

Type: Integer

Required: No

Type

Indicates whether the user requires a password to authenticate.

Type: String

Valid Values: `password` | `no-password`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

AvailabilityZone

Describes an Availability Zone in which the cluster is launched.

Contents

Note

In the following list, the required parameters are described first.

Name

The name of the Availability Zone.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

CacheCluster

Contains all of the attributes of a specific cluster.

Contents

Note

In the following list, the required parameters are described first.

ARN

The ARN (Amazon Resource Name) of the cache cluster.

Type: String

Required: No

AtRestEncryptionEnabled

A flag that enables encryption at-rest when set to `true`.

You cannot modify the value of `AtRestEncryptionEnabled` after the cluster is created. To enable at-rest encryption on a cluster you must set `AtRestEncryptionEnabled` to `true` when you create a cluster.

Required: Only available when creating a replication group in an Amazon VPC using redis version 3.2.6, 4.x or later.

Default: `false`

Type: Boolean

Required: No

AuthTokenEnabled

A flag that enables using an `AuthToken` (password) when issuing Redis commands.

Default: `false`

Type: Boolean

Required: No

AuthTokenLastModifiedDate

The date the auth token was last modified

Type: Timestamp

Required: No

AutoMinorVersionUpgrade

This parameter is currently disabled.

Type: Boolean

Required: No

CacheClusterCreateTime

The date and time when the cluster was created.

Type: Timestamp

Required: No

CacheClusterId

The user-supplied identifier of the cluster. This identifier is a unique key that identifies a cluster.

Type: String

Required: No

CacheClusterStatus

The current state of this cluster, one of the following values: `available`, `creating`, `deleted`, `deleting`, `incompatible-network`, `modifying`, `rebooting cluster nodes`, `restore-failed`, or `snapshotting`.

Type: String

Required: No

CacheNodes.CacheNode.N

A list of cache nodes that are members of the cluster.

Type: Array of [CacheNode](#) (p. 232) objects

Required: No

CacheNodeType

The name of the compute and memory capacity node type for the cluster.

The following node types are supported by ElastiCache. Generally speaking, the current generation types provide more memory and computational power at lower cost when compared to their equivalent previous generation counterparts.

- General purpose:
 - Current generation:

M6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

`cache.m6g.large`, `cache.m6g.xlarge`, `cache.m6g.2xlarge`, `cache.m6g.4xlarge`,
`cache.m6g.8xlarge`, `cache.m6g.12xlarge`, `cache.m6g.16xlarge`

Note

For region availability, see [Supported Node Types](#)

M5 node types: `cache.m5.large`, `cache.m5.xlarge`, `cache.m5.2xlarge`,
`cache.m5.4xlarge`, `cache.m5.12xlarge`, `cache.m5.24xlarge`

M4 node types: `cache.m4.large`, `cache.m4.xlarge`, `cache.m4.2xlarge`,
`cache.m4.4xlarge`, `cache.m4.10xlarge`

T3 node types: `cache.t3.micro`, `cache.t3.small`, `cache.t3.medium`

T2 node types: `cache.t2.micro`, `cache.t2.small`, `cache.t2.medium`

- Previous generation: (not recommended)

T1 node types: `cache.t1.micro`

M1 node types: `cache.m1.small`, `cache.m1.medium`, `cache.m1.large`,
`cache.m1.xlarge`

M3 node types: `cache.m3.medium`, `cache.m3.large`, `cache.m3.xlarge`, `cache.m3.2xlarge`

- Compute optimized:
 - Previous generation: (not recommended)

C1 node types: `cache.c1.xlarge`

- Memory optimized:
 - Current generation:

R6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

`cache.r6g.large`, `cache.r6g.xlarge`, `cache.r6g.2xlarge`, `cache.r6g.4xlarge`, `cache.r6g.8xlarge`, `cache.r6g.12xlarge`, `cache.r6g.16xlarge`

Note

For region availability, see [Supported Node Types](#)

R5 node types: `cache.r5.large`, `cache.r5.xlarge`, `cache.r5.2xlarge`, `cache.r5.4xlarge`, `cache.r5.12xlarge`, `cache.r5.24xlarge`

R4 node types: `cache.r4.large`, `cache.r4.xlarge`, `cache.r4.2xlarge`, `cache.r4.4xlarge`, `cache.r4.8xlarge`, `cache.r4.16xlarge`

- Previous generation: (not recommended)

M2 node types: `cache.m2.xlarge`, `cache.m2.2xlarge`, `cache.m2.4xlarge`

R3 node types: `cache.r3.large`, `cache.r3.xlarge`, `cache.r3.2xlarge`, `cache.r3.4xlarge`, `cache.r3.8xlarge`

Additional node type info

- All current generation instance types are created in Amazon VPC by default.
- Redis append-only files (AOF) are not supported for T1 or T2 instances.
- Redis Multi-AZ with automatic failover is not supported on T1 instances.
- Redis configuration variables `appendonly` and `appendfsync` are not supported on Redis version 2.8.22 and later.

Type: String

Required: No

CacheParameterGroup

Status of the cache parameter group.

Type: [CacheParameterGroupStatus](#) (p. 242) object

Required: No

CacheSecurityGroups.CacheSecurityGroup.N

A list of cache security group elements, composed of name and status sub-elements.

Type: Array of [CacheSecurityGroupMembership](#) (p. 245) objects

Required: No

CacheSubnetGroupName

The name of the cache subnet group associated with the cluster.

Type: String

Required: No

ClientDownloadLandingPage

The URL of the web page where you can download the latest ElastiCache client library.

Type: String

Required: No

ConfigurationEndpoint

Represents a Memcached cluster endpoint which can be used by an application to connect to any node in the cluster. The configuration endpoint will always have `.cfg` in it.

Example: `mem-3.9dvc4r.cfg.usw2.cache.amazonaws.com:11211`

Type: [Endpoint \(p. 254\)](#) object

Required: No

Engine

The name of the cache engine (`memcached` or `redis`) to be used for this cluster.

Type: String

Required: No

EngineVersion

The version of the cache engine that is used in this cluster.

Type: String

Required: No

LogDeliveryConfigurations.LogDeliveryConfiguration.N

Returns the destination, format and type of the logs.

Type: Array of [LogDeliveryConfiguration \(p. 266\)](#) objects

Required: No

NotificationConfiguration

Describes a notification topic and its status. Notification topics are used for publishing ElastiCache events to subscribers using Amazon Simple Notification Service (SNS).

Type: [NotificationConfiguration \(p. 281\)](#) object

Required: No

NumCacheNodes

The number of cache nodes in the cluster.

For clusters running Redis, this value must be 1. For clusters running Memcached, this value must be between 1 and 40.

Type: Integer

Required: No

PendingModifiedValues

A group of settings that are applied to the cluster in the future, or that are currently being applied.

Type: [PendingModifiedValues](#) (p. 286) object

Required: No

PreferredAvailabilityZone

The name of the Availability Zone in which the cluster is located or "Multiple" if the cache nodes are located in different Availability Zones.

Type: String

Required: No

PreferredMaintenanceWindow

Specifies the weekly time range during which maintenance on the cluster is performed. It is specified as a range in the format ddd:hh24:mi-ddd:hh24:mi (24H Clock UTC). The minimum maintenance window is a 60 minute period.

Valid values for ddd are:

- sun
- mon
- tue
- wed
- thu
- fri
- sat

Example: sun:23:00-mon:01:30

Type: String

Required: No

PreferredOutpostArn

The outpost ARN in which the cache cluster is created.

Type: String

Required: No

ReplicationGroupId

The replication group to which this cluster belongs. If this field is empty, the cluster is not associated with any replication group.

Type: String

Required: No

ReplicationGroupLogDeliveryEnabled

A boolean value indicating whether log delivery is enabled for the replication group.

Type: Boolean

Required: No

SecurityGroups.member.N

A list of VPC Security Groups associated with the cluster.

Type: Array of [SecurityGroupMembership](#) (p. 307) objects

Required: No

SnapshotRetentionLimit

The number of days for which ElastiCache retains automatic cluster snapshots before deleting them. For example, if you set `SnapshotRetentionLimit` to 5, a snapshot that was taken today is retained for 5 days before being deleted.

Important

If the value of `SnapshotRetentionLimit` is set to zero (0), backups are turned off.

Type: Integer

Required: No

SnapshotWindow

The daily time range (in UTC) during which ElastiCache begins taking a daily snapshot of your cluster.

Example: 05:00–09:00

Type: String

Required: No

TransitEncryptionEnabled

A flag that enables in-transit encryption when set to `true`.

You cannot modify the value of `TransitEncryptionEnabled` after the cluster is created. To enable in-transit encryption on a cluster you must set `TransitEncryptionEnabled` to `true` when you create a cluster.

Required: Only available when creating a replication group in an Amazon VPC using redis version 3.2.6, 4.x or later.

Default: `false`

Type: Boolean

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

CacheEngineVersion

Provides all of the details about a particular cache engine version.

Contents

Note

In the following list, the required parameters are described first.

CacheEngineDescription

The description of the cache engine.

Type: String

Required: No

CacheEngineVersionDescription

The description of the cache engine version.

Type: String

Required: No

CacheParameterGroupFamily

The name of the cache parameter group family associated with this cache engine.

Valid values are: `memcached1.4` | `memcached1.5` | `memcached1.6` | `redis2.6` | `redis2.8` | `redis3.2` | `redis4.0` | `redis5.0` | `redis6.x` |

Type: String

Required: No

Engine

The name of the cache engine.

Type: String

Required: No

EngineVersion

The version number of the cache engine.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

CacheNode

Represents an individual cache node within a cluster. Each cache node runs its own instance of the cluster's protocol-compliant caching software - either Memcached or Redis.

The following node types are supported by ElastiCache. Generally speaking, the current generation types provide more memory and computational power at lower cost when compared to their equivalent previous generation counterparts.

- General purpose:
 - Current generation:

M6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

cache.m6g.large, cache.m6g.xlarge, cache.m6g.2xlarge, cache.m6g.4xlarge, cache.m6g.8xlarge, cache.m6g.12xlarge, cache.m6g.16xlarge

Note

For region availability, see [Supported Node Types](#)

M5 node types: cache.m5.large, cache.m5.xlarge, cache.m5.2xlarge, cache.m5.4xlarge, cache.m5.12xlarge, cache.m5.24xlarge

M4 node types: cache.m4.large, cache.m4.xlarge, cache.m4.2xlarge, cache.m4.4xlarge, cache.m4.10xlarge

T3 node types: cache.t3.micro, cache.t3.small, cache.t3.medium

T2 node types: cache.t2.micro, cache.t2.small, cache.t2.medium

- Previous generation: (not recommended)

T1 node types: cache.t1.micro

M1 node types: cache.m1.small, cache.m1.medium, cache.m1.large, cache.m1.xlarge

M3 node types: cache.m3.medium, cache.m3.large, cache.m3.xlarge, cache.m3.2xlarge

- Compute optimized:
 - Previous generation: (not recommended)

C1 node types: cache.c1.xlarge

- Memory optimized:
 - Current generation:

R6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

cache.r6g.large, cache.r6g.xlarge, cache.r6g.2xlarge, cache.r6g.4xlarge, cache.r6g.8xlarge, cache.r6g.12xlarge, cache.r6g.16xlarge

Note

For region availability, see [Supported Node Types](#)

R5 node types: cache.r5.large, cache.r5.xlarge, cache.r5.2xlarge, cache.r5.4xlarge, cache.r5.12xlarge, cache.r5.24xlarge

R4 node types: cache.r4.large, cache.r4.xlarge, cache.r4.2xlarge, cache.r4.4xlarge, cache.r4.8xlarge, cache.r4.16xlarge

- Previous generation: (not recommended)

M2 node types: `cache.m2.xlarge`, `cache.m2.2xlarge`, `cache.m2.4xlarge`

R3 node types: `cache.r3.large`, `cache.r3.xlarge`, `cache.r3.2xlarge`, `cache.r3.4xlarge`, `cache.r3.8xlarge`

Additional node type info

- All current generation instance types are created in Amazon VPC by default.
- Redis append-only files (AOF) are not supported for T1 or T2 instances.
- Redis Multi-AZ with automatic failover is not supported on T1 instances.
- Redis configuration variables `appendonly` and `appendfsync` are not supported on Redis version 2.8.22 and later.

Contents

Note

In the following list, the required parameters are described first.

CacheNodeCreateTime

The date and time when the cache node was created.

Type: Timestamp

Required: No

CacheNodeId

The cache node identifier. A node ID is a numeric identifier (0001, 0002, etc.). The combination of cluster ID and node ID uniquely identifies every cache node used in a customer's Amazon account.

Type: String

Required: No

CacheNodeStatus

The current state of this cache node, one of the following values: `available`, `creating`, `rebooting`, or `deleting`.

Type: String

Required: No

CustomerAvailabilityZone

The Availability Zone where this node was created and now resides.

Type: String

Required: No

CustomerOutpostArn

The customer outpost ARN of the cache node.

Type: String

Required: No

Endpoint

The hostname for connecting to this cache node.

Type: [Endpoint](#) (p. 254) object

Required: No

ParameterGroupStatus

The status of the parameter group applied to this cache node.

Type: String

Required: No

SourceCacheNodeId

The ID of the primary node to which this read replica node is synchronized. If this field is empty, this node is not associated with a primary cluster.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

CacheNodeTypeSpecificParameter

A parameter that has a different value for each cache node type it is applied to. For example, in a Redis cluster, a `cache.m1.large` cache node type would have a larger `maxmemory` value than a `cache.m1.small` type.

Contents

Note

In the following list, the required parameters are described first.

AllowedValues

The valid range of values for the parameter.

Type: String

Required: No

CacheNodeTypeSpecificValues.CacheNodeTypeSpecificValue.N

A list of cache node types and their corresponding values for this parameter.

Type: Array of [CacheNodeTypeSpecificValue](#) (p. 237) objects

Required: No

ChangeType

Indicates whether a change to the parameter is applied immediately or requires a reboot for the change to be applied. You can force a reboot or wait until the next maintenance window's reboot. For more information, see [Rebooting a Cluster](#).

Type: String

Valid Values: `immediate` | `requires-reboot`

Required: No

DataType

The valid data type for the parameter.

Type: String

Required: No

Description

A description of the parameter.

Type: String

Required: No

IsModifiable

Indicates whether (`true`) or not (`false`) the parameter can be modified. Some parameters have security or operational implications that prevent them from being changed.

Type: Boolean

Required: No

MinimumEngineVersion

The earliest cache engine version to which the parameter can apply.

Type: String

Required: No

ParameterName

The name of the parameter.

Type: String

Required: No

Source

The source of the parameter value.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

CacheNodeTypeSpecificValue

A value that applies only to a certain cache node type.

Contents

Note

In the following list, the required parameters are described first.

CacheNodeType

The cache node type for which this value applies.

Type: String

Required: No

Value

The value for the cache node type.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

CacheNodeUpdateStatus

The status of the service update on the cache node

Contents

Note

In the following list, the required parameters are described first.

CacheNodeId

The node ID of the cache cluster

Type: String

Required: No

NodeDeletionDate

The deletion date of the node

Type: Timestamp

Required: No

NodeUpdateEndDate

The end date of the update for a node

Type: Timestamp

Required: No

NodeUpdateInitiatedBy

Reflects whether the update was initiated by the customer or automatically applied

Type: String

Valid Values: `system` | `customer`

Required: No

NodeUpdateInitiatedDate

The date when the update is triggered

Type: Timestamp

Required: No

NodeUpdateStartDate

The start date of the update for a node

Type: Timestamp

Required: No

NodeUpdateStatus

The update status of the node

Type: String

Valid Values: not-applied | waiting-to-start | in-progress | stopping | stopped
| complete

Required: No

NodeUpdateStatusModifiedDate

The date when the NodeUpdateStatus was last modified>

Type: Timestamp

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

CacheParameterGroup

Represents the output of a `CreateCacheParameterGroup` operation.

Contents

Note

In the following list, the required parameters are described first.

ARN

The ARN (Amazon Resource Name) of the cache parameter group.

Type: String

Required: No

CacheParameterGroupFamily

The name of the cache parameter group family that this cache parameter group is compatible with.

Valid values are: `memcached1.4` | `memcached1.5` | `memcached1.6` | `redis2.6` | `redis2.8` | `redis3.2` | `redis4.0` | `redis5.0` | `redis6.x` |

Type: String

Required: No

CacheParameterGroupName

The name of the cache parameter group.

Type: String

Required: No

Description

The description for this cache parameter group.

Type: String

Required: No

IsGlobal

Indicates whether the parameter group is associated with a Global datastore

Type: Boolean

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

CacheParameterGroupStatus

Status of the cache parameter group.

Contents

Note

In the following list, the required parameters are described first.

CacheNodeIdsToReboot.CacheNodeId.N

A list of the cache node IDs which need to be rebooted for parameter changes to be applied. A node ID is a numeric identifier (0001, 0002, etc.).

Type: Array of strings

Required: No

CacheParameterGroupName

The name of the cache parameter group.

Type: String

Required: No

ParameterApplyStatus

The status of parameter updates.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

CacheSecurityGroup

Represents the output of one of the following operations:

- `AuthorizeCacheSecurityGroupIngress`
- `CreateCacheSecurityGroup`
- `RevokeCacheSecurityGroupIngress`

Contents

Note

In the following list, the required parameters are described first.

ARN

The ARN of the cache security group,

Type: String

Required: No

CacheSecurityGroupName

The name of the cache security group.

Type: String

Required: No

Description

The description of the cache security group.

Type: String

Required: No

EC2SecurityGroups.EC2SecurityGroup.N

A list of Amazon EC2 security groups that are associated with this cache security group.

Type: Array of [EC2SecurityGroup](#) (p. 253) objects

Required: No

OwnerId

The Amazon account ID of the cache security group owner.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)

- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

CacheSecurityGroupMembership

Represents a cluster's status within a particular cache security group.

Contents

Note

In the following list, the required parameters are described first.

CacheSecurityGroupName

The name of the cache security group.

Type: String

Required: No

Status

The membership status in the cache security group. The status changes when a cache security group is modified, or when the cache security groups assigned to a cluster are modified.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

CacheSubnetGroup

Represents the output of one of the following operations:

- `CreateCacheSubnetGroup`
- `ModifyCacheSubnetGroup`

Contents

Note

In the following list, the required parameters are described first.

ARN

The ARN (Amazon Resource Name) of the cache subnet group.

Type: String

Required: No

CacheSubnetGroupDescription

The description of the cache subnet group.

Type: String

Required: No

CacheSubnetGroupName

The name of the cache subnet group.

Type: String

Required: No

Subnets.Subnet.N

A list of subnets associated with the cache subnet group.

Type: Array of [Subnet](#) (p. 317) objects

Required: No

VpcId

The Amazon Virtual Private Cloud identifier (VPC ID) of the cache subnet group.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)

- [AWS SDK for Ruby V3](#)

CloudWatchLogsDestinationDetails

The configuration details of the CloudWatch Logs destination.

Contents

Note

In the following list, the required parameters are described first.

LogGroup

The name of the CloudWatch Logs log group.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ConfigureShard

Node group (shard) configuration options when adding or removing replicas. Each node group (shard) configuration has the following members: `NodeGroupId`, `NewReplicaCount`, and `PreferredAvailabilityZones`.

Contents

Note

In the following list, the required parameters are described first.

NewReplicaCount

The number of replicas you want in this node group at the end of this operation. The maximum value for `NewReplicaCount` is 5. The minimum value depends upon the type of Redis replication group you are working with.

The minimum number of replicas in a shard or replication group is:

- Redis (cluster mode disabled)
 - If Multi-AZ: 1
 - If Multi-AZ: 0
- Redis (cluster mode enabled): 0 (though you will not be able to failover to a replica if your primary node fails)

Type: Integer

Required: Yes

NodeGroupId

The 4-digit id for the node group you are configuring. For Redis (cluster mode disabled) replication groups, the node group id is always 0001. To find a Redis (cluster mode enabled)'s node group's (shard's) id, see [Finding a Shard's Id](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 4.

Pattern: `\d+`

Required: Yes

PreferredAvailabilityZones.PreferredAvailabilityZone.N

A list of `PreferredAvailabilityZone` strings that specify which availability zones the replication group's nodes are to be in. The number of `PreferredAvailabilityZone` values must equal the value of `NewReplicaCount` plus 1 to account for the primary node. If this member of `ReplicaConfiguration` is omitted, ElastiCache for Redis selects the availability zone for each of the replicas.

Type: Array of strings

Required: No

PreferredOutpostArns.PreferredOutpostArn.N

The outpost ARNs in which the cache cluster is created.

Type: Array of strings

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

CustomerNodeEndpoint

The endpoint from which data should be migrated.

Contents

Note

In the following list, the required parameters are described first.

Address

The address of the node endpoint

Type: String

Required: No

Port

The port of the node endpoint

Type: Integer

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

DestinationDetails

Configuration details of either a CloudWatch Logs destination or Kinesis Data Firehose destination.

Contents

Note

In the following list, the required parameters are described first.

CloudWatchLogsDetails

The configuration details of the CloudWatch Logs destination.

Type: [CloudWatchLogsDestinationDetails](#) (p. 248) object

Required: No

KinesisFirehoseDetails

The configuration details of the Kinesis Data Firehose destination.

Type: [KinesisFirehoseDestinationDetails](#) (p. 265) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

EC2SecurityGroup

Provides ownership and status information for an Amazon EC2 security group.

Contents

Note

In the following list, the required parameters are described first.

EC2SecurityGroupName

The name of the Amazon EC2 security group.

Type: String

Required: No

EC2SecurityGroupOwnerId

The Amazon account ID of the Amazon EC2 security group owner.

Type: String

Required: No

Status

The status of the Amazon EC2 security group.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

Endpoint

Represents the information required for client programs to connect to a cache node.

Contents

Note

In the following list, the required parameters are described first.

Address

The DNS hostname of the cache node.

Type: String

Required: No

Port

The port number that the cache engine is listening on.

Type: Integer

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

EngineDefaults

Represents the output of a `DescribeEngineDefaultParameters` operation.

Contents

Note

In the following list, the required parameters are described first.

CacheNodeTypeSpecificParameters.CacheNodeTypeSpecificParameter.N

A list of parameters specific to a particular cache node type. Each element in the list contains detailed information about one parameter.

Type: Array of [CacheNodeTypeSpecificParameter](#) (p. 235) objects

Required: No

CacheParameterGroupFamily

Specifies the name of the cache parameter group family to which the engine default parameters apply.

Valid values are: `memcached1.4` | `memcached1.5` | `memcached1.6` | `redis2.6` | `redis2.8` | `redis3.2` | `redis4.0` | `redis5.0` | `redis6.x` |

Type: String

Required: No

Marker

Provides an identifier to allow retrieval of paginated results.

Type: String

Required: No

Parameters.Parameter.N

Contains a list of engine default parameters.

Type: Array of [Parameter](#) (p. 282) objects

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

Event

Represents a single occurrence of something interesting within the system. Some examples of events are creating a cluster, adding or removing a cache node, or rebooting a node.

Contents

Note

In the following list, the required parameters are described first.

Date

The date and time when the event occurred.

Type: Timestamp

Required: No

Message

The text of the event.

Type: String

Required: No

SourceIdentifier

The identifier for the source of the event. For example, if the event occurred at the cluster level, the identifier would be the name of the cluster.

Type: String

Required: No

SourceType

Specifies the origin of this event - a cluster, a parameter group, a security group, etc.

Type: String

Valid Values: `cache-cluster` | `cache-parameter-group` | `cache-security-group` | `cache-subnet-group` | `replication-group` | `user` | `user-group`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

Filter

Used to streamline results of a search based on the property being filtered.

Contents

Note

In the following list, the required parameters are described first.

Name

The property being filtered. For example, UserId.

Type: String

Pattern: `.*\S.*`

Required: Yes

Values.member.N

The property values to filter on. For example, "user-123".

Type: Array of strings

Array Members: Minimum number of 1 item.

Pattern: `.*\S.*`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

GlobalNodeGroup

Indicates the slot configuration and global identifier for a slice group.

Contents

Note

In the following list, the required parameters are described first.

GlobalNodeGroupId

The name of the global node group

Type: String

Required: No

Slots

The keyspace for this node group

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

GlobalReplicationGroup

Consists of a primary cluster that accepts writes and an associated secondary cluster that resides in a different Amazon region. The secondary cluster accepts only reads. The primary cluster automatically replicates updates to the secondary cluster.

- The **GlobalReplicationGroupIdSuffix** represents the name of the Global datastore, which is what you use to associate a secondary cluster.

Contents

Note

In the following list, the required parameters are described first.

ARN

The ARN (Amazon Resource Name) of the global replication group.

Type: String

Required: No

AtRestEncryptionEnabled

A flag that enables encryption at rest when set to `true`.

You cannot modify the value of `AtRestEncryptionEnabled` after the replication group is created. To enable encryption at rest on a replication group you must set `AtRestEncryptionEnabled` to `true` when you create the replication group.

Required: Only available when creating a replication group in an Amazon VPC using redis version 3.2.6, 4.x or later.

Type: Boolean

Required: No

AuthTokenEnabled

A flag that enables using an `AuthToken` (password) when issuing Redis commands.

Default: `false`

Type: Boolean

Required: No

CacheNodeType

The cache node type of the Global datastore

Type: String

Required: No

ClusterEnabled

A flag that indicates whether the Global datastore is cluster enabled.

Type: Boolean

Required: No

Engine

The ElastiCache engine. For Redis only.

Type: String

Required: No

EngineVersion

The ElastiCache Redis engine version.

Type: String

Required: No

GlobalNodeGroups.GlobalNodeGroup.N

Indicates the slot configuration and global identifier for each slice group.

Type: Array of [GlobalNodeGroup](#) (p. 258) objects

Required: No

GlobalReplicationGroupDescription

The optional description of the Global datastore

Type: String

Required: No

GlobalReplicationGroupId

The name of the Global datastore

Type: String

Required: No

Members.GlobalReplicationGroupMember.N

The replication groups that comprise the Global datastore.

Type: Array of [GlobalReplicationGroupMember](#) (p. 263) objects

Required: No

Status

The status of the Global datastore

Type: String

Required: No

TransitEncryptionEnabled

A flag that enables in-transit encryption when set to true. You cannot modify the value of `TransitEncryptionEnabled` after the cluster is created. To enable in-transit encryption on a cluster you must set `TransitEncryptionEnabled` to true when you create a cluster.

Required: Only available when creating a replication group in an Amazon VPC using redis version 3.2.6, 4.x or later.

Type: Boolean

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

GlobalReplicationGroupInfo

The name of the Global datastore and role of this replication group in the Global datastore.

Contents

Note

In the following list, the required parameters are described first.

GlobalReplicationGroupId

The name of the Global datastore

Type: String

Required: No

GlobalReplicationGroupMemberRole

The role of the replication group in a Global datastore. Can be primary or secondary.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

GlobalReplicationGroupMember

A member of a Global datastore. It contains the Replication Group Id, the Amazon region and the role of the replication group.

Contents

Note

In the following list, the required parameters are described first.

AutomaticFailover

Indicates whether automatic failover is enabled for the replication group.

Type: String

Valid Values: `enabled` | `disabled` | `enabling` | `disabling`

Required: No

ReplicationGroupId

The replication group id of the Global datastore member.

Type: String

Required: No

ReplicationGroupRegion

The Amazon region of the Global datastore member.

Type: String

Required: No

Role

Indicates the role of the replication group, primary or secondary.

Type: String

Required: No

Status

The status of the membership of the replication group.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

KinesisFirehoseDestinationDetails

The configuration details of the Kinesis Data Firehose destination.

Contents

Note

In the following list, the required parameters are described first.

DeliveryStream

The name of the Kinesis Data Firehose delivery stream.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

LogDeliveryConfiguration

Returns the destination, format and type of the logs.

Contents

Note

In the following list, the required parameters are described first.

DestinationDetails

Configuration details of either a CloudWatch Logs destination or Kinesis Data Firehose destination.

Type: [DestinationDetails](#) (p. 252) object

Required: No

DestinationType

Returns the destination type, either `cloudwatch-logs` or `kinesis-firehose`.

Type: String

Valid Values: `cloudwatch-logs` | `kinesis-firehose`

Required: No

LogFormat

Returns the log format, either JSON or TEXT.

Type: String

Valid Values: `text` | `json`

Required: No

LogType

Refers to [slow-log](#).

Type: String

Valid Values: `slow-log`

Required: No

Message

Returns an error message for the log delivery configuration.

Type: String

Required: No

Status

Returns the log delivery configuration status. Values are one of `enabling` | `disabling` | `modifying` | `active` | `error`

Type: String

Valid Values: `active` | `enabling` | `modifying` | `disabling` | `error`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

LogDeliveryConfigurationRequest

Specifies the destination, format and type of the logs.

Contents

Note

In the following list, the required parameters are described first.

DestinationDetails

Configuration details of either a CloudWatch Logs destination or Kinesis Data Firehose destination.

Type: [DestinationDetails](#) (p. 252) object

Required: No

DestinationType

Specify either `cloudwatch-logs` or `kinesis-firehose` as the destination type.

Type: String

Valid Values: `cloudwatch-logs` | `kinesis-firehose`

Required: No

Enabled

Specify if log delivery is enabled. Default `true`.

Type: Boolean

Required: No

LogFormat

Specifies either JSON or TEXT

Type: String

Valid Values: `text` | `json`

Required: No

LogType

Refers to [slow-log](#).

Type: String

Valid Values: `slow-log`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)

- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

NodeGroup

Represents a collection of cache nodes in a replication group. One node in the node group is the read/write primary node. All the other nodes are read-only Replica nodes.

Contents

Note

In the following list, the required parameters are described first.

NodeGroupId

The identifier for the node group (shard). A Redis (cluster mode disabled) replication group contains only 1 node group; therefore, the node group ID is 0001. A Redis (cluster mode enabled) replication group contains 1 to 90 node groups numbered 0001 to 0090. Optionally, the user can provide the id for a node group.

Type: String

Required: No

NodeGroupMembers.NodeGroupMember.N

A list containing information about individual nodes within the node group (shard).

Type: Array of [NodeGroupMember](#) (p. 274) objects

Required: No

PrimaryEndpoint

The endpoint of the primary node in this node group (shard).

Type: [Endpoint](#) (p. 254) object

Required: No

ReaderEndpoint

The endpoint of the replica nodes in this node group (shard).

Type: [Endpoint](#) (p. 254) object

Required: No

Slots

The keyspace for this node group (shard).

Type: String

Required: No

Status

The current state of this replication group - creating, available, modifying, deleting.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

NodeGroupConfiguration

Node group (shard) configuration options. Each node group (shard) configuration has the following: Slots, PrimaryAvailabilityZone, ReplicaAvailabilityZones, ReplicaCount.

Contents

Note

In the following list, the required parameters are described first.

NodeGroupId

Either the ElastiCache for Redis supplied 4-digit id or a user supplied id for the node group these configuration values apply to.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 4.

Pattern: \d+

Required: No

PrimaryAvailabilityZone

The Availability Zone where the primary node of this node group (shard) is launched.

Type: String

Required: No

PrimaryOutpostArn

The outpost ARN of the primary node.

Type: String

Required: No

ReplicaAvailabilityZones.AvailabilityZone.N

A list of Availability Zones to be used for the read replicas. The number of Availability Zones in this list must match the value of ReplicaCount or ReplicasPerNodeGroup if not specified.

Type: Array of strings

Required: No

ReplicaCount

The number of read replica nodes in this node group (shard).

Type: Integer

Required: No

ReplicaOutpostArns.OutpostArn.N

The outpost ARN of the node replicas.

Type: Array of strings

Required: No

Slots

A string that specifies the keyspaces for a particular node group. Keyspaces range from 0 to 16,383. The string is in the format `startkey-endkey`.

Example: "0-3999"

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

NodeGroupMember

Represents a single node within a node group (shard).

Contents

Note

In the following list, the required parameters are described first.

CacheClusterId

The ID of the cluster to which the node belongs.

Type: String

Required: No

CacheNodeId

The ID of the node within its cluster. A node ID is a numeric identifier (0001, 0002, etc.).

Type: String

Required: No

CurrentRole

The role that is currently assigned to the node - `primary` or `replica`. This member is only applicable for Redis (cluster mode disabled) replication groups.

Type: String

Required: No

PreferredAvailabilityZone

The name of the Availability Zone in which the node is located.

Type: String

Required: No

PreferredOutpostArn

The outpost ARN of the node group member.

Type: String

Required: No

ReadEndpoint

The information required for client programs to connect to a node for read operations. The read endpoint is only applicable on Redis (cluster mode disabled) clusters.

Type: [Endpoint](#) (p. 254) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

NodeGroupMemberUpdateStatus

The status of the service update on the node group member

Contents

Note

In the following list, the required parameters are described first.

CacheClusterId

The cache cluster ID

Type: String

Required: No

CacheNodeId

The node ID of the cache cluster

Type: String

Required: No

NodeDeletionDate

The deletion date of the node

Type: Timestamp

Required: No

NodeUpdateEndDate

The end date of the update for a node

Type: Timestamp

Required: No

NodeUpdateInitiatedBy

Reflects whether the update was initiated by the customer or automatically applied

Type: String

Valid Values: `system` | `customer`

Required: No

NodeUpdateInitiatedDate

The date when the update is triggered

Type: Timestamp

Required: No

NodeUpdateStartDate

The start date of the update for a node

Type: Timestamp

Required: No

NodeUpdateStatus

The update status of the node

Type: String

Valid Values: not-applied | waiting-to-start | in-progress | stopping | stopped
| complete

Required: No

NodeUpdateStatusModifiedDate

The date when the NodeUpdateStatus was last modified

Type: Timestamp

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

NodeGroupUpdateStatus

The status of the service update on the node group

Contents

Note

In the following list, the required parameters are described first.

NodeGroupId

The ID of the node group

Type: String

Required: No

NodeGroupMemberUpdateStatus.NodeGroupMemberUpdateStatus.N

The status of the service update on the node group member

Type: Array of [NodeGroupMemberUpdateStatus](#) (p. 276) objects

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

NodeSnapshot

Represents an individual cache node in a snapshot of a cluster.

Contents

Note

In the following list, the required parameters are described first.

CacheClusterId

A unique identifier for the source cluster.

Type: String

Required: No

CacheNodeCreateTime

The date and time when the cache node was created in the source cluster.

Type: Timestamp

Required: No

CacheNodeId

The cache node identifier for the node in the source cluster.

Type: String

Required: No

CacheSize

The size of the cache on the source cache node.

Type: String

Required: No

NodeGroupConfiguration

The configuration for the source node group (shard).

Type: [NodeGroupConfiguration](#) (p. 272) object

Required: No

NodeGroupId

A unique identifier for the source node group (shard).

Type: String

Required: No

SnapshotCreateTime

The date and time when the source node's metadata and cache data set was obtained for the snapshot.

Type: Timestamp

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

NotificationConfiguration

Describes a notification topic and its status. Notification topics are used for publishing ElastiCache events to subscribers using Amazon Simple Notification Service (SNS).

Contents

Note

In the following list, the required parameters are described first.

TopicArn

The Amazon Resource Name (ARN) that identifies the topic.

Type: String

Required: No

TopicStatus

The current state of the topic.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

Parameter

Describes an individual setting that controls some aspect of ElastiCache behavior.

Contents

Note

In the following list, the required parameters are described first.

AllowedValues

The valid range of values for the parameter.

Type: String

Required: No

ChangeType

Indicates whether a change to the parameter is applied immediately or requires a reboot for the change to be applied. You can force a reboot or wait until the next maintenance window's reboot. For more information, see [Rebooting a Cluster](#).

Type: String

Valid Values: `immediate` | `requires-reboot`

Required: No

DataType

The valid data type for the parameter.

Type: String

Required: No

Description

A description of the parameter.

Type: String

Required: No

IsModifiable

Indicates whether (`true`) or not (`false`) the parameter can be modified. Some parameters have security or operational implications that prevent them from being changed.

Type: Boolean

Required: No

MinimumEngineVersion

The earliest cache engine version to which the parameter can apply.

Type: String

Required: No

ParameterName

The name of the parameter.

Type: String

Required: No

ParameterValue

The value of the parameter.

Type: String

Required: No

Source

The source of the parameter.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ParameterNameValue

Describes a name-value pair that is used to update the value of a parameter.

Contents

Note

In the following list, the required parameters are described first.

ParameterName

The name of the parameter.

Type: String

Required: No

ParameterValue

The value of the parameter.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

PendingLogDeliveryConfiguration

The log delivery configurations being modified

Contents

Note

In the following list, the required parameters are described first.

DestinationDetails

Configuration details of either a CloudWatch Logs destination or Kinesis Data Firehose destination.

Type: [DestinationDetails](#) (p. 252) object

Required: No

DestinationType

Returns the destination type, either CloudWatch Logs or Kinesis Data Firehose.

Type: String

Valid Values: `cloudwatch-logs` | `kinesis-firehose`

Required: No

LogFormat

Returns the log format, either JSON or TEXT

Type: String

Valid Values: `text` | `json`

Required: No

LogType

Refers to [slow-log](#).

Type: String

Valid Values: `slow-log`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

PendingModifiedValues

A group of settings that are applied to the cluster in the future, or that are currently being applied.

Contents

Note

In the following list, the required parameters are described first.

AuthTokenStatus

The auth token status

Type: String

Valid Values: `SETTING` | `ROTATING`

Required: No

CacheNodeIdsToRemove.CacheNodeId.N

A list of cache node IDs that are being removed (or will be removed) from the cluster. A node ID is a 4-digit numeric identifier (0001, 0002, etc.).

Type: Array of strings

Required: No

CacheNodeType

The cache node type that this cluster or replication group is scaled to.

Type: String

Required: No

EngineVersion

The new cache engine version that the cluster runs.

Type: String

Required: No

NumCacheNodes

The new number of cache nodes for the cluster.

For clusters running Redis, this value must be 1. For clusters running Memcached, this value must be between 1 and 40.

Type: Integer

Required: No

PendingLogDeliveryConfiguration.member.N

The log delivery configurations being modified

Type: Array of [PendingLogDeliveryConfiguration](#) (p. 285) objects

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ProcessedUpdateAction

Update action that has been processed for the corresponding apply/stop request

Contents

Note

In the following list, the required parameters are described first.

CacheClusterId

The ID of the cache cluster

Type: String

Required: No

ReplicationGroupId

The ID of the replication group

Type: String

Required: No

ServiceUpdateName

The unique ID of the service update

Type: String

Required: No

UpdateActionStatus

The status of the update action on the Redis cluster

Type: String

Valid Values: not-applied | waiting-to-start | in-progress | stopping | stopped
| complete | scheduling | scheduled | not-applicable

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

RecurringCharge

Contains the specific price and frequency of a recurring charges for a reserved cache node, or for a reserved cache node offering.

Contents

Note

In the following list, the required parameters are described first.

RecurringChargeAmount

The monetary amount of the recurring charge.

Type: Double

Required: No

RecurringChargeFrequency

The frequency of the recurring charge.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

RegionalConfiguration

A list of the replication groups

Contents

Note

In the following list, the required parameters are described first.

ReplicationGroupId

The name of the secondary cluster

Type: String

Required: Yes

ReplicationGroupRegion

The Amazon region where the cluster is stored

Type: String

Required: Yes

ReshardingConfiguration.ReshardingConfiguration.N

A list of `PreferredAvailabilityZones` objects that specifies the configuration of a node group in the resharded cluster.

Type: Array of [ReshardingConfiguration](#) (p. 305) objects

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ReplicationGroup

Contains all of the attributes of a specific Redis replication group.

Contents

Note

In the following list, the required parameters are described first.

ARN

The ARN (Amazon Resource Name) of the replication group.

Type: String

Required: No

AtRestEncryptionEnabled

A flag that enables encryption at-rest when set to `true`.

You cannot modify the value of `AtRestEncryptionEnabled` after the cluster is created. To enable encryption at-rest on a cluster you must set `AtRestEncryptionEnabled` to `true` when you create a cluster.

Required: Only available when creating a replication group in an Amazon VPC using redis version `3.2.6`, `4.x` or later.

Default: `false`

Type: Boolean

Required: No

AuthTokenEnabled

A flag that enables using an `AuthToken` (password) when issuing Redis commands.

Default: `false`

Type: Boolean

Required: No

AuthTokenLastModifiedDate

The date the auth token was last modified

Type: Timestamp

Required: No

AutomaticFailover

Indicates the status of automatic failover for this Redis replication group.

Type: String

Valid Values: `enabled` | `disabled` | `enabling` | `disabling`

Required: No

CacheNodeType

The name of the compute and memory capacity node type for each node in the replication group.

Type: String

Required: No

ClusterEnabled

A flag indicating whether or not this replication group is cluster enabled; i.e., whether its data can be partitioned across multiple shards (API/CLI: node groups).

Valid values: `true` | `false`

Type: Boolean

Required: No

ConfigurationEndpoint

The configuration endpoint for this replication group. Use the configuration endpoint to connect to this replication group.

Type: [Endpoint](#) (p. 254) object

Required: No

Description

The user supplied description of the replication group.

Type: String

Required: No

GlobalReplicationGroupInfo

The name of the Global datastore and role of this replication group in the Global datastore.

Type: [GlobalReplicationGroupInfo](#) (p. 262) object

Required: No

KmsKeyId

The ID of the KMS key used to encrypt the disk in the cluster.

Type: String

Required: No

LogDeliveryConfigurations.LogDeliveryConfiguration.N

Returns the destination, format and type of the logs.

Type: Array of [LogDeliveryConfiguration](#) (p. 266) objects

Required: No

MemberClusters.ClusterId.N

The names of all the cache clusters that are part of this replication group.

Type: Array of strings

Required: No

MemberClustersOutpostArns.ReplicationGroupOutpostArn.N

The outpost ARNs of the replication group's member clusters.

Type: Array of strings

Required: No

MultiAZ

A flag indicating if you have Multi-AZ enabled to enhance fault tolerance. For more information, see [Minimizing Downtime: Multi-AZ](#)

Type: String

Valid Values: `enabled` | `disabled`

Required: No

NodeGroups.NodeGroup.N

A list of node groups in this replication group. For Redis (cluster mode disabled) replication groups, this is a single-element list. For Redis (cluster mode enabled) replication groups, the list contains an entry for each node group (shard).

Type: Array of [NodeGroup](#) (p. 270) objects

Required: No

PendingModifiedValues

A group of settings to be applied to the replication group, either immediately or during the next maintenance window.

Type: [ReplicationGroupPendingModifiedValues](#) (p. 296) object

Required: No

ReplicationGroupCreateTime

The date and time when the cluster was created.

Type: Timestamp

Required: No

ReplicationGroupId

The identifier for the replication group.

Type: String

Required: No

SnapshotRetentionLimit

The number of days for which ElastiCache retains automatic cluster snapshots before deleting them. For example, if you set `SnapshotRetentionLimit` to 5, a snapshot that was taken today is retained for 5 days before being deleted.

Important

If the value of `SnapshotRetentionLimit` is set to zero (0), backups are turned off.

Type: Integer

Required: No

SnapshottingClusterId

The cluster ID that is used as the daily snapshot source for the replication group.

Type: String

Required: No

SnapshotWindow

The daily time range (in UTC) during which ElastiCache begins taking a daily snapshot of your node group (shard).

Example: 05:00-09:00

If you do not specify this parameter, ElastiCache automatically chooses an appropriate time range.

Note

This parameter is only valid if the `Engine` parameter is `redis`.

Type: String

Required: No

Status

The current state of this replication group - `creating`, `available`, `modifying`, `deleting`, `create-failed`, `snapshotting`.

Type: String

Required: No

TransitEncryptionEnabled

A flag that enables in-transit encryption when set to `true`.

You cannot modify the value of `TransitEncryptionEnabled` after the cluster is created. To enable in-transit encryption on a cluster you must set `TransitEncryptionEnabled` to `true` when you create a cluster.

Required: Only available when creating a replication group in an Amazon VPC using redis version 3.2.6, 4.x or later.

Default: `false`

Type: Boolean

Required: No

UserGroupIds.member.N

The ID of the user group associated to the replication group.

Type: Array of strings

Length Constraints: Minimum length of 1.

Pattern: `[a-zA-Z][a-zA-Z0-9\-\]*`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)

- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ReplicationGroupPendingModifiedValues

The settings to be applied to the Redis replication group, either immediately or during the next maintenance window.

Contents

Note

In the following list, the required parameters are described first.

AuthTokenStatus

The auth token status

Type: String

Valid Values: `SETTING` | `ROTATING`

Required: No

AutomaticFailoverStatus

Indicates the status of automatic failover for this Redis replication group.

Type: String

Valid Values: `enabled` | `disabled`

Required: No

PendingLogDeliveryConfiguration.member.N

The log delivery configurations being modified

Type: Array of [PendingLogDeliveryConfiguration](#) (p. 285) objects

Required: No

PrimaryClusterId

The primary cluster ID that is applied immediately (if `--apply-immediately` was specified), or during the next maintenance window.

Type: String

Required: No

Resharding

The status of an online resharding operation.

Type: [ReshardingStatus](#) (p. 306) object

Required: No

UserGroups

The user group being modified.

Type: [UserGroupsUpdateStatus](#) (p. 330) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ReservedCacheNode

Represents the output of a `PurchaseReservedCacheNodesOffering` operation.

Contents

Note

In the following list, the required parameters are described first.

CacheNodeCount

The number of cache nodes that have been reserved.

Type: Integer

Required: No

CacheNodeType

The cache node type for the reserved cache nodes.

The following node types are supported by ElastiCache. Generally speaking, the current generation types provide more memory and computational power at lower cost when compared to their equivalent previous generation counterparts.

- General purpose:
 - Current generation:

M6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

`cache.m6g.large`, `cache.m6g.xlarge`, `cache.m6g.2xlarge`, `cache.m6g.4xlarge`,
`cache.m6g.8xlarge`, `cache.m6g.12xlarge`, `cache.m6g.16xlarge`

Note

For region availability, see [Supported Node Types](#)

M5 node types: `cache.m5.large`, `cache.m5.xlarge`, `cache.m5.2xlarge`,
`cache.m5.4xlarge`, `cache.m5.12xlarge`, `cache.m5.24xlarge`

M4 node types: `cache.m4.large`, `cache.m4.xlarge`, `cache.m4.2xlarge`,
`cache.m4.4xlarge`, `cache.m4.10xlarge`

T3 node types: `cache.t3.micro`, `cache.t3.small`, `cache.t3.medium`

T2 node types: `cache.t2.micro`, `cache.t2.small`, `cache.t2.medium`

- Previous generation: (not recommended)

T1 node types: `cache.t1.micro`

M1 node types: `cache.m1.small`, `cache.m1.medium`, `cache.m1.large`,
`cache.m1.xlarge`

M3 node types: `cache.m3.medium`, `cache.m3.large`, `cache.m3.xlarge`,
`cache.m3.2xlarge`

- Compute optimized:
 - Previous generation: (not recommended)

C1 node types: `cache.c1.xlarge`

- Memory optimized:
- Current generation:

R6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

`cache.r6g.large`, `cache.r6g.xlarge`, `cache.r6g.2xlarge`, `cache.r6g.4xlarge`,
`cache.r6g.8xlarge`, `cache.r6g.12xlarge`, `cache.r6g.16xlarge`

Note

For region availability, see [Supported Node Types](#)

R5 node types: `cache.r5.large`, `cache.r5.xlarge`, `cache.r5.2xlarge`,
`cache.r5.4xlarge`, `cache.r5.12xlarge`, `cache.r5.24xlarge`

R4 node types: `cache.r4.large`, `cache.r4.xlarge`, `cache.r4.2xlarge`,
`cache.r4.4xlarge`, `cache.r4.8xlarge`, `cache.r4.16xlarge`

- Previous generation: (not recommended)

M2 node types: `cache.m2.xlarge`, `cache.m2.2xlarge`, `cache.m2.4xlarge`

R3 node types: `cache.r3.large`, `cache.r3.xlarge`, `cache.r3.2xlarge`,
`cache.r3.4xlarge`, `cache.r3.8xlarge`

Additional node type info

- All current generation instance types are created in Amazon VPC by default.
- Redis append-only files (AOF) are not supported for T1 or T2 instances.
- Redis Multi-AZ with automatic failover is not supported on T1 instances.
- Redis configuration variables `appendonly` and `appendfsync` are not supported on Redis version 2.8.22 and later.

Type: String

Required: No

Duration

The duration of the reservation in seconds.

Type: Integer

Required: No

FixedPrice

The fixed price charged for this reserved cache node.

Type: Double

Required: No

OfferingType

The offering type of this reserved cache node.

Type: String

Required: No

ProductDescription

The description of the reserved cache node.

Type: String

Required: No

RecurringCharges.RecurringCharge.N

The recurring price charged to run this reserved cache node.

Type: Array of [RecurringCharge](#) (p. 289) objects

Required: No

ReservationARN

The Amazon Resource Name (ARN) of the reserved cache node.

Example: `arn:aws:elasticache:us-east-1:123456789012:reserved-instance:ri-2017-03-27-08-33-25-582`

Type: String

Required: No

ReservedCacheNodeId

The unique identifier for the reservation.

Type: String

Required: No

ReservedCacheNodesOfferingId

The offering identifier.

Type: String

Required: No

StartTime

The time the reservation started.

Type: Timestamp

Required: No

State

The state of the reserved cache node.

Type: String

Required: No

UsagePrice

The hourly price charged for this reserved cache node.

Type: Double

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ReservedCacheNodesOffering

Describes all of the attributes of a reserved cache node offering.

Contents

Note

In the following list, the required parameters are described first.

CacheNodeType

The cache node type for the reserved cache node.

The following node types are supported by ElastiCache. Generally speaking, the current generation types provide more memory and computational power at lower cost when compared to their equivalent previous generation counterparts.

- General purpose:
 - Current generation:

M6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

`cache.m6g.large`, `cache.m6g.xlarge`, `cache.m6g.2xlarge`, `cache.m6g.4xlarge`,
`cache.m6g.8xlarge`, `cache.m6g.12xlarge`, `cache.m6g.16xlarge`

Note

For region availability, see [Supported Node Types](#)

M5 node types: `cache.m5.large`, `cache.m5.xlarge`, `cache.m5.2xlarge`,
`cache.m5.4xlarge`, `cache.m5.12xlarge`, `cache.m5.24xlarge`

M4 node types: `cache.m4.large`, `cache.m4.xlarge`, `cache.m4.2xlarge`,
`cache.m4.4xlarge`, `cache.m4.10xlarge`

T3 node types: `cache.t3.micro`, `cache.t3.small`, `cache.t3.medium`

T2 node types: `cache.t2.micro`, `cache.t2.small`, `cache.t2.medium`

- Previous generation: (not recommended)

T1 node types: `cache.t1.micro`

M1 node types: `cache.m1.small`, `cache.m1.medium`, `cache.m1.large`,
`cache.m1.xlarge`

M3 node types: `cache.m3.medium`, `cache.m3.large`, `cache.m3.xlarge`,
`cache.m3.2xlarge`

- Compute optimized:
 - Previous generation: (not recommended)

C1 node types: `cache.c1.xlarge`

- Memory optimized:
 - Current generation:

R6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

`cache.r6g.large`, `cache.r6g.xlarge`, `cache.r6g.2xlarge`, `cache.r6g.4xlarge`,
`cache.r6g.8xlarge`, `cache.r6g.12xlarge`, `cache.r6g.16xlarge`

Note

For region availability, see [Supported Node Types](#)

R5 node types: `cache.r5.large`, `cache.r5.xlarge`, `cache.r5.2xlarge`, `cache.r5.4xlarge`, `cache.r5.12xlarge`, `cache.r5.24xlarge`

R4 node types: `cache.r4.large`, `cache.r4.xlarge`, `cache.r4.2xlarge`, `cache.r4.4xlarge`, `cache.r4.8xlarge`, `cache.r4.16xlarge`

- Previous generation: (not recommended)

M2 node types: `cache.m2.xlarge`, `cache.m2.2xlarge`, `cache.m2.4xlarge`

R3 node types: `cache.r3.large`, `cache.r3.xlarge`, `cache.r3.2xlarge`, `cache.r3.4xlarge`, `cache.r3.8xlarge`

Additional node type info

- All current generation instance types are created in Amazon VPC by default.
- Redis append-only files (AOF) are not supported for T1 or T2 instances.
- Redis Multi-AZ with automatic failover is not supported on T1 instances.
- Redis configuration variables `appendonly` and `appendfsync` are not supported on Redis version 2.8.22 and later.

Type: String

Required: No

Duration

The duration of the offering. in seconds.

Type: Integer

Required: No

FixedPrice

The fixed price charged for this offering.

Type: Double

Required: No

OfferingType

The offering type.

Type: String

Required: No

ProductDescription

The cache engine used by the offering.

Type: String

Required: No

RecurringCharges.RecurringCharge.N

The recurring price charged to run this reserved cache node.

Type: Array of [RecurringCharge](#) (p. 289) objects

Required: No

ReservedCacheNodesOfferingId

A unique identifier for the reserved cache node offering.

Type: String

Required: No

UsagePrice

The hourly price charged for this offering.

Type: Double

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ReshardingConfiguration

A list of `PreferredAvailabilityZones` objects that specifies the configuration of a node group in the resharded cluster.

Contents

Note

In the following list, the required parameters are described first.

NodeGroupId

Either the ElastiCache for Redis supplied 4-digit id or a user supplied id for the node group these configuration values apply to.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 4.

Pattern: `\d+`

Required: No

PreferredAvailabilityZones.AvailabilityZone.N

A list of preferred availability zones for the nodes in this cluster.

Type: Array of strings

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ReshardingStatus

The status of an online resharding operation.

Contents

Note

In the following list, the required parameters are described first.

SlotMigration

Represents the progress of an online resharding operation.

Type: [SlotMigration](#) (p. 310) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

SecurityGroupMembership

Represents a single cache security group and its status.

Contents

Note

In the following list, the required parameters are described first.

SecurityGroupId

The identifier of the cache security group.

Type: String

Required: No

Status

The status of the cache security group membership. The status changes whenever a cache security group is modified, or when the cache security groups assigned to a cluster are modified.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ServiceUpdate

An update that you can apply to your Redis clusters.

Contents

Note

In the following list, the required parameters are described first.

AutoUpdateAfterRecommendedApplyByDate

Indicates whether the service update will be automatically applied once the recommended apply-by date has expired.

Type: Boolean

Required: No

Engine

The ElastiCache engine to which the update applies. Either Redis or Memcached

Type: String

Required: No

EngineVersion

The ElastiCache engine version to which the update applies. Either Redis or Memcached engine version

Type: String

Required: No

EstimatedUpdateTime

The estimated length of time the service update will take

Type: String

Required: No

ServiceUpdateDescription

Provides details of the service update

Type: String

Required: No

ServiceUpdateEndDate

The date after which the service update is no longer available

Type: Timestamp

Required: No

ServiceUpdateName

The unique ID of the service update

Type: String

Required: No

ServiceUpdateRecommendedApplyByDate

The recommendend date to apply the service update in order to ensure compliance. For information on compliance, see [Self-Service Security Updates for Compliance](#).

Type: Timestamp

Required: No

ServiceUpdateReleaseDate

The date when the service update is initially available

Type: Timestamp

Required: No

ServiceUpdateSeverity

The severity of the service update

Type: String

Valid Values: `critical` | `important` | `medium` | `low`

Required: No

ServiceUpdateStatus

The status of the service update

Type: String

Valid Values: `available` | `cancelled` | `expired`

Required: No

ServiceUpdateType

Reflects the nature of the service update

Type: String

Valid Values: `security-update`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

SlotMigration

Represents the progress of an online resharding operation.

Contents

Note

In the following list, the required parameters are described first.

ProgressPercentage

The percentage of the slot migration that is complete.

Type: Double

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

Snapshot

Represents a copy of an entire Redis cluster as of the time when the snapshot was taken.

Contents

Note

In the following list, the required parameters are described first.

ARN

The ARN (Amazon Resource Name) of the snapshot.

Type: String

Required: No

AutomaticFailover

Indicates the status of automatic failover for the source Redis replication group.

Type: String

Valid Values: `enabled` | `disabled` | `enabling` | `disabling`

Required: No

AutoMinorVersionUpgrade

This parameter is currently disabled.

Type: Boolean

Required: No

CacheClusterCreateTime

The date and time when the source cluster was created.

Type: Timestamp

Required: No

CacheClusterId

The user-supplied identifier of the source cluster.

Type: String

Required: No

CacheNodeType

The name of the compute and memory capacity node type for the source cluster.

The following node types are supported by ElastiCache. Generally speaking, the current generation types provide more memory and computational power at lower cost when compared to their equivalent previous generation counterparts.

- General purpose:
 - Current generation:

M6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

cache.m6g.large, cache.m6g.xlarge, cache.m6g.2xlarge, cache.m6g.4xlarge, cache.m6g.8xlarge, cache.m6g.12xlarge, cache.m6g.16xlarge

Note

For region availability, see [Supported Node Types](#)

M5 node types: cache.m5.large, cache.m5.xlarge, cache.m5.2xlarge, cache.m5.4xlarge, cache.m5.12xlarge, cache.m5.24xlarge

M4 node types: cache.m4.large, cache.m4.xlarge, cache.m4.2xlarge, cache.m4.4xlarge, cache.m4.10xlarge

T3 node types: cache.t3.micro, cache.t3.small, cache.t3.medium

T2 node types: cache.t2.micro, cache.t2.small, cache.t2.medium

- Previous generation: (not recommended)

T1 node types: cache.t1.micro

M1 node types: cache.m1.small, cache.m1.medium, cache.m1.large, cache.m1.xlarge

M3 node types: cache.m3.medium, cache.m3.large, cache.m3.xlarge, cache.m3.2xlarge

- Compute optimized:
 - Previous generation: (not recommended)

C1 node types: cache.c1.xlarge

- Memory optimized:
 - Current generation:

R6g node types (available only for Redis engine version 5.0.6 onward and for Memcached engine version 1.5.16 onward).

cache.r6g.large, cache.r6g.xlarge, cache.r6g.2xlarge, cache.r6g.4xlarge, cache.r6g.8xlarge, cache.r6g.12xlarge, cache.r6g.16xlarge

Note

For region availability, see [Supported Node Types](#)

R5 node types: cache.r5.large, cache.r5.xlarge, cache.r5.2xlarge, cache.r5.4xlarge, cache.r5.12xlarge, cache.r5.24xlarge

R4 node types: cache.r4.large, cache.r4.xlarge, cache.r4.2xlarge, cache.r4.4xlarge, cache.r4.8xlarge, cache.r4.16xlarge

- Previous generation: (not recommended)

M2 node types: cache.m2.xlarge, cache.m2.2xlarge, cache.m2.4xlarge

R3 node types: cache.r3.large, cache.r3.xlarge, cache.r3.2xlarge, cache.r3.4xlarge, cache.r3.8xlarge

Additional node type info

- All current generation instance types are created in Amazon VPC by default.
- Redis append-only files (AOF) are not supported for T1 or T2 instances.

- Redis Multi-AZ with automatic failover is not supported on T1 instances.
- Redis configuration variables `appendonly` and `appendfsync` are not supported on Redis version 2.8.22 and later.

Type: String

Required: No

CacheParameterGroupName

The cache parameter group that is associated with the source cluster.

Type: String

Required: No

CacheSubnetGroupName

The name of the cache subnet group associated with the source cluster.

Type: String

Required: No

Engine

The name of the cache engine (`memcached` or `redis`) used by the source cluster.

Type: String

Required: No

EngineVersion

The version of the cache engine version that is used by the source cluster.

Type: String

Required: No

KmsKeyId

The ID of the KMS key used to encrypt the snapshot.

Type: String

Required: No

NodeSnapshots.NodeSnapshot.N

A list of the cache nodes in the source cluster.

Type: Array of [NodeSnapshot](#) (p. 279) objects

Required: No

NumCacheNodes

The number of cache nodes in the source cluster.

For clusters running Redis, this value must be 1. For clusters running Memcached, this value must be between 1 and 40.

Type: Integer

Required: No

NumNodeGroups

The number of node groups (shards) in this snapshot. When restoring from a snapshot, the number of node groups (shards) in the snapshot and in the restored replication group must be the same.

Type: Integer

Required: No

Port

The port number used by each cache nodes in the source cluster.

Type: Integer

Required: No

PreferredAvailabilityZone

The name of the Availability Zone in which the source cluster is located.

Type: String

Required: No

PreferredMaintenanceWindow

Specifies the weekly time range during which maintenance on the cluster is performed. It is specified as a range in the format ddd:hh24:mi-ddd:hh24:mi (24H Clock UTC). The minimum maintenance window is a 60 minute period.

Valid values for ddd are:

- sun
- mon
- tue
- wed
- thu
- fri
- sat

Example: sun:23:00-mon:01:30

Type: String

Required: No

PreferredOutpostArn

The ARN (Amazon Resource Name) of the preferred outpost.

Type: String

Required: No

ReplicationGroupDescription

A description of the source replication group.

Type: String

Required: No

ReplicationGroupId

The unique identifier of the source replication group.

Type: String

Required: No

SnapshotName

The name of a snapshot. For an automatic snapshot, the name is system-generated. For a manual snapshot, this is the user-provided name.

Type: String

Required: No

SnapshotRetentionLimit

For an automatic snapshot, the number of days for which ElastiCache retains the snapshot before deleting it.

For manual snapshots, this field reflects the `SnapshotRetentionLimit` for the source cluster when the snapshot was created. This field is otherwise ignored: Manual snapshots do not expire, and can only be deleted using the `DeleteSnapshot` operation.

Important If the value of `SnapshotRetentionLimit` is set to zero (0), backups are turned off.

Type: Integer

Required: No

SnapshotSource

Indicates whether the snapshot is from an automatic backup (`automated`) or was created manually (`manual`).

Type: String

Required: No

SnapshotStatus

The status of the snapshot. Valid values: `creating` | `available` | `restoring` | `copying` | `deleting`.

Type: String

Required: No

SnapshotWindow

The daily time range during which ElastiCache takes daily snapshots of the source cluster.

Type: String

Required: No

TopicArn

The Amazon Resource Name (ARN) for the topic used by the source cluster for publishing notifications.

Type: String

Required: No

VpcId

The Amazon Virtual Private Cloud identifier (VPC ID) of the cache subnet group for the source cluster.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

Subnet

Represents the subnet associated with a cluster. This parameter refers to subnets defined in Amazon Virtual Private Cloud (Amazon VPC) and used with ElastiCache.

Contents

Note

In the following list, the required parameters are described first.

SubnetAvailabilityZone

The Availability Zone associated with the subnet.

Type: [AvailabilityZone](#) (p. 223) object

Required: No

SubnetIdentifier

The unique identifier for the subnet.

Type: String

Required: No

SubnetOutpost

The outpost ARN of the subnet.

Type: [SubnetOutpost](#) (p. 318) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

SubnetOutpost

The ID of the outpost subnet.

Contents

Note

In the following list, the required parameters are described first.

SubnetOutpostArn

The outpost ARN of the subnet.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

Tag

A tag that can be added to an ElastiCache cluster or replication group. Tags are composed of a Key/Value pair. You can use tags to categorize and track all your ElastiCache resources, with the exception of global replication group. When you add or remove tags on replication groups, those actions will be replicated to all nodes in the replication group. A tag with a null Value is permitted.

Contents

Note

In the following list, the required parameters are described first.

Key

The key for the tag. May not be null.

Type: String

Required: No

Value

The tag's value. May be null.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

TimeRangeFilter

Filters update actions from the service updates that are in available status during the time range.

Contents

Note

In the following list, the required parameters are described first.

EndTime

The end time of the time range filter

Type: Timestamp

Required: No

StartTime

The start time of the time range filter

Type: Timestamp

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

UnprocessedUpdateAction

Update action that has failed to be processed for the corresponding apply/stop request

Contents

Note

In the following list, the required parameters are described first.

CacheClusterId

The ID of the cache cluster

Type: String

Required: No

ErrorMessage

The error message that describes the reason the request was not processed

Type: String

Required: No

ErrorType

The error type for requests that are not processed

Type: String

Required: No

ReplicationGroupId

The replication group ID

Type: String

Required: No

ServiceUpdateName

The unique ID of the service update

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

UpdateAction

The status of the service update for a specific replication group

Contents

Note

In the following list, the required parameters are described first.

CacheClusterId

The ID of the cache cluster

Type: String

Required: No

CacheNodeUpdateStatus.CacheNodeUpdateStatus.N

The status of the service update on the cache node

Type: Array of [CacheNodeUpdateStatus](#) (p. 238) objects

Required: No

Engine

The Elasticache engine to which the update applies. Either Redis or Memcached

Type: String

Required: No

EstimatedUpdateTime

The estimated length of time for the update to complete

Type: String

Required: No

NodeGroupUpdateStatus.NodeGroupUpdateStatus.N

The status of the service update on the node group

Type: Array of [NodeGroupUpdateStatus](#) (p. 278) objects

Required: No

NodesUpdated

The progress of the service update on the replication group

Type: String

Required: No

ReplicationGroupId

The ID of the replication group

Type: String

Required: No

ServiceUpdateName

The unique ID of the service update

Type: String

Required: No

ServiceUpdateRecommendedApplyByDate

The recommended date to apply the service update to ensure compliance. For information on compliance, see [Self-Service Security Updates for Compliance](#).

Type: Timestamp

Required: No

ServiceUpdateReleaseDate

The date the update is first available

Type: Timestamp

Required: No

ServiceUpdateSeverity

The severity of the service update

Type: String

Valid Values: `critical` | `important` | `medium` | `low`

Required: No

ServiceUpdateStatus

The status of the service update

Type: String

Valid Values: `available` | `cancelled` | `expired`

Required: No

ServiceUpdateType

Reflects the nature of the service update

Type: String

Valid Values: `security-update`

Required: No

SlaMet

If yes, all nodes in the replication group have been updated by the recommended apply-by date. If no, at least one node in the replication group have not been updated by the recommended apply-by date. If N/A, the replication group was created after the recommended apply-by date.

Type: String

Valid Values: `yes` | `no` | `n/a`

Required: No

UpdateActionAvailableDate

The date that the service update is available to a replication group

Type: Timestamp

Required: No

UpdateActionStatus

The status of the update action

Type: String

Valid Values: not-applied | waiting-to-start | in-progress | stopping | stopped
| complete | scheduling | scheduled | not-applicable

Required: No

UpdateActionStatusModifiedDate

The date when the UpdateActionStatus was last modified

Type: Timestamp

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

User

Contents

Note

In the following list, the required parameters are described first.

AccessString

Access permissions string used for this user.

Type: String

Required: No

ARN

The Amazon Resource Name (ARN) of the user.

Type: String

Required: No

Authentication

Denotes whether the user requires a password to authenticate.

Type: [Authentication](#) (p. 222) object

Required: No

Engine

The current supported value is Redis.

Type: String

Pattern: [a-zA-Z]*

Required: No

Status

Indicates the user status. Can be "active", "modifying" or "deleting".

Type: String

Required: No

UserGroupIds.member.N

Returns a list of the user group IDs the user belongs to.

Type: Array of strings

Length Constraints: Minimum length of 1.

Pattern: [a-zA-Z][a-zA-Z0-9\-*]

Required: No

UserId

The ID of the user.

Type: String

Required: No

UserName

The username of the user.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

UserGroup

Contents

Note

In the following list, the required parameters are described first.

ARN

The Amazon Resource Name (ARN) of the user group.

Type: String

Required: No

Engine

The current supported value is Redis.

Type: String

Pattern: [a-zA-Z]*

Required: No

PendingChanges

A list of updates being applied to the user group.

Type: [UserGroupPendingChanges](#) (p. 329) object

Required: No

ReplicationGroups.member.N

A list of replication groups that the user group can access.

Type: Array of strings

Required: No

Status

Indicates user group status. Can be "creating", "active", "modifying", "deleting".

Type: String

Required: No

UserGroupId

The ID of the user group.

Type: String

Required: No

UserIds.member.N

The list of user IDs that belong to the user group.

Type: Array of strings

Length Constraints: Minimum length of 1.

Pattern: `[a-zA-Z][a-zA-Z0-9\-*]`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

UserGroupPendingChanges

Returns the updates being applied to the user group.

Contents

Note

In the following list, the required parameters are described first.

UserIdsToAdd.member.N

The list of user IDs to add.

Type: Array of strings

Length Constraints: Minimum length of 1.

Pattern: `[a-zA-Z][a-zA-Z0-9\-\]*`

Required: No

UserIdsToRemove.member.N

The list of user IDs to remove.

Type: Array of strings

Length Constraints: Minimum length of 1.

Pattern: `[a-zA-Z][a-zA-Z0-9\-\]*`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

UserGroupsUpdateStatus

The status of the user group update.

Contents

Note

In the following list, the required parameters are described first.

UserGroupIdsToAdd.member.N

The ID of the user group to add.

Type: Array of strings

Length Constraints: Minimum length of 1.

Pattern: `[a-zA-Z][a-zA-Z0-9\-\-]*`

Required: No

UserGroupIdsToRemove.member.N

The ID of the user group to remove.

Type: Array of strings

Length Constraints: Minimum length of 1.

Pattern: `[a-zA-Z][a-zA-Z0-9\-\-]*`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

Common Parameters

The following list contains the parameters that all actions use for signing Signature Version 4 requests with a query string. Any action-specific parameters are listed in the topic for that action. For more information about Signature Version 4, see [Signature Version 4 Signing Process](#) in the *Amazon Web Services General Reference*.

Action

The action to be performed.

Type: string

Required: Yes

Version

The API version that the request is written for, expressed in the format YYYY-MM-DD.

Type: string

Required: Yes

X-Amz-Algorithm

The hash algorithm that you used to create the request signature.

Condition: Specify this parameter when you include authentication information in a query string instead of in the HTTP authorization header.

Type: string

Valid Values: `AWS4-HMAC-SHA256`

Required: Conditional

X-Amz-Credential

The credential scope value, which is a string that includes your access key, the date, the region you are targeting, the service you are requesting, and a termination string ("aws4_request"). The value is expressed in the following format: `access_key/YYYYMMDD/region/service/aws4_request`.

For more information, see [Task 2: Create a String to Sign for Signature Version 4](#) in the *Amazon Web Services General Reference*.

Condition: Specify this parameter when you include authentication information in a query string instead of in the HTTP authorization header.

Type: string

Required: Conditional

X-Amz-Date

The date that is used to create the signature. The format must be ISO 8601 basic format (YYYYMMDD'THHMMSS'Z'). For example, the following date time is a valid X-Amz-Date value: `20120325T120000Z`.

Condition: X-Amz-Date is optional for all requests; it can be used to override the date used for signing requests. If the Date header is specified in the ISO 8601 basic format, X-Amz-Date is

not required. When X-Amz-Date is used, it always overrides the value of the Date header. For more information, see [Handling Dates in Signature Version 4](#) in the *Amazon Web Services General Reference*.

Type: string

Required: Conditional

X-Amz-Security-Token

The temporary security token that was obtained through a call to AWS Security Token Service (AWS STS). For a list of services that support temporary security credentials from AWS Security Token Service, go to [AWS Services That Work with IAM](#) in the *IAM User Guide*.

Condition: If you're using temporary security credentials from the AWS Security Token Service, you must include the security token.

Type: string

Required: Conditional

X-Amz-Signature

Specifies the hex-encoded signature that was calculated from the string to sign and the derived signing key.

Condition: Specify this parameter when you include authentication information in a query string instead of in the HTTP authorization header.

Type: string

Required: Conditional

X-Amz-SignedHeaders

Specifies all the HTTP headers that were included as part of the canonical request. For more information about specifying signed headers, see [Task 1: Create a Canonical Request For Signature Version 4](#) in the *Amazon Web Services General Reference*.

Condition: Specify this parameter when you include authentication information in a query string instead of in the HTTP authorization header.

Type: string

Required: Conditional

Common Errors

This section lists the errors common to the API actions of all AWS services. For errors specific to an API action for this service, see the topic for that API action.

AccessDeniedException

You do not have sufficient access to perform this action.

HTTP Status Code: 400

IncompleteSignature

The request signature does not conform to AWS standards.

HTTP Status Code: 400

InternalFailure

The request processing has failed because of an unknown error, exception or failure.

HTTP Status Code: 500

InvalidAction

The action or operation requested is invalid. Verify that the action is typed correctly.

HTTP Status Code: 400

InvalidClientTokenId

The X.509 certificate or AWS access key ID provided does not exist in our records.

HTTP Status Code: 403

InvalidParameterCombination

Parameters that must not be used together were used together.

HTTP Status Code: 400

InvalidParameterValue

An invalid or out-of-range value was supplied for the input parameter.

HTTP Status Code: 400

InvalidQueryParameter

The AWS query string is malformed or does not adhere to AWS standards.

HTTP Status Code: 400

MalformedQueryString

The query string contains a syntax error.

HTTP Status Code: 404

MissingAction

The request is missing an action or a required parameter.

HTTP Status Code: 400

MissingAuthenticationToken

The request must contain either a valid (registered) AWS access key ID or X.509 certificate.

HTTP Status Code: 403

MissingParameter

A required parameter for the specified action is not supplied.

HTTP Status Code: 400

NotAuthorized

You do not have permission to perform this action.

HTTP Status Code: 400

OptInRequired

The AWS access key ID needs a subscription for the service.

HTTP Status Code: 403

RequestExpired

The request reached the service more than 15 minutes after the date stamp on the request or more than 15 minutes after the request expiration date (such as for pre-signed URLs), or the date stamp on the request is more than 15 minutes in the future.

HTTP Status Code: 400

ServiceUnavailable

The request has failed due to a temporary failure of the server.

HTTP Status Code: 503

ThrottlingException

The request was denied due to request throttling.

HTTP Status Code: 400

ValidationError

The input fails to satisfy the constraints specified by an AWS service.

HTTP Status Code: 400