
Amazon Kinesis Data Firehose

API Reference

API Version 2015-08-04



Amazon Kinesis Data Firehose: API Reference

Copyright © Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

Table of Contents

Welcome	1
Actions	2
CreateDeliveryStream	3
Request Syntax	4
Request Parameters	10
Response Syntax	12
Response Elements	12
Errors	12
Examples	13
See Also	13
DeleteDeliveryStream	15
Request Syntax	15
Request Parameters	15
Response Elements	15
Errors	16
Examples	16
See Also	16
DescribeDeliveryStream	18
Request Syntax	18
Request Parameters	18
Response Syntax	19
Response Elements	25
Errors	25
Examples	26
See Also	27
ListDeliveryStreams	28
Request Syntax	28
Request Parameters	28
Response Syntax	29
Response Elements	29
Errors	29
Examples	29
See Also	30
ListTagsForDeliveryStream	31
Request Syntax	31
Request Parameters	31
Response Syntax	31
Response Elements	32
Errors	32
Examples	32
See Also	33
PutRecord	34
Request Syntax	34
Request Parameters	34
Response Syntax	35
Response Elements	35
Errors	35
Examples	36
See Also	36
PutRecordBatch	38
Request Syntax	38
Request Parameters	39
Response Syntax	39
Response Elements	39

Errors	40
Examples	40
See Also	41
StartDeliveryStreamEncryption	43
Request Syntax	43
Request Parameters	43
Response Elements	44
Errors	44
Examples	44
See Also	45
StopDeliveryStreamEncryption	46
Request Syntax	46
Request Parameters	46
Response Elements	46
Errors	46
Examples	47
See Also	47
TagDeliveryStream	49
Request Syntax	49
Request Parameters	49
Response Elements	49
Errors	49
Examples	50
See Also	51
UntagDeliveryStream	52
Request Syntax	52
Request Parameters	52
Response Elements	52
Errors	52
Examples	53
See Also	53
UpdateDestination	55
Request Syntax	55
Request Parameters	61
Response Elements	62
Errors	63
Examples	63
See Also	64
Data Types	65
BufferingHints	67
Contents	67
See Also	67
CloudWatchLoggingOptions	68
Contents	68
See Also	68
CopyCommand	69
Contents	69
See Also	69
DataFormatConversionConfiguration	71
Contents	71
See Also	71
DeliveryStreamDescription	72
Contents	72
See Also	74
DeliveryStreamEncryptionConfiguration	75
Contents	75
See Also	75

DeliveryStreamEncryptionConfigurationInput	77
Contents	77
See Also	77
Deserializer	79
Contents	79
See Also	79
DestinationDescription	80
Contents	80
See Also	81
DynamicPartitioningConfiguration	82
Contents	82
See Also	82
ElasticsearchBufferingHints	83
Contents	83
See Also	83
ElasticsearchDestinationConfiguration	84
Contents	84
See Also	86
ElasticsearchDestinationDescription	87
Contents	87
See Also	89
ElasticsearchDestinationUpdate	90
Contents	90
See Also	92
ElasticsearchRetryOptions	93
Contents	93
See Also	93
EncryptionConfiguration	94
Contents	94
See Also	94
ExtendedS3DestinationConfiguration	95
Contents	95
See Also	97
ExtendedS3DestinationDescription	98
Contents	98
See Also	100
ExtendedS3DestinationUpdate	101
Contents	101
See Also	103
FailureDescription	104
Contents	104
See Also	104
HiveJsonSerDe	105
Contents	105
See Also	105
HttpEndpointBufferingHints	106
Contents	106
See Also	106
HttpEndpointCommonAttribute	107
Contents	107
See Also	107
HttpEndpointConfiguration	108
Contents	108
See Also	108
HttpEndpointDescription	109
Contents	109
See Also	109

HttpEndpointDestinationConfiguration	110
Contents	110
See Also	111
HttpEndpointDestinationDescription	112
Contents	112
See Also	113
HttpEndpointDestinationUpdate	114
Contents	114
See Also	115
HttpEndpointRequestConfiguration	116
Contents	116
See Also	116
HttpEndpointRetryOptions	117
Contents	117
See Also	117
InputFormatConfiguration	118
Contents	118
See Also	118
KinesisStreamSourceConfiguration	119
Contents	119
See Also	119
KinesisStreamSourceDescription	120
Contents	120
See Also	120
KMSEncryptionConfig	121
Contents	121
See Also	121
OpenXJsonSerDe	122
Contents	122
See Also	122
OrcSerDe	124
Contents	124
See Also	125
OutputFormatConfiguration	127
Contents	127
See Also	127
ParquetSerDe	128
Contents	128
See Also	129
ProcessingConfiguration	130
Contents	130
See Also	130
Processor	131
Contents	131
See Also	131
ProcessorParameter	132
Contents	132
See Also	132
PutRecordBatchResponseEntry	133
Contents	133
See Also	133
Record	134
Contents	134
See Also	134
RedshiftDestinationConfiguration	135
Contents	135
See Also	136

RedshiftDestinationDescription	138
Contents	138
See Also	139
RedshiftDestinationUpdate	140
Contents	140
See Also	141
RedshiftRetryOptions	143
Contents	143
See Also	143
RetryOptions	144
Contents	144
See Also	144
S3DestinationConfiguration	145
Contents	145
See Also	146
S3DestinationDescription	147
Contents	147
See Also	148
S3DestinationUpdate	149
Contents	149
See Also	150
SchemaConfiguration	151
Contents	151
See Also	152
Serializer	153
Contents	153
See Also	153
SourceDescription	154
Contents	154
See Also	154
SplunkDestinationConfiguration	155
Contents	155
See Also	156
SplunkDestinationDescription	157
Contents	157
See Also	158
SplunkDestinationUpdate	159
Contents	159
See Also	160
SplunkRetryOptions	161
Contents	161
See Also	161
Tag	162
Contents	162
See Also	162
VpcConfiguration	163
Contents	163
See Also	164
VpcConfigurationDescription	165
Contents	165
See Also	166
Common Errors	167

Welcome

Amazon Kinesis Data Firehose is a fully managed service that delivers real-time streaming data to destinations such as Amazon Simple Storage Service (Amazon S3), Amazon Elasticsearch Service (Amazon ES), Amazon Redshift, and Splunk.

This document was last published on October 6, 2021.

Actions

The following actions are supported:

- [CreateDeliveryStream](#) (p. 3)
- [DeleteDeliveryStream](#) (p. 15)
- [DescribeDeliveryStream](#) (p. 18)
- [ListDeliveryStreams](#) (p. 28)
- [ListTagsForDeliveryStream](#) (p. 31)
- [PutRecord](#) (p. 34)
- [PutRecordBatch](#) (p. 38)
- [StartDeliveryStreamEncryption](#) (p. 43)
- [StopDeliveryStreamEncryption](#) (p. 46)
- [TagDeliveryStream](#) (p. 49)
- [UntagDeliveryStream](#) (p. 52)
- [UpdateDestination](#) (p. 55)

CreateDeliveryStream

Creates a Kinesis Data Firehose delivery stream.

By default, you can create up to 50 delivery streams per AWS Region.

This is an asynchronous operation that immediately returns. The initial status of the delivery stream is `CREATING`. After the delivery stream is created, its status is `ACTIVE` and it now accepts data. If the delivery stream creation fails, the status transitions to `CREATING_FAILED`. Attempts to send data to a delivery stream that is not in the `ACTIVE` state cause an exception. To check the state of a delivery stream, use [DescribeDeliveryStream](#) (p. 18).

If the status of a delivery stream is `CREATING_FAILED`, this status doesn't change, and you can't invoke `CreateDeliveryStream` again on it. However, you can invoke the [DeleteDeliveryStream](#) (p. 15) operation to delete it.

A Kinesis Data Firehose delivery stream can be configured to receive records directly from providers using [PutRecord](#) (p. 34) or [PutRecordBatch](#) (p. 38), or it can be configured to use an existing Kinesis stream as its source. To specify a Kinesis data stream as input, set the `DeliveryStreamType` parameter to `KinesisStreamAsSource`, and provide the Kinesis stream Amazon Resource Name (ARN) and role ARN in the `KinesisStreamSourceConfiguration` parameter.

To create a delivery stream with server-side encryption (SSE) enabled, include [DeliveryStreamEncryptionConfigurationInput](#) (p. 77) in your request. This is optional. You can also invoke [StartDeliveryStreamEncryption](#) (p. 43) to turn on SSE for an existing delivery stream that doesn't have SSE enabled.

A delivery stream is configured with a single destination: Amazon S3, Amazon ES, Amazon Redshift, or Splunk. You must specify only one of the following destination configuration parameters: `ExtendedS3DestinationConfiguration`, `S3DestinationConfiguration`, `ElasticsearchDestinationConfiguration`, `RedshiftDestinationConfiguration`, or `SplunkDestinationConfiguration`.

When you specify `S3DestinationConfiguration`, you can also provide the following optional values: `BufferingHints`, `EncryptionConfiguration`, and `CompressionFormat`. By default, if no `BufferingHints` value is provided, Kinesis Data Firehose buffers data up to 5 MB or for 5 minutes, whichever condition is satisfied first. `BufferingHints` is a hint, so there are some cases where the service cannot adhere to these conditions strictly. For example, record boundaries might be such that the size is a little over or under the configured buffering size. By default, no encryption is performed. We strongly recommend that you enable encryption to ensure secure data storage in Amazon S3.

A few notes about Amazon Redshift as a destination:

- An Amazon Redshift destination requires an S3 bucket as intermediate location. Kinesis Data Firehose first delivers data to Amazon S3 and then uses `COPY` syntax to load data into an Amazon Redshift table. This is specified in the `RedshiftDestinationConfiguration.S3Configuration` parameter.
- The compression formats `SNAPPY` or `ZIP` cannot be specified in `RedshiftDestinationConfiguration.S3Configuration` because the Amazon Redshift `COPY` operation that reads from the S3 bucket doesn't support these compression formats.
- We strongly recommend that you use the user name and password you provide exclusively with Kinesis Data Firehose, and that the permissions for the account are restricted for Amazon Redshift `INSERT` permissions.

Kinesis Data Firehose assumes the IAM role that is configured as part of the destination. The role should allow the Kinesis Data Firehose principal to assume the role, and the role should have permissions that

allow the service to deliver the data. For more information, see [Grant Kinesis Data Firehose Access to an Amazon S3 Destination](#) in the *Amazon Kinesis Data Firehose Developer Guide*.

Request Syntax

```
{
  "DeliveryStreamEncryptionConfigurationInput": {
    "KeyARN": "string",
    "KeyType": "string"
  },
  "DeliveryStreamName": "string",
  "DeliveryStreamType": "string",
  "ElasticsearchDestinationConfiguration": {
    "BufferingHints": {
      "IntervalInSeconds": number,
      "SizeInMBs": number
    },
    "CloudWatchLoggingOptions": {
      "Enabled": boolean,
      "LogGroupName": "string",
      "LogStreamName": "string"
    },
    "ClusterEndpoint": "string",
    "DomainARN": "string",
    "IndexName": "string",
    "IndexRotationPeriod": "string",
    "ProcessingConfiguration": {
      "Enabled": boolean,
      "Processors": [
        {
          "Parameters": [
            {
              "ParameterName": "string",
              "ParameterValue": "string"
            }
          ],
          "Type": "string"
        }
      ]
    },
    "RetryOptions": {
      "DurationInSeconds": number
    },
    "RoleARN": "string",
    "S3BackupMode": "string",
    "S3Configuration": {
      "BucketARN": "string",
      "BufferingHints": {
        "IntervalInSeconds": number,
        "SizeInMBs": number
      },
      "CloudWatchLoggingOptions": {
        "Enabled": boolean,
        "LogGroupName": "string",
        "LogStreamName": "string"
      },
      "CompressionFormat": "string",
      "EncryptionConfiguration": {
        "KMSEncryptionConfig": {
          "AWSKMSKeyARN": "string"
        },
        "NoEncryptionConfig": "string"
      },
      "ErrorOutputPrefix": "string",
```

```
    "Prefix": "string",
    "RoleARN": "string"
  },
  "TypeName": "string",
  "VpcConfiguration": {
    "RoleARN": "string",
    "SecurityGroupIds": [ "string" ],
    "SubnetIds": [ "string" ]
  }
},
"ExtendedS3DestinationConfiguration": {
  "BucketARN": "string",
  "BufferingHints": {
    "IntervalInSeconds": number,
    "SizeInMBs": number
  },
  "CloudWatchLoggingOptions": {
    "Enabled": boolean,
    "LogGroupName": "string",
    "LogStreamName": "string"
  },
  "CompressionFormat": "string",
  "DataFormatConversionConfiguration": {
    "Enabled": boolean,
    "InputFormatConfiguration": {
      "Deserializer": {
        "HiveJsonSerDe": {
          "TimestampFormats": [ "string" ]
        },
        "OpenXJsonSerDe": {
          "CaseInsensitive": boolean,
          "ColumnToJsonKeyMappings": {
            "string" : "string"
          },
          "ConvertDotsInJsonKeysToUnderscores": boolean
        }
      }
    },
    "OutputFormatConfiguration": {
      "Serializer": {
        "OrcSerDe": {
          "BlockSizeBytes": number,
          "BloomFilterColumns": [ "string" ],
          "BloomFilterFalsePositiveProbability": number,
          "Compression": "string",
          "DictionaryKeyThreshold": number,
          "EnablePadding": boolean,
          "FormatVersion": "string",
          "PaddingTolerance": number,
          "RowIndexStride": number,
          "StripeSizeBytes": number
        },
        "ParquetSerDe": {
          "BlockSizeBytes": number,
          "Compression": "string",
          "EnableDictionaryCompression": boolean,
          "MaxPaddingBytes": number,
          "PageSizeBytes": number,
          "WriterVersion": "string"
        }
      }
    }
  },
  "SchemaConfiguration": {
    "CatalogId": "string",
    "DatabaseName": "string",
    "Region": "string",
```

```
        "RoleARN": "string",
        "TableName": "string",
        "VersionId": "string"
    },
    "DynamicPartitioningConfiguration": {
        "Enabled": boolean,
        "RetryOptions": {
            "DurationInSeconds": number
        }
    },
    "EncryptionConfiguration": {
        "KMSEncryptionConfig": {
            "AWSKMSKeyARN": "string"
        },
        "NoEncryptionConfig": "string"
    },
    "ErrorOutputPrefix": "string",
    "Prefix": "string",
    "ProcessingConfiguration": {
        "Enabled": boolean,
        "Processors": [
            {
                "Parameters": [
                    {
                        "ParameterName": "string",
                        "ParameterValue": "string"
                    }
                ],
                "Type": "string"
            }
        ]
    },
    "RoleARN": "string",
    "S3BackupConfiguration": {
        "BucketARN": "string",
        "BufferingHints": {
            "IntervalInSeconds": number,
            "SizeInMBs": number
        },
        "CloudWatchLoggingOptions": {
            "Enabled": boolean,
            "LogGroupName": "string",
            "LogStreamName": "string"
        },
        "CompressionFormat": "string",
        "EncryptionConfiguration": {
            "KMSEncryptionConfig": {
                "AWSKMSKeyARN": "string"
            },
            "NoEncryptionConfig": "string"
        },
        "ErrorOutputPrefix": "string",
        "Prefix": "string",
        "RoleARN": "string"
    },
    "S3BackupMode": "string"
},
"HttpEndpointDestinationConfiguration": {
    "BufferingHints": {
        "IntervalInSeconds": number,
        "SizeInMBs": number
    },
    "CloudWatchLoggingOptions": {
        "Enabled": boolean,
        "LogGroupName": "string",
```

```
    "LogStreamName": "string"
  },
  "EndpointConfiguration": {
    "AccessKey": "string",
    "Name": "string",
    "Url": "string"
  },
  "ProcessingConfiguration": {
    "Enabled": boolean,
    "Processors": [
      {
        "Parameters": [
          {
            "ParameterName": "string",
            "ParameterValue": "string"
          }
        ],
        "Type": "string"
      }
    ]
  },
  "RequestConfiguration": {
    "CommonAttributes": [
      {
        "AttributeName": "string",
        "AttributeValue": "string"
      }
    ],
    "ContentEncoding": "string"
  },
  "RetryOptions": {
    "DurationInSeconds": number
  },
  "RoleARN": "string",
  "S3BackupMode": "string",
  "S3Configuration": {
    "BucketARN": "string",
    "BufferingHints": {
      "IntervalInSeconds": number,
      "SizeInMBs": number
    },
    "CloudWatchLoggingOptions": {
      "Enabled": boolean,
      "LogGroupName": "string",
      "LogStreamName": "string"
    },
    "CompressionFormat": "string",
    "EncryptionConfiguration": {
      "KMSEncryptionConfig": {
        "AWSKMSKeyARN": "string"
      },
      "NoEncryptionConfig": "string"
    },
    "ErrorOutputPrefix": "string",
    "Prefix": "string",
    "RoleARN": "string"
  }
},
"KinesisStreamSourceConfiguration": {
  "KinesisStreamARN": "string",
  "RoleARN": "string"
},
"RedshiftDestinationConfiguration": {
  "CloudWatchLoggingOptions": {
    "Enabled": boolean,
    "LogGroupName": "string",
```

```
    "LogStreamName": "string"
  },
  "ClusterJDBCURL": "string",
  "CopyCommand": {
    "CopyOptions": "string",
    "DataTableColumns": "string",
    "DataTableName": "string"
  },
  "Password": "string",
  "ProcessingConfiguration": {
    "Enabled": boolean,
    "Processors": [
      {
        "Parameters": [
          {
            "ParameterName": "string",
            "ParameterValue": "string"
          }
        ],
        "Type": "string"
      }
    ]
  },
  "RetryOptions": {
    "DurationInSeconds": number
  },
  "RoleARN": "string",
  "S3BackupConfiguration": {
    "BucketARN": "string",
    "BufferingHints": {
      "IntervalInSeconds": number,
      "SizeInMBs": number
    },
    "CloudWatchLoggingOptions": {
      "Enabled": boolean,
      "LogGroupName": "string",
      "LogStreamName": "string"
    },
    "CompressionFormat": "string",
    "EncryptionConfiguration": {
      "KMSEncryptionConfig": {
        "AWSKMSKeyARN": "string"
      },
      "NoEncryptionConfig": "string"
    },
    "ErrorOutputPrefix": "string",
    "Prefix": "string",
    "RoleARN": "string"
  },
  "S3BackupMode": "string",
  "S3Configuration": {
    "BucketARN": "string",
    "BufferingHints": {
      "IntervalInSeconds": number,
      "SizeInMBs": number
    },
    "CloudWatchLoggingOptions": {
      "Enabled": boolean,
      "LogGroupName": "string",
      "LogStreamName": "string"
    },
    "CompressionFormat": "string",
    "EncryptionConfiguration": {
      "KMSEncryptionConfig": {
        "AWSKMSKeyARN": "string"
      }
    }
  },
```

```

        "NoEncryptionConfig": "string"
    },
    "ErrorOutputPrefix": "string",
    "Prefix": "string",
    "RoleARN": "string"
},
"Username": "string"
},
"S3DestinationConfiguration": {
    "BucketARN": "string",
    "BufferingHints": {
        "IntervalInSeconds": number,
        "SizeInMBs": number
    },
    "CloudWatchLoggingOptions": {
        "Enabled": boolean,
        "LogGroupName": "string",
        "LogStreamName": "string"
    },
    "CompressionFormat": "string",
    "EncryptionConfiguration": {
        "KMSEncryptionConfig": {
            "AWSKMSKeyARN": "string"
        },
        "NoEncryptionConfig": "string"
    },
    "ErrorOutputPrefix": "string",
    "Prefix": "string",
    "RoleARN": "string"
},
"SplunkDestinationConfiguration": {
    "CloudWatchLoggingOptions": {
        "Enabled": boolean,
        "LogGroupName": "string",
        "LogStreamName": "string"
    },
    "HECAcknowledgmentTimeoutInSeconds": number,
    "HECEndpoint": "string",
    "HECEndpointType": "string",
    "HECToken": "string",
    "ProcessingConfiguration": {
        "Enabled": boolean,
        "Processors": [
            {
                "Parameters": [
                    {
                        "ParameterName": "string",
                        "ParameterValue": "string"
                    }
                ],
                "Type": "string"
            }
        ]
    },
    "RetryOptions": {
        "DurationInSeconds": number
    },
    "S3BackupMode": "string",
    "S3Configuration": {
        "BucketARN": "string",
        "BufferingHints": {
            "IntervalInSeconds": number,
            "SizeInMBs": number
        },
        "CloudWatchLoggingOptions": {
            "Enabled": boolean,

```



```
    "LogGroupName": "string",
    "LogStreamName": "string"
  },
  "CompressionFormat": "string",
  "EncryptionConfiguration": {
    "KMSEncryptionConfig": {
      "AWSKMSKeyARN": "string"
    },
    "NoEncryptionConfig": "string"
  },
  "ErrorOutputPrefix": "string",
  "Prefix": "string",
  "RoleARN": "string"
}
},
"Tags": [
  {
    "Key": "string",
    "Value": "string"
  }
]
}
```

Request Parameters

The request accepts the following data in JSON format.

DeliveryStreamEncryptionConfigurationInput (p. 4)

Used to specify the type and Amazon Resource Name (ARN) of the KMS key needed for Server-Side Encryption (SSE).

Type: [DeliveryStreamEncryptionConfigurationInput](#) (p. 77) object

Required: No

DeliveryStreamName (p. 4)

The name of the delivery stream. This name must be unique per AWS account in the same AWS Region. If the delivery streams are in different accounts or different Regions, you can have multiple delivery streams with the same name.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 64.

Pattern: [a-zA-Z0-9_.-]+

Required: Yes

DeliveryStreamType (p. 4)

The delivery stream type. This parameter can be one of the following values:

- `DirectPut`: Provider applications access the delivery stream directly.
- `KinesisStreamAsSource`: The delivery stream uses a Kinesis data stream as a source.

Type: String

Valid Values: `DirectPut` | `KinesisStreamAsSource`

Required: No

ElasticsearchDestinationConfiguration (p. 4)

The destination in Amazon ES. You can specify only one destination.

Type: [ElasticsearchDestinationConfiguration \(p. 84\)](#) object

Required: No

ExtendedS3DestinationConfiguration (p. 4)

The destination in Amazon S3. You can specify only one destination.

Type: [ExtendedS3DestinationConfiguration \(p. 95\)](#) object

Required: No

HttpEndpointDestinationConfiguration (p. 4)

Enables configuring Kinesis Firehose to deliver data to any HTTP endpoint destination. You can specify only one destination.

Type: [HttpEndpointDestinationConfiguration \(p. 110\)](#) object

Required: No

KinesisStreamSourceConfiguration (p. 4)

When a Kinesis data stream is used as the source for the delivery stream, a [KinesisStreamSourceConfiguration \(p. 119\)](#) containing the Kinesis data stream Amazon Resource Name (ARN) and the role ARN for the source stream.

Type: [KinesisStreamSourceConfiguration \(p. 119\)](#) object

Required: No

RedshiftDestinationConfiguration (p. 4)

The destination in Amazon Redshift. You can specify only one destination.

Type: [RedshiftDestinationConfiguration \(p. 135\)](#) object

Required: No

S3DestinationConfiguration (p. 4)

[Deprecated] The destination in Amazon S3. You can specify only one destination.

Type: [S3DestinationConfiguration \(p. 145\)](#) object

Required: No

SplunkDestinationConfiguration (p. 4)

The destination in Splunk. You can specify only one destination.

Type: [SplunkDestinationConfiguration \(p. 155\)](#) object

Required: No

Tags (p. 4)

A set of tags to assign to the delivery stream. A tag is a key-value pair that you can define and assign to AWS resources. Tags are metadata. For example, you can add friendly names and descriptions or other types of information that can help you distinguish the delivery stream. For more information about tags, see [Using Cost Allocation Tags](#) in the AWS Billing and Cost Management User Guide.

You can specify up to 50 tags when creating a delivery stream.

Type: Array of [Tag](#) (p. 162) objects

Array Members: Minimum number of 1 item. Maximum number of 50 items.

Required: No

Response Syntax

```
{  
  "DeliveryStreamARN": "string"  
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

DeliveryStreamARN (p. 12)

The ARN of the delivery stream.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: arn:.*

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 167).

InvalidArgumentException

The specified input parameter has a value that is not valid.

HTTP Status Code: 400

InvalidKMSResourceException

Kinesis Data Firehose throws this exception when an attempt to put records or to start or stop delivery stream encryption fails. This happens when the KMS service throws one of the following exception types: `AccessDeniedException`, `InvalidStateException`, `DisabledException`, or `NotFoundException`.

HTTP Status Code: 400

LimitExceededException

You have already reached the limit for a requested resource.

HTTP Status Code: 400

ResourceInUseException

The resource is already in use and not available for this operation.

HTTP Status Code: 400

Examples

Example

The following JSON example creates a delivery stream named `exampleStreamName` with an Amazon S3 destination. To use this example, first replace the placeholders for the `RoleARN` and `BucketARN` keys with valid strings. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Sample Request

```
POST / HTTP/1.1
Host: firehose.<region>.<domain>
Content-Length: <PayloadSizeBytes>
User-Agent: <UserAgentString>
Content-Type: application/x-amz-json-1.1
Authorization: <AuthParams>
Connection: Keep-Alive
X-Amz-Date: <Date>
X-Amz-Target: Firehose_20150804.CreateDeliveryStream
{
  "DeliveryStreamName": "exampleStreamName",
  "S3DestinationConfiguration": {
    "RoleARN": "insert-role-ARN",
    "BucketARN": "insert-bucket-ARN",
    "BufferingHints": {
      "SizeInMBs": 3,
      "IntervalInSeconds": 60
    },
    "CompressionFormat": "ZIP"
  }
}
```

Sample Response

```
HTTP/1.1 200 OK
x-amzn-RequestId: <RequestId>
Content-Type: application/x-amz-json-1.1
Content-Length: <PayloadSizeBytes>
Date: <Date>
{
  "DeliveryStreamARN": "arn:aws:firehose:us-east-1:814985986679:deliverystream/
exampleStreamName"
}
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)

- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DeleteDeliveryStream

Deletes a delivery stream and its data.

To check the state of a delivery stream, use [DescribeDeliveryStream \(p. 18\)](#). You can delete a delivery stream only if it is in one of the following states: `ACTIVE`, `DELETING`, `CREATING_FAILED`, or `DELETING_FAILED`. You can't delete a delivery stream that is in the `CREATING` state. While the deletion request is in process, the delivery stream is in the `DELETING` state.

While the delivery stream is in the `DELETING` state, the service might continue to accept records, but it doesn't make any guarantees with respect to delivering the data. Therefore, as a best practice, first stop any applications that are sending records before you delete a delivery stream.

Request Syntax

```
{  
  "AllowForceDelete": boolean,  
  "DeliveryStreamName": "string"  
}
```

Request Parameters

The request accepts the following data in JSON format.

[AllowForceDelete \(p. 15\)](#)

Set this to true if you want to delete the delivery stream even if Kinesis Data Firehose is unable to retire the grant for the CMK. Kinesis Data Firehose might be unable to retire the grant due to a customer error, such as when the CMK or the grant are in an invalid state. If you force deletion, you can then use the [RevokeGrant](#) operation to revoke the grant you gave to Kinesis Data Firehose. If a failure to retire the grant happens due to an AWS KMS issue, Kinesis Data Firehose keeps retrying the delete operation.

The default value is false.

Type: Boolean

Required: No

[DeliveryStreamName \(p. 15\)](#)

The name of the delivery stream.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 64.

Pattern: `[a-zA-Z0-9_.-]+`

Required: Yes

Response Elements

If the action is successful, the service sends back an HTTP 200 response with an empty HTTP body.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 167\)](#).

ResourceInUseException

The resource is already in use and not available for this operation.

HTTP Status Code: 400

ResourceNotFoundException

The specified resource could not be found.

HTTP Status Code: 400

Examples

Example

The following JSON example deletes a delivery stream named `exampleStreamName`.

Sample Request

```
POST / HTTP/1.1
Host: firehose.<region>.<domain>
Content-Length: <PayloadSizeBytes>
User-Agent: <UserAgentString>
Content-Type: application/x-amz-json-1.1
Authorization: <AuthParams>
Connection: Keep-Alive
X-Amz-Date: <Date>
X-Amz-Target: Firehose_20150804.DeleteDeliveryStream
{
  "DeliveryStreamName": "exampleStreamName"
}
```

Sample Response

```
HTTP/1.1 200 OK
x-amzn-RequestId: <RequestId>
Content-Type: application/x-amz-json-1.1
Content-Length: <PayloadSizeBytes>
Date: <Date>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)

- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeDeliveryStream

Describes the specified delivery stream and its status. For example, after your delivery stream is created, call `DescribeDeliveryStream` to see whether the delivery stream is `ACTIVE` and therefore ready for data to be sent to it.

If the status of a delivery stream is `CREATING_FAILED`, this status doesn't change, and you can't invoke [CreateDeliveryStream](#) (p. 3) again on it. However, you can invoke the [DeleteDeliveryStream](#) (p. 15) operation to delete it. If the status is `DELETING_FAILED`, you can force deletion by invoking [DeleteDeliveryStream](#) (p. 15) again but with [DeleteDeliveryStream:AllowForceDelete](#) (p. 15) set to true.

Request Syntax

```
{
  "DeliveryStreamName": "string",
  "ExclusiveStartDestinationId": "string",
  "Limit": number
}
```

Request Parameters

The request accepts the following data in JSON format.

[DeliveryStreamName](#) (p. 18)

The name of the delivery stream.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 64.

Pattern: [a-zA-Z0-9_.-]+

Required: Yes

[ExclusiveStartDestinationId](#) (p. 18)

The ID of the destination to start returning the destination information. Kinesis Data Firehose supports one destination per delivery stream.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: [a-zA-Z0-9-]+

Required: No

[Limit](#) (p. 18)

The limit on the number of destinations to return. You can have one destination per delivery stream.

Type: Integer

Valid Range: Minimum value of 1. Maximum value of 10000.

Required: No

Response Syntax

```
{
  "DeliveryStreamDescription": {
    "CreateTimestamp": number,
    "DeliveryStreamARN": "string",
    "DeliveryStreamEncryptionConfiguration": {
      "FailureDescription": {
        "Details": "string",
        "Type": "string"
      },
      "KeyARN": "string",
      "KeyType": "string",
      "Status": "string"
    },
    "DeliveryStreamName": "string",
    "DeliveryStreamStatus": "string",
    "DeliveryStreamType": "string",
    "Destinations": [
      {
        "DestinationId": "string",
        "ElasticsearchDestinationDescription": {
          "BufferingHints": {
            "IntervalInSeconds": number,
            "SizeInMBs": number
          },
          "CloudWatchLoggingOptions": {
            "Enabled": boolean,
            "LogGroupName": "string",
            "LogStreamName": "string"
          },
          "ClusterEndpoint": "string",
          "DomainARN": "string",
          "IndexName": "string",
          "IndexRotationPeriod": "string",
          "ProcessingConfiguration": {
            "Enabled": boolean,
            "Processors": [
              {
                "Parameters": [
                  {
                    "ParameterName": "string",
                    "ParameterValue": "string"
                  }
                ],
                "Type": "string"
              }
            ]
          }
        },
        "Type": "string"
      }
    ],
    "RetryOptions": {
      "DurationInSeconds": number
    },
    "RoleARN": "string",
    "S3BackupMode": "string",
    "S3DestinationDescription": {
      "BucketARN": "string",
      "BufferingHints": {
        "IntervalInSeconds": number,
        "SizeInMBs": number
      },
      "CloudWatchLoggingOptions": {
        "Enabled": boolean,
        "LogGroupName": "string",
        "LogStreamName": "string"
      }
    }
  }
}
```

```
    },
    "CompressionFormat": "string",
    "EncryptionConfiguration": {
      "KMSEncryptionConfig": {
        "AWSKMSKeyARN": "string"
      },
      "NoEncryptionConfig": "string"
    },
    "ErrorOutputPrefix": "string",
    "Prefix": "string",
    "RoleARN": "string"
  },
  "TypeName": "string",
  "VpcConfigurationDescription": {
    "RoleARN": "string",
    "SecurityGroupIds": [ "string" ],
    "SubnetIds": [ "string" ],
    "VpcId": "string"
  }
},
"ExtendedS3DestinationDescription": {
  "BucketARN": "string",
  "BufferingHints": {
    "IntervalInSeconds": number,
    "SizeInMBs": number
  },
  "CloudWatchLoggingOptions": {
    "Enabled": boolean,
    "LogGroupName": "string",
    "LogStreamName": "string"
  },
  "CompressionFormat": "string",
  "DataFormatConversionConfiguration": {
    "Enabled": boolean,
    "InputFormatConfiguration": {
      "Deserializer": {
        "HiveJsonSerDe": {
          "TimestampFormats": [ "string" ]
        },
        "OpenXJsonSerDe": {
          "CaseInsensitive": boolean,
          "ColumnToJsonKeyMappings": {
            "string": "string"
          },
          "ConvertDotsInJsonKeysToUnderscores": boolean
        }
      }
    }
  },
  "OutputFormatConfiguration": {
    "Serializer": {
      "OrcSerDe": {
        "BlockSizeBytes": number,
        "BloomFilterColumns": [ "string" ],
        "BloomFilterFalsePositiveProbability": number,
        "Compression": "string",
        "DictionaryKeyThreshold": number,
        "EnablePadding": boolean,
        "FormatVersion": "string",
        "PaddingTolerance": number,
        "RowIndexStride": number,
        "StripeSizeBytes": number
      },
      "ParquetSerDe": {
        "BlockSizeBytes": number,
        "Compression": "string",
        "EnableDictionaryCompression": boolean,
```

```
        "MaxPaddingBytes": number,
        "PageSizeBytes": number,
        "WriterVersion": "string"
    }
}
},
"SchemaConfiguration": {
    "CatalogId": "string",
    "DatabaseName": "string",
    "Region": "string",
    "RoleARN": "string",
    "TableName": "string",
    "VersionId": "string"
}
},
"DynamicPartitioningConfiguration": {
    "Enabled": boolean,
    "RetryOptions": {
        "DurationInSeconds": number
    }
}
},
"EncryptionConfiguration": {
    "KMSEncryptionConfig": {
        "AWSKMSKeyARN": "string"
    },
    "NoEncryptionConfig": "string"
}
},
"ErrorOutputPrefix": "string",
"Prefix": "string",
"ProcessingConfiguration": {
    "Enabled": boolean,
    "Processors": [
        {
            "Parameters": [
                {
                    "ParameterName": "string",
                    "ParameterValue": "string"
                }
            ],
            "Type": "string"
        }
    ]
}
},
"RoleARN": "string",
"S3BackupDescription": {
    "BucketARN": "string",
    "BufferingHints": {
        "IntervalInSeconds": number,
        "SizeInMBs": number
    },
    "CloudWatchLoggingOptions": {
        "Enabled": boolean,
        "LogGroupName": "string",
        "LogStreamName": "string"
    },
    "CompressionFormat": "string",
    "EncryptionConfiguration": {
        "KMSEncryptionConfig": {
            "AWSKMSKeyARN": "string"
        },
        "NoEncryptionConfig": "string"
    },
    "ErrorOutputPrefix": "string",
    "Prefix": "string",
    "RoleARN": "string"
}
},
```

```
    "S3BackupMode": "string"
  },
  "HttpEndpointDestinationDescription": {
    "BufferingHints": {
      "IntervalInSeconds": number,
      "SizeInMBs": number
    },
    "CloudWatchLoggingOptions": {
      "Enabled": boolean,
      "LogGroupName": "string",
      "LogStreamName": "string"
    },
    "EndpointConfiguration": {
      "Name": "string",
      "Url": "string"
    },
    "ProcessingConfiguration": {
      "Enabled": boolean,
      "Processors": [
        {
          "Parameters": [
            {
              "ParameterName": "string",
              "ParameterValue": "string"
            }
          ],
          "Type": "string"
        }
      ]
    },
    "RequestConfiguration": {
      "CommonAttributes": [
        {
          "AttributeName": "string",
          "AttributeValue": "string"
        }
      ],
      "ContentEncoding": "string"
    },
    "RetryOptions": {
      "DurationInSeconds": number
    },
    "RoleARN": "string",
    "S3BackupMode": "string",
    "S3DestinationDescription": {
      "BucketARN": "string",
      "BufferingHints": {
        "IntervalInSeconds": number,
        "SizeInMBs": number
      },
      "CloudWatchLoggingOptions": {
        "Enabled": boolean,
        "LogGroupName": "string",
        "LogStreamName": "string"
      },
      "CompressionFormat": "string",
      "EncryptionConfiguration": {
        "KMSEncryptionConfig": {
          "AWSKMSKeyARN": "string"
        },
        "NoEncryptionConfig": "string"
      },
      "ErrorOutputPrefix": "string",
      "Prefix": "string",
      "RoleARN": "string"
    }
  }
}
```

```
},
"RedshiftDestinationDescription": {
  "CloudWatchLoggingOptions": {
    "Enabled": boolean,
    "LogGroupName": "string",
    "LogStreamName": "string"
  },
  "ClusterJDBCURL": "string",
  "CopyCommand": {
    "CopyOptions": "string",
    "DataTableColumns": "string",
    "DataTableName": "string"
  },
  "ProcessingConfiguration": {
    "Enabled": boolean,
    "Processors": [
      {
        "Parameters": [
          {
            "ParameterName": "string",
            "ParameterValue": "string"
          }
        ],
        "Type": "string"
      }
    ]
  },
  "RetryOptions": {
    "DurationInSeconds": number
  },
  "RoleARN": "string",
  "S3BackupDescription": {
    "BucketARN": "string",
    "BufferingHints": {
      "IntervalInSeconds": number,
      "SizeInMBs": number
    },
    "CloudWatchLoggingOptions": {
      "Enabled": boolean,
      "LogGroupName": "string",
      "LogStreamName": "string"
    },
    "CompressionFormat": "string",
    "EncryptionConfiguration": {
      "KMSEncryptionConfig": {
        "AWSKMSKeyARN": "string"
      },
      "NoEncryptionConfig": "string"
    },
    "ErrorOutputPrefix": "string",
    "Prefix": "string",
    "RoleARN": "string"
  },
  "S3BackupMode": "string",
  "S3DestinationDescription": {
    "BucketARN": "string",
    "BufferingHints": {
      "IntervalInSeconds": number,
      "SizeInMBs": number
    },
    "CloudWatchLoggingOptions": {
      "Enabled": boolean,
      "LogGroupName": "string",
      "LogStreamName": "string"
    },
    "CompressionFormat": "string",
```

```
    "EncryptionConfiguration": {
      "KMSEncryptionConfig": {
        "AWSKMSKeyARN": "string"
      },
      "NoEncryptionConfig": "string"
    },
    "ErrorOutputPrefix": "string",
    "Prefix": "string",
    "RoleARN": "string"
  },
  "Username": "string"
},
"S3DestinationDescription": {
  "BucketARN": "string",
  "BufferingHints": {
    "IntervalInSeconds": number,
    "SizeInMBs": number
  },
  "CloudWatchLoggingOptions": {
    "Enabled": boolean,
    "LogGroupName": "string",
    "LogStreamName": "string"
  },
  "CompressionFormat": "string",
  "EncryptionConfiguration": {
    "KMSEncryptionConfig": {
      "AWSKMSKeyARN": "string"
    },
    "NoEncryptionConfig": "string"
  },
  "ErrorOutputPrefix": "string",
  "Prefix": "string",
  "RoleARN": "string"
},
"SplunkDestinationDescription": {
  "CloudWatchLoggingOptions": {
    "Enabled": boolean,
    "LogGroupName": "string",
    "LogStreamName": "string"
  },
  "HECAcknowledgmentTimeoutInSeconds": number,
  "HECEndpoint": "string",
  "HECEndpointType": "string",
  "HECToken": "string",
  "ProcessingConfiguration": {
    "Enabled": boolean,
    "Processors": [
      {
        "Parameters": [
          {
            "ParameterName": "string",
            "ParameterValue": "string"
          }
        ],
        "Type": "string"
      }
    ]
  },
  "RetryOptions": {
    "DurationInSeconds": number
  },
  "S3BackupMode": "string",
  "S3DestinationDescription": {
    "BucketARN": "string",
    "BufferingHints": {
      "IntervalInSeconds": number,
```

```
        "SizeInMBs": number
      },
      "CloudWatchLoggingOptions": {
        "Enabled": boolean,
        "LogGroupName": "string",
        "LogStreamName": "string"
      },
      "CompressionFormat": "string",
      "EncryptionConfiguration": {
        "KMSEncryptionConfig": {
          "AWSKMSKeyARN": "string"
        },
        "NoEncryptionConfig": "string"
      },
      "ErrorOutputPrefix": "string",
      "Prefix": "string",
      "RoleARN": "string"
    }
  }
},
"FailureDescription": {
  "Details": "string",
  "Type": "string"
},
"HasMoreDestinations": boolean,
"LastUpdateTimestamp": number,
"Source": {
  "KinesisStreamSourceDescription": {
    "DeliveryStartTimestamp": number,
    "KinesisStreamARN": "string",
    "RoleARN": "string"
  }
},
"VersionId": "string"
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

DeliveryStreamDescription (p. 19)

Information about the delivery stream.

Type: [DeliveryStreamDescription](#) (p. 72) object

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 167).

ResourceNotFoundException

The specified resource could not be found.

HTTP Status Code: 400

Examples

Example

The following JSON example describes a delivery stream.

Sample Request

```
POST / HTTP/1.1
Host: firehose.<region>.<domain>
Content-Length: <PayloadSizeBytes>
User-Agent: <UserAgentString>
Content-Type: application/x-amz-json-1.1
Authorization: <AuthParams>
Connection: Keep-Alive
X-Amz-Date: <Date>
X-Amz-Target: Firehose_20150804.DescribeDeliveryStream
{
  "DeliveryStreamName": "exampleStreamName"
}
```

Sample Response

```
HTTP/1.1 200 OK
x-amzn-RequestId: <RequestId>
Content-Type: application/x-amz-json-1.1
Content-Length: <PayloadSizeBytes>
Date: <Date>
{
  "DeliveryStreamDescription": {
    "DeliveryStreamType": "DirectPut",
    "HasMoreDestinations": false,
    "VersionId": "1",
    "CreateTimestamp": 1517595920.596,
    "DeliveryStreamARN": "arn:aws:firehose:us-east-1:111222333444:deliverystream/
exampleStreamName",
    "DeliveryStreamStatus": "ACTIVE",
    "DeliveryStreamName": "exampleStreamName",
    "DeliveryStreamEncryptionConfiguration": {
      "Status": "DISABLED"
    },
    "Destinations": [
      {
        "DestinationId": "destinationId-0000000000001",
        "ExtendedS3DestinationDescription": {
          "RoleARN": "arn:aws:iam::111222333444:role/exampleStreamName",
          "Prefix": "",
          "BufferingHints": {
            "IntervalInSeconds": 60,
            "SizeInMBs": 1
          },
          "EncryptionConfiguration": {
            "NoEncryptionConfig": "NoEncryption"
          },
          "CompressionFormat": "UNCOMPRESSED",
          "S3BackupMode": "Disabled",
          "CloudWatchLoggingOptions": {
            "Enabled": true,
            "LogStreamName": "S3Delivery",
            "LogGroupName": "/aws/kinesisfirehose/exampleStreamName"
          }
        }
      }
    ]
  }
}
```

```
    "BucketARN": "arn:aws:s3:::somebucket",
    "ProcessingConfiguration": {
      "Enabled": false,
      "Processors": []
    }
  },
  "S3DestinationDescription": {
    "RoleARN": "arn:aws:iam::111222333444:role/exampleStreamName",
    "Prefix": "",
    "BufferingHints": {
      "IntervalInSeconds": 60,
      "SizeInMBs": 1
    },
    "EncryptionConfiguration": {
      "NoEncryptionConfig": "NoEncryption"
    },
    "CompressionFormat": "UNCOMPRESSED",
    "CloudWatchLoggingOptions": {
      "Enabled": true,
      "LogStreamName": "S3Delivery",
      "LogGroupName": "/aws/kinesisfirehose/exampleStreamName"
    },
    "BucketARN": "arn:aws:s3:::somebucket"
  }
}
]
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ListDeliveryStreams

Lists your delivery streams in alphabetical order of their names.

The number of delivery streams might be too large to return using a single call to `ListDeliveryStreams`. You can limit the number of delivery streams returned, using the `Limit` parameter. To determine whether there are more delivery streams to list, check the value of `HasMoreDeliveryStreams` in the output. If there are more delivery streams to list, you can request them by calling this operation again and setting the `ExclusiveStartDeliveryStreamName` parameter to the name of the last delivery stream returned in the last call.

Request Syntax

```
{  
  "DeliveryStreamType": "string",  
  "ExclusiveStartDeliveryStreamName": "string",  
  "Limit": number  
}
```

Request Parameters

The request accepts the following data in JSON format.

DeliveryStreamType (p. 28)

The delivery stream type. This can be one of the following values:

- `DirectPut`: Provider applications access the delivery stream directly.
- `KinesisStreamAsSource`: The delivery stream uses a Kinesis data stream as a source.

This parameter is optional. If this parameter is omitted, delivery streams of all types are returned.

Type: String

Valid Values: `DirectPut` | `KinesisStreamAsSource`

Required: No

ExclusiveStartDeliveryStreamName (p. 28)

The list of delivery streams returned by this call to `ListDeliveryStreams` will start with the delivery stream whose name comes alphabetically immediately after the name you specify in `ExclusiveStartDeliveryStreamName`.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 64.

Pattern: `[a-zA-Z0-9_.-]+`

Required: No

Limit (p. 28)

The maximum number of delivery streams to list. The default value is 10.

Type: Integer

Valid Range: Minimum value of 1. Maximum value of 10000.

Required: No

Response Syntax

```
{
  "DeliveryStreamNames": [ "string" ],
  "HasMoreDeliveryStreams": boolean
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

DeliveryStreamNames (p. 29)

The names of the delivery streams.

Type: Array of strings

Length Constraints: Minimum length of 1. Maximum length of 64.

Pattern: [a-zA-Z0-9_.-]+

HasMoreDeliveryStreams (p. 29)

Indicates whether there are more delivery streams available to list.

Type: Boolean

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 167).

Examples

Example

The following JSON example lists up to three delivery streams configured for `DirectPut`. The response indicates that there are more `DirectPut` delivery streams to be listed. To list the remaining `DirectPut` delivery streams, set `DeliveryStreamType` to `DirectPut` and `ExclusiveStartDeliveryStreamName` to `last_stream_in_first_listing`, and then run the operation again.

Sample Request

```
POST / HTTP/1.1
Host: firehose.<region>.<domain>
Content-Length: <PayloadSizeBytes>
User-Agent: <UserAgentString>
Content-Type: application/x-amz-json-1.1
Authorization: <AuthParams>
Connection: Keep-Alive
X-Amz-Date: <Date>
```

```
X-Amz-Target: Firehose_20150804.ListDeliveryStreams
{
  "DeliveryStreamType": "DirectPut",
  "Limit": 3
}
```

Sample Response

```
HTTP/1.1 200 OK
x-amzn-RequestId: <RequestId>
Content-Type: application/x-amz-json-1.1
Content-Length: <PayloadSizeBytes>
Date: <Date>
{
  "DeliveryStreamNames": [
    "some_delivery_stream",
    "another_example_delivery_stream",
    "last_stream_in_first_listing"
  ],
  "HasMoreDeliveryStreams": true
}
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ListTagsForDeliveryStream

Lists the tags for the specified delivery stream. This operation has a limit of five transactions per second per account.

Request Syntax

```
{
  "DeliveryStreamName": "string",
  "ExclusiveStartTagKey": "string",
  "Limit": number
}
```

Request Parameters

The request accepts the following data in JSON format.

DeliveryStreamName (p. 31)

The name of the delivery stream whose tags you want to list.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 64.

Pattern: [a-zA-Z0-9_.-]+

Required: Yes

ExclusiveStartTagKey (p. 31)

The key to use as the starting point for the list of tags. If you set this parameter, `ListTagsForDeliveryStream` gets all tags that occur after `ExclusiveStartTagKey`.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 128.

Pattern: ^(? !aws:)[\p{L}\p{Z}\p{N}_. : \/+ - @%]*\$

Required: No

Limit (p. 31)

The number of tags to return. If this number is less than the total number of tags associated with the delivery stream, `HasMoreTags` is set to `true` in the response. To list additional tags, set `ExclusiveStartTagKey` to the last key in the response.

Type: Integer

Valid Range: Minimum value of 1. Maximum value of 50.

Required: No

Response Syntax

```
{
```

```
"HasMoreTags": boolean,
"Tags": [
  {
    "Key": "string",
    "Value": "string"
  }
]
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

HasMoreTags (p. 31)

If this is `true` in the response, more tags are available. To list the remaining tags, set `ExclusiveStartTagKey` to the key of the last tag returned and call `ListTagsForDeliveryStream` again.

Type: Boolean

Tags (p. 31)

A list of tags associated with `DeliveryStreamName`, starting with the first tag after `ExclusiveStartTagKey` and up to the specified `Limit`.

Type: Array of [Tag \(p. 162\)](#) objects

Array Members: Minimum number of 0 items. Maximum number of 50 items.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 167\)](#).

InvalidArgumentException

The specified input parameter has a value that is not valid.

HTTP Status Code: 400

LimitExceededException

You have already reached the limit for a requested resource.

HTTP Status Code: 400

ResourceNotFoundException

The specified resource could not be found.

HTTP Status Code: 400

Examples

To list the tags for a stream

The following JSON example lists the tags for the specified delivery stream.

Sample Request

```
POST / HTTP/1.1
Host: firehose.<region>.<domain>
Content-Length: <PayloadSizeBytes>
User-Agent: <UserAgentString>
Content-Type: application/x-amz-json-1.1
Authorization: <AuthParams>
Connection: Keep-Alive
X-Amz-Date: <Date>
X-Amz-Target: Firehose_20150804.ListTagsForDeliveryStream
{
  "DeliveryStreamName": "exampleDeliveryStreamName"
}
```

Sample Response

```
HTTP/1.1 200 OK
x-amzn-RequestId: <RequestId>
Content-Type: application/x-amz-json-1.1
Content-Length: <PayloadSizeBytes>
Date: <Date>
{
  "HasMoreTags": "false",
  "Tags" : [
    {
      "Key": "Project",
      "Value": "myProject"
    },
    {
      "Key": "Environment",
      "Value": "Production"
    }
  ]
}
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

PutRecord

Writes a single data record into an Amazon Kinesis Data Firehose delivery stream. To write multiple data records into a delivery stream, use [PutRecordBatch](#) (p. 38). Applications using these operations are referred to as producers.

By default, each delivery stream can take in up to 2,000 transactions per second, 5,000 records per second, or 5 MB per second. If you use [PutRecord](#) (p. 34) and [PutRecordBatch](#) (p. 38), the limits are an aggregate across these two operations for each delivery stream. For more information about limits and how to request an increase, see [Amazon Kinesis Data Firehose Limits](#).

You must specify the name of the delivery stream and the data record when using [PutRecord](#) (p. 34). The data record consists of a data blob that can be up to 1,000 KiB in size, and any kind of data. For example, it can be a segment from a log file, geographic location data, website clickstream data, and so on.

Kinesis Data Firehose buffers records before delivering them to the destination. To disambiguate the data blobs at the destination, a common solution is to use delimiters in the data, such as a newline (\n) or some other character unique within the data. This allows the consumer application to parse individual data items when reading the data from the destination.

The `PutRecord` operation returns a `RecordId`, which is a unique string assigned to each record. Producer applications can use this ID for purposes such as auditability and investigation.

If the `PutRecord` operation throws a `ServiceUnavailableException`, back off and retry. If the exception persists, it is possible that the throughput limits have been exceeded for the delivery stream.

Data records sent to Kinesis Data Firehose are stored for 24 hours from the time they are added to a delivery stream as it tries to send the records to the destination. If the destination is unreachable for more than 24 hours, the data is no longer available.

Important

Don't concatenate two or more base64 strings to form the data fields of your records. Instead, concatenate the raw data, then perform base64 encoding.

Request Syntax

```
{
  "DeliveryStreamName": "string",
  "Record": {
    "Data": blob
  }
}
```

Request Parameters

The request accepts the following data in JSON format.

[DeliveryStreamName](#) (p. 34)

The name of the delivery stream.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 64.

Pattern: [a-zA-Z0-9_.-]+

Required: Yes

Record (p. 34)

The record.

Type: [Record \(p. 134\)](#) object

Required: Yes

Response Syntax

```
{  
  "Encrypted": boolean,  
  "RecordId": "string"  
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

Encrypted (p. 35)

Indicates whether server-side encryption (SSE) was enabled during this operation.

Type: Boolean

RecordId (p. 35)

The ID of the record.

Type: String

Length Constraints: Minimum length of 1.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 167\)](#).

InvalidArgumentException

The specified input parameter has a value that is not valid.

HTTP Status Code: 400

InvalidKMSResourceException

Kinesis Data Firehose throws this exception when an attempt to put records or to start or stop delivery stream encryption fails. This happens when the KMS service throws one of the following exception types: `AccessDeniedException`, `InvalidStateException`, `DisabledException`, or `NotFoundException`.

HTTP Status Code: 400

ResourceNotFoundException

The specified resource could not be found.

HTTP Status Code: 400

ServiceUnavailableException

The service is unavailable. Back off and retry the operation. If you continue to see the exception, throughput limits for the delivery stream may have been exceeded. For more information about limits and how to request an increase, see [Amazon Kinesis Data Firehose Limits](#).

HTTP Status Code: 500

Examples

Example

The following JSON puts a record in the delivery stream named `some_delivery_stream`:

Sample Request

```
POST / HTTP/1.1
Host: firehose.<region>.<domain>
Content-Length: <PayloadSizeBytes>
User-Agent: <UserAgentString>
Content-Type: application/x-amz-json-1.1
Authorization: <AuthParams>
Connection: Keep-Alive
X-Amz-Date: <Date>
X-Amz-Target: Firehose_20150804.PutRecord
{
  "DeliveryStreamName": "some_delivery_stream",
  "Record": {
    "Data": "... "
  }
}
```

Sample Response

```
HTTP/1.1 200 OK
x-amzn-RequestId: <RequestId>
Content-Type: application/x-amz-json-1.1
Content-Length: <PayloadSizeBytes>
Date: <Date>
{
  "RecordId": "CGojNMJq3msHbGoc+lmGSpifCmfFm7lFhuts//4Ft6sFVokyE6t+5ioEAjNm+sgQ6iVf/
YePEXBK6epIW4QeXqJp2xsbfZUNXsfOY1QrYXgRBCKznkjMMTP0BqJGOBm3fB//dHgEE0XDtc4wW065i/
tJyYI1Vy8qn8FMhpKZuh5bvG482XkkBxFmMGnhPTQwQ4A1IOP0sE0X99YnBK8RECdeQ2zxynvZ"
}
```

Example

The following example shows how you can use the [AWS CLI](#) to put a record in a delivery stream.

```
aws firehose put-record --delivery-stream-name mystream --record="{\"Data\\\":\\\"1\\\"}"
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

PutRecordBatch

Writes multiple data records into a delivery stream in a single call, which can achieve higher throughput per producer than when writing single records. To write single data records into a delivery stream, use [PutRecord](#) (p. 34). Applications using these operations are referred to as producers.

For information about service quota, see [Amazon Kinesis Data Firehose Quota](#).

Each [PutRecordBatch](#) (p. 38) request supports up to 500 records. Each record in the request can be as large as 1,000 KB (before base64 encoding), up to a limit of 4 MB for the entire request. These limits cannot be changed.

You must specify the name of the delivery stream and the data record when using [PutRecord](#) (p. 34). The data record consists of a data blob that can be up to 1,000 KB in size, and any kind of data. For example, it could be a segment from a log file, geographic location data, website clickstream data, and so on.

Kinesis Data Firehose buffers records before delivering them to the destination. To disambiguate the data blobs at the destination, a common solution is to use delimiters in the data, such as a newline (\n) or some other character unique within the data. This allows the consumer application to parse individual data items when reading the data from the destination.

The [PutRecordBatch](#) (p. 38) response includes a count of failed records, `FailedPutCount`, and an array of responses, `RequestResponses`. Even if the [PutRecordBatch](#) (p. 38) call succeeds, the value of `FailedPutCount` may be greater than 0, indicating that there are records for which the operation didn't succeed. Each entry in the `RequestResponses` array provides additional information about the processed record. It directly correlates with a record in the request array using the same ordering, from the top to the bottom. The response array always includes the same number of records as the request array. `RequestResponses` includes both successfully and unsuccessfully processed records. Kinesis Data Firehose tries to process all records in each [PutRecordBatch](#) (p. 38) request. A single record failure does not stop the processing of subsequent records.

A successfully processed record includes a `RecordId` value, which is unique for the record. An unsuccessfully processed record includes `ErrorCode` and `ErrorMessage` values. `ErrorCode` reflects the type of error, and is one of the following values: `ServiceUnavailableException` or `InternalFailure`. `ErrorMessage` provides more detailed information about the error.

If there is an internal server error or a timeout, the write might have completed or it might have failed. If `FailedPutCount` is greater than 0, retry the request, resending only those records that might have failed processing. This minimizes the possible duplicate records and also reduces the total bytes sent (and corresponding charges). We recommend that you handle any duplicates at the destination.

If [PutRecordBatch](#) (p. 38) throws `ServiceUnavailableException`, back off and retry. If the exception persists, it is possible that the throughput limits have been exceeded for the delivery stream.

Data records sent to Kinesis Data Firehose are stored for 24 hours from the time they are added to a delivery stream as it attempts to send the records to the destination. If the destination is unreachable for more than 24 hours, the data is no longer available.

Important

Don't concatenate two or more base64 strings to form the data fields of your records. Instead, concatenate the raw data, then perform base64 encoding.

Request Syntax

```
{  
  "DeliveryStreamName": "string",
```

```
"Records": [  
  {  
    "Data": blob  
  }  
]
```

Request Parameters

The request accepts the following data in JSON format.

DeliveryStreamName (p. 38)

The name of the delivery stream.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 64.

Pattern: [a-zA-Z0-9_.-]+

Required: Yes

Records (p. 38)

One or more records.

Type: Array of [Record](#) (p. 134) objects

Array Members: Minimum number of 1 item. Maximum number of 500 items.

Required: Yes

Response Syntax

```
{  
  "Encrypted": boolean,  
  "FailedPutCount": number,  
  "RequestResponses": [  
    {  
      "ErrorCode": "string",  
      "ErrorMessage": "string",  
      "RecordId": "string"  
    }  
  ]  
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

Encrypted (p. 39)

Indicates whether server-side encryption (SSE) was enabled during this operation.

Type: Boolean

FailedPutCount (p. 39)

The number of records that might have failed processing. This number might be greater than 0 even if the [PutRecordBatch \(p. 38\)](#) call succeeds. Check `FailedPutCount` to determine whether there are records that you need to resend.

Type: Integer

Valid Range: Minimum value of 0.

RequestResponses (p. 39)

The results array. For each record, the index of the response element is the same as the index used in the request array.

Type: Array of [PutRecordBatchResponseEntry \(p. 133\)](#) objects

Array Members: Minimum number of 1 item. Maximum number of 500 items.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 167\)](#).

InvalidArgumentException

The specified input parameter has a value that is not valid.

HTTP Status Code: 400

InvalidKMSResourceException

Kinesis Data Firehose throws this exception when an attempt to put records or to start or stop delivery stream encryption fails. This happens when the KMS service throws one of the following exception types: `AccessDeniedException`, `InvalidStateException`, `DisabledException`, or `NotFoundException`.

HTTP Status Code: 400

ResourceNotFoundException

The specified resource could not be found.

HTTP Status Code: 400

ServiceUnavailableException

The service is unavailable. Back off and retry the operation. If you continue to see the exception, throughput limits for the delivery stream may have been exceeded. For more information about limits and how to request an increase, see [Amazon Kinesis Data Firehose Limits](#).

HTTP Status Code: 500

Examples

Example

The following JSON puts two records in the delivery stream named `some_delivery_stream`:

Sample Request

```
POST / HTTP/1.1
```

```
Host: firehose.<region>.<domain>
Content-Length: <PayloadSizeBytes>
User-Agent: <UserAgentString>
Content-Type: application/x-amz-json-1.1
Authorization: <AuthParams>
Connection: Keep-Alive
X-Amz-Date: <Date>
X-Amz-Target: Firehose_20150804.PutRecordBatch
{
  "DeliveryStreamName": "some_delivery_stream",
  "Records": [
    {
      "Data": "Some data blob."
    },
    {
      "Data": "Another blob of data."
    }
  ]
}
```

Sample Response

```
HTTP/1.1 200 OK
x-amzn-RequestId: <RequestId>
Content-Type: application/x-amz-json-1.1
Content-Length: <PayloadSizeBytes>
Date: <Date>
{
  "FailedPutCount": 0,
  "RequestResponses": [
    {
      "RecordId": "AJJBALlfiFN9HyhPj6Dc+XqcRlTjqyIbr927TsEmWWpN39EK/
JbRTbXDZFNCCrWIs/4YUd3gjYUs6giZn76Tui4vv0ljDOMFJvLzqLuupu93RscZRmpWlCP8DeiFzJJGvqIr1LRE/
MDozYenOz+v1ZIqEjECvfMwIz6silvDaGuYtIhVzwd8yWadGFCX4ODes6W07W9W7Q5ViMPUKSt6F0nn7OhVkJ3/"
    },
    {
      "RecordId": "goGaFS919Mmv7lYET0oMaw+UL9iFpzi100o
+csoIc3lSmvkqpwzQuT0RPqZ7QqfR1FJ
+HxJciW/8paFMWPByJ6qVDhiE7TtJxARKaP4YDccvWHXhD7x6Y4bc9AHZ0uzy
+BEuTTnH5Zsip1IrEPgu8lj8a7fd6sluTe/kZ54GVG+Sc105IrxIu08kPbRiUTWL
+Wg0A3HzOQeePqokRKKqmoXLG0pzx/8O"
    }
  ]
}
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

StartDeliveryStreamEncryption

Enables server-side encryption (SSE) for the delivery stream.

This operation is asynchronous. It returns immediately. When you invoke it, Kinesis Data Firehose first sets the encryption status of the stream to `ENABLING`, and then to `ENABLED`. The encryption status of a delivery stream is the `Status` property in [DeliveryStreamEncryptionConfiguration](#) (p. 75). If the operation fails, the encryption status changes to `ENABLING_FAILED`. You can continue to read and write data to your delivery stream while the encryption status is `ENABLING`, but the data is not encrypted. It can take up to 5 seconds after the encryption status changes to `ENABLED` before all records written to the delivery stream are encrypted. To find out whether a record or a batch of records was encrypted, check the response elements [PutRecord:Encrypted](#) (p. 35) and [PutRecordBatch:Encrypted](#) (p. 39), respectively.

To check the encryption status of a delivery stream, use [DescribeDeliveryStream](#) (p. 18).

Even if encryption is currently enabled for a delivery stream, you can still invoke this operation on it to change the ARN of the CMK or both its type and ARN. If you invoke this method to change the CMK, and the old CMK is of type `CUSTOMER_MANAGED_CMK`, Kinesis Data Firehose schedules the grant it had on the old CMK for retirement. If the new CMK is of type `CUSTOMER_MANAGED_CMK`, Kinesis Data Firehose creates a grant that enables it to use the new CMK to encrypt and decrypt data and to manage the grant.

If a delivery stream already has encryption enabled and then you invoke this operation to change the ARN of the CMK or both its type and ARN and you get `ENABLING_FAILED`, this only means that the attempt to change the CMK failed. In this case, encryption remains enabled with the old CMK.

If the encryption status of your delivery stream is `ENABLING_FAILED`, you can invoke this operation again with a valid CMK. The CMK must be enabled and the key policy mustn't explicitly deny the permission for Kinesis Data Firehose to invoke KMS encrypt and decrypt operations.

You can enable SSE for a delivery stream only if it's a delivery stream that uses `DirectPut` as its source.

The `StartDeliveryStreamEncryption` and `StopDeliveryStreamEncryption` operations have a combined limit of 25 calls per delivery stream per 24 hours. For example, you reach the limit if you call `StartDeliveryStreamEncryption` 13 times and `StopDeliveryStreamEncryption` 12 times for the same delivery stream in a 24-hour period.

Request Syntax

```
{
  "DeliveryStreamEncryptionConfigurationInput": {
    "KeyARN": "string",
    "KeyType": "string"
  },
  "DeliveryStreamName": "string"
}
```

Request Parameters

The request accepts the following data in JSON format.

DeliveryStreamEncryptionConfigurationInput (p. 43)

Used to specify the type and Amazon Resource Name (ARN) of the KMS key needed for Server-Side Encryption (SSE).

Type: [DeliveryStreamEncryptionConfigurationInput](#) (p. 77) object

Required: No

DeliveryStreamName (p. 43)

The name of the delivery stream for which you want to enable server-side encryption (SSE).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 64.

Pattern: [a-zA-Z0-9_.-]+

Required: Yes

Response Elements

If the action is successful, the service sends back an HTTP 200 response with an empty HTTP body.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 167\)](#).

InvalidArgumentException

The specified input parameter has a value that is not valid.

HTTP Status Code: 400

InvalidKMSResourceException

Kinesis Data Firehose throws this exception when an attempt to put records or to start or stop delivery stream encryption fails. This happens when the KMS service throws one of the following exception types: `AccessDeniedException`, `InvalidStateException`, `DisabledException`, or `NotFoundException`.

HTTP Status Code: 400

LimitExceededException

You have already reached the limit for a requested resource.

HTTP Status Code: 400

ResourceInUseException

The resource is already in use and not available for this operation.

HTTP Status Code: 400

ResourceNotFoundException

The specified resource could not be found.

HTTP Status Code: 400

Examples

To start server-side encryption for a stream

The following JSON example starts server-side encryption (SSE) for the specified stream.

Sample Request

```
POST / HTTP/1.1
Host: firehose.<region>.<domain>
Content-Length: <PayloadSizeBytes>
User-Agent: <UserAgentString>
Content-Type: application/x-amz-json-1.1
Authorization: <AuthParams>
Connection: Keep-Alive
X-Amz-Date: <Date>
X-Amz-Target: Firehose_20150804.StartDeliveryStreamEncryption
{
  "DeliveryStreamName": "exampleDeliveryStreamName"
}
```

Sample Response

```
HTTP/1.1 200 OK
x-amzn-RequestId: <RequestId>
Content-Type: application/x-amz-json-1.1
Content-Length: <PayloadSizeBytes>
Date: <Date>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

StopDeliveryStreamEncryption

Disables server-side encryption (SSE) for the delivery stream.

This operation is asynchronous. It returns immediately. When you invoke it, Kinesis Data Firehose first sets the encryption status of the stream to `DISABLING`, and then to `DISABLED`. You can continue to read and write data to your stream while its status is `DISABLING`. It can take up to 5 seconds after the encryption status changes to `DISABLED` before all records written to the delivery stream are no longer subject to encryption. To find out whether a record or a batch of records was encrypted, check the response elements [PutRecord:Encrypted](#) (p. 35) and [PutRecordBatch:Encrypted](#) (p. 39), respectively.

To check the encryption state of a delivery stream, use [DescribeDeliveryStream](#) (p. 18).

If SSE is enabled using a customer managed CMK and then you invoke `StopDeliveryStreamEncryption`, Kinesis Data Firehose schedules the related KMS grant for retirement and then retires it after it ensures that it is finished delivering records to the destination.

The `StartDeliveryStreamEncryption` and `StopDeliveryStreamEncryption` operations have a combined limit of 25 calls per delivery stream per 24 hours. For example, you reach the limit if you call `StartDeliveryStreamEncryption` 13 times and `StopDeliveryStreamEncryption` 12 times for the same delivery stream in a 24-hour period.

Request Syntax

```
{  
  "DeliveryStreamName": "string"  
}
```

Request Parameters

The request accepts the following data in JSON format.

DeliveryStreamName (p. 46)

The name of the delivery stream for which you want to disable server-side encryption (SSE).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 64.

Pattern: [a-zA-Z0-9_.-]+

Required: Yes

Response Elements

If the action is successful, the service sends back an HTTP 200 response with an empty HTTP body.

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 167).

InvalidArgumentException

The specified input parameter has a value that is not valid.

HTTP Status Code: 400

LimitExceededException

You have already reached the limit for a requested resource.

HTTP Status Code: 400

ResourceInUseException

The resource is already in use and not available for this operation.

HTTP Status Code: 400

ResourceNotFoundException

The specified resource could not be found.

HTTP Status Code: 400

Examples

To stop server-side encryption for a stream

The following JSON example stops server-side encryption (SSE) for the specified stream.

Sample Request

```
POST / HTTP/1.1
Host: firehose.<region>.<domain>
Content-Length: <PayloadSizeBytes>
User-Agent: <UserAgentString>
Content-Type: application/x-amz-json-1.1
Authorization: <AuthParams>
Connection: Keep-Alive
X-Amz-Date: <Date>
X-Amz-Target: Firehose_20150804.StopDeliveryStreamEncryption
{
  "DeliveryStreamName": "exampleDeliveryStreamName"
}
```

Sample Response

```
HTTP/1.1 200 OK
x-amzn-RequestId: <RequestId>
Content-Type: application/x-amz-json-1.1
Content-Length: <PayloadSizeBytes>
Date: <Date>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)

- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

TagDeliveryStream

Adds or updates tags for the specified delivery stream. A tag is a key-value pair that you can define and assign to AWS resources. If you specify a tag that already exists, the tag value is replaced with the value that you specify in the request. Tags are metadata. For example, you can add friendly names and descriptions or other types of information that can help you distinguish the delivery stream. For more information about tags, see [Using Cost Allocation Tags](#) in the *AWS Billing and Cost Management User Guide*.

Each delivery stream can have up to 50 tags.

This operation has a limit of five transactions per second per account.

Request Syntax

```
{
  "DeliveryStreamName": "string",
  "Tags": [
    {
      "Key": "string",
      "Value": "string"
    }
  ]
}
```

Request Parameters

The request accepts the following data in JSON format.

DeliveryStreamName (p. 49)

The name of the delivery stream to which you want to add the tags.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 64.

Pattern: [a-zA-Z0-9_.-]+

Required: Yes

Tags (p. 49)

A set of key-value pairs to use to create the tags.

Type: Array of [Tag](#) (p. 162) objects

Array Members: Minimum number of 1 item. Maximum number of 50 items.

Required: Yes

Response Elements

If the action is successful, the service sends back an HTTP 200 response with an empty HTTP body.

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 167).

InvalidArgumentException

The specified input parameter has a value that is not valid.

HTTP Status Code: 400

LimitExceededException

You have already reached the limit for a requested resource.

HTTP Status Code: 400

ResourceInUseException

The resource is already in use and not available for this operation.

HTTP Status Code: 400

ResourceNotFoundException

The specified resource could not be found.

HTTP Status Code: 400

Examples

To add tags to a stream

The following JSON example adds two tags to the specified stream.

Sample Request

```
POST / HTTP/1.1
Host: firehose.<region>.<domain>
Content-Length: <PayloadSizeBytes>
User-Agent: <UserAgentString>
Content-Type: application/x-amz-json-1.1
Authorization: <AuthParams>
Connection: Keep-Alive
X-Amz-Date: <Date>
X-Amz-Target: Firehose_20150804.TagDeliveryStream
{
  "DeliveryStreamName": "exampleDeliveryStreamName",
  "Tags": [
    {
      "Key": "Project",
      "Value": "myProject"
    },
    {
      "Key": "Environment",
      "Value": "Production"
    }
  ]
}
```

Sample Response

```
HTTP/1.1 200 OK
x-amzn-RequestId: <RequestId>
Content-Type: application/x-amz-json-1.1
Content-Length: <PayloadSizeBytes>
```

Date: <Date>

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

UntagDeliveryStream

Removes tags from the specified delivery stream. Removed tags are deleted, and you can't recover them after this operation successfully completes.

If you specify a tag that doesn't exist, the operation ignores it.

This operation has a limit of five transactions per second per account.

Request Syntax

```
{
  "DeliveryStreamName": "string",
  "TagKeys": [ "string" ]
}
```

Request Parameters

The request accepts the following data in JSON format.

DeliveryStreamName (p. 52)

The name of the delivery stream.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 64.

Pattern: [a-zA-Z0-9_.-]+

Required: Yes

TagKeys (p. 52)

A list of tag keys. Each corresponding tag is removed from the delivery stream.

Type: Array of strings

Array Members: Minimum number of 1 item. Maximum number of 50 items.

Length Constraints: Minimum length of 1. Maximum length of 128.

Pattern: ^(? !aws:)[\p{L}\p{Z}\p{N}_.:/=+\-@%]*\$

Required: Yes

Response Elements

If the action is successful, the service sends back an HTTP 200 response with an empty HTTP body.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 167\)](#).

InvalidArgumentException

The specified input parameter has a value that is not valid.

HTTP Status Code: 400

LimitExceededException

You have already reached the limit for a requested resource.

HTTP Status Code: 400

ResourceInUseException

The resource is already in use and not available for this operation.

HTTP Status Code: 400

ResourceNotFoundException

The specified resource could not be found.

HTTP Status Code: 400

Examples

To remove tags from a stream

The following JSON example removes the specified tag from the specified stream.

Sample Request

```
POST / HTTP/1.1
Host: firehose.<region>.<domain>
Content-Length: <PayloadSizeBytes>
User-Agent: <UserAgentString>
Content-Type: application/x-amz-json-1.1
Authorization: <AuthParams>
Connection: Keep-Alive
X-Amz-Date: <Date>
X-Amz-Target: Firehose_20150804.UntagDeliveryStream
{
  "DeliveryStreamName": "exampleDeliveryStreamName",
  "TagKeys": ["Project", "Environment"]
}
```

Sample Response

```
HTTP/1.1 200 OK
x-amzn-RequestId: <RequestId>
Content-Type: application/x-amz-json-1.1
Content-Length: <PayloadSizeBytes>
Date: <Date>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)

- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

UpdateDestination

Updates the specified destination of the specified delivery stream.

Use this operation to change the destination type (for example, to replace the Amazon S3 destination with Amazon Redshift) or change the parameters associated with a destination (for example, to change the bucket name of the Amazon S3 destination). The update might not occur immediately. The target delivery stream remains active while the configurations are updated, so data writes to the delivery stream can continue during this process. The updated configurations are usually effective within a few minutes.

Switching between Amazon ES and other services is not supported. For an Amazon ES destination, you can only update to another Amazon ES destination.

If the destination type is the same, Kinesis Data Firehose merges the configuration parameters specified with the destination configuration that already exists on the delivery stream. If any of the parameters are not specified in the call, the existing values are retained. For example, in the Amazon S3 destination, if [EncryptionConfiguration](#) (p. 94) is not specified, then the existing `EncryptionConfiguration` is maintained on the destination.

If the destination type is not the same, for example, changing the destination from Amazon S3 to Amazon Redshift, Kinesis Data Firehose does not merge any parameters. In this case, all parameters must be specified.

Kinesis Data Firehose uses `CurrentDeliveryStreamVersionId` to avoid race conditions and conflicting merges. This is a required field, and the service updates the configuration only if the existing configuration has a version ID that matches. After the update is applied successfully, the version ID is updated, and can be retrieved using [DescribeDeliveryStream](#) (p. 18). Use the new version ID to set `CurrentDeliveryStreamVersionId` in the next call.

Request Syntax

```
{
  "CurrentDeliveryStreamVersionId": "string",
  "DeliveryStreamName": "string",
  "DestinationId": "string",
  "ElasticsearchDestinationUpdate": {
    "BufferingHints": {
      "IntervalInSeconds": number,
      "SizeInMBs": number
    },
    "CloudWatchLoggingOptions": {
      "Enabled": boolean,
      "LogGroupName": "string",
      "LogStreamName": "string"
    },
    "ClusterEndpoint": "string",
    "DomainARN": "string",
    "IndexName": "string",
    "IndexRotationPeriod": "string",
    "ProcessingConfiguration": {
      "Enabled": boolean,
      "Processors": [
        {
          "Parameters": [
            {
              "ParameterName": "string",
              "ParameterValue": "string"
            }
          ]
        }
      ]
    }
  }
}
```

```

        ],
        "Type": "string"
    }
]
},
"RetryOptions": {
    "DurationInSeconds": number
},
"RoleARN": "string",
"S3Update": {
    "BucketARN": "string",
    "BufferingHints": {
        "IntervalInSeconds": number,
        "SizeInMBs": number
    },
    "CloudWatchLoggingOptions": {
        "Enabled": boolean,
        "LogGroupName": "string",
        "LogStreamName": "string"
    },
    "CompressionFormat": "string",
    "EncryptionConfiguration": {
        "KMSEncryptionConfig": {
            "AWSKMSKeyARN": "string"
        },
        "NoEncryptionConfig": "string"
    },
    "ErrorOutputPrefix": "string",
    "Prefix": "string",
    "RoleARN": "string"
},
"TypeName": "string"
},
"ExtendedS3DestinationUpdate": {
    "BucketARN": "string",
    "BufferingHints": {
        "IntervalInSeconds": number,
        "SizeInMBs": number
    },
    "CloudWatchLoggingOptions": {
        "Enabled": boolean,
        "LogGroupName": "string",
        "LogStreamName": "string"
    },
    "CompressionFormat": "string",
    "DataFormatConversionConfiguration": {
        "Enabled": boolean,
        "InputFormatConfiguration": {
            "Deserializer": {
                "HiveJsonSerDe": {
                    "TimestampFormats": [ "string" ]
                },
                "OpenXJsonSerDe": {
                    "CaseInsensitive": boolean,
                    "ColumnToJsonKeyMappings": {
                        "string": "string"
                    },
                    "ConvertDotsInJsonKeysToUnderscores": boolean
                }
            }
        },
        "OutputFormatConfiguration": {
            "Serializer": {
                "OrcSerDe": {
                    "BlockSizeBytes": number,
                    "BloomFilterColumns": [ "string" ],

```

```

        "BloomFilterFalsePositiveProbability": number,
        "Compression": "string",
        "DictionaryKeyThreshold": number,
        "EnablePadding": boolean,
        "FormatVersion": "string",
        "PaddingTolerance": number,
        "RowIndexStride": number,
        "StripeSizeBytes": number
    },
    "ParquetSerDe": {
        "BlockSizeBytes": number,
        "Compression": "string",
        "EnableDictionaryCompression": boolean,
        "MaxPaddingBytes": number,
        "PageSizeBytes": number,
        "WriterVersion": "string"
    }
},
"SchemaConfiguration": {
    "CatalogId": "string",
    "DatabaseName": "string",
    "Region": "string",
    "RoleARN": "string",
    "TableName": "string",
    "VersionId": "string"
}
},
"DynamicPartitioningConfiguration": {
    "Enabled": boolean,
    "RetryOptions": {
        "DurationInSeconds": number
    }
}
},
"EncryptionConfiguration": {
    "KMSEncryptionConfig": {
        "AWSKMSKeyARN": "string"
    },
    "NoEncryptionConfig": "string"
}
},
"ErrorOutputPrefix": "string",
"Prefix": "string",
"ProcessingConfiguration": {
    "Enabled": boolean,
    "Processors": [
        {
            "Parameters": [
                {
                    "ParameterName": "string",
                    "ParameterValue": "string"
                }
            ],
            "Type": "string"
        }
    ]
}
},
"RoleARN": "string",
"S3BackupMode": "string",
"S3BackupUpdate": {
    "BucketARN": "string",
    "BufferingHints": {
        "IntervalInSeconds": number,
        "SizeInMBs": number
    }
}
},
"CloudWatchLoggingOptions": {
    "Enabled": boolean,

```



```

        "LogGroupName": "string",
        "LogStreamName": "string"
    },
    "CompressionFormat": "string",
    "EncryptionConfiguration": {
        "KMSEncryptionConfig": {
            "AWSKMSKeyARN": "string"
        },
        "NoEncryptionConfig": "string"
    },
    "ErrorOutputPrefix": "string",
    "Prefix": "string",
    "RoleARN": "string"
}
},
"HttpEndpointDestinationUpdate": {
    "BufferingHints": {
        "IntervalInSeconds": number,
        "SizeInMBs": number
    },
    "CloudWatchLoggingOptions": {
        "Enabled": boolean,
        "LogGroupName": "string",
        "LogStreamName": "string"
    },
    "EndpointConfiguration": {
        "AccessKey": "string",
        "Name": "string",
        "Url": "string"
    },
    "ProcessingConfiguration": {
        "Enabled": boolean,
        "Processors": [
            {
                "Parameters": [
                    {
                        "ParameterName": "string",
                        "ParameterValue": "string"
                    }
                ],
                "Type": "string"
            }
        ]
    },
    "RequestConfiguration": {
        "CommonAttributes": [
            {
                "AttributeName": "string",
                "AttributeValue": "string"
            }
        ],
        "ContentEncoding": "string"
    },
    "RetryOptions": {
        "DurationInSeconds": number
    },
    "RoleARN": "string",
    "S3BackupMode": "string",
    "S3Update": {
        "BucketARN": "string",
        "BufferingHints": {
            "IntervalInSeconds": number,
            "SizeInMBs": number
        },
        "CloudWatchLoggingOptions": {
            "Enabled": boolean,

```

```
        "LogGroupName": "string",
        "LogStreamName": "string"
    },
    "CompressionFormat": "string",
    "EncryptionConfiguration": {
        "KMSEncryptionConfig": {
            "AWSKMSKeyARN": "string"
        },
        "NoEncryptionConfig": "string"
    },
    "ErrorOutputPrefix": "string",
    "Prefix": "string",
    "RoleARN": "string"
}
},
"RedshiftDestinationUpdate": {
    "CloudWatchLoggingOptions": {
        "Enabled": boolean,
        "LogGroupName": "string",
        "LogStreamName": "string"
    },
    "ClusterJDBCURL": "string",
    "CopyCommand": {
        "CopyOptions": "string",
        "DataTableColumns": "string",
        "DataTableName": "string"
    },
    "Password": "string",
    "ProcessingConfiguration": {
        "Enabled": boolean,
        "Processors": [
            {
                "Parameters": [
                    {
                        "ParameterName": "string",
                        "ParameterValue": "string"
                    }
                ],
                "Type": "string"
            }
        ]
    },
    "RetryOptions": {
        "DurationInSeconds": number
    },
    "RoleARN": "string",
    "S3BackupMode": "string",
    "S3BackupUpdate": {
        "BucketARN": "string",
        "BufferingHints": {
            "IntervalInSeconds": number,
            "SizeInMBs": number
        },
        "CloudWatchLoggingOptions": {
            "Enabled": boolean,
            "LogGroupName": "string",
            "LogStreamName": "string"
        },
        "CompressionFormat": "string",
        "EncryptionConfiguration": {
            "KMSEncryptionConfig": {
                "AWSKMSKeyARN": "string"
            },
            "NoEncryptionConfig": "string"
        },
        "ErrorOutputPrefix": "string",
```

```

    "Prefix": "string",
    "RoleARN": "string"
  },
  "S3Update": {
    "BucketARN": "string",
    "BufferingHints": {
      "IntervalInSeconds": number,
      "SizeInMBs": number
    },
    "CloudWatchLoggingOptions": {
      "Enabled": boolean,
      "LogGroupName": "string",
      "LogStreamName": "string"
    },
    "CompressionFormat": "string",
    "EncryptionConfiguration": {
      "KMSEncryptionConfig": {
        "AWSKMSKeyARN": "string"
      },
      "NoEncryptionConfig": "string"
    },
    "ErrorOutputPrefix": "string",
    "Prefix": "string",
    "RoleARN": "string"
  },
  "Username": "string"
},
"S3DestinationUpdate": {
  "BucketARN": "string",
  "BufferingHints": {
    "IntervalInSeconds": number,
    "SizeInMBs": number
  },
  "CloudWatchLoggingOptions": {
    "Enabled": boolean,
    "LogGroupName": "string",
    "LogStreamName": "string"
  },
  "CompressionFormat": "string",
  "EncryptionConfiguration": {
    "KMSEncryptionConfig": {
      "AWSKMSKeyARN": "string"
    },
    "NoEncryptionConfig": "string"
  },
  "ErrorOutputPrefix": "string",
  "Prefix": "string",
  "RoleARN": "string"
},
"SplunkDestinationUpdate": {
  "CloudWatchLoggingOptions": {
    "Enabled": boolean,
    "LogGroupName": "string",
    "LogStreamName": "string"
  },
  "HECAcknowledgmentTimeoutInSeconds": number,
  "HECEndpoint": "string",
  "HECEndpointType": "string",
  "HECToken": "string",
  "ProcessingConfiguration": {
    "Enabled": boolean,
    "Processors": [
      {
        "Parameters": [
          {
            "ParameterName": "string",

```

```
        "ParameterValue": "string"
      }
    ],
    "Type": "string"
  }
]
},
"RetryOptions": {
  "DurationInSeconds": number
},
"S3BackupMode": "string",
"S3Update": {
  "BucketARN": "string",
  "BufferingHints": {
    "IntervalInSeconds": number,
    "SizeInMBs": number
  },
  "CloudWatchLoggingOptions": {
    "Enabled": boolean,
    "LogGroupName": "string",
    "LogStreamName": "string"
  },
  "CompressionFormat": "string",
  "EncryptionConfiguration": {
    "KMSEncryptionConfig": {
      "AWSKMSKeyARN": "string"
    },
    "NoEncryptionConfig": "string"
  },
  "ErrorOutputPrefix": "string",
  "Prefix": "string",
  "RoleARN": "string"
}
}
```

Request Parameters

The request accepts the following data in JSON format.

CurrentDeliveryStreamVersionId (p. 55)

Obtain this value from the `VersionId` result of [DeliveryStreamDescription](#) (p. 72). This value is required, and helps the service perform conditional operations. For example, if there is an interleaving update and this value is null, then the update destination fails. After the update is successful, the `VersionId` value is updated. The service then performs a merge of the old configuration with the new configuration.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 50.

Pattern: [0-9]+

Required: Yes

DeliveryStreamName (p. 55)

The name of the delivery stream.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 64.

Pattern: [a-zA-Z0-9_.-]+

Required: Yes

DestinationId (p. 55)

The ID of the destination.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: [a-zA-Z0-9-]+

Required: Yes

ElasticsearchDestinationUpdate (p. 55)

Describes an update for a destination in Amazon ES.

Type: [ElasticsearchDestinationUpdate \(p. 90\)](#) object

Required: No

ExtendedS3DestinationUpdate (p. 55)

Describes an update for a destination in Amazon S3.

Type: [ExtendedS3DestinationUpdate \(p. 101\)](#) object

Required: No

HttpEndpointDestinationUpdate (p. 55)

Describes an update to the specified HTTP endpoint destination.

Type: [HttpEndpointDestinationUpdate \(p. 114\)](#) object

Required: No

RedshiftDestinationUpdate (p. 55)

Describes an update for a destination in Amazon Redshift.

Type: [RedshiftDestinationUpdate \(p. 140\)](#) object

Required: No

S3DestinationUpdate (p. 55)

[Deprecated] Describes an update for a destination in Amazon S3.

Type: [S3DestinationUpdate \(p. 149\)](#) object

Required: No

SplunkDestinationUpdate (p. 55)

Describes an update for a destination in Splunk.

Type: [SplunkDestinationUpdate \(p. 159\)](#) object

Required: No

Response Elements

If the action is successful, the service sends back an HTTP 200 response with an empty HTTP body.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 167\)](#).

ConcurrentModificationException

Another modification has already happened. Fetch `VersionId` again and use it to update the destination.

HTTP Status Code: 400

InvalidArgumentException

The specified input parameter has a value that is not valid.

HTTP Status Code: 400

ResourceInUseException

The resource is already in use and not available for this operation.

HTTP Status Code: 400

ResourceNotFoundException

The specified resource could not be found.

HTTP Status Code: 400

Examples

Example

The following JSON updates the buffering hints for the destination to 150 seconds and 2 MiB. You can obtain the other values that you need to use in this example from the response to a [DescribeDeliveryStream \(p. 18\)](#) invocation for the delivery stream in question.

Sample Request

```
POST / HTTP/1.1
Host: firehose.<region>.<domain>
Content-Length: <PayloadSizeBytes>
User-Agent: <UserAgentString>
Content-Type: application/x-amz-json-1.1
Authorization: <AuthParams>
Connection: Keep-Alive
X-Amz-Date: <Date>
X-Amz-Target: Firehose_20150804.UpdateDestination
{
  "CurrentDeliveryStreamVersionId": "1",
  "DeliveryStreamName": "exampleStreamName",
  "DestinationId": "destinationId-000000000001",
  "ExtendedS3DestinationUpdate": {
    "BucketARN": "arn:aws:s3:::somebucket",
    "BufferingHints": {
      "IntervalInSeconds": 150,
      "SizeInMBs": 2
    },
    "RoleARN": "arn:aws:iam::111222333444:role/exampleStreamName"
  }
}
```

```
}
```

Sample Response

```
HTTP/1.1 200 OK
x-amzn-RequestId: <RequestId>
Content-Type: application/x-amz-json-1.1
Content-Length: <PayloadSizeBytes>
Date: <Date>
```

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

Data Types

The Amazon Kinesis Firehose API contains several data types that various actions use. This section describes each data type in detail.

Note

The order of each element in a data type structure is not guaranteed. Applications should not assume a particular order.

The following data types are supported:

- [BufferingHints](#) (p. 67)
- [CloudWatchLoggingOptions](#) (p. 68)
- [CopyCommand](#) (p. 69)
- [DataFormatConversionConfiguration](#) (p. 71)
- [DeliveryStreamDescription](#) (p. 72)
- [DeliveryStreamEncryptionConfiguration](#) (p. 75)
- [DeliveryStreamEncryptionConfigurationInput](#) (p. 77)
- [Deserializer](#) (p. 79)
- [DestinationDescription](#) (p. 80)
- [DynamicPartitioningConfiguration](#) (p. 82)
- [ElasticsearchBufferingHints](#) (p. 83)
- [ElasticsearchDestinationConfiguration](#) (p. 84)
- [ElasticsearchDestinationDescription](#) (p. 87)
- [ElasticsearchDestinationUpdate](#) (p. 90)
- [ElasticsearchRetryOptions](#) (p. 93)
- [EncryptionConfiguration](#) (p. 94)
- [ExtendedS3DestinationConfiguration](#) (p. 95)
- [ExtendedS3DestinationDescription](#) (p. 98)
- [ExtendedS3DestinationUpdate](#) (p. 101)
- [FailureDescription](#) (p. 104)
- [HiveJsonSerDe](#) (p. 105)
- [HttpEndpointBufferingHints](#) (p. 106)
- [HttpEndpointCommonAttribute](#) (p. 107)
- [HttpEndpointConfiguration](#) (p. 108)
- [HttpEndpointDescription](#) (p. 109)
- [HttpEndpointDestinationConfiguration](#) (p. 110)
- [HttpEndpointDestinationDescription](#) (p. 112)
- [HttpEndpointDestinationUpdate](#) (p. 114)
- [HttpEndpointRequestConfiguration](#) (p. 116)
- [HttpEndpointRetryOptions](#) (p. 117)
- [InputFormatConfiguration](#) (p. 118)
- [KinesisStreamSourceConfiguration](#) (p. 119)
- [KinesisStreamSourceDescription](#) (p. 120)
- [KMSEncryptionConfig](#) (p. 121)
- [OpenXJsonSerDe](#) (p. 122)

- [OrcSerDe](#) (p. 124)
- [OutputFormatConfiguration](#) (p. 127)
- [ParquetSerDe](#) (p. 128)
- [ProcessingConfiguration](#) (p. 130)
- [Processor](#) (p. 131)
- [ProcessorParameter](#) (p. 132)
- [PutRecordBatchResponseEntry](#) (p. 133)
- [Record](#) (p. 134)
- [RedshiftDestinationConfiguration](#) (p. 135)
- [RedshiftDestinationDescription](#) (p. 138)
- [RedshiftDestinationUpdate](#) (p. 140)
- [RedshiftRetryOptions](#) (p. 143)
- [RetryOptions](#) (p. 144)
- [S3DestinationConfiguration](#) (p. 145)
- [S3DestinationDescription](#) (p. 147)
- [S3DestinationUpdate](#) (p. 149)
- [SchemaConfiguration](#) (p. 151)
- [Serializer](#) (p. 153)
- [SourceDescription](#) (p. 154)
- [SplunkDestinationConfiguration](#) (p. 155)
- [SplunkDestinationDescription](#) (p. 157)
- [SplunkDestinationUpdate](#) (p. 159)
- [SplunkRetryOptions](#) (p. 161)
- [Tag](#) (p. 162)
- [VpcConfiguration](#) (p. 163)
- [VpcConfigurationDescription](#) (p. 165)

BufferingHints

Describes hints for the buffering to perform before delivering data to the destination. These options are treated as hints, and therefore Kinesis Data Firehose might choose to use different values when it is optimal. The `SizeInMBs` and `IntervalInSeconds` parameters are optional. However, if specify a value for one of them, you must also provide a value for the other.

Contents

IntervalInSeconds

Buffer incoming data for the specified period of time, in seconds, before delivering it to the destination. The default value is 300. This parameter is optional but if you specify a value for it, you must also specify a value for `SizeInMBs`, and vice versa.

Type: Integer

Valid Range: Minimum value of 60. Maximum value of 900.

Required: No

SizeInMBs

Buffer incoming data to the specified size, in MiBs, before delivering it to the destination. The default value is 5. This parameter is optional but if you specify a value for it, you must also specify a value for `IntervalInSeconds`, and vice versa.

We recommend setting this parameter to a value greater than the amount of data you typically ingest into the delivery stream in 10 seconds. For example, if you typically ingest data at 1 MiB/sec, the value should be 10 MiB or higher.

Type: Integer

Valid Range: Minimum value of 1. Maximum value of 128.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

CloudWatchLoggingOptions

Describes the Amazon CloudWatch logging options for your delivery stream.

Contents

Enabled

Enables or disables CloudWatch logging.

Type: Boolean

Required: No

LogGroupName

The CloudWatch group name for logging. This value is required if CloudWatch logging is enabled.

Type: String

Length Constraints: Minimum length of 0. Maximum length of 512.

Pattern: [\ . \ _ / # A - Z a - z 0 - 9] *

Required: No

LogStreamName

The CloudWatch log stream name for logging. This value is required if CloudWatch logging is enabled.

Type: String

Length Constraints: Minimum length of 0. Maximum length of 512.

Pattern: [^ : *] *

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

CopyCommand

Describes a COPY command for Amazon Redshift.

Contents

CopyOptions

Optional parameters to use with the Amazon Redshift COPY command. For more information, see the "Optional Parameters" section of [Amazon Redshift COPY command](#). Some possible examples that would apply to Kinesis Data Firehose are as follows:

`delimiter '\t' lzop;` - fields are delimited with "\t" (TAB character) and compressed using lzop.

`delimiter '|' -` fields are delimited with "|" (this is the default delimiter).

`delimiter '|' escape -` the delimiter should be escaped.

`fixedwidth 'venueid:3,venueid:25,venueid:12,venueid:2,venueid:6' -` fields are fixed width in the source, with each width specified after every column in the table.

`JSON 's3://mybucket/jsonpaths.txt' -` data is in JSON format, and the path specified is the format of the data.

For more examples, see [Amazon Redshift COPY command examples](#).

Type: String

Length Constraints: Minimum length of 0. Maximum length of 204800.

Pattern: . *

Required: No

DataTableColumns

A comma-separated list of column names.

Type: String

Length Constraints: Minimum length of 0. Maximum length of 204800.

Pattern: . *

Required: No

DataTableName

The name of the target table. The table must already exist in the database.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: . *

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

DataFormatConversionConfiguration

Specifies that you want Kinesis Data Firehose to convert data from the JSON format to the Parquet or ORC format before writing it to Amazon S3. Kinesis Data Firehose uses the serializer and deserializer that you specify, in addition to the column information from the AWS Glue table, to deserialize your input data from JSON and then serialize it to the Parquet or ORC format. For more information, see [Kinesis Data Firehose Record Format Conversion](#).

Contents

Enabled

Defaults to `true`. Set it to `false` if you want to disable format conversion while preserving the configuration details.

Type: Boolean

Required: No

InputFormatConfiguration

Specifies the deserializer that you want Kinesis Data Firehose to use to convert the format of your data from JSON. This parameter is required if `Enabled` is set to `true`.

Type: [InputFormatConfiguration](#) (p. 118) object

Required: No

OutputFormatConfiguration

Specifies the serializer that you want Kinesis Data Firehose to use to convert the format of your data to the Parquet or ORC format. This parameter is required if `Enabled` is set to `true`.

Type: [OutputFormatConfiguration](#) (p. 127) object

Required: No

SchemaConfiguration

Specifies the AWS Glue Data Catalog table that contains the column information. This parameter is required if `Enabled` is set to `true`.

Type: [SchemaConfiguration](#) (p. 151) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

DeliveryStreamDescription

Contains information about a delivery stream.

Contents

CreateTimestamp

The date and time that the delivery stream was created.

Type: Timestamp

Required: No

DeliveryStreamARN

The Amazon Resource Name (ARN) of the delivery stream. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: Yes

DeliveryStreamEncryptionConfiguration

Indicates the server-side encryption (SSE) status for the delivery stream.

Type: [DeliveryStreamEncryptionConfiguration](#) (p. 75) object

Required: No

DeliveryStreamName

The name of the delivery stream.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 64.

Pattern: `[a-zA-Z0-9_.-]+`

Required: Yes

DeliveryStreamStatus

The status of the delivery stream. If the status of a delivery stream is `CREATING_FAILED`, this status doesn't change, and you can't invoke `CreateDeliveryStream` again on it. However, you can invoke the [DeleteDeliveryStream](#) (p. 15) operation to delete it.

Type: String

Valid Values: `CREATING` | `CREATING_FAILED` | `DELETING` | `DELETING_FAILED` | `ACTIVE`

Required: Yes

DeliveryStreamType

The delivery stream type. This can be one of the following values:

- `DirectPut`: Provider applications access the delivery stream directly.
- `KinesisStreamAsSource`: The delivery stream uses a Kinesis data stream as a source.

Type: String

Valid Values: `DirectPut` | `KinesisStreamAsSource`

Required: Yes

Destinations

The destinations.

Type: Array of [DestinationDescription](#) (p. 80) objects

Required: Yes

FailureDescription

Provides details in case one of the following operations fails due to an error related to KMS: [CreateDeliveryStream](#) (p. 3), [DeleteDeliveryStream](#) (p. 15), [StartDeliveryStreamEncryption](#) (p. 43), [StopDeliveryStreamEncryption](#) (p. 46).

Type: [FailureDescription](#) (p. 104) object

Required: No

HasMoreDestinations

Indicates whether there are more destinations available to list.

Type: Boolean

Required: Yes

LastUpdateTimestamp

The date and time that the delivery stream was last updated.

Type: Timestamp

Required: No

Source

If the `DeliveryStreamType` parameter is `KinesisStreamAsSource`, a [SourceDescription](#) (p. 154) object describing the source Kinesis data stream.

Type: [SourceDescription](#) (p. 154) object

Required: No

VersionId

Each time the destination is updated for a delivery stream, the version ID is changed, and the current version ID is required when updating the destination. This is so that the service knows it is applying the changes to the correct version of the delivery stream.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 50.

Pattern: `[0-9]+`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

DeliveryStreamEncryptionConfiguration

Contains information about the server-side encryption (SSE) status for the delivery stream, the type customer master key (CMK) in use, if any, and the ARN of the CMK. You can get `DeliveryStreamEncryptionConfiguration` by invoking the [DescribeDeliveryStream](#) (p. 18) operation.

Contents

FailureDescription

Provides details in case one of the following operations fails due to an error related to KMS: [CreateDeliveryStream](#) (p. 3), [DeleteDeliveryStream](#) (p. 15), [StartDeliveryStreamEncryption](#) (p. 43), [StopDeliveryStreamEncryption](#) (p. 46).

Type: [FailureDescription](#) (p. 104) object

Required: No

KeyARN

If `KeyType` is `CUSTOMER_MANAGED_CMK`, this field contains the ARN of the customer managed CMK. If `KeyType` is `AWS_OWNED_CMK`, `DeliveryStreamEncryptionConfiguration` doesn't contain a value for `KeyARN`.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: No

KeyType

Indicates the type of customer master key (CMK) that is used for encryption. The default setting is `AWS_OWNED_CMK`. For more information about CMKs, see [Customer Master Keys \(CMKs\)](#).

Type: String

Valid Values: `AWS_OWNED_CMK` | `CUSTOMER_MANAGED_CMK`

Required: No

Status

This is the server-side encryption (SSE) status for the delivery stream. For a full description of the different values of this status, see [StartDeliveryStreamEncryption](#) (p. 43) and [StopDeliveryStreamEncryption](#) (p. 46). If this status is `ENABLING_FAILED` or `DISABLING_FAILED`, it is the status of the most recent attempt to enable or disable SSE, respectively.

Type: String

Valid Values: `ENABLED` | `ENABLING` | `ENABLING_FAILED` | `DISABLED` | `DISABLING` | `DISABLING_FAILED`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

DeliveryStreamEncryptionConfigurationInput

Specifies the type and Amazon Resource Name (ARN) of the CMK to use for Server-Side Encryption (SSE).

Contents

KeyARN

If you set `KeyType` to `CUSTOMER_MANAGED_CMK`, you must specify the Amazon Resource Name (ARN) of the CMK. If you set `KeyType` to `AWS_OWNED_CMK`, Kinesis Data Firehose uses a service-account CMK.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: No

KeyType

Indicates the type of customer master key (CMK) to use for encryption. The default setting is `AWS_OWNED_CMK`. For more information about CMKs, see [Customer Master Keys \(CMKs\)](#). When you invoke [CreateDeliveryStream](#) (p. 3) or [StartDeliveryStreamEncryption](#) (p. 43) with `KeyType` set to `CUSTOMER_MANAGED_CMK`, Kinesis Data Firehose invokes the Amazon KMS operation [CreateGrant](#) to create a grant that allows the Kinesis Data Firehose service to use the customer managed CMK to perform encryption and decryption. Kinesis Data Firehose manages that grant.

When you invoke [StartDeliveryStreamEncryption](#) (p. 43) to change the CMK for a delivery stream that is encrypted with a customer managed CMK, Kinesis Data Firehose schedules the grant it had on the old CMK for retirement.

You can use a CMK of type `CUSTOMER_MANAGED_CMK` to encrypt up to 500 delivery streams. If a [CreateDeliveryStream](#) (p. 3) or [StartDeliveryStreamEncryption](#) (p. 43) operation exceeds this limit, Kinesis Data Firehose throws a `LimitExceededException`.

Important

To encrypt your delivery stream, use symmetric CMKs. Kinesis Data Firehose doesn't support asymmetric CMKs. For information about symmetric and asymmetric CMKs, see [About Symmetric and Asymmetric CMKs](#) in the AWS Key Management Service developer guide.

Type: String

Valid Values: `AWS_OWNED_CMK` | `CUSTOMER_MANAGED_CMK`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

Deserializer

The deserializer you want Kinesis Data Firehose to use for converting the input data from JSON. Kinesis Data Firehose then serializes the data to its final format using the [Serializer](#) (p. 153). Kinesis Data Firehose supports two types of deserializers: the [Apache Hive JSON SerDe](#) and the [OpenX JSON SerDe](#).

Contents

HiveJsonSerDe

The native Hive / HCatalog JsonSerDe. Used by Kinesis Data Firehose for deserializing data, which means converting it from the JSON format in preparation for serializing it to the Parquet or ORC format. This is one of two deserializers you can choose, depending on which one offers the functionality you need. The other option is the OpenX SerDe.

Type: [HiveJsonSerDe](#) (p. 105) object

Required: No

OpenXJsonSerDe

The OpenX SerDe. Used by Kinesis Data Firehose for deserializing data, which means converting it from the JSON format in preparation for serializing it to the Parquet or ORC format. This is one of two deserializers you can choose, depending on which one offers the functionality you need. The other option is the native Hive / HCatalog JsonSerDe.

Type: [OpenXJsonSerDe](#) (p. 122) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

DestinationDescription

Describes the destination for a delivery stream.

Contents

DestinationId

The ID of the destination.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: [a-zA-Z0-9-]+

Required: Yes

ElasticsearchDestinationDescription

The destination in Amazon ES.

Type: [ElasticsearchDestinationDescription](#) (p. 87) object

Required: No

ExtendedS3DestinationDescription

The destination in Amazon S3.

Type: [ExtendedS3DestinationDescription](#) (p. 98) object

Required: No

HttpEndpointDestinationDescription

Describes the specified HTTP endpoint destination.

Type: [HttpEndpointDestinationDescription](#) (p. 112) object

Required: No

RedshiftDestinationDescription

The destination in Amazon Redshift.

Type: [RedshiftDestinationDescription](#) (p. 138) object

Required: No

S3DestinationDescription

[Deprecated] The destination in Amazon S3.

Type: [S3DestinationDescription](#) (p. 147) object

Required: No

SplunkDestinationDescription

The destination in Splunk.

Type: [SplunkDestinationDescription](#) (p. 157) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

DynamicPartitioningConfiguration

The configuration of the dynamic partitioning mechanism that creates smaller data sets from the streaming data by partitioning it based on partition keys. Currently, dynamic partitioning is only supported for Amazon S3 destinations.

Contents

Enabled

Specifies that the dynamic partitioning is enabled for this Kinesis Data Firehose delivery stream.

Type: Boolean

Required: No

RetryOptions

The retry behavior in case Kinesis Data Firehose is unable to deliver data to an Amazon S3 prefix.

Type: [RetryOptions](#) (p. 144) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ElasticsearchBufferingHints

Describes the buffering to perform before delivering data to the Amazon ES destination.

Contents

IntervalInSeconds

Buffer incoming data for the specified period of time, in seconds, before delivering it to the destination. The default value is 300 (5 minutes).

Type: Integer

Valid Range: Minimum value of 60. Maximum value of 900.

Required: No

SizeInMBs

Buffer incoming data to the specified size, in MBs, before delivering it to the destination. The default value is 5.

We recommend setting this parameter to a value greater than the amount of data you typically ingest into the delivery stream in 10 seconds. For example, if you typically ingest data at 1 MB/sec, the value should be 10 MB or higher.

Type: Integer

Valid Range: Minimum value of 1. Maximum value of 100.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ElasticsearchDestinationConfiguration

Describes the configuration of a destination in Amazon ES.

Contents

BufferingHints

The buffering options. If no value is specified, the default values for `ElasticsearchBufferingHints` are used.

Type: [ElasticsearchBufferingHints](#) (p. 83) object

Required: No

CloudWatchLoggingOptions

The Amazon CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

ClusterEndpoint

The endpoint to use when communicating with the cluster. Specify either this `ClusterEndpoint` or the `DomainARN` field.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `https:.*`

Required: No

DomainARN

The ARN of the Amazon ES domain. The IAM role must have permissions for `DescribeElasticsearchDomain`, `DescribeElasticsearchDomains`, and `DescribeElasticsearchDomainConfig` after assuming the role specified in **RoleARN**. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Specify either `ClusterEndpoint` or `DomainARN`.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: No

IndexName

The Elasticsearch index name.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 80.

Pattern: `.*`

Required: Yes

IndexRotationPeriod

The Elasticsearch index rotation period. Index rotation appends a timestamp to the `IndexName` to facilitate the expiration of old data. For more information, see [Index Rotation for the Amazon ES Destination](#). The default value is `OneDay`.

Type: String

Valid Values: `NoRotation` | `OneHour` | `OneDay` | `OneWeek` | `OneMonth`

Required: No

ProcessingConfiguration

The data processing configuration.

Type: [ProcessingConfiguration](#) (p. 130) object

Required: No

RetryOptions

The retry behavior in case Kinesis Data Firehose is unable to deliver documents to Amazon ES. The default value is 300 (5 minutes).

Type: [ElasticsearchRetryOptions](#) (p. 93) object

Required: No

RoleARN

The Amazon Resource Name (ARN) of the IAM role to be assumed by Kinesis Data Firehose for calling the Amazon ES Configuration API and for indexing documents. For more information, see [Grant Kinesis Data Firehose Access to an Amazon S3 Destination](#) and [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: Yes

S3BackupMode

Defines how documents should be delivered to Amazon S3. When it is set to `FailedDocumentsOnly`, Kinesis Data Firehose writes any documents that could not be indexed to the configured Amazon S3 destination, with `elasticsearch-failed/` appended to the key prefix. When set to `AllDocuments`, Kinesis Data Firehose delivers all incoming records to Amazon S3, and also writes failed documents with `elasticsearch-failed/` appended to the prefix. For more information, see [Amazon S3 Backup for the Amazon ES Destination](#). Default value is `FailedDocumentsOnly`.

You can't change this backup mode after you create the delivery stream.

Type: String

Valid Values: `FailedDocumentsOnly` | `AllDocuments`

Required: No

S3Configuration

The configuration for the backup Amazon S3 location.

Type: [S3DestinationConfiguration](#) (p. 145) object

Required: Yes

TypeName

The Elasticsearch type name. For Elasticsearch 6.x, there can be only one type per index. If you try to specify a new type for an existing index that already has another type, Kinesis Data Firehose returns an error during run time.

For Elasticsearch 7.x, don't specify a `TypeName`.

Type: String

Length Constraints: Minimum length of 0. Maximum length of 100.

Pattern: `. *`

Required: No

VpcConfiguration

The details of the VPC of the Amazon ES destination.

Type: [VpcConfiguration](#) (p. 163) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ElasticsearchDestinationDescription

The destination description in Amazon ES.

Contents

BufferingHints

The buffering options.

Type: [ElasticsearchBufferingHints](#) (p. 83) object

Required: No

CloudWatchLoggingOptions

The Amazon CloudWatch logging options.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

ClusterEndpoint

The endpoint to use when communicating with the cluster. Kinesis Data Firehose uses either this `ClusterEndpoint` or the `DomainARN` field to send data to Amazon ES.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `https:.*`

Required: No

DomainARN

The ARN of the Amazon ES domain. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Kinesis Data Firehose uses either `ClusterEndpoint` or `DomainARN` to send data to Amazon ES.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: No

IndexName

The Elasticsearch index name.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 80.

Pattern: `.*`

Required: No

IndexRotationPeriod

The Elasticsearch index rotation period

Type: String

Valid Values: `NoRotation` | `OneHour` | `OneDay` | `OneWeek` | `OneMonth`

Required: No

ProcessingConfiguration

The data processing configuration.

Type: [ProcessingConfiguration](#) (p. 130) object

Required: No

RetryOptions

The Amazon ES retry options.

Type: [ElasticsearchRetryOptions](#) (p. 93) object

Required: No

RoleARN

The Amazon Resource Name (ARN) of the AWS credentials. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: No

S3BackupMode

The Amazon S3 backup mode.

Type: String

Valid Values: `FailedDocumentsOnly` | `AllDocuments`

Required: No

S3DestinationDescription

The Amazon S3 destination.

Type: [S3DestinationDescription](#) (p. 147) object

Required: No

TypeName

The Elasticsearch type name. This applies to Elasticsearch 6.x and lower versions. For Elasticsearch 7.x, there's no value for `TypeName`.

Type: String

Length Constraints: Minimum length of 0. Maximum length of 100.

Pattern: `.*`

Required: No

VpcConfigurationDescription

The details of the VPC of the Amazon ES destination.

Type: [VpcConfigurationDescription](#) (p. 165) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ElasticsearchDestinationUpdate

Describes an update for a destination in Amazon ES.

Contents

BufferingHints

The buffering options. If no value is specified, `ElasticsearchBufferingHints` object default values are used.

Type: [ElasticsearchBufferingHints](#) (p. 83) object

Required: No

CloudWatchLoggingOptions

The CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

ClusterEndpoint

The endpoint to use when communicating with the cluster. Specify either this `ClusterEndpoint` or the `DomainARN` field.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `https:.*`

Required: No

DomainARN

The ARN of the Amazon ES domain. The IAM role must have permissions for `DescribeElasticsearchDomain`, `DescribeElasticsearchDomains`, and `DescribeElasticsearchDomainConfig` after assuming the IAM role specified in `RoleARN`. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Specify either `ClusterEndpoint` or `DomainARN`.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: No

IndexName

The Elasticsearch index name.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 80.

Pattern: `.*`

Required: No

IndexRotationPeriod

The Elasticsearch index rotation period. Index rotation appends a timestamp to `IndexName` to facilitate the expiration of old data. For more information, see [Index Rotation for the Amazon ES Destination](#). Default value is `OneDay`.

Type: String

Valid Values: `NoRotation` | `OneHour` | `OneDay` | `OneWeek` | `OneMonth`

Required: No

ProcessingConfiguration

The data processing configuration.

Type: [ProcessingConfiguration](#) (p. 130) object

Required: No

RetryOptions

The retry behavior in case Kinesis Data Firehose is unable to deliver documents to Amazon ES. The default value is 300 (5 minutes).

Type: [ElasticsearchRetryOptions](#) (p. 93) object

Required: No

RoleARN

The Amazon Resource Name (ARN) of the IAM role to be assumed by Kinesis Data Firehose for calling the Amazon ES Configuration API and for indexing documents. For more information, see [Grant Kinesis Data Firehose Access to an Amazon S3 Destination](#) and [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: No

S3Update

The Amazon S3 destination.

Type: [S3DestinationUpdate](#) (p. 149) object

Required: No

TypeName

The Elasticsearch type name. For Elasticsearch 6.x, there can be only one type per index. If you try to specify a new type for an existing index that already has another type, Kinesis Data Firehose returns an error during runtime.

If you upgrade Elasticsearch from 6.x to 7.x and don't update your delivery stream, Kinesis Data Firehose still delivers data to Elasticsearch with the old index name and type name. If you want to update your delivery stream with a new index name, provide an empty string for `TypeName`.

Type: String

Length Constraints: Minimum length of 0. Maximum length of 100.

Pattern: . *

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ElasticsearchRetryOptions

Configures retry behavior in case Kinesis Data Firehose is unable to deliver documents to Amazon ES.

Contents

DurationInSeconds

After an initial failure to deliver to Amazon ES, the total amount of time during which Kinesis Data Firehose retries delivery (including the first attempt). After this time has elapsed, the failed documents are written to Amazon S3. Default value is 300 seconds (5 minutes). A value of 0 (zero) results in no retries.

Type: Integer

Valid Range: Minimum value of 0. Maximum value of 7200.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

EncryptionConfiguration

Describes the encryption for a destination in Amazon S3.

Contents

KMSEncryptionConfig

The encryption key.

Type: [KMSEncryptionConfig](#) (p. 121) object

Required: No

NoEncryptionConfig

Specifically override existing encryption information to ensure that no encryption is used.

Type: String

Valid Values: `NoEncryption`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ExtendedS3DestinationConfiguration

Describes the configuration of a destination in Amazon S3.

Contents

BucketARN

The ARN of the S3 bucket. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: `arn:.*`

Required: Yes

BufferingHints

The buffering option.

Type: [BufferingHints](#) (p. 67) object

Required: No

CloudWatchLoggingOptions

The Amazon CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

CompressionFormat

The compression format. If no value is specified, the default is UNCOMPRESSED.

Type: String

Valid Values: `UNCOMPRESSED` | `GZIP` | `ZIP` | `Snappy` | `HADOOP_SNAPPY`

Required: No

DataFormatConversionConfiguration

The serializer, deserializer, and schema for converting data from the JSON format to the Parquet or ORC format before writing it to Amazon S3.

Type: [DataFormatConversionConfiguration](#) (p. 71) object

Required: No

DynamicPartitioningConfiguration

The configuration of the dynamic partitioning mechanism that creates smaller data sets from the streaming data by partitioning it based on partition keys. Currently, dynamic partitioning is only supported for Amazon S3 destinations.

Type: [DynamicPartitioningConfiguration](#) (p. 82) object

Required: No

EncryptionConfiguration

The encryption configuration. If no value is specified, the default is no encryption.

Type: [EncryptionConfiguration](#) (p. 94) object

Required: No

ErrorOutputPrefix

A prefix that Kinesis Data Firehose evaluates and adds to failed records before writing them to S3. This prefix appears immediately following the bucket name. For information about how to specify this prefix, see [Custom Prefixes for Amazon S3 Objects](#).

Type: String

Length Constraints: Minimum length of 0. Maximum length of 1024.

Pattern: . *

Required: No

Prefix

The "YYYY/MM/DD/HH" time format prefix is automatically used for delivered Amazon S3 files. You can also specify a custom prefix, as described in [Custom Prefixes for Amazon S3 Objects](#).

Type: String

Length Constraints: Minimum length of 0. Maximum length of 1024.

Pattern: . *

Required: No

ProcessingConfiguration

The data processing configuration.

Type: [ProcessingConfiguration](#) (p. 130) object

Required: No

RoleARN

The Amazon Resource Name (ARN) of the AWS credentials. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: arn:.*

Required: Yes

S3BackupConfiguration

The configuration for backup in Amazon S3.

Type: [S3DestinationConfiguration](#) (p. 145) object

Required: No

S3BackupMode

The Amazon S3 backup mode. After you create a delivery stream, you can update it to enable Amazon S3 backup if it is disabled. If backup is enabled, you can't update the delivery stream to disable it.

Type: String

Valid Values: `Disabled` | `Enabled`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ExtendedS3DestinationDescription

Describes a destination in Amazon S3.

Contents

BucketARN

The ARN of the S3 bucket. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: arn:.*

Required: Yes

BufferingHints

The buffering option.

Type: [BufferingHints](#) (p. 67) object

Required: Yes

CloudWatchLoggingOptions

The Amazon CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

CompressionFormat

The compression format. If no value is specified, the default is UNCOMPRESSED.

Type: String

Valid Values: UNCOMPRESSED | GZIP | ZIP | Snappy | HADOOP_SNAPPY

Required: Yes

DataFormatConversionConfiguration

The serializer, deserializer, and schema for converting data from the JSON format to the Parquet or ORC format before writing it to Amazon S3.

Type: [DataFormatConversionConfiguration](#) (p. 71) object

Required: No

DynamicPartitioningConfiguration

The configuration of the dynamic partitioning mechanism that creates smaller data sets from the streaming data by partitioning it based on partition keys. Currently, dynamic partitioning is only supported for Amazon S3 destinations.

Type: [DynamicPartitioningConfiguration](#) (p. 82) object

Required: No

EncryptionConfiguration

The encryption configuration. If no value is specified, the default is no encryption.

Type: [EncryptionConfiguration](#) (p. 94) object

Required: Yes

ErrorOutputPrefix

A prefix that Kinesis Data Firehose evaluates and adds to failed records before writing them to S3. This prefix appears immediately following the bucket name. For information about how to specify this prefix, see [Custom Prefixes for Amazon S3 Objects](#).

Type: String

Length Constraints: Minimum length of 0. Maximum length of 1024.

Pattern: . *

Required: No

Prefix

The "YYYY/MM/DD/HH" time format prefix is automatically used for delivered Amazon S3 files. You can also specify a custom prefix, as described in [Custom Prefixes for Amazon S3 Objects](#).

Type: String

Length Constraints: Minimum length of 0. Maximum length of 1024.

Pattern: . *

Required: No

ProcessingConfiguration

The data processing configuration.

Type: [ProcessingConfiguration](#) (p. 130) object

Required: No

RoleARN

The Amazon Resource Name (ARN) of the AWS credentials. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: arn:.*

Required: Yes

S3BackupDescription

The configuration for backup in Amazon S3.

Type: [S3DestinationDescription](#) (p. 147) object

Required: No

S3BackupMode

The Amazon S3 backup mode.

Type: String

Valid Values: Disabled | Enabled

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ExtendedS3DestinationUpdate

Describes an update for a destination in Amazon S3.

Contents

BucketARN

The ARN of the S3 bucket. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: arn:.*

Required: No

BufferingHints

The buffering option.

Type: [BufferingHints](#) (p. 67) object

Required: No

CloudWatchLoggingOptions

The Amazon CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

CompressionFormat

The compression format. If no value is specified, the default is UNCOMPRESSED.

Type: String

Valid Values: UNCOMPRESSED | GZIP | ZIP | Snappy | HADOOP_SNAPPY

Required: No

DataFormatConversionConfiguration

The serializer, deserializer, and schema for converting data from the JSON format to the Parquet or ORC format before writing it to Amazon S3.

Type: [DataFormatConversionConfiguration](#) (p. 71) object

Required: No

DynamicPartitioningConfiguration

The configuration of the dynamic partitioning mechanism that creates smaller data sets from the streaming data by partitioning it based on partition keys. Currently, dynamic partitioning is only supported for Amazon S3 destinations.

Type: [DynamicPartitioningConfiguration](#) (p. 82) object

Required: No

EncryptionConfiguration

The encryption configuration. If no value is specified, the default is no encryption.

Type: [EncryptionConfiguration](#) (p. 94) object

Required: No

ErrorOutputPrefix

A prefix that Kinesis Data Firehose evaluates and adds to failed records before writing them to S3. This prefix appears immediately following the bucket name. For information about how to specify this prefix, see [Custom Prefixes for Amazon S3 Objects](#).

Type: String

Length Constraints: Minimum length of 0. Maximum length of 1024.

Pattern: . *

Required: No

Prefix

The "YYYY/MM/DD/HH" time format prefix is automatically used for delivered Amazon S3 files. You can also specify a custom prefix, as described in [Custom Prefixes for Amazon S3 Objects](#).

Type: String

Length Constraints: Minimum length of 0. Maximum length of 1024.

Pattern: . *

Required: No

ProcessingConfiguration

The data processing configuration.

Type: [ProcessingConfiguration](#) (p. 130) object

Required: No

RoleARN

The Amazon Resource Name (ARN) of the AWS credentials. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: arn:.*

Required: No

S3BackupMode

You can update a delivery stream to enable Amazon S3 backup if it is disabled. If backup is enabled, you can't update the delivery stream to disable it.

Type: String

Valid Values: Disabled | Enabled

Required: No

S3BackupUpdate

The Amazon S3 destination for backup.

Type: [S3DestinationUpdate](#) (p. 149) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

FailureDescription

Provides details in case one of the following operations fails due to an error related to KMS: [CreateDeliveryStream](#) (p. 3), [DeleteDeliveryStream](#) (p. 15), [StartDeliveryStreamEncryption](#) (p. 43), [StopDeliveryStreamEncryption](#) (p. 46).

Contents

Details

A message providing details about the error that caused the failure.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1024.

Pattern: `^(?!\\s*$).+`

Required: Yes

Type

The type of error that caused the failure.

Type: String

Valid Values: RETIRE_KMS_GRANT_FAILED | CREATE_KMS_GRANT_FAILED | KMS_ACCESS_DENIED | DISABLED_KMS_KEY | INVALID_KMS_KEY | KMS_KEY_NOT_FOUND | KMS_OPT_IN_REQUIRED | CREATE_ENI_FAILED | DELETE_ENI_FAILED | SUBNET_NOT_FOUND | SECURITY_GROUP_NOT_FOUND | ENI_ACCESS_DENIED | SUBNET_ACCESS_DENIED | SECURITY_GROUP_ACCESS_DENIED | UNKNOWN_ERROR

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

HiveJsonSerDe

The native Hive / HCatalog JsonSerDe. Used by Kinesis Data Firehose for deserializing data, which means converting it from the JSON format in preparation for serializing it to the Parquet or ORC format. This is one of two deserializers you can choose, depending on which one offers the functionality you need. The other option is the OpenX SerDe.

Contents

TimestampFormats

Indicates how you want Kinesis Data Firehose to parse the date and timestamps that may be present in your input data JSON. To specify these format strings, follow the pattern syntax of JodaTime's `DateTimeFormat` format strings. For more information, see [Class DateTimeFormat](#). You can also use the special value `millis` to parse timestamps in epoch milliseconds. If you don't specify a format, Kinesis Data Firehose uses `java.sql.Timestamp::valueOf` by default.

Type: Array of strings

Length Constraints: Minimum length of 1. Maximum length of 1024.

Pattern: `^(?!\\s*$).+`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

HttpEndpointBufferingHints

Describes the buffering options that can be applied before data is delivered to the HTTP endpoint destination. Kinesis Data Firehose treats these options as hints, and it might choose to use more optimal values. The `SizeInMBs` and `IntervalInSeconds` parameters are optional. However, if specify a value for one of them, you must also provide a value for the other.

Contents

IntervalInSeconds

Buffer incoming data for the specified period of time, in seconds, before delivering it to the destination. The default value is 300 (5 minutes).

Type: Integer

Valid Range: Minimum value of 60. Maximum value of 900.

Required: No

SizeInMBs

Buffer incoming data to the specified size, in MBs, before delivering it to the destination. The default value is 5.

We recommend setting this parameter to a value greater than the amount of data you typically ingest into the delivery stream in 10 seconds. For example, if you typically ingest data at 1 MB/sec, the value should be 10 MB or higher.

Type: Integer

Valid Range: Minimum value of 1. Maximum value of 64.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

HttpEndpointCommonAttribute

Describes the metadata that's delivered to the specified HTTP endpoint destination.

Contents

AttributeName

The name of the HTTP endpoint common attribute.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 256.

Pattern: `^(?!\\s*$)\\.+`

Required: Yes

AttributeValue

The value of the HTTP endpoint common attribute.

Type: String

Length Constraints: Minimum length of 0. Maximum length of 1024.

Pattern: `.*`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

HttpEndpointConfiguration

Describes the configuration of the HTTP endpoint to which Kinesis Firehose delivers data.

Contents

AccessKey

The access key required for Kinesis Firehose to authenticate with the HTTP endpoint selected as the destination.

Type: String

Length Constraints: Minimum length of 0. Maximum length of 4096.

Pattern: . *

Required: No

Name

The name of the HTTP endpoint selected as the destination.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 256.

Pattern: ^(?!\s*\$) . +

Required: No

Url

The URL of the HTTP endpoint selected as the destination.

Important

If you choose an HTTP endpoint as your destination, review and follow the instructions in the [Appendix - HTTP Endpoint Delivery Request and Response Specifications](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: https:// . *

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

HttpEndpointDescription

Describes the HTTP endpoint selected as the destination.

Contents

Name

The name of the HTTP endpoint selected as the destination.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 256.

Pattern: `^(?!\\s*$).+`

Required: No

Url

The URL of the HTTP endpoint selected as the destination.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `https://.*`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

HttpEndpointDestinationConfiguration

Describes the configuration of the HTTP endpoint destination.

Contents

BufferingHints

The buffering options that can be used before data is delivered to the specified destination. Kinesis Data Firehose treats these options as hints, and it might choose to use more optimal values. The `SizeInMBs` and `IntervalInSeconds` parameters are optional. However, if you specify a value for one of them, you must also provide a value for the other.

Type: [HttpEndpointBufferingHints](#) (p. 106) object

Required: No

CloudWatchLoggingOptions

Describes the Amazon CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

EndpointConfiguration

The configuration of the HTTP endpoint selected as the destination.

Type: [HttpEndpointConfiguration](#) (p. 108) object

Required: Yes

ProcessingConfiguration

Describes a data processing configuration.

Type: [ProcessingConfiguration](#) (p. 130) object

Required: No

RequestConfiguration

The configuration of the request sent to the HTTP endpoint specified as the destination.

Type: [HttpEndpointRequestConfiguration](#) (p. 116) object

Required: No

RetryOptions

Describes the retry behavior in case Kinesis Data Firehose is unable to deliver data to the specified HTTP endpoint destination, or if it doesn't receive a valid acknowledgment of receipt from the specified HTTP endpoint destination.

Type: [HttpEndpointRetryOptions](#) (p. 117) object

Required: No

RoleARN

Kinesis Data Firehose uses this IAM role for all the permissions that the delivery stream needs.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: No

S3BackupMode

Describes the S3 bucket backup options for the data that Kinesis Data Firehose delivers to the HTTP endpoint destination. You can back up all documents (`AllData`) or only the documents that Kinesis Data Firehose could not deliver to the specified HTTP endpoint destination (`FailedDataOnly`).

Type: String

Valid Values: `FailedDataOnly` | `AllData`

Required: No

S3Configuration

Describes the configuration of a destination in Amazon S3.

Type: [S3DestinationConfiguration](#) (p. 145) object

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

HttpEndpointDestinationDescription

Describes the HTTP endpoint destination.

Contents

BufferingHints

Describes buffering options that can be applied to the data before it is delivered to the HTTPS endpoint destination. Kinesis Data Firehose treats these options as hints, and it might choose to use more optimal values. The `SizeInMBs` and `IntervalInSeconds` parameters are optional. However, if specify a value for one of them, you must also provide a value for the other.

Type: [HttpEndpointBufferingHints](#) (p. 106) object

Required: No

CloudWatchLoggingOptions

Describes the Amazon CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

EndpointConfiguration

The configuration of the specified HTTP endpoint destination.

Type: [HttpEndpointDescription](#) (p. 109) object

Required: No

ProcessingConfiguration

Describes a data processing configuration.

Type: [ProcessingConfiguration](#) (p. 130) object

Required: No

RequestConfiguration

The configuration of request sent to the HTTP endpoint specified as the destination.

Type: [HttpEndpointRequestConfiguration](#) (p. 116) object

Required: No

RetryOptions

Describes the retry behavior in case Kinesis Data Firehose is unable to deliver data to the specified HTTP endpoint destination, or if it doesn't receive a valid acknowledgment of receipt from the specified HTTP endpoint destination.

Type: [HttpEndpointRetryOptions](#) (p. 117) object

Required: No

RoleARN

Kinesis Data Firehose uses this IAM role for all the permissions that the delivery stream needs.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: No

S3BackupMode

Describes the S3 bucket backup options for the data that Kinesis Firehose delivers to the HTTP endpoint destination. You can back up all documents (`AllData`) or only the documents that Kinesis Data Firehose could not deliver to the specified HTTP endpoint destination (`FailedDataOnly`).

Type: String

Valid Values: `FailedDataOnly` | `AllData`

Required: No

S3DestinationDescription

Describes a destination in Amazon S3.

Type: [S3DestinationDescription](#) (p. 147) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

HttpEndpointDestinationUpdate

Updates the specified HTTP endpoint destination.

Contents

BufferingHints

Describes buffering options that can be applied to the data before it is delivered to the HTTPS endpoint destination. Kinesis Data Firehose treats these options as hints, and it might choose to use more optimal values. The `SizeInMBs` and `IntervalInSeconds` parameters are optional. However, if specify a value for one of them, you must also provide a value for the other.

Type: [HttpEndpointBufferingHints](#) (p. 106) object

Required: No

CloudWatchLoggingOptions

Describes the Amazon CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

EndpointConfiguration

Describes the configuration of the HTTP endpoint destination.

Type: [HttpEndpointConfiguration](#) (p. 108) object

Required: No

ProcessingConfiguration

Describes a data processing configuration.

Type: [ProcessingConfiguration](#) (p. 130) object

Required: No

RequestConfiguration

The configuration of the request sent to the HTTP endpoint specified as the destination.

Type: [HttpEndpointRequestConfiguration](#) (p. 116) object

Required: No

RetryOptions

Describes the retry behavior in case Kinesis Data Firehose is unable to deliver data to the specified HTTP endpoint destination, or if it doesn't receive a valid acknowledgment of receipt from the specified HTTP endpoint destination.

Type: [HttpEndpointRetryOptions](#) (p. 117) object

Required: No

RoleARN

Kinesis Data Firehose uses this IAM role for all the permissions that the delivery stream needs.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: No

S3BackupMode

Describes the S3 bucket backup options for the data that Kinesis Firehose delivers to the HTTP endpoint destination. You can back up all documents (`AllData`) or only the documents that Kinesis Data Firehose could not deliver to the specified HTTP endpoint destination (`FailedDataOnly`).

Type: String

Valid Values: `FailedDataOnly` | `AllData`

Required: No

S3Update

Describes an update for a destination in Amazon S3.

Type: [S3DestinationUpdate](#) (p. 149) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

HttpEndpointRequestConfiguration

The configuration of the HTTP endpoint request.

Contents

CommonAttributes

Describes the metadata sent to the HTTP endpoint destination.

Type: Array of [HttpEndpointCommonAttribute](#) (p. 107) objects

Array Members: Minimum number of 0 items. Maximum number of 50 items.

Required: No

ContentEncoding

Kinesis Data Firehose uses the content encoding to compress the body of a request before sending the request to the destination. For more information, see [Content-Encoding](#) in MDN Web Docs, the official Mozilla documentation.

Type: String

Valid Values: `NONE` | `GZIP`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

HttpEndpointRetryOptions

Describes the retry behavior in case Kinesis Data Firehose is unable to deliver data to the specified HTTP endpoint destination, or if it doesn't receive a valid acknowledgment of receipt from the specified HTTP endpoint destination.

Contents

DurationInSeconds

The total amount of time that Kinesis Data Firehose spends on retries. This duration starts after the initial attempt to send data to the custom destination via HTTPS endpoint fails. It doesn't include the periods during which Kinesis Data Firehose waits for acknowledgment from the specified destination after each attempt.

Type: Integer

Valid Range: Minimum value of 0. Maximum value of 7200.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

InputFormatConfiguration

Specifies the deserializer you want to use to convert the format of the input data. This parameter is required if `Enabled` is set to `true`.

Contents

Deserializer

Specifies which deserializer to use. You can choose either the Apache Hive JSON SerDe or the OpenX JSON SerDe. If both are non-null, the server rejects the request.

Type: [Deserializer](#) (p. 79) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

KinesisStreamSourceConfiguration

The stream and role Amazon Resource Names (ARNs) for a Kinesis data stream used as the source for a delivery stream.

Contents

KinesisStreamARN

The ARN of the source Kinesis data stream. For more information, see [Amazon Kinesis Data Streams ARN Format](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: Yes

RoleARN

The ARN of the role that provides access to the source Kinesis data stream. For more information, see [AWS Identity and Access Management \(IAM\) ARN Format](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

KinesisStreamSourceDescription

Details about a Kinesis data stream used as the source for a Kinesis Data Firehose delivery stream.

Contents

DeliveryStartTimestamp

Kinesis Data Firehose starts retrieving records from the Kinesis data stream starting with this timestamp.

Type: Timestamp

Required: No

KinesisStreamARN

The Amazon Resource Name (ARN) of the source Kinesis data stream. For more information, see [Amazon Kinesis Data Streams ARN Format](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: No

RoleARN

The ARN of the role used by the source Kinesis data stream. For more information, see [AWS Identity and Access Management \(IAM\) ARN Format](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

KMSEncryptionConfig

Describes an encryption key for a destination in Amazon S3.

Contents

AWSKMSKeyARN

The Amazon Resource Name (ARN) of the encryption key. Must belong to the same AWS Region as the destination Amazon S3 bucket. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

OpenXJsonSerDe

The OpenX SerDe. Used by Kinesis Data Firehose for deserializing data, which means converting it from the JSON format in preparation for serializing it to the Parquet or ORC format. This is one of two deserializers you can choose, depending on which one offers the functionality you need. The other option is the native Hive / HCatalog JsonSerDe.

Contents

CaseInsensitive

When set to `true`, which is the default, Kinesis Data Firehose converts JSON keys to lowercase before deserializing them.

Type: Boolean

Required: No

ColumnToJsonKeyMappings

Maps column names to JSON keys that aren't identical to the column names. This is useful when the JSON contains keys that are Hive keywords. For example, `timestamp` is a Hive keyword. If you have a JSON key named `timestamp`, set this parameter to `{"ts": "timestamp"}` to map this key to a column named `ts`.

Type: String to string map

Key Length Constraints: Minimum length of 1. Maximum length of 1024.

Key Pattern: `^\S+$`

Value Length Constraints: Minimum length of 1. Maximum length of 1024.

Value Pattern: `^(?!\\s*$)\\.+`

Required: No

ConvertDotsInJsonKeysToUnderscores

When set to `true`, specifies that the names of the keys include dots and that you want Kinesis Data Firehose to replace them with underscores. This is useful because Apache Hive does not allow dots in column names. For example, if the JSON contains a key whose name is `"a.b"`, you can define the column name to be `"a_b"` when using this option.

The default is `false`.

Type: Boolean

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

OrcSerDe

A serializer to use for converting data to the ORC format before storing it in Amazon S3. For more information, see [Apache ORC](#).

Contents

BlockSizeBytes

The Hadoop Distributed File System (HDFS) block size. This is useful if you intend to copy the data from Amazon S3 to HDFS before querying. The default is 256 MiB and the minimum is 64 MiB. Kinesis Data Firehose uses this value for padding calculations.

Type: Integer

Valid Range: Minimum value of 67108864.

Required: No

BloomFilterColumns

The column names for which you want Kinesis Data Firehose to create bloom filters. The default is `null`.

Type: Array of strings

Length Constraints: Minimum length of 1. Maximum length of 1024.

Pattern: `^\S+$`

Required: No

BloomFilterFalsePositiveProbability

The Bloom filter false positive probability (FPP). The lower the FPP, the bigger the Bloom filter. The default value is 0.05, the minimum is 0, and the maximum is 1.

Type: Double

Valid Range: Minimum value of 0. Maximum value of 1.

Required: No

Compression

The compression code to use over data blocks. The default is `SNAPPY`.

Type: String

Valid Values: `NONE` | `ZLIB` | `SNAPPY`

Required: No

DictionaryKeyThreshold

Represents the fraction of the total number of non-null rows. To turn off dictionary encoding, set this fraction to a number that is less than the number of distinct keys in a dictionary. To always use dictionary encoding, set this threshold to 1.

Type: Double

Valid Range: Minimum value of 0. Maximum value of 1.

Required: No

EnablePadding

Set this to `true` to indicate that you want stripes to be padded to the HDFS block boundaries. This is useful if you intend to copy the data from Amazon S3 to HDFS before querying. The default is `false`.

Type: Boolean

Required: No

FormatVersion

The version of the file to write. The possible values are `v0_11` and `v0_12`. The default is `v0_12`.

Type: String

Valid Values: `v0_11` | `v0_12`

Required: No

PaddingTolerance

A number between 0 and 1 that defines the tolerance for block padding as a decimal fraction of stripe size. The default value is 0.05, which means 5 percent of stripe size.

For the default values of 64 MiB ORC stripes and 256 MiB HDFS blocks, the default block padding tolerance of 5 percent reserves a maximum of 3.2 MiB for padding within the 256 MiB block. In such a case, if the available size within the block is more than 3.2 MiB, a new, smaller stripe is inserted to fit within that space. This ensures that no stripe crosses block boundaries and causes remote reads within a node-local task.

Kinesis Data Firehose ignores this parameter when [OrcSerDe:EnablePadding](#) (p. 125) is `false`.

Type: Double

Valid Range: Minimum value of 0. Maximum value of 1.

Required: No

RowIndexStride

The number of rows between index entries. The default is 10,000 and the minimum is 1,000.

Type: Integer

Valid Range: Minimum value of 1000.

Required: No

StripeSizeBytes

The number of bytes in each stripe. The default is 64 MiB and the minimum is 8 MiB.

Type: Integer

Valid Range: Minimum value of 8388608.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

OutputFormatConfiguration

Specifies the serializer that you want Kinesis Data Firehose to use to convert the format of your data before it writes it to Amazon S3. This parameter is required if `Enabled` is set to `true`.

Contents

Serializer

Specifies which serializer to use. You can choose either the ORC SerDe or the Parquet SerDe. If both are non-null, the server rejects the request.

Type: [Serializer](#) (p. 153) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ParquetSerDe

A serializer to use for converting data to the Parquet format before storing it in Amazon S3. For more information, see [Apache Parquet](#).

Contents

BlockSizeBytes

The Hadoop Distributed File System (HDFS) block size. This is useful if you intend to copy the data from Amazon S3 to HDFS before querying. The default is 256 MiB and the minimum is 64 MiB. Kinesis Data Firehose uses this value for padding calculations.

Type: Integer

Valid Range: Minimum value of 67108864.

Required: No

Compression

The compression code to use over data blocks. The possible values are UNCOMPRESSED, SNAPPY, and GZIP, with the default being SNAPPY. Use SNAPPY for higher decompression speed. Use GZIP if the compression ratio is more important than speed.

Type: String

Valid Values: UNCOMPRESSED | GZIP | SNAPPY

Required: No

EnableDictionaryCompression

Indicates whether to enable dictionary compression.

Type: Boolean

Required: No

MaxPaddingBytes

The maximum amount of padding to apply. This is useful if you intend to copy the data from Amazon S3 to HDFS before querying. The default is 0.

Type: Integer

Valid Range: Minimum value of 0.

Required: No

PageSizeBytes

The Parquet page size. Column chunks are divided into pages. A page is conceptually an indivisible unit (in terms of compression and encoding). The minimum value is 64 KiB and the default is 1 MiB.

Type: Integer

Valid Range: Minimum value of 65536.

Required: No

WriterVersion

Indicates the version of row format to output. The possible values are v1 and v2. The default is v1.

Type: String

Valid Values: `V1` | `V2`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ProcessingConfiguration

Describes a data processing configuration.

Contents

Enabled

Enables or disables data processing.

Type: Boolean

Required: No

Processors

The data processors.

Type: Array of [Processor](#) (p. 131) objects

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

Processor

Describes a data processor.

Contents

Parameters

The processor parameters.

Type: Array of [ProcessorParameter](#) (p. 132) objects

Required: No

Type

The type of processor.

Type: String

Valid Values: `RecordDeAggregation` | `Lambda` | `MetadataExtraction` | `AppendDelimiterToRecord`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

ProcessorParameter

Describes the processor parameter.

Contents

ParameterName

The name of the parameter. Currently the following default values are supported: 3 for `NumberOfRetries`, 60 for the `BufferIntervalInSeconds`, and 3 for the `BufferSizeInMBs`.

Type: String

Valid Values: `LambdaArn` | `NumberOfRetries` | `MetadataExtractionQuery` | `JsonParsingEngine` | `RoleArn` | `BufferSizeInMBs` | `BufferIntervalInSeconds` | `SubRecordType` | `Delimiter`

Required: Yes

ParameterValue

The parameter value.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 5120.

Pattern: `^(?!\\s*$).+`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

PutRecordBatchResponseEntry

Contains the result for an individual record from a [PutRecordBatch](#) (p. 38) request. If the record is successfully added to your delivery stream, it receives a record ID. If the record fails to be added to your delivery stream, the result includes an error code and an error message.

Contents

ErrorCode

The error code for an individual record result.

Type: String

Required: No

ErrorMessage

The error message for an individual record result.

Type: String

Required: No

RecordId

The ID of the record.

Type: String

Length Constraints: Minimum length of 1.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

Record

The unit of data in a delivery stream.

Contents

Data

The data blob, which is base64-encoded when the blob is serialized. The maximum size of the data blob, before base64-encoding, is 1,000 KiB.

Type: Base64-encoded binary data object

Length Constraints: Minimum length of 0. Maximum length of 1024000.

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

RedshiftDestinationConfiguration

Describes the configuration of a destination in Amazon Redshift.

Contents

CloudWatchLoggingOptions

The CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

ClusterJDBCURL

The database connection string.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: jdbc:(redshift|postgresql)://((?!-)[A-Za-z0-9-]{1,63}(?<!--)\.)+redshift\.([a-zA-Z0-9\.]++):\d{1,5}/[a-zA-Z0-9_-\$-]+

Required: Yes

CopyCommand

The COPY command.

Type: [CopyCommand](#) (p. 69) object

Required: Yes

Password

The user password.

Type: String

Length Constraints: Minimum length of 6. Maximum length of 512.

Pattern: .*

Required: Yes

ProcessingConfiguration

The data processing configuration.

Type: [ProcessingConfiguration](#) (p. 130) object

Required: No

RetryOptions

The retry behavior in case Kinesis Data Firehose is unable to deliver documents to Amazon Redshift. Default value is 3600 (60 minutes).

Type: [RedshiftRetryOptions](#) (p. 143) object

Required: No

RoleARN

The Amazon Resource Name (ARN) of the AWS credentials. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: Yes

S3BackupConfiguration

The configuration for backup in Amazon S3.

Type: [S3DestinationConfiguration](#) (p. 145) object

Required: No

S3BackupMode

The Amazon S3 backup mode. After you create a delivery stream, you can update it to enable Amazon S3 backup if it is disabled. If backup is enabled, you can't update the delivery stream to disable it.

Type: String

Valid Values: `Disabled` | `Enabled`

Required: No

S3Configuration

The configuration for the intermediate Amazon S3 location from which Amazon Redshift obtains data. Restrictions are described in the topic for [CreateDeliveryStream](#) (p. 3).

The compression formats `SNAPPY` or `ZIP` cannot be specified in `RedshiftDestinationConfiguration.S3Configuration` because the Amazon Redshift `COPY` operation that reads from the S3 bucket doesn't support these compression formats.

Type: [S3DestinationConfiguration](#) (p. 145) object

Required: Yes

Username

The name of the user.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `.*`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)

- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

RedshiftDestinationDescription

Describes a destination in Amazon Redshift.

Contents

CloudWatchLoggingOptions

The Amazon CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

ClusterJDBCURL

The database connection string.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: jdbc:(redshift|postgresql)://((?!-)[A-Za-z0-9-]{1,63}(?<!--)\.)*redshift\.([a-zA-Z0-9\.]*)?:\d{1,5}/[a-zA-Z0-9_-\$-]+

Required: Yes

CopyCommand

The COPY command.

Type: [CopyCommand](#) (p. 69) object

Required: Yes

ProcessingConfiguration

The data processing configuration.

Type: [ProcessingConfiguration](#) (p. 130) object

Required: No

RetryOptions

The retry behavior in case Kinesis Data Firehose is unable to deliver documents to Amazon Redshift. Default value is 3600 (60 minutes).

Type: [RedshiftRetryOptions](#) (p. 143) object

Required: No

RoleARN

The Amazon Resource Name (ARN) of the AWS credentials. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: arn:.*

Required: Yes

S3BackupDescription

The configuration for backup in Amazon S3.

Type: [S3DestinationDescription](#) (p. 147) object

Required: No

S3BackupMode

The Amazon S3 backup mode.

Type: String

Valid Values: `Disabled` | `Enabled`

Required: No

S3DestinationDescription

The Amazon S3 destination.

Type: [S3DestinationDescription](#) (p. 147) object

Required: Yes

Username

The name of the user.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `. *`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

RedshiftDestinationUpdate

Describes an update for a destination in Amazon Redshift.

Contents

CloudWatchLoggingOptions

The Amazon CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

ClusterJDBCURL

The database connection string.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: jdbc:(redshift|postgresql)://((?!-)[A-Za-z0-9-]{1,63}(?<!--)\.)+redshift\.([a-zA-Z0-9\.]++):\d{1,5}/[a-zA-Z0-9_-\$-]+

Required: No

CopyCommand

The COPY command.

Type: [CopyCommand](#) (p. 69) object

Required: No

Password

The user password.

Type: String

Length Constraints: Minimum length of 6. Maximum length of 512.

Pattern: .*

Required: No

ProcessingConfiguration

The data processing configuration.

Type: [ProcessingConfiguration](#) (p. 130) object

Required: No

RetryOptions

The retry behavior in case Kinesis Data Firehose is unable to deliver documents to Amazon Redshift. Default value is 3600 (60 minutes).

Type: [RedshiftRetryOptions](#) (p. 143) object

Required: No

RoleARN

The Amazon Resource Name (ARN) of the AWS credentials. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: No

S3BackupMode

You can update a delivery stream to enable Amazon S3 backup if it is disabled. If backup is enabled, you can't update the delivery stream to disable it.

Type: String

Valid Values: `Disabled` | `Enabled`

Required: No

S3BackupUpdate

The Amazon S3 destination for backup.

Type: [S3DestinationUpdate](#) (p. 149) object

Required: No

S3Update

The Amazon S3 destination.

The compression formats `SNAPPY` or `ZIP` cannot be specified in `RedshiftDestinationUpdate.S3Update` because the Amazon Redshift `COPY` operation that reads from the S3 bucket doesn't support these compression formats.

Type: [S3DestinationUpdate](#) (p. 149) object

Required: No

Username

The name of the user.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `.*`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)

- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

RedshiftRetryOptions

Configures retry behavior in case Kinesis Data Firehose is unable to deliver documents to Amazon Redshift.

Contents

DurationInSeconds

The length of time during which Kinesis Data Firehose retries delivery after a failure, starting from the initial request and including the first attempt. The default value is 3600 seconds (60 minutes). Kinesis Data Firehose does not retry if the value of `DurationInSeconds` is 0 (zero) or if the first delivery attempt takes longer than the current value.

Type: Integer

Valid Range: Minimum value of 0. Maximum value of 7200.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

RetryOptions

The retry behavior in case Kinesis Data Firehose is unable to deliver data to an Amazon S3 prefix.

Contents

DurationInSeconds

The period of time during which Kinesis Data Firehose retries to deliver data to the specified Amazon S3 prefix.

Type: Integer

Valid Range: Minimum value of 0. Maximum value of 7200.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

S3DestinationConfiguration

Describes the configuration of a destination in Amazon S3.

Contents

BucketARN

The ARN of the S3 bucket. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: `arn:.*`

Required: Yes

BufferingHints

The buffering option. If no value is specified, `BufferingHints` object default values are used.

Type: [BufferingHints](#) (p. 67) object

Required: No

CloudWatchLoggingOptions

The CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

CompressionFormat

The compression format. If no value is specified, the default is `UNCOMPRESSED`.

The compression formats `SNAPPY` or `ZIP` cannot be specified for Amazon Redshift destinations because they are not supported by the Amazon Redshift `COPY` operation that reads from the S3 bucket.

Type: String

Valid Values: `UNCOMPRESSED` | `GZIP` | `ZIP` | `Snappy` | `HADOOP_SNAPPY`

Required: No

EncryptionConfiguration

The encryption configuration. If no value is specified, the default is no encryption.

Type: [EncryptionConfiguration](#) (p. 94) object

Required: No

ErrorOutputPrefix

A prefix that Kinesis Data Firehose evaluates and adds to failed records before writing them to S3. This prefix appears immediately following the bucket name. For information about how to specify this prefix, see [Custom Prefixes for Amazon S3 Objects](#).

Type: String

Length Constraints: Minimum length of 0. Maximum length of 1024.

Pattern: .*

Required: No

Prefix

The "YYYY/MM/DD/HH" time format prefix is automatically used for delivered Amazon S3 files. You can also specify a custom prefix, as described in [Custom Prefixes for Amazon S3 Objects](#).

Type: String

Length Constraints: Minimum length of 0. Maximum length of 1024.

Pattern: .*

Required: No

RoleARN

The Amazon Resource Name (ARN) of the AWS credentials. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: arn:.*

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

S3DestinationDescription

Describes a destination in Amazon S3.

Contents

BucketARN

The ARN of the S3 bucket. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: `arn:.*`

Required: Yes

BufferingHints

The buffering option. If no value is specified, `BufferingHints` object default values are used.

Type: [BufferingHints](#) (p. 67) object

Required: Yes

CloudWatchLoggingOptions

The Amazon CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

CompressionFormat

The compression format. If no value is specified, the default is `UNCOMPRESSED`.

Type: String

Valid Values: `UNCOMPRESSED` | `GZIP` | `ZIP` | `Snappy` | `HADOOP_SNAPPY`

Required: Yes

EncryptionConfiguration

The encryption configuration. If no value is specified, the default is no encryption.

Type: [EncryptionConfiguration](#) (p. 94) object

Required: Yes

ErrorOutputPrefix

A prefix that Kinesis Data Firehose evaluates and adds to failed records before writing them to S3. This prefix appears immediately following the bucket name. For information about how to specify this prefix, see [Custom Prefixes for Amazon S3 Objects](#).

Type: String

Length Constraints: Minimum length of 0. Maximum length of 1024.

Pattern: .*

Required: No

Prefix

The "YYYY/MM/DD/HH" time format prefix is automatically used for delivered Amazon S3 files. You can also specify a custom prefix, as described in [Custom Prefixes for Amazon S3 Objects](#).

Type: String

Length Constraints: Minimum length of 0. Maximum length of 1024.

Pattern: .*

Required: No

RoleARN

The Amazon Resource Name (ARN) of the AWS credentials. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: arn:.*

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

S3DestinationUpdate

Describes an update for a destination in Amazon S3.

Contents

BucketARN

The ARN of the S3 bucket. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: `arn:.*`

Required: No

BufferingHints

The buffering option. If no value is specified, `BufferingHints` object default values are used.

Type: [BufferingHints](#) (p. 67) object

Required: No

CloudWatchLoggingOptions

The CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

CompressionFormat

The compression format. If no value is specified, the default is `UNCOMPRESSED`.

The compression formats `SNAPPY` or `ZIP` cannot be specified for Amazon Redshift destinations because they are not supported by the Amazon Redshift `COPY` operation that reads from the S3 bucket.

Type: String

Valid Values: `UNCOMPRESSED` | `GZIP` | `ZIP` | `Snappy` | `HADOOP_SNAPPY`

Required: No

EncryptionConfiguration

The encryption configuration. If no value is specified, the default is no encryption.

Type: [EncryptionConfiguration](#) (p. 94) object

Required: No

ErrorOutputPrefix

A prefix that Kinesis Data Firehose evaluates and adds to failed records before writing them to S3. This prefix appears immediately following the bucket name. For information about how to specify this prefix, see [Custom Prefixes for Amazon S3 Objects](#).

Type: String

Length Constraints: Minimum length of 0. Maximum length of 1024.

Pattern: .*

Required: No

Prefix

The "YYYY/MM/DD/HH" time format prefix is automatically used for delivered Amazon S3 files. You can also specify a custom prefix, as described in [Custom Prefixes for Amazon S3 Objects](#).

Type: String

Length Constraints: Minimum length of 0. Maximum length of 1024.

Pattern: .*

Required: No

RoleARN

The Amazon Resource Name (ARN) of the AWS credentials. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: arn:.*

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

SchemaConfiguration

Specifies the schema to which you want Kinesis Data Firehose to configure your data before it writes it to Amazon S3. This parameter is required if `Enabled` is set to `true`.

Contents

CatalogId

The ID of the AWS Glue Data Catalog. If you don't supply this, the AWS account ID is used by default.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1024.

Pattern: `^\S+$`

Required: No

DatabaseName

Specifies the name of the AWS Glue database that contains the schema for the output data.

Important

If the `SchemaConfiguration` request parameter is used as part of invoking the `CreateDeliveryStream` API, then the `DatabaseName` property is required and its value must be specified.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1024.

Pattern: `^\S+$`

Required: No

Region

If you don't specify an AWS Region, the default is the current Region.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1024.

Pattern: `^\S+$`

Required: No

RoleARN

The role that Kinesis Data Firehose can use to access AWS Glue. This role must be in the same account you use for Kinesis Data Firehose. Cross-account roles aren't allowed.

Important

If the `SchemaConfiguration` request parameter is used as part of invoking the `CreateDeliveryStream` API, then the `RoleARN` property is required and its value must be specified.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1024.

Pattern: `^\S+$`

Required: No

TableName

Specifies the AWS Glue table that contains the column information that constitutes your data schema.

Important

If the `SchemaConfiguration` request parameter is used as part of invoking the `CreateDeliveryStream` API, then the `TableName` property is required and its value must be specified.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1024.

Pattern: `^\S+$`

Required: No

VersionId

Specifies the table version for the output data schema. If you don't specify this version ID, or if you set it to `LATEST`, Kinesis Data Firehose uses the most recent version. This means that any updates to the table are automatically picked up.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1024.

Pattern: `^\S+$`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

Serializer

The serializer that you want Kinesis Data Firehose to use to convert data to the target format before writing it to Amazon S3. Kinesis Data Firehose supports two types of serializers: the [ORC SerDe](#) and the [Parquet SerDe](#).

Contents

OrcSerDe

A serializer to use for converting data to the ORC format before storing it in Amazon S3. For more information, see [Apache ORC](#).

Type: [OrcSerDe](#) (p. 124) object

Required: No

ParquetSerDe

A serializer to use for converting data to the Parquet format before storing it in Amazon S3. For more information, see [Apache Parquet](#).

Type: [ParquetSerDe](#) (p. 128) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

SourceDescription

Details about a Kinesis data stream used as the source for a Kinesis Data Firehose delivery stream.

Contents

KinesisStreamSourceDescription

The [KinesisStreamSourceDescription](#) (p. 120) value for the source Kinesis data stream.

Type: [KinesisStreamSourceDescription](#) (p. 120) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

SplunkDestinationConfiguration

Describes the configuration of a destination in Splunk.

Contents

CloudWatchLoggingOptions

The Amazon CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

HECAcknowledgmentTimeoutInSeconds

The amount of time that Kinesis Data Firehose waits to receive an acknowledgment from Splunk after it sends it data. At the end of the timeout period, Kinesis Data Firehose either tries to send the data again or considers it an error, based on your retry settings.

Type: Integer

Valid Range: Minimum value of 180. Maximum value of 600.

Required: No

HECEndpoint

The HTTP Event Collector (HEC) endpoint to which Kinesis Data Firehose sends your data.

Type: String

Length Constraints: Minimum length of 0. Maximum length of 2048.

Pattern: . *

Required: Yes

HECEndpointType

This type can be either "Raw" or "Event."

Type: String

Valid Values: Raw | Event

Required: Yes

HECToken

This is a GUID that you obtain from your Splunk cluster when you create a new HEC endpoint.

Type: String

Length Constraints: Minimum length of 0. Maximum length of 2048.

Pattern: . *

Required: Yes

ProcessingConfiguration

The data processing configuration.

Type: [ProcessingConfiguration](#) (p. 130) object

Required: No

RetryOptions

The retry behavior in case Kinesis Data Firehose is unable to deliver data to Splunk, or if it doesn't receive an acknowledgment of receipt from Splunk.

Type: [SplunkRetryOptions](#) (p. 161) object

Required: No

S3BackupMode

Defines how documents should be delivered to Amazon S3. When set to `FailedEventsOnly`, Kinesis Data Firehose writes any data that could not be indexed to the configured Amazon S3 destination. When set to `AllEvents`, Kinesis Data Firehose delivers all incoming records to Amazon S3, and also writes failed documents to Amazon S3. The default value is `FailedEventsOnly`.

You can update this backup mode from `FailedEventsOnly` to `AllEvents`. You can't update it from `AllEvents` to `FailedEventsOnly`.

Type: String

Valid Values: `FailedEventsOnly` | `AllEvents`

Required: No

S3Configuration

The configuration for the backup Amazon S3 location.

Type: [S3DestinationConfiguration](#) (p. 145) object

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

SplunkDestinationDescription

Describes a destination in Splunk.

Contents

CloudWatchLoggingOptions

The Amazon CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

HECAcknowledgmentTimeoutInSeconds

The amount of time that Kinesis Data Firehose waits to receive an acknowledgment from Splunk after it sends it data. At the end of the timeout period, Kinesis Data Firehose either tries to send the data again or considers it an error, based on your retry settings.

Type: Integer

Valid Range: Minimum value of 180. Maximum value of 600.

Required: No

HECEndpoint

The HTTP Event Collector (HEC) endpoint to which Kinesis Data Firehose sends your data.

Type: String

Length Constraints: Minimum length of 0. Maximum length of 2048.

Pattern: . *

Required: No

HECEndpointType

This type can be either "Raw" or "Event."

Type: String

Valid Values: Raw | Event

Required: No

HECToken

A GUID you obtain from your Splunk cluster when you create a new HEC endpoint.

Type: String

Length Constraints: Minimum length of 0. Maximum length of 2048.

Pattern: . *

Required: No

ProcessingConfiguration

The data processing configuration.

Type: [ProcessingConfiguration](#) (p. 130) object

Required: No

RetryOptions

The retry behavior in case Kinesis Data Firehose is unable to deliver data to Splunk or if it doesn't receive an acknowledgment of receipt from Splunk.

Type: [SplunkRetryOptions](#) (p. 161) object

Required: No

S3BackupMode

Defines how documents should be delivered to Amazon S3. When set to `FailedDocumentsOnly`, Kinesis Data Firehose writes any data that could not be indexed to the configured Amazon S3 destination. When set to `AllDocuments`, Kinesis Data Firehose delivers all incoming records to Amazon S3, and also writes failed documents to Amazon S3. Default value is `FailedDocumentsOnly`.

Type: String

Valid Values: `FailedEventsOnly` | `AllEvents`

Required: No

S3DestinationDescription

The Amazon S3 destination.>

Type: [S3DestinationDescription](#) (p. 147) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

SplunkDestinationUpdate

Describes an update for a destination in Splunk.

Contents

CloudWatchLoggingOptions

The Amazon CloudWatch logging options for your delivery stream.

Type: [CloudWatchLoggingOptions](#) (p. 68) object

Required: No

HECAcknowledgmentTimeoutInSeconds

The amount of time that Kinesis Data Firehose waits to receive an acknowledgment from Splunk after it sends data. At the end of the timeout period, Kinesis Data Firehose either tries to send the data again or considers it an error, based on your retry settings.

Type: Integer

Valid Range: Minimum value of 180. Maximum value of 600.

Required: No

HECEndpoint

The HTTP Event Collector (HEC) endpoint to which Kinesis Data Firehose sends your data.

Type: String

Length Constraints: Minimum length of 0. Maximum length of 2048.

Pattern: . *

Required: No

HECEndpointType

This type can be either "Raw" or "Event."

Type: String

Valid Values: Raw | Event

Required: No

HECToken

A GUID that you obtain from your Splunk cluster when you create a new HEC endpoint.

Type: String

Length Constraints: Minimum length of 0. Maximum length of 2048.

Pattern: . *

Required: No

ProcessingConfiguration

The data processing configuration.

Type: [ProcessingConfiguration](#) (p. 130) object

Required: No

RetryOptions

The retry behavior in case Kinesis Data Firehose is unable to deliver data to Splunk or if it doesn't receive an acknowledgment of receipt from Splunk.

Type: [SplunkRetryOptions](#) (p. 161) object

Required: No

S3BackupMode

Specifies how you want Kinesis Data Firehose to back up documents to Amazon S3. When set to `FailedDocumentsOnly`, Kinesis Data Firehose writes any data that could not be indexed to the configured Amazon S3 destination. When set to `AllEvents`, Kinesis Data Firehose delivers all incoming records to Amazon S3, and also writes failed documents to Amazon S3. The default value is `FailedEventsOnly`.

You can update this backup mode from `FailedEventsOnly` to `AllEvents`. You can't update it from `AllEvents` to `FailedEventsOnly`.

Type: String

Valid Values: `FailedEventsOnly` | `AllEvents`

Required: No

S3Update

Your update to the configuration of the backup Amazon S3 location.

Type: [S3DestinationUpdate](#) (p. 149) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

SplunkRetryOptions

Configures retry behavior in case Kinesis Data Firehose is unable to deliver documents to Splunk, or if it doesn't receive an acknowledgment from Splunk.

Contents

DurationInSeconds

The total amount of time that Kinesis Data Firehose spends on retries. This duration starts after the initial attempt to send data to Splunk fails. It doesn't include the periods during which Kinesis Data Firehose waits for acknowledgment from Splunk after each attempt.

Type: Integer

Valid Range: Minimum value of 0. Maximum value of 7200.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

Tag

Metadata that you can assign to a delivery stream, consisting of a key-value pair.

Contents

Key

A unique identifier for the tag. Maximum length: 128 characters. Valid characters: Unicode letters, digits, white space, _ . / = + - % @

Type: String

Length Constraints: Minimum length of 1. Maximum length of 128.

Pattern: `^(?!aws:)[\p{L}\p{Z}\p{N}_.:/=+\-@%]*$`

Required: Yes

Value

An optional string, which you can use to describe or define the tag. Maximum length: 256 characters. Valid characters: Unicode letters, digits, white space, _ . / = + - % @

Type: String

Length Constraints: Minimum length of 0. Maximum length of 256.

Pattern: `^[\p{L}\p{Z}\p{N}_.:/=+\-@%]*$`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

VpcConfiguration

The details of the VPC of the Amazon ES destination.

Contents

RoleARN

The ARN of the IAM role that you want the delivery stream to use to create endpoints in the destination VPC. You can use your existing Kinesis Data Firehose delivery role or you can specify a new role. In either case, make sure that the role trusts the Kinesis Data Firehose service principal and that it grants the following permissions:

- `ec2:DescribeVpcs`
- `ec2:DescribeVpcAttribute`
- `ec2:DescribeSubnets`
- `ec2:DescribeSecurityGroups`
- `ec2:DescribeNetworkInterfaces`
- `ec2:CreateNetworkInterface`
- `ec2:CreateNetworkInterfacePermission`
- `ec2>DeleteNetworkInterface`

If you revoke these permissions after you create the delivery stream, Kinesis Data Firehose can't scale out by creating more ENIs when necessary. You might therefore see a degradation in performance.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: Yes

SecurityGroupIds

The IDs of the security groups that you want Kinesis Data Firehose to use when it creates ENIs in the VPC of the Amazon ES destination. You can use the same security group that the Amazon ES domain uses or different ones. If you specify different security groups here, ensure that they allow outbound HTTPS traffic to the Amazon ES domain's security group. Also ensure that the Amazon ES domain's security group allows HTTPS traffic from the security groups specified here. If you use the same security group for both your delivery stream and the Amazon ES domain, make sure the security group inbound rule allows HTTPS traffic. For more information about security group rules, see [Security group rules](#) in the Amazon VPC documentation.

Type: Array of strings

Array Members: Minimum number of 1 item. Maximum number of 5 items.

Length Constraints: Minimum length of 1. Maximum length of 1024.

Pattern: `^\S+$`

Required: Yes

SubnetIds

The IDs of the subnets that you want Kinesis Data Firehose to use to create ENIs in the VPC of the Amazon ES destination. Make sure that the routing tables and inbound and outbound rules

allow traffic to flow from the subnets whose IDs are specified here to the subnets that have the destination Amazon ES endpoints. Kinesis Data Firehose creates at least one ENI in each of the subnets that are specified here. Do not delete or modify these ENIs.

The number of ENIs that Kinesis Data Firehose creates in the subnets specified here scales up and down automatically based on throughput. To enable Kinesis Data Firehose to scale up the number of ENIs to match throughput, ensure that you have sufficient quota. To help you calculate the quota you need, assume that Kinesis Data Firehose can create up to three ENIs for this delivery stream for each of the subnets specified here. For more information about ENI quota, see [Network Interfaces](#) in the Amazon VPC Quotas topic.

Type: Array of strings

Array Members: Minimum number of 1 item. Maximum number of 16 items.

Length Constraints: Minimum length of 1. Maximum length of 1024.

Pattern: `^\S+$`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

VpcConfigurationDescription

The details of the VPC of the Amazon ES destination.

Contents

RoleARN

The ARN of the IAM role that the delivery stream uses to create endpoints in the destination VPC. You can use your existing Kinesis Data Firehose delivery role or you can specify a new role. In either case, make sure that the role trusts the Kinesis Data Firehose service principal and that it grants the following permissions:

- `ec2:DescribeVpcs`
- `ec2:DescribeVpcAttribute`
- `ec2:DescribeSubnets`
- `ec2:DescribeSecurityGroups`
- `ec2:DescribeNetworkInterfaces`
- `ec2:CreateNetworkInterface`
- `ec2:CreateNetworkInterfacePermission`
- `ec2>DeleteNetworkInterface`

If you revoke these permissions after you create the delivery stream, Kinesis Data Firehose can't scale out by creating more ENIs when necessary. You might therefore see a degradation in performance.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 512.

Pattern: `arn:.*`

Required: Yes

SecurityGroupIds

The IDs of the security groups that Kinesis Data Firehose uses when it creates ENIs in the VPC of the Amazon ES destination. You can use the same security group that the Amazon ES domain uses or different ones. If you specify different security groups, ensure that they allow outbound HTTPS traffic to the Amazon ES domain's security group. Also ensure that the Amazon ES domain's security group allows HTTPS traffic from the security groups specified here. If you use the same security group for both your delivery stream and the Amazon ES domain, make sure the security group inbound rule allows HTTPS traffic. For more information about security group rules, see [Security group rules](#) in the Amazon VPC documentation.

Type: Array of strings

Array Members: Minimum number of 1 item. Maximum number of 5 items.

Length Constraints: Minimum length of 1. Maximum length of 1024.

Pattern: `^\S+$`

Required: Yes

SubnetIds

The IDs of the subnets that Kinesis Data Firehose uses to create ENIs in the VPC of the Amazon ES destination. Make sure that the routing tables and inbound and outbound rules allow traffic to flow

from the subnets whose IDs are specified here to the subnets that have the destination Amazon ES endpoints. Kinesis Data Firehose creates at least one ENI in each of the subnets that are specified here. Do not delete or modify these ENIs.

The number of ENIs that Kinesis Data Firehose creates in the subnets specified here scales up and down automatically based on throughput. To enable Kinesis Data Firehose to scale up the number of ENIs to match throughput, ensure that you have sufficient quota. To help you calculate the quota you need, assume that Kinesis Data Firehose can create up to three ENIs for this delivery stream for each of the subnets specified here. For more information about ENI quota, see [Network Interfaces](#) in the Amazon VPC Quotas topic.

Type: Array of strings

Array Members: Minimum number of 1 item. Maximum number of 16 items.

Length Constraints: Minimum length of 1. Maximum length of 1024.

Pattern: `^\S+$`

Required: Yes

VpcId

The ID of the Amazon ES destination's VPC.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1024.

Pattern: `^\S+$`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java V2](#)
- [AWS SDK for Ruby V3](#)

Common Errors

This section lists the errors common to the API actions of all AWS services. For errors specific to an API action for this service, see the topic for that API action.

AccessDeniedException

You do not have sufficient access to perform this action.

HTTP Status Code: 400

IncompleteSignature

The request signature does not conform to AWS standards.

HTTP Status Code: 400

InternalFailure

The request processing has failed because of an unknown error, exception or failure.

HTTP Status Code: 500

InvalidAction

The action or operation requested is invalid. Verify that the action is typed correctly.

HTTP Status Code: 400

InvalidClientTokenId

The X.509 certificate or AWS access key ID provided does not exist in our records.

HTTP Status Code: 403

InvalidParameterCombination

Parameters that must not be used together were used together.

HTTP Status Code: 400

InvalidParameterValue

An invalid or out-of-range value was supplied for the input parameter.

HTTP Status Code: 400

InvalidQueryParameter

The AWS query string is malformed or does not adhere to AWS standards.

HTTP Status Code: 400

MalformedQueryString

The query string contains a syntax error.

HTTP Status Code: 404

MissingAction

The request is missing an action or a required parameter.

HTTP Status Code: 400

MissingAuthenticationToken

The request must contain either a valid (registered) AWS access key ID or X.509 certificate.

HTTP Status Code: 403

MissingParameter

A required parameter for the specified action is not supplied.

HTTP Status Code: 400

NotAuthorized

You do not have permission to perform this action.

HTTP Status Code: 400

OptInRequired

The AWS access key ID needs a subscription for the service.

HTTP Status Code: 403

RequestExpired

The request reached the service more than 15 minutes after the date stamp on the request or more than 15 minutes after the request expiration date (such as for pre-signed URLs), or the date stamp on the request is more than 15 minutes in the future.

HTTP Status Code: 400

ServiceUnavailable

The request has failed due to a temporary failure of the server.

HTTP Status Code: 503

ThrottlingException

The request was denied due to request throttling.

HTTP Status Code: 400

ValidationError

The input fails to satisfy the constraints specified by an AWS service.

HTTP Status Code: 400